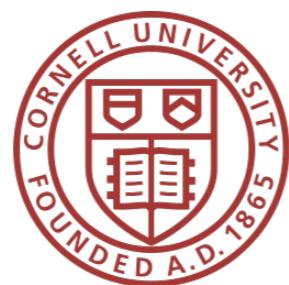


Query Optimization for Data Analysis

Immanuel Trummer



Cornell University



Astronomy

Biology

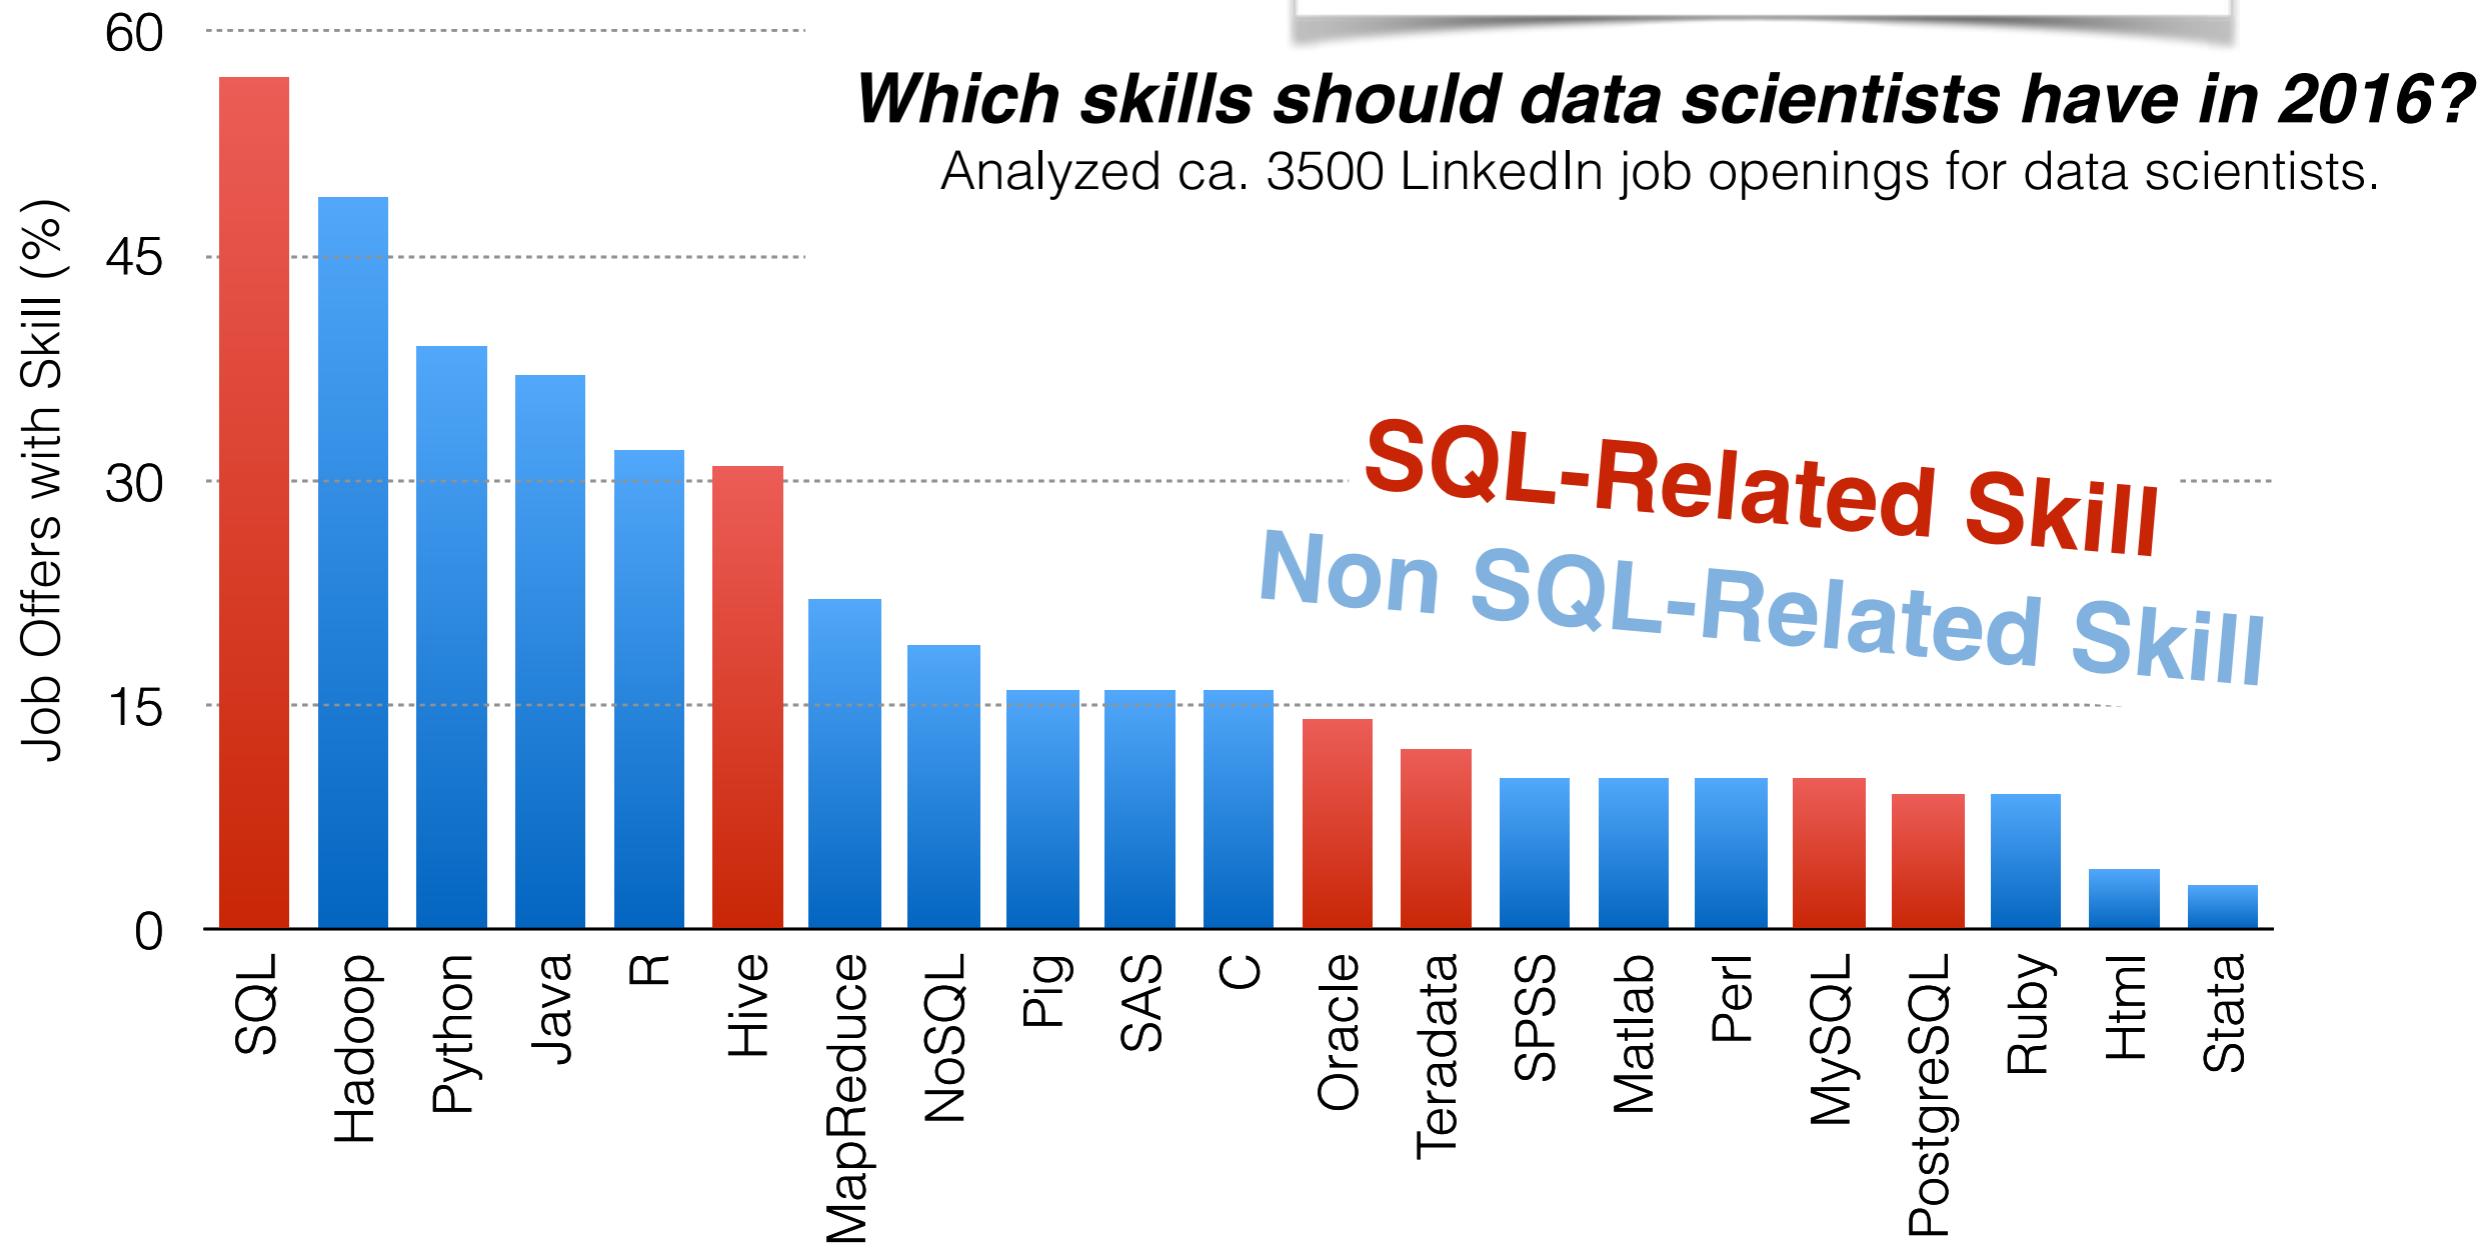
Physics

Data Analysis

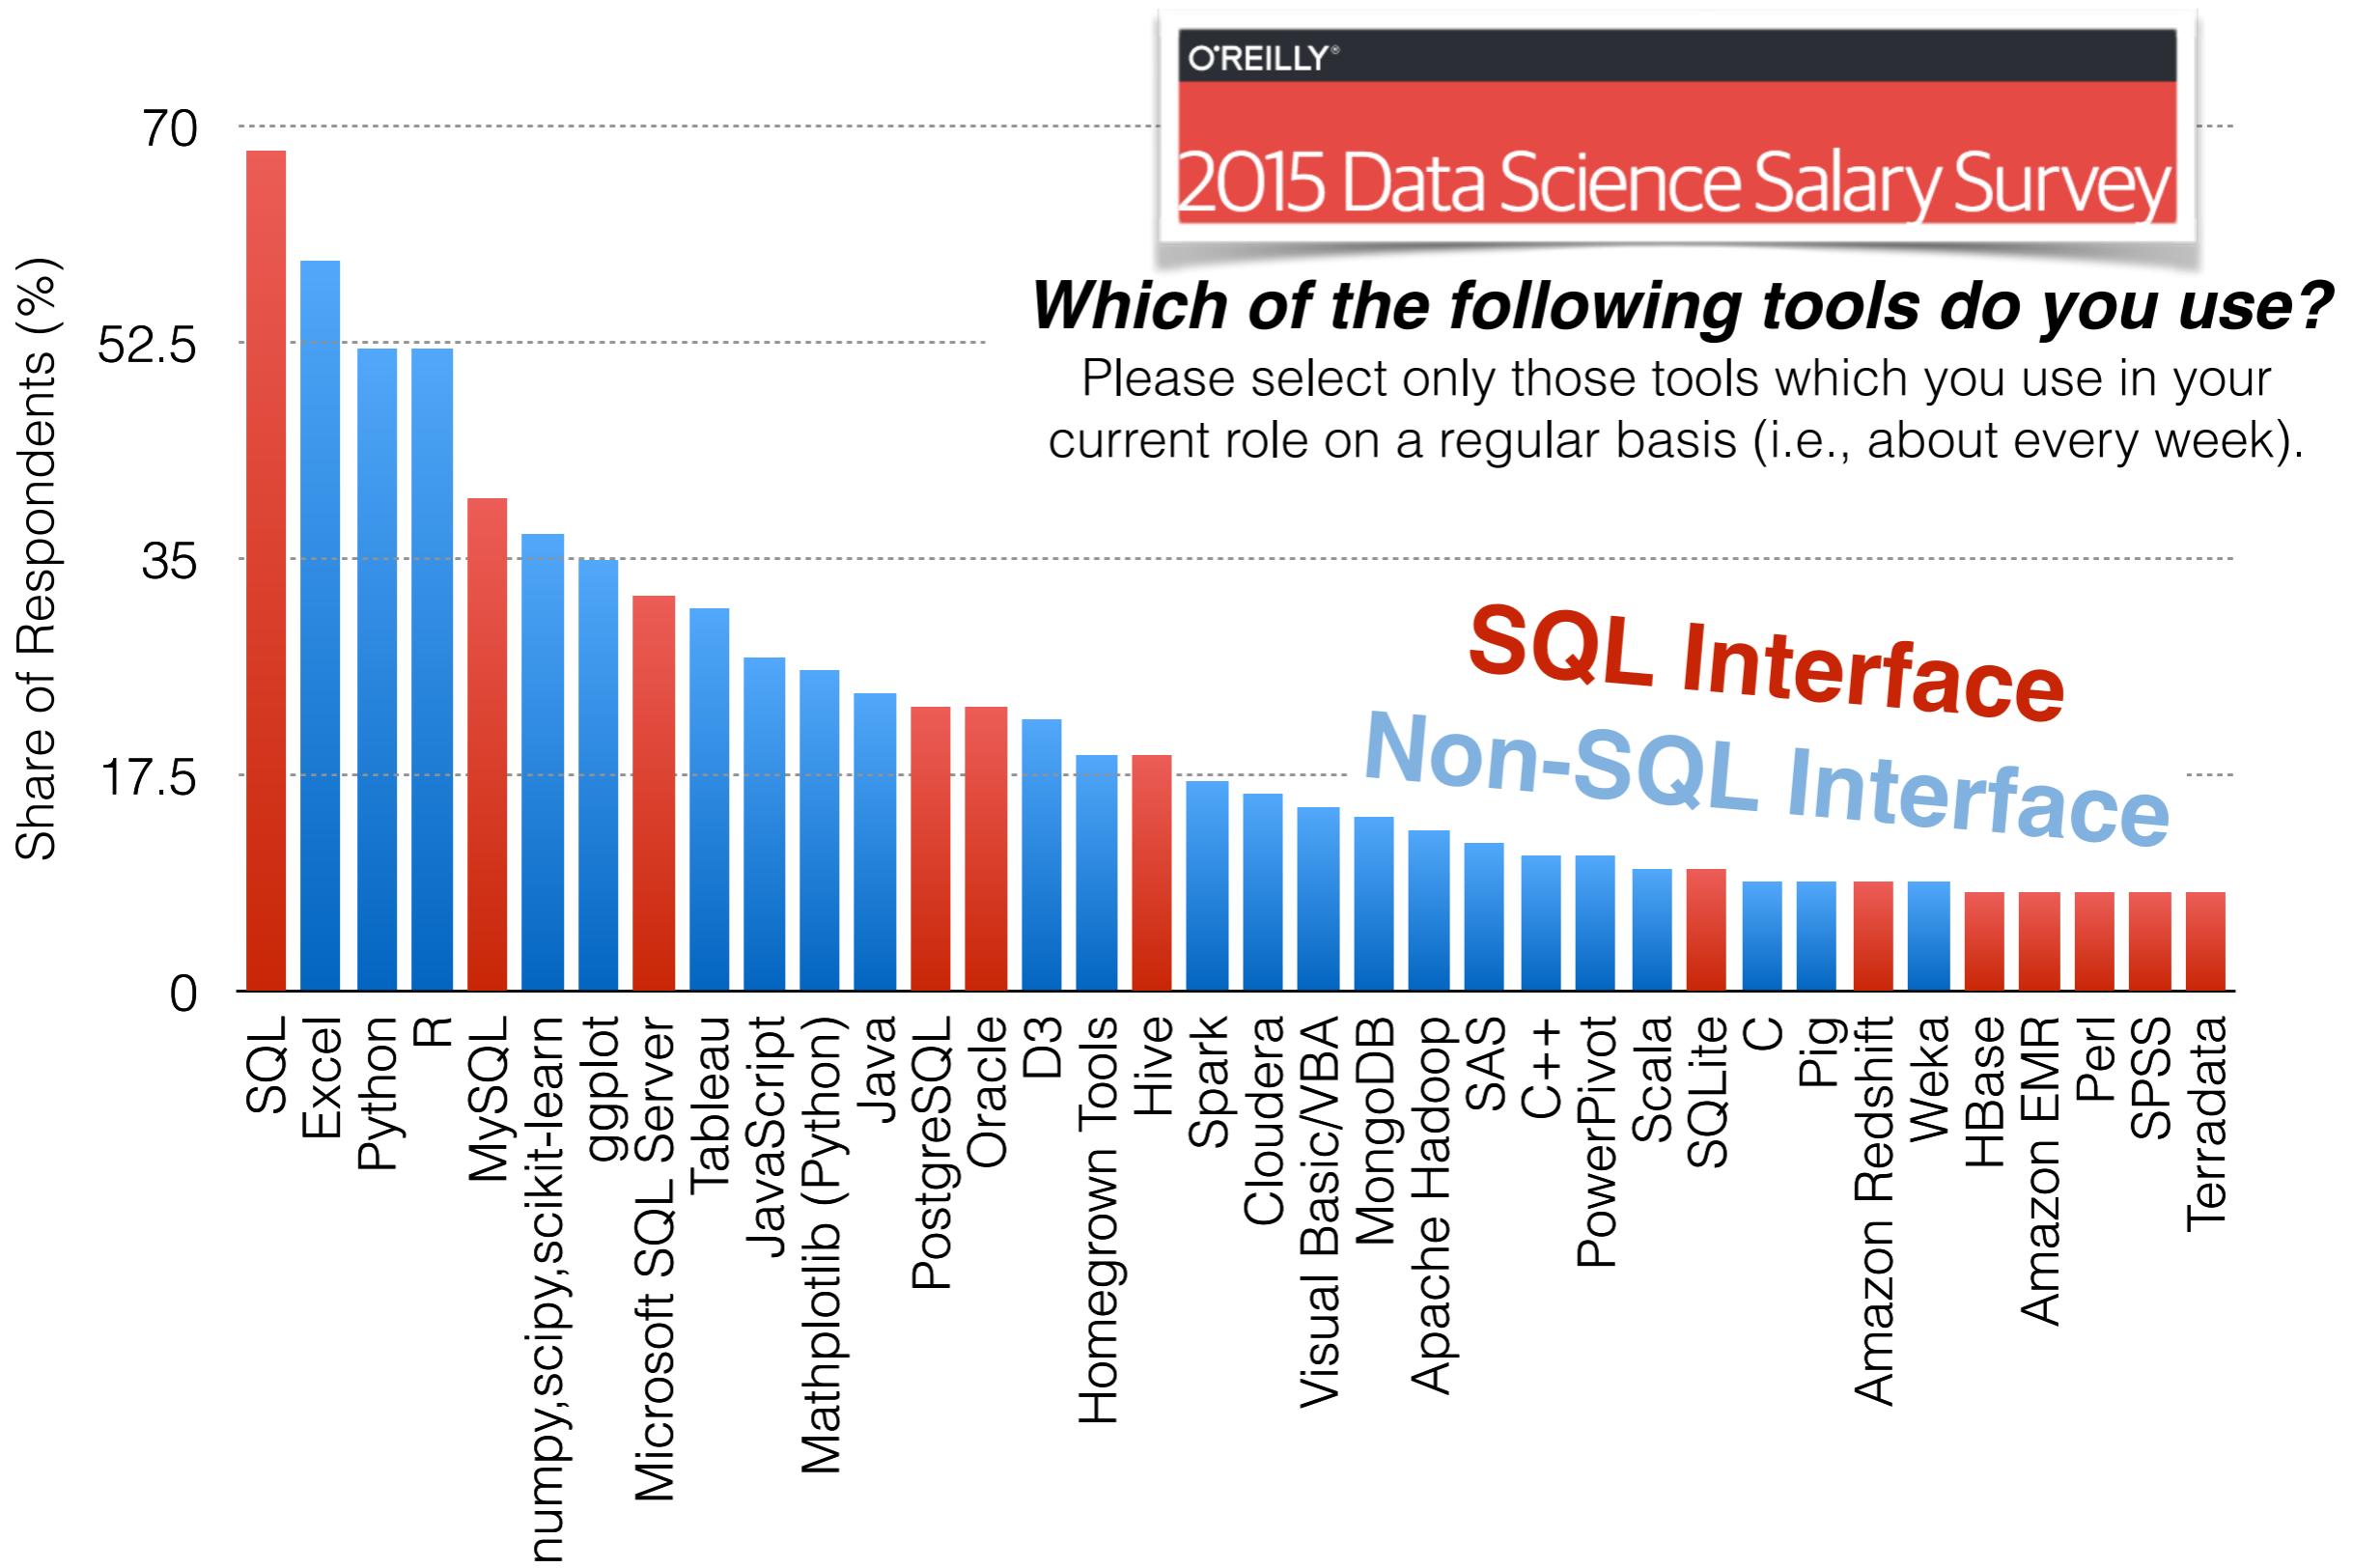
Retail

Politics

Popular Data Analysis Skills



Popular Data Analysis Tools



Classical Query Optimization

Tuning Knobs

*Operation order
Operators*

Problem Model

Query



**Query
Optimization**

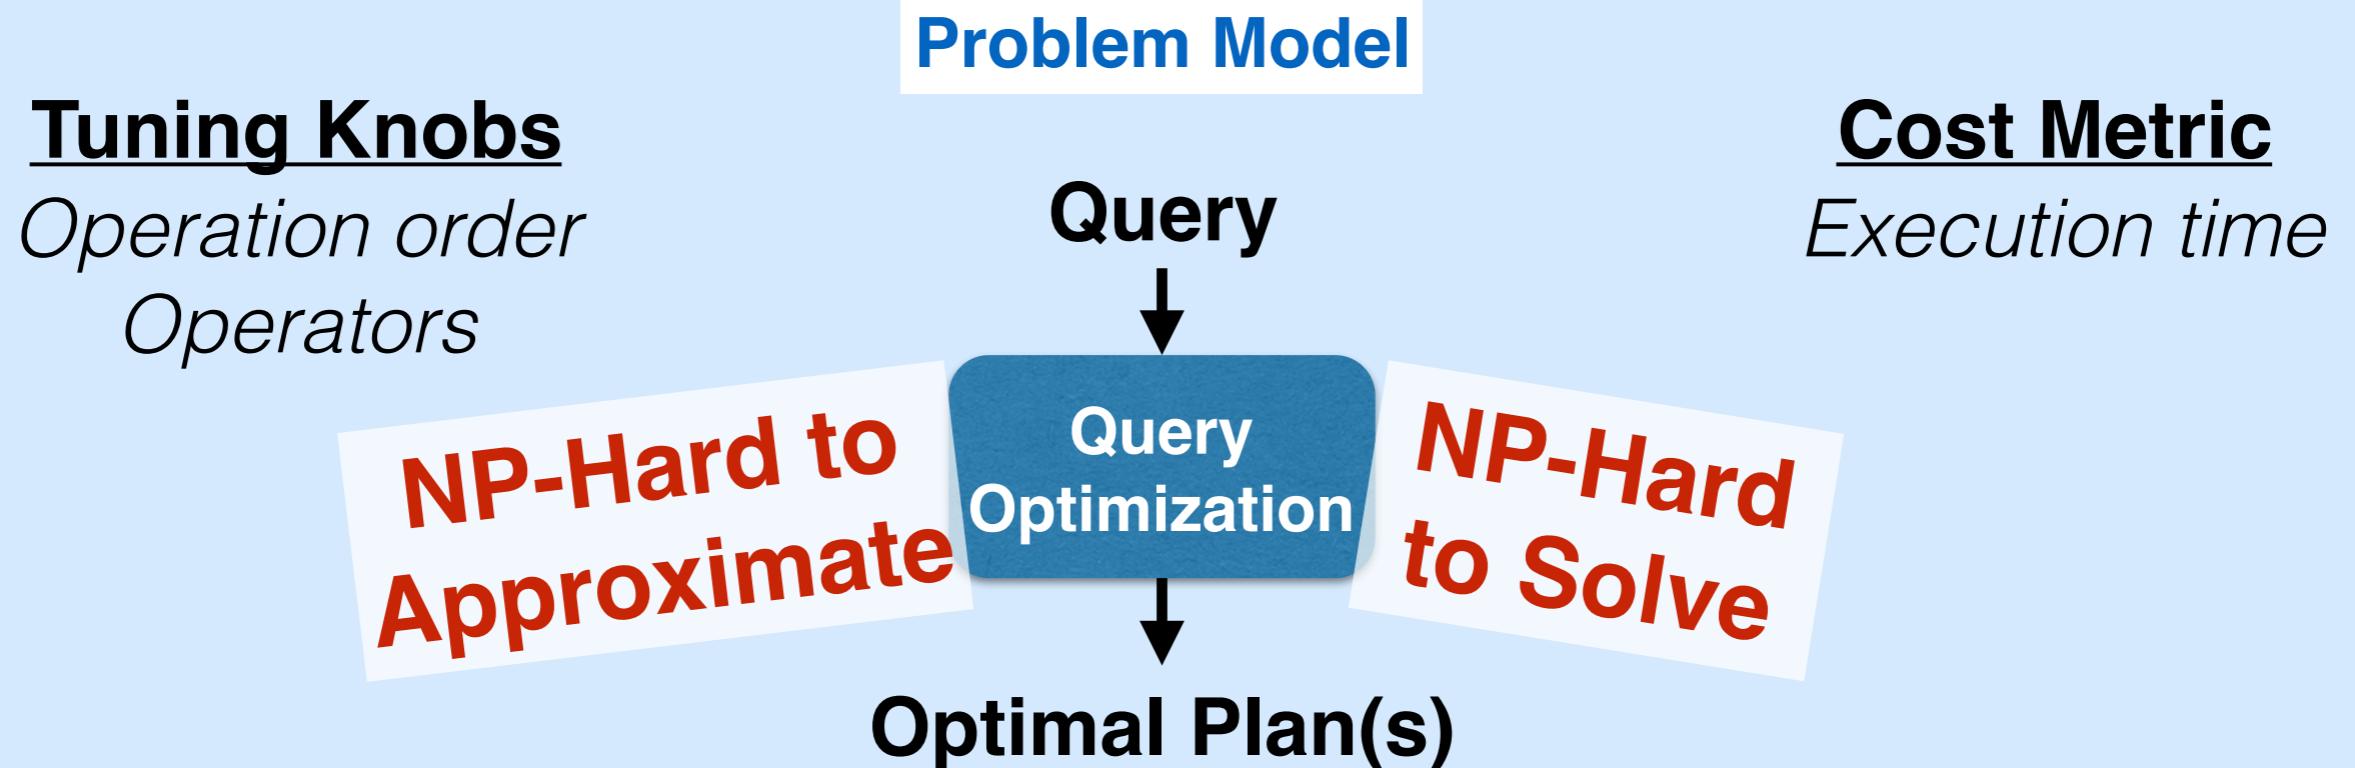


Cost Metric

Execution time

Optimal Plan(s)

Classical Query Optimization



Classical Query Optimization

Fixed processing platform,
few operators available



Context in 1979

Optimization Context

Problem Model

Tuning Knobs

Operation order

Operators

Cost Metric

Execution time

Query



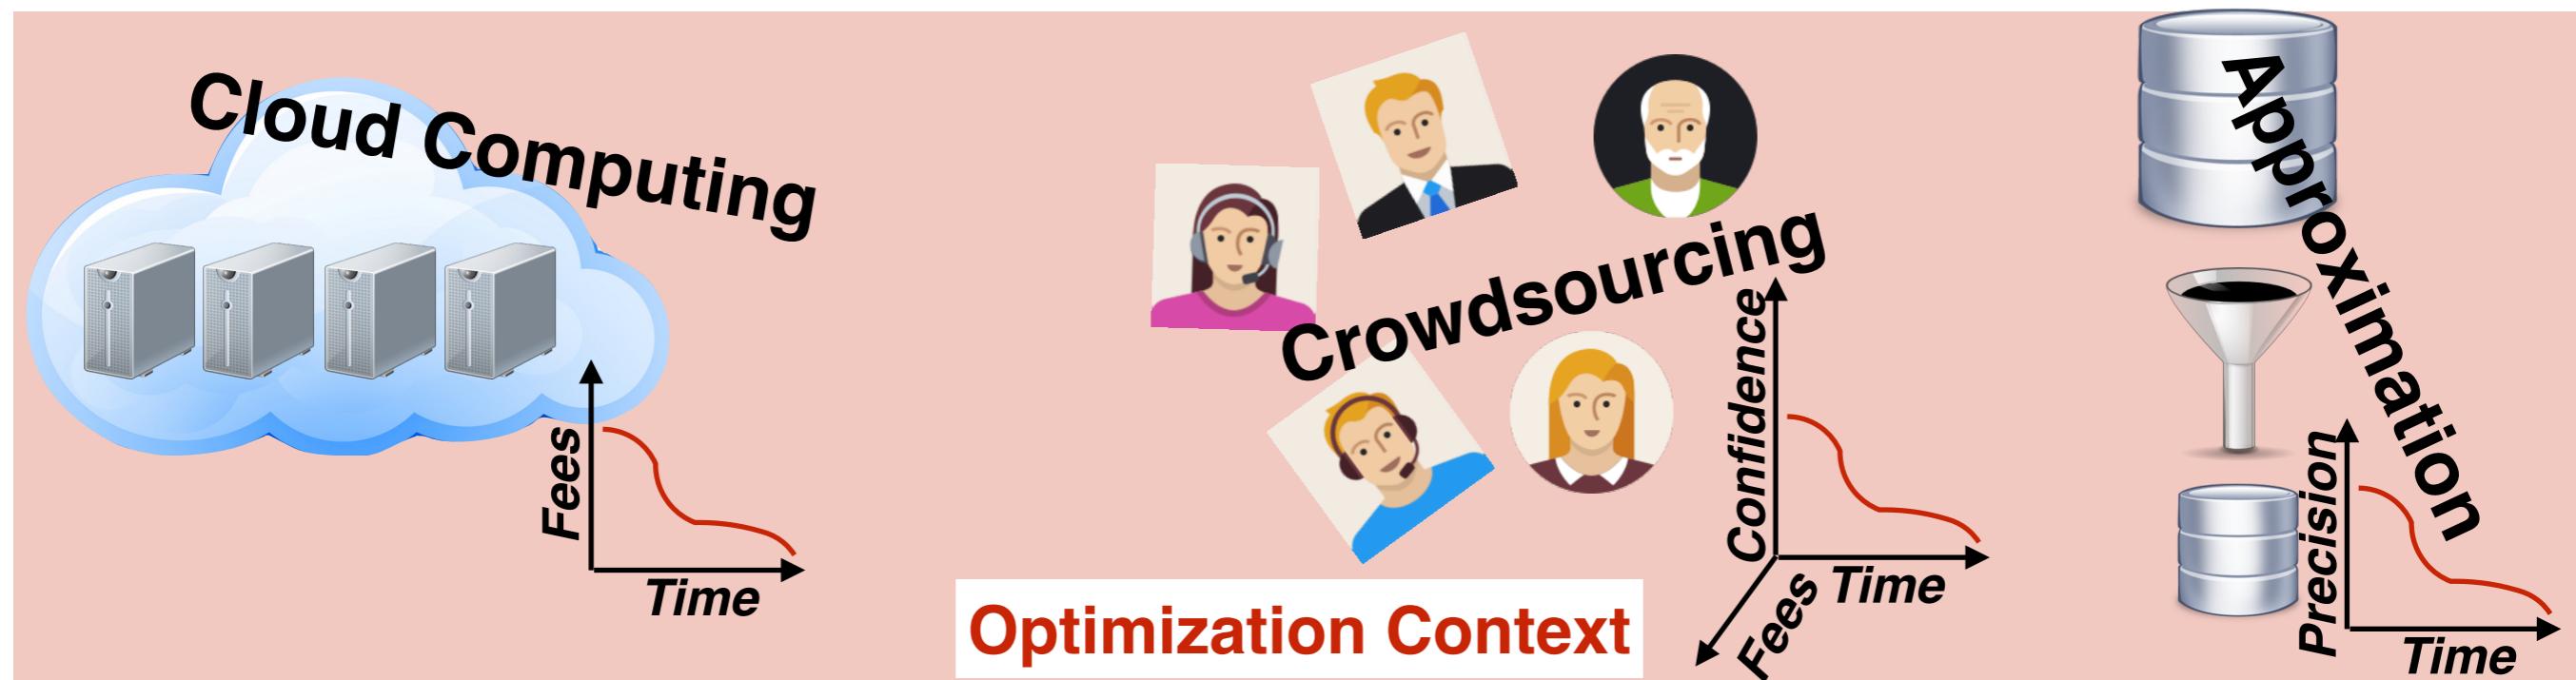
Query
Optimization

NP-Hard to
Solve

Optimal Plan(s)

NP-Hard to
Approximate

Modern Query Optimization



Problem Model

Tuning Knobs

*Operation order
Operators*

**NP-Hard to
Approximate**

Query

↓
Query
Optimization

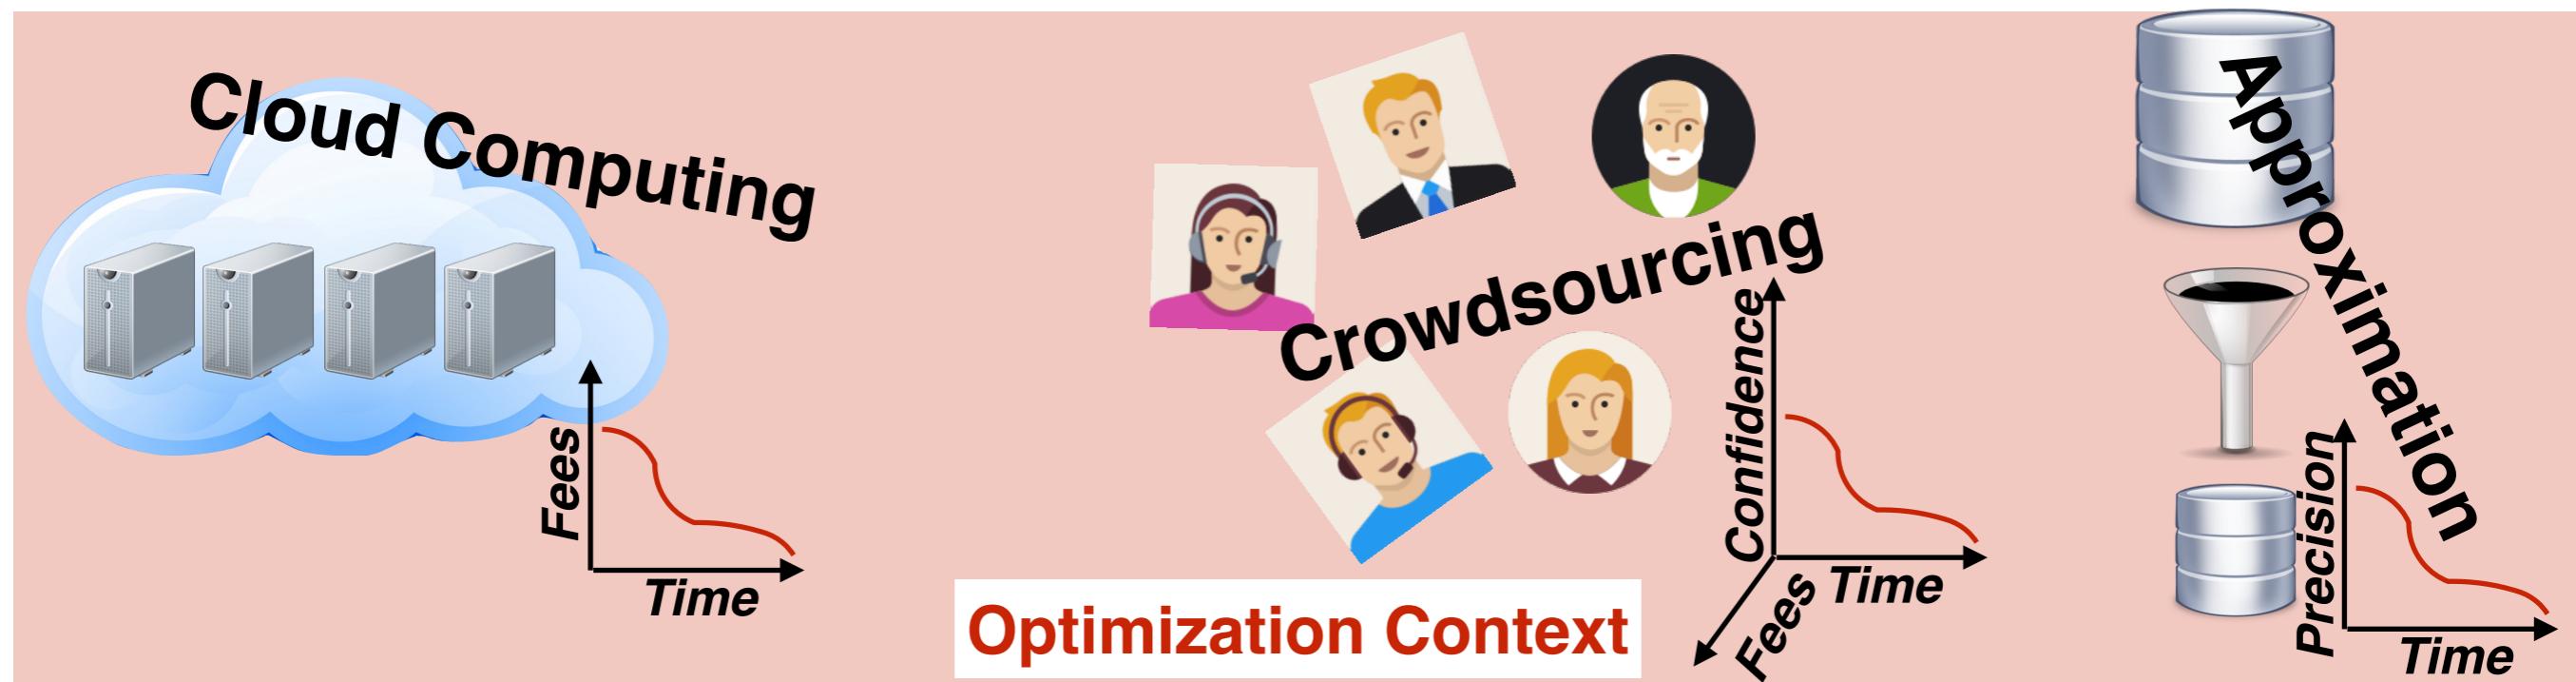
Optimal Plan(s)

Cost Metric

Execution time

**NP-Hard
to Solve**

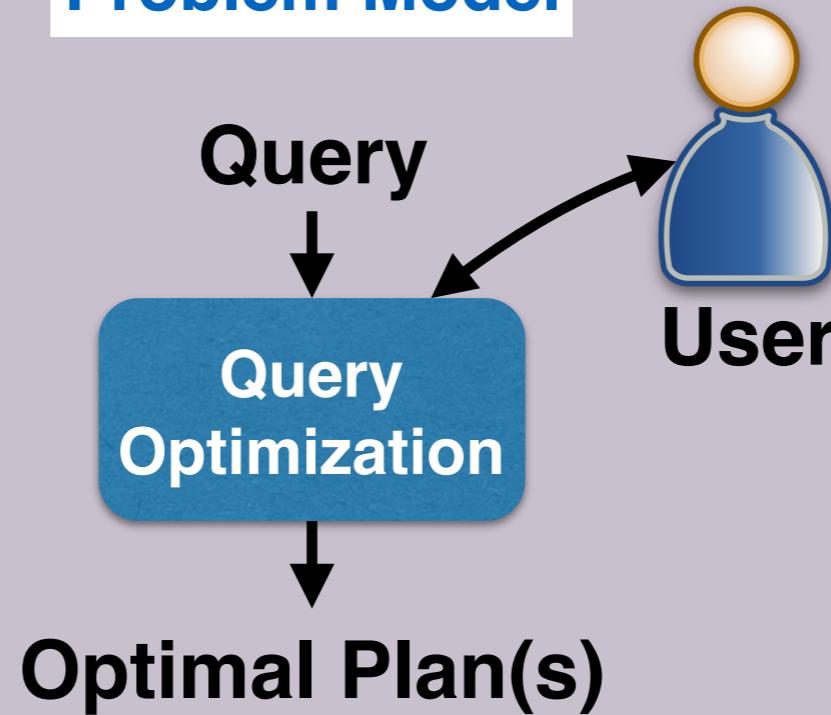
Modern Query Optimization



Problem Model

Tuning Knobs

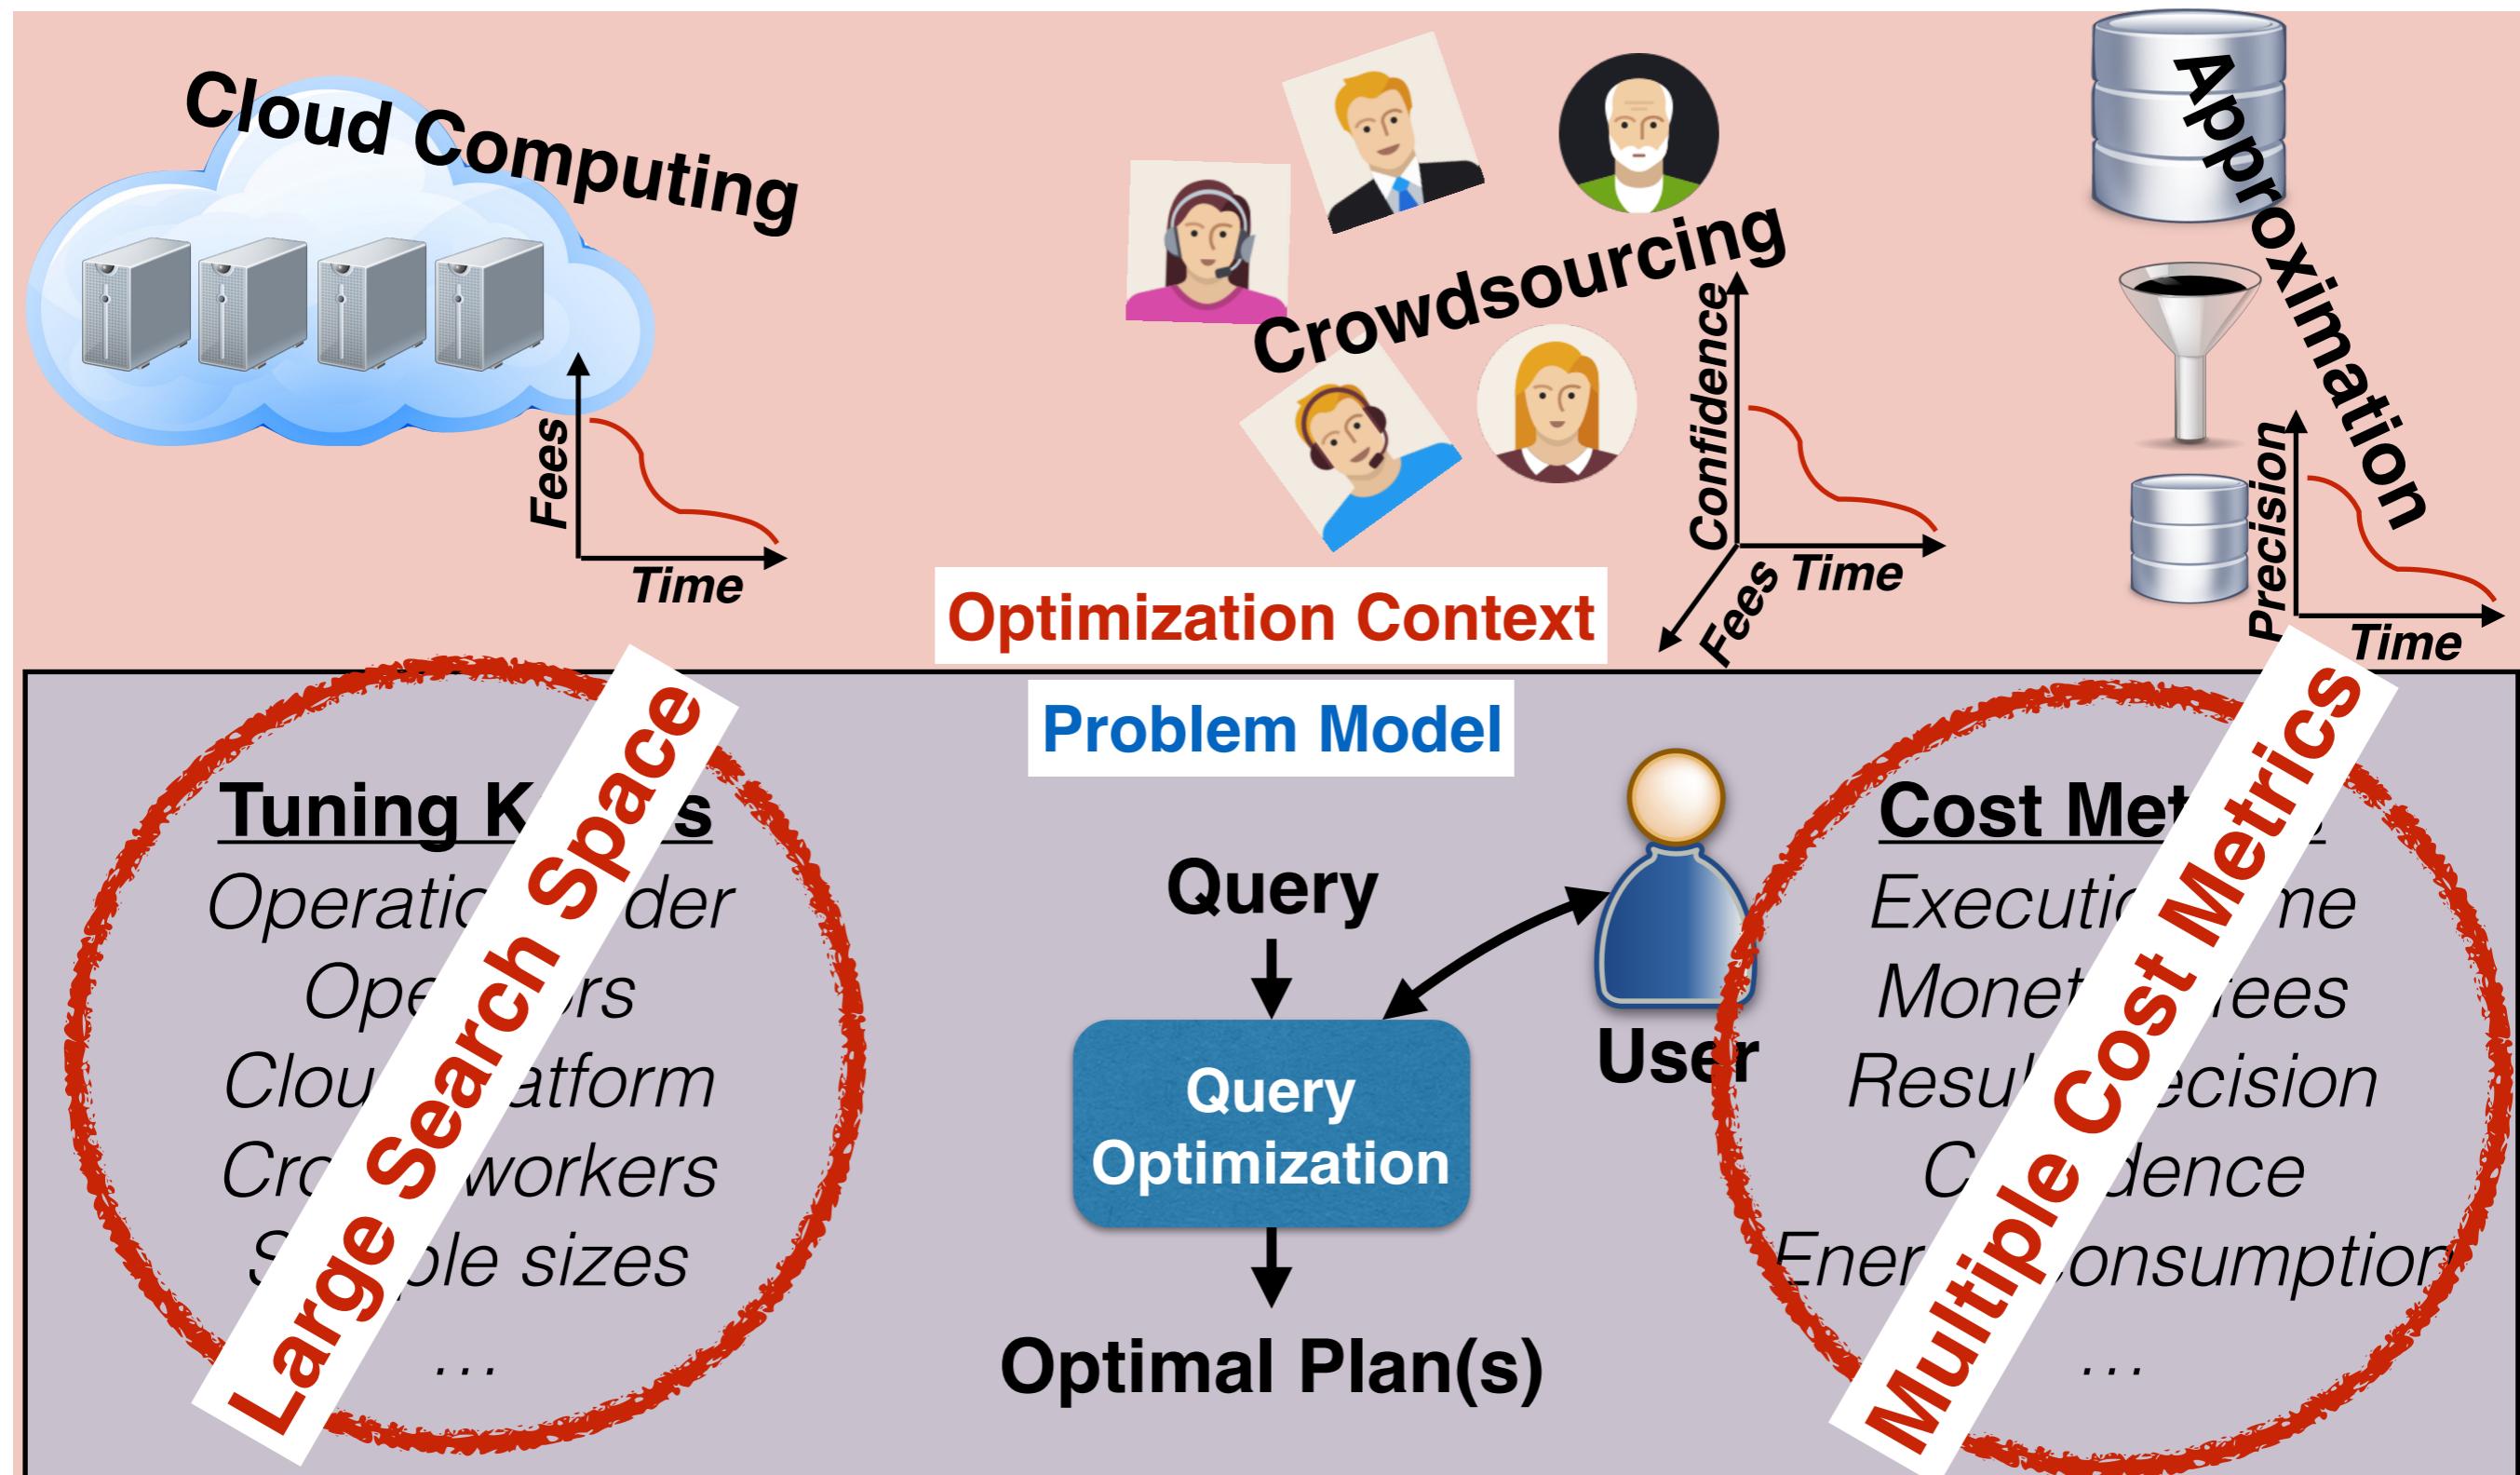
*Operation order
Operators
Cloud platform
Crowd workers
Sample sizes
...*



Cost Metrics

*Execution time
Monetary fees
Result precision
Confidence
Energy consumption
...*

Modern Query Optimization



Example: Google

Query

Types ⚕ Properties ⚕ Entities ⚕ Opinions ⚕ Models ...

Example: Google

Query

Types ⋈ Properties ⋈ Entities ⋈ Opinions ⋈ Models ...

Query Plan



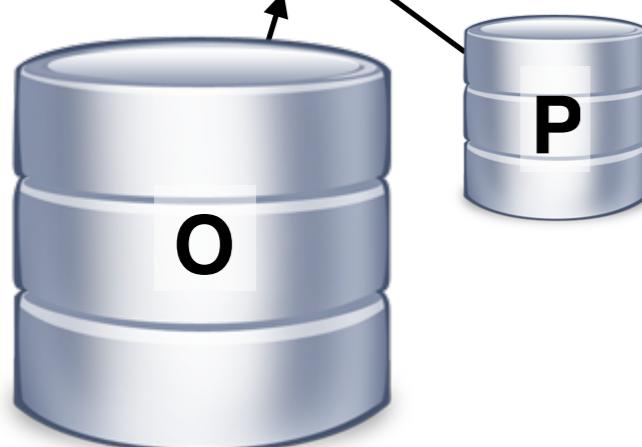
Example: Google

Query

Types \bowtie Properties \bowtie Entities \bowtie Opinions \bowtie Models ...

Query Plan

O \bowtie P



Example: Google

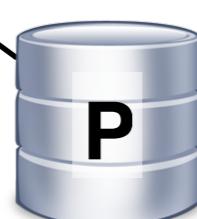
Query

Types \bowtie Properties \bowtie Entities \bowtie Opinions \bowtie Models ...

Query Plan

O \bowtie P \bowtie T

O \bowtie P



Example: Google

Query

Types \bowtie Properties \bowtie Entities \bowtie Opinions \bowtie Models ...

Query Plan

O \bowtie P \bowtie T \bowtie E

O \bowtie P \bowtie T

O \bowtie P



Example: Google

Query

Types \bowtie Properties \bowtie Entities \bowtie Opinions \bowtie Models ...

Query Plan

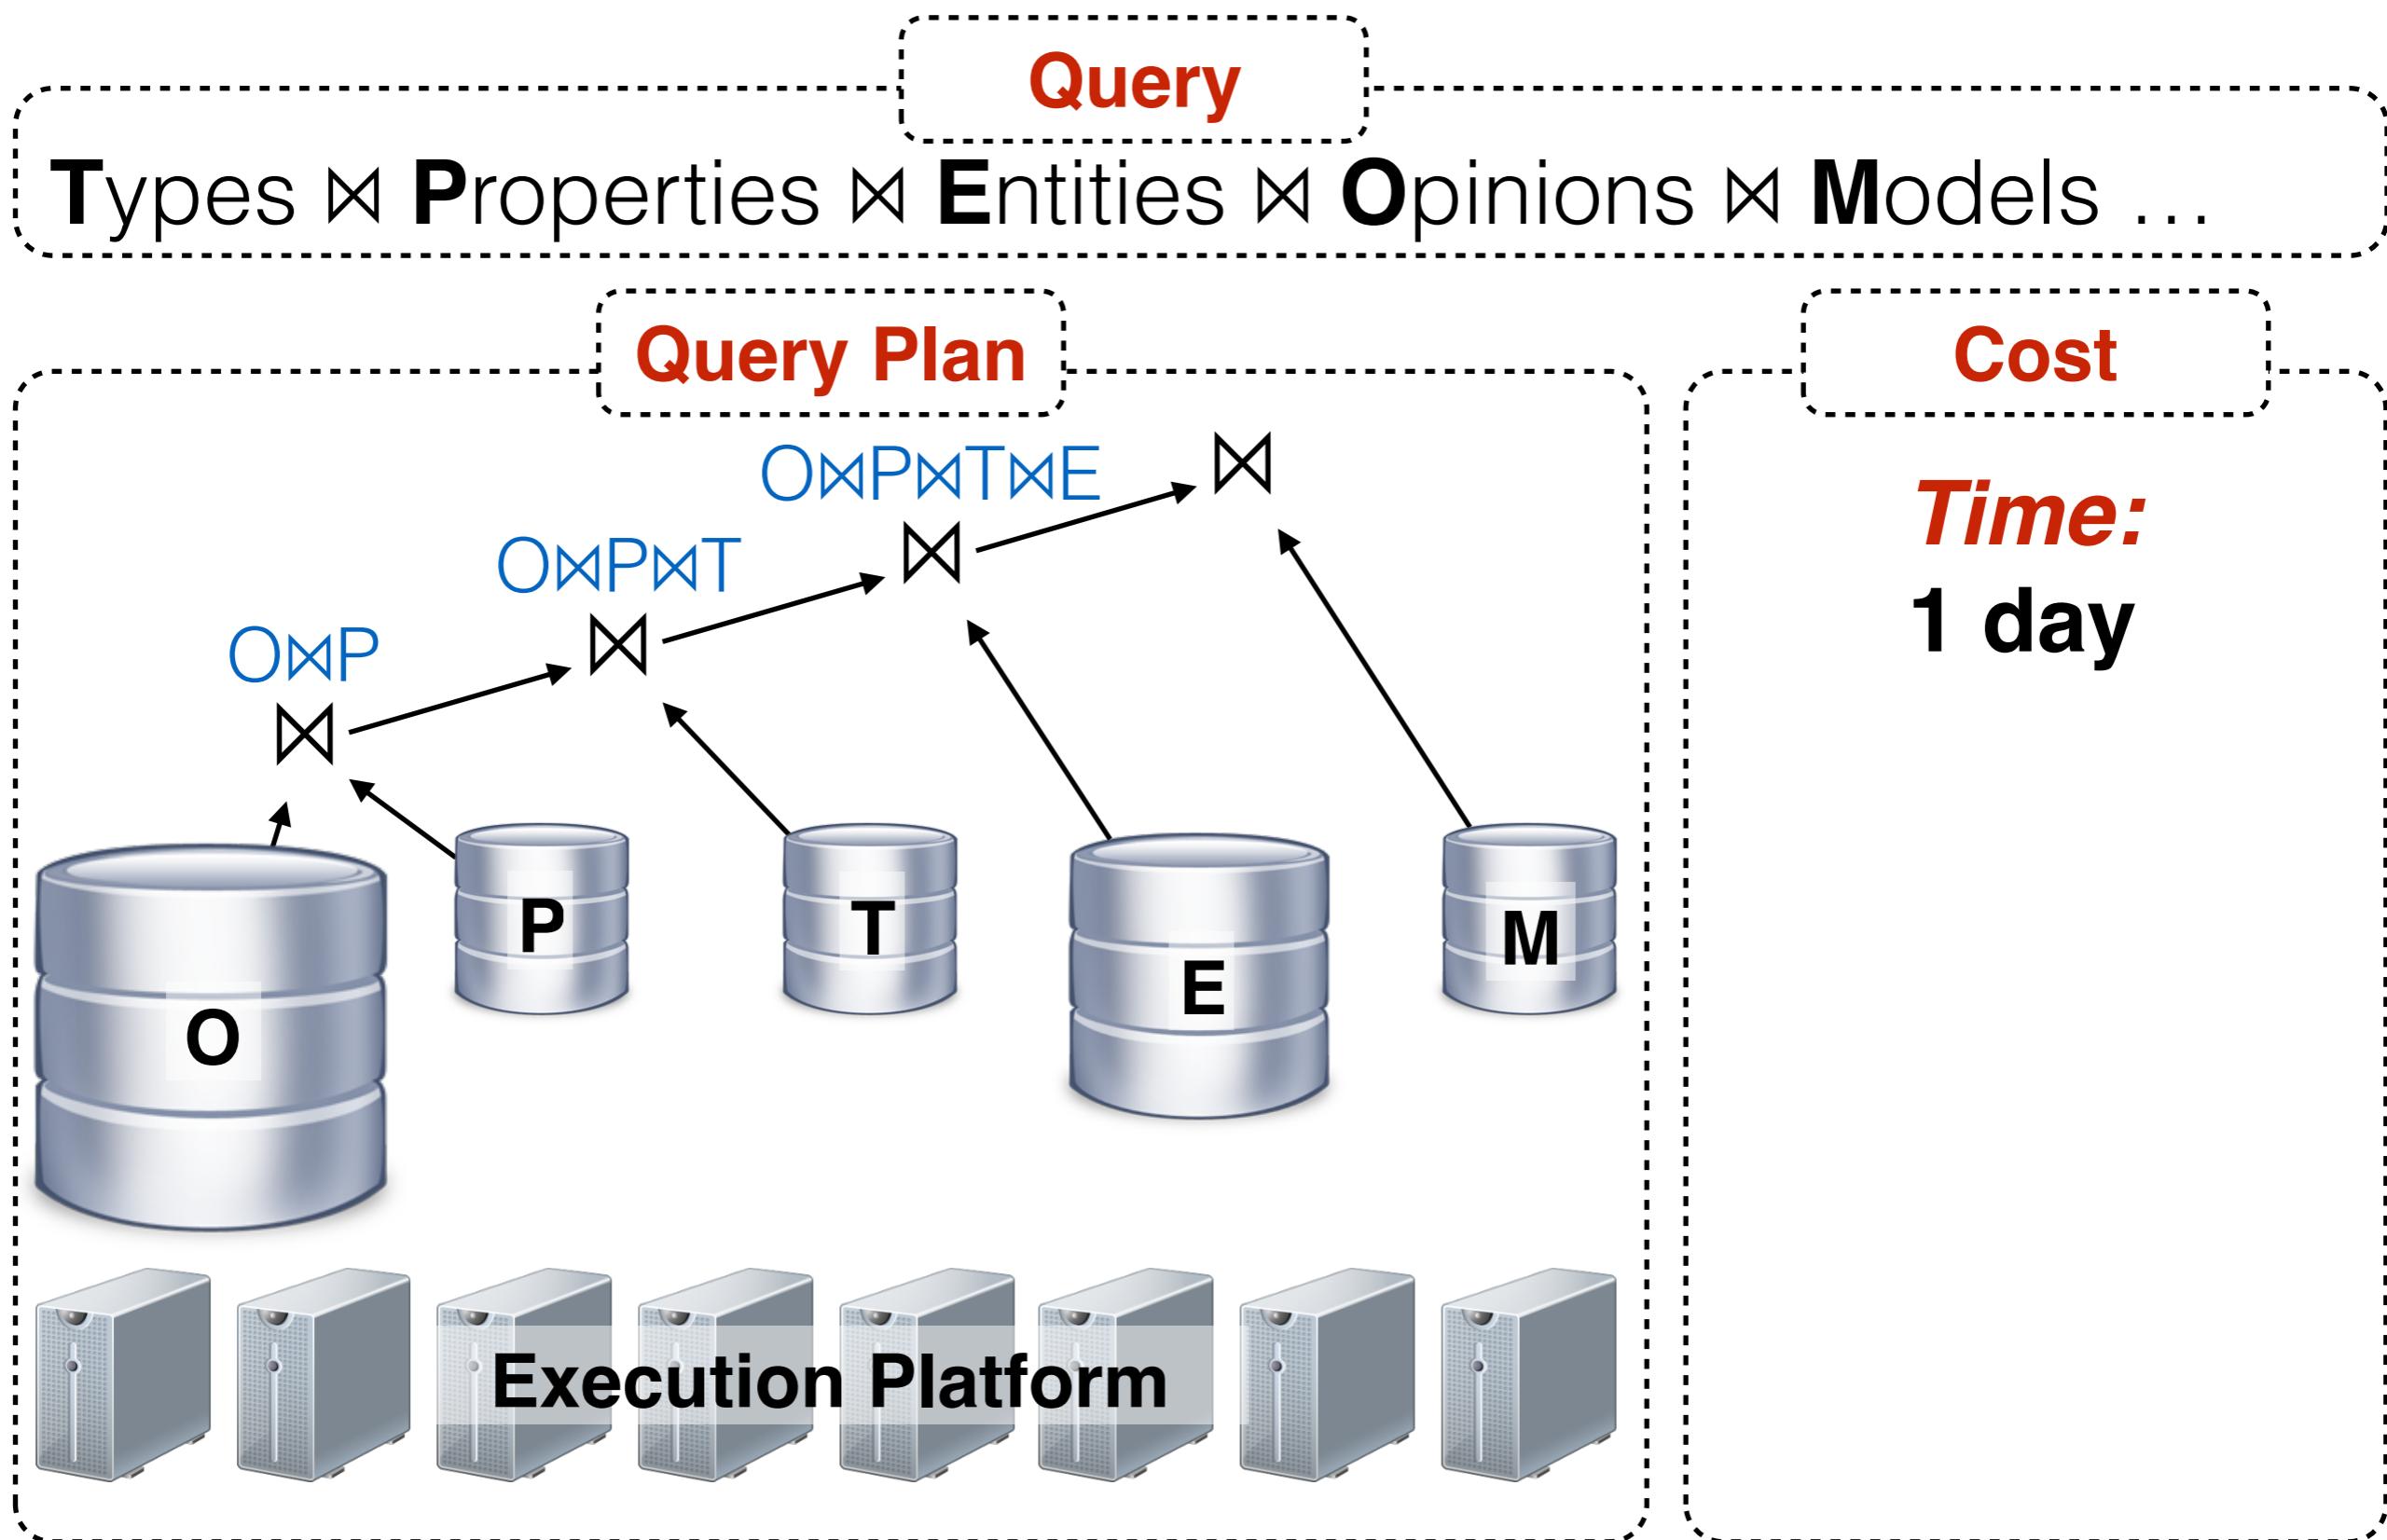
O \bowtie P \bowtie T \bowtie E

O \bowtie P \bowtie T

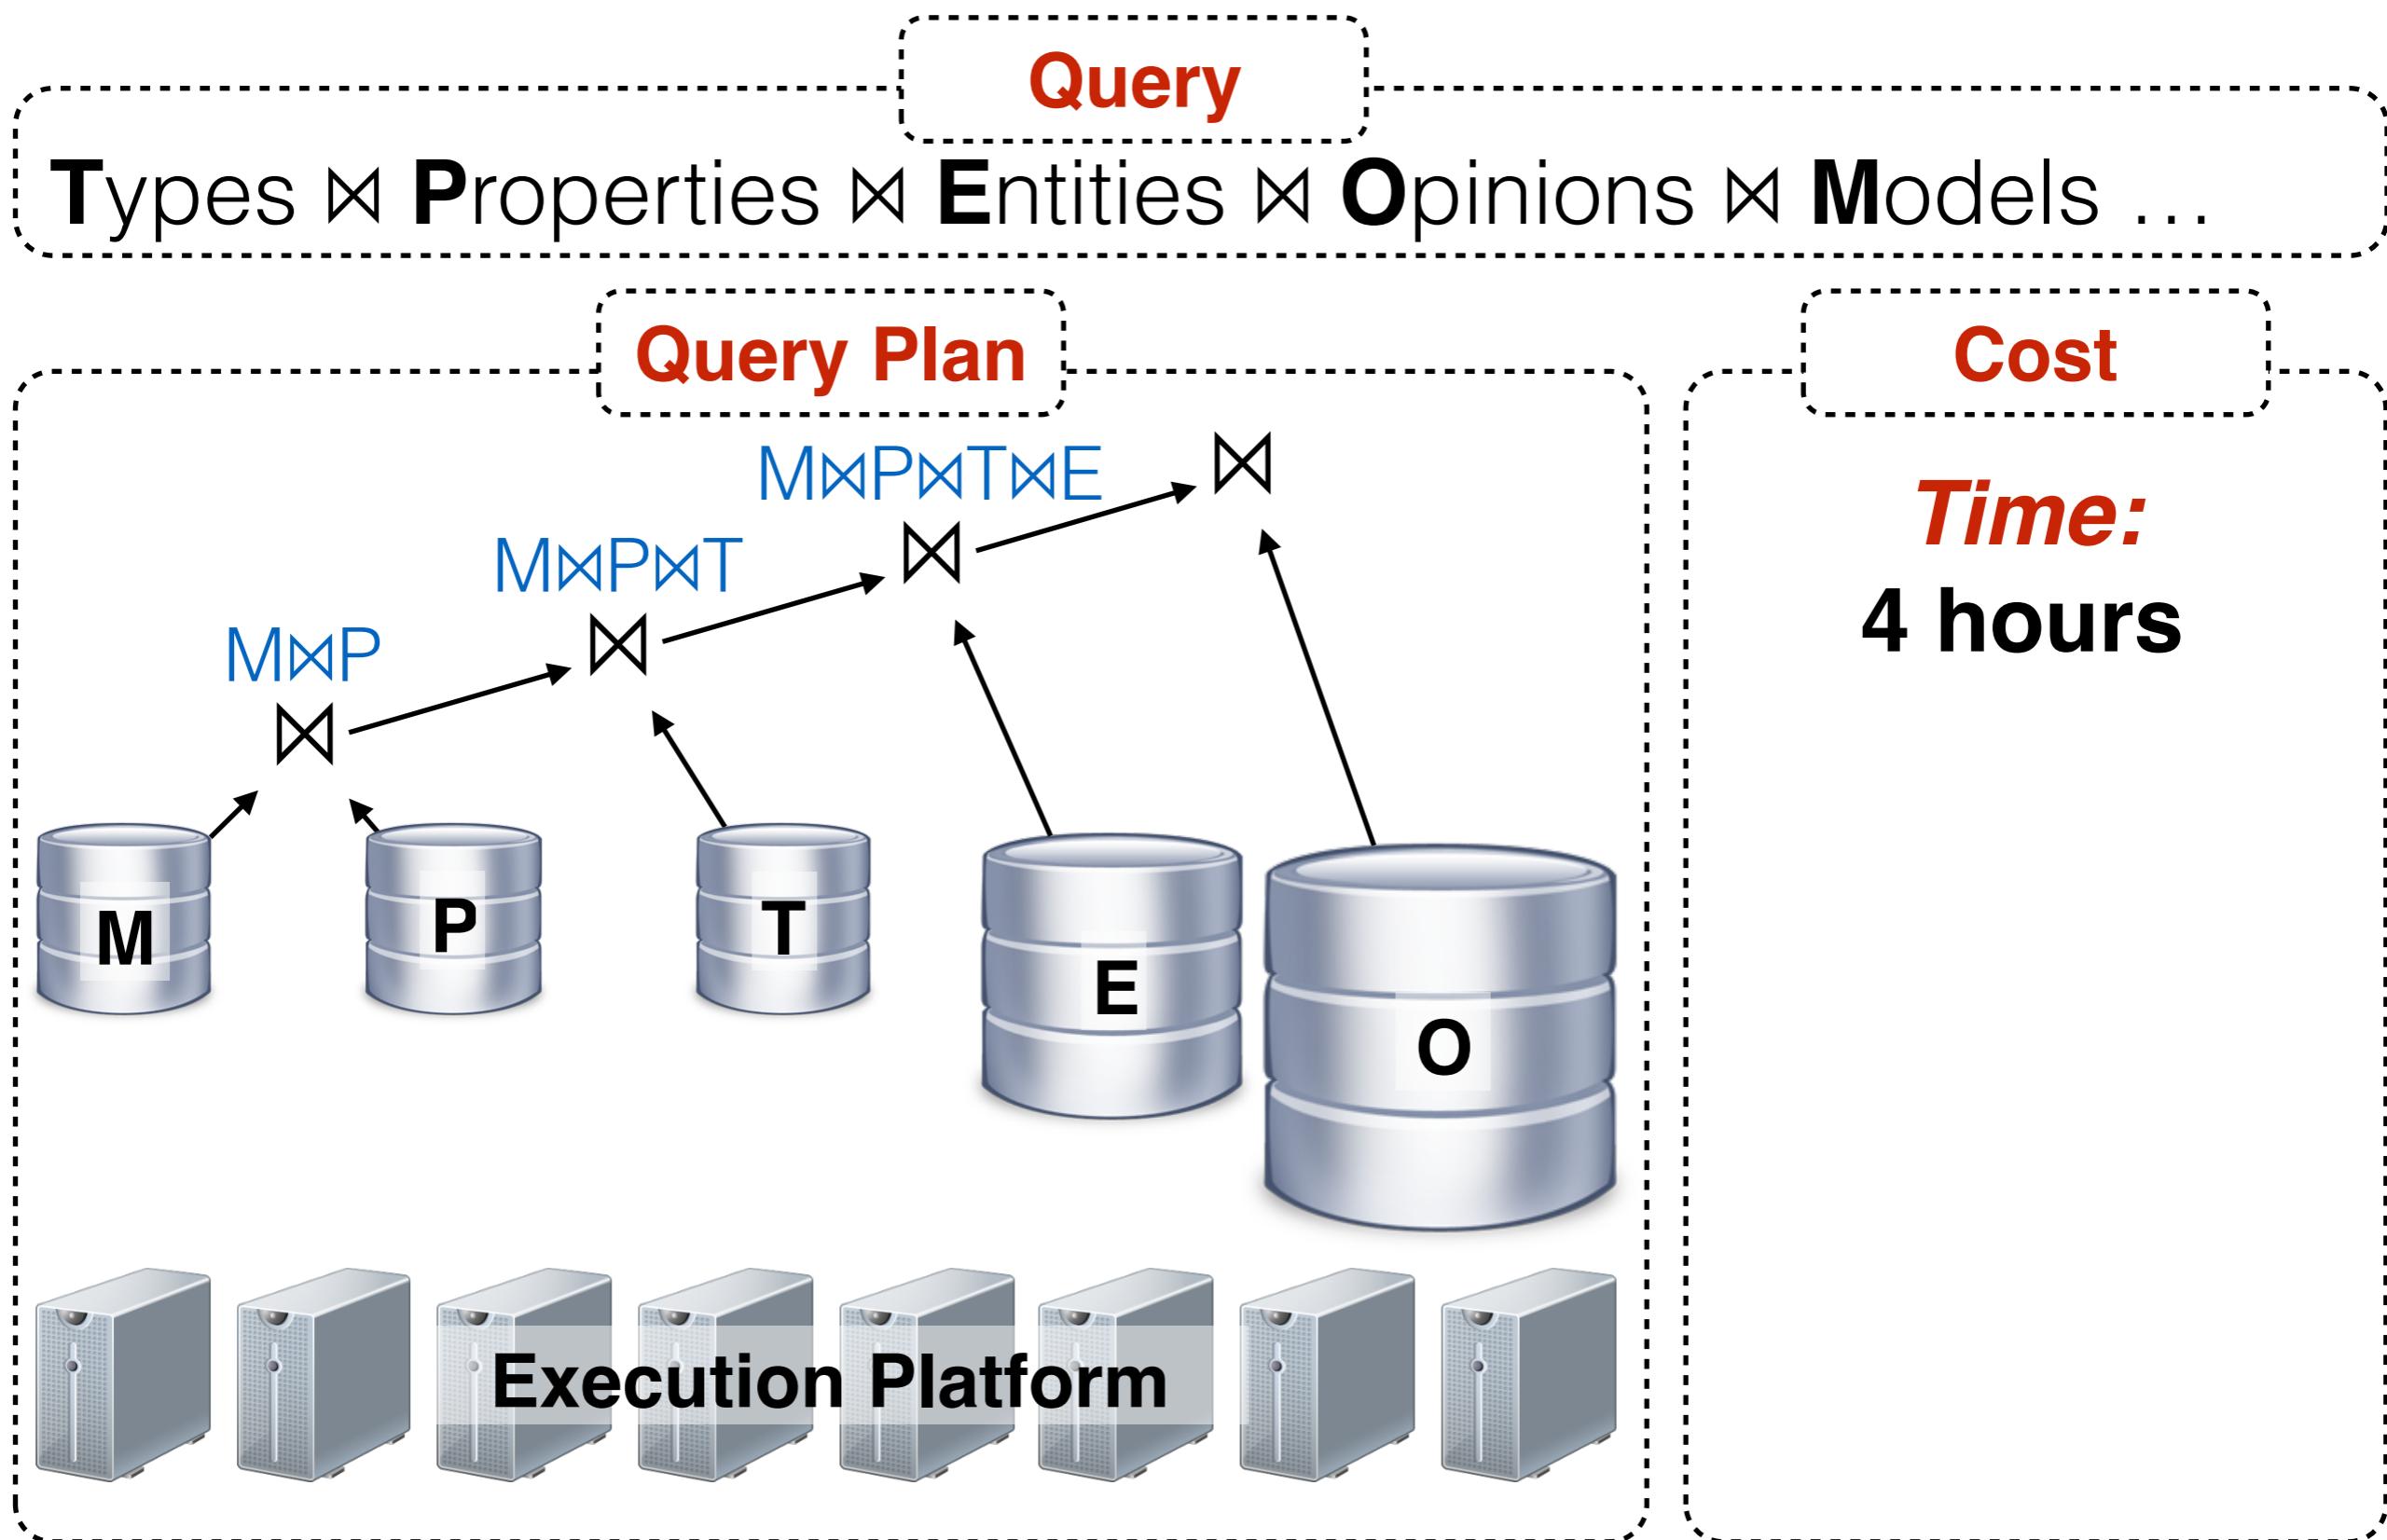
O \bowtie P



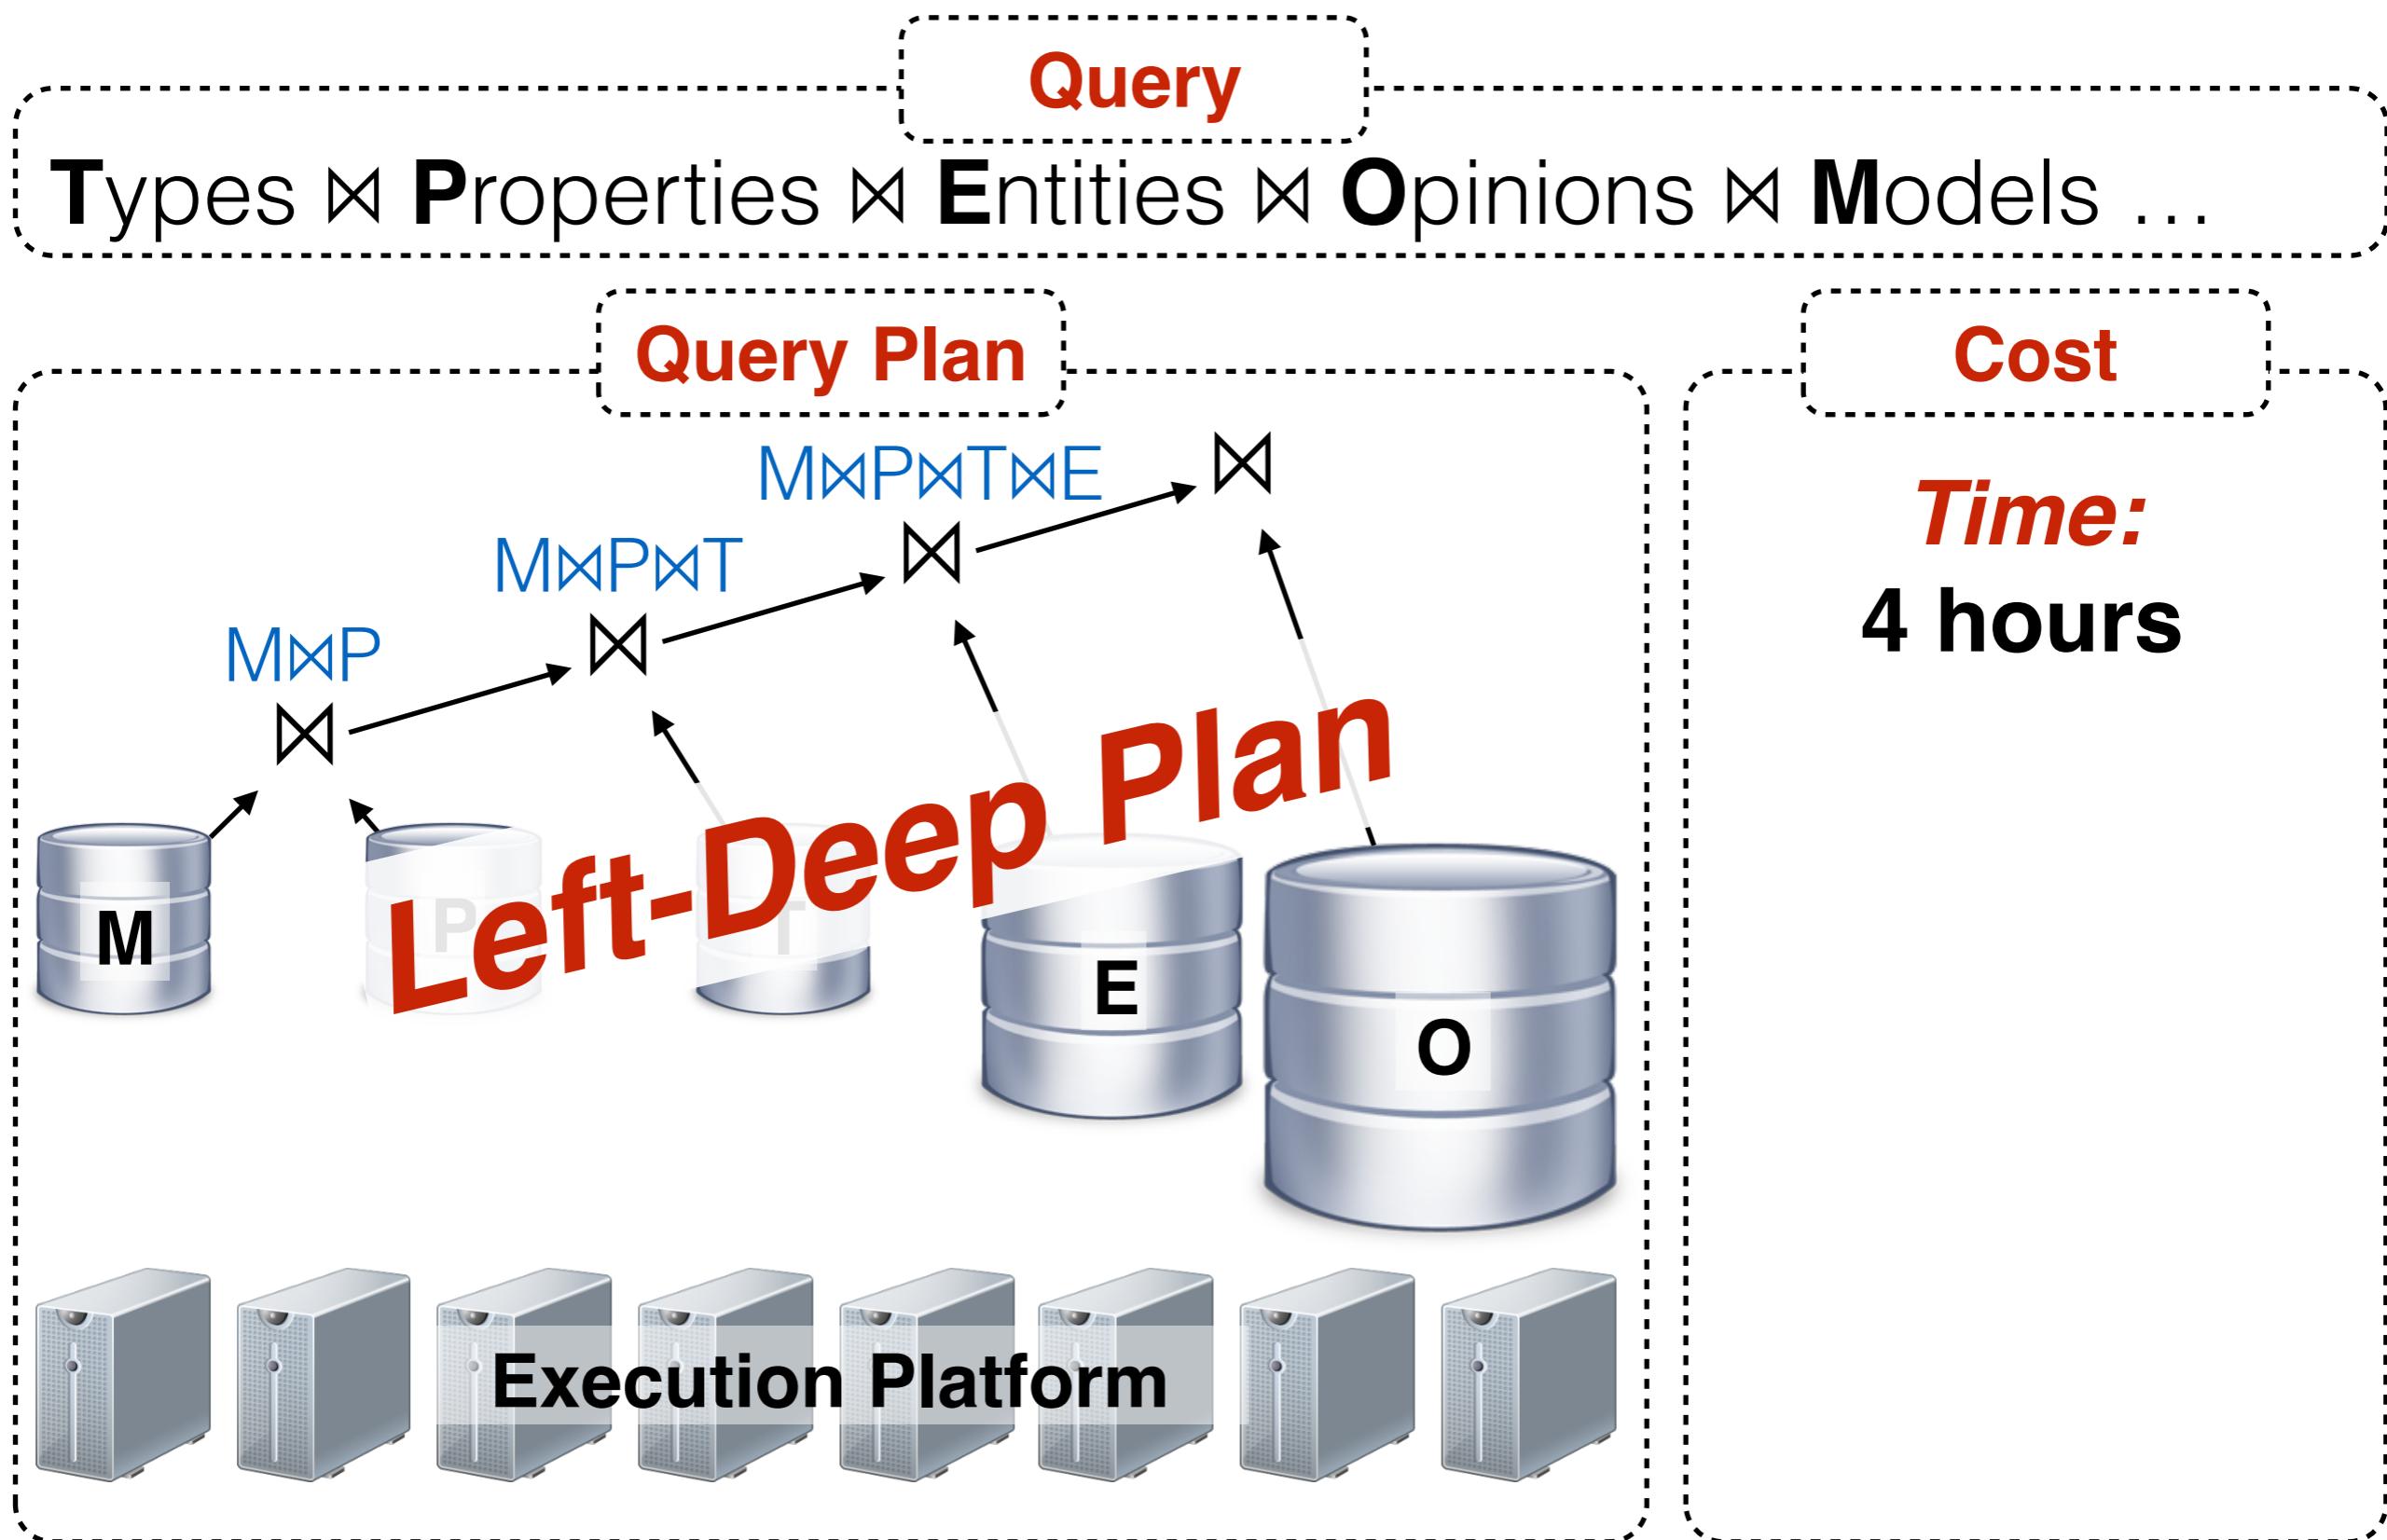
Example: Google



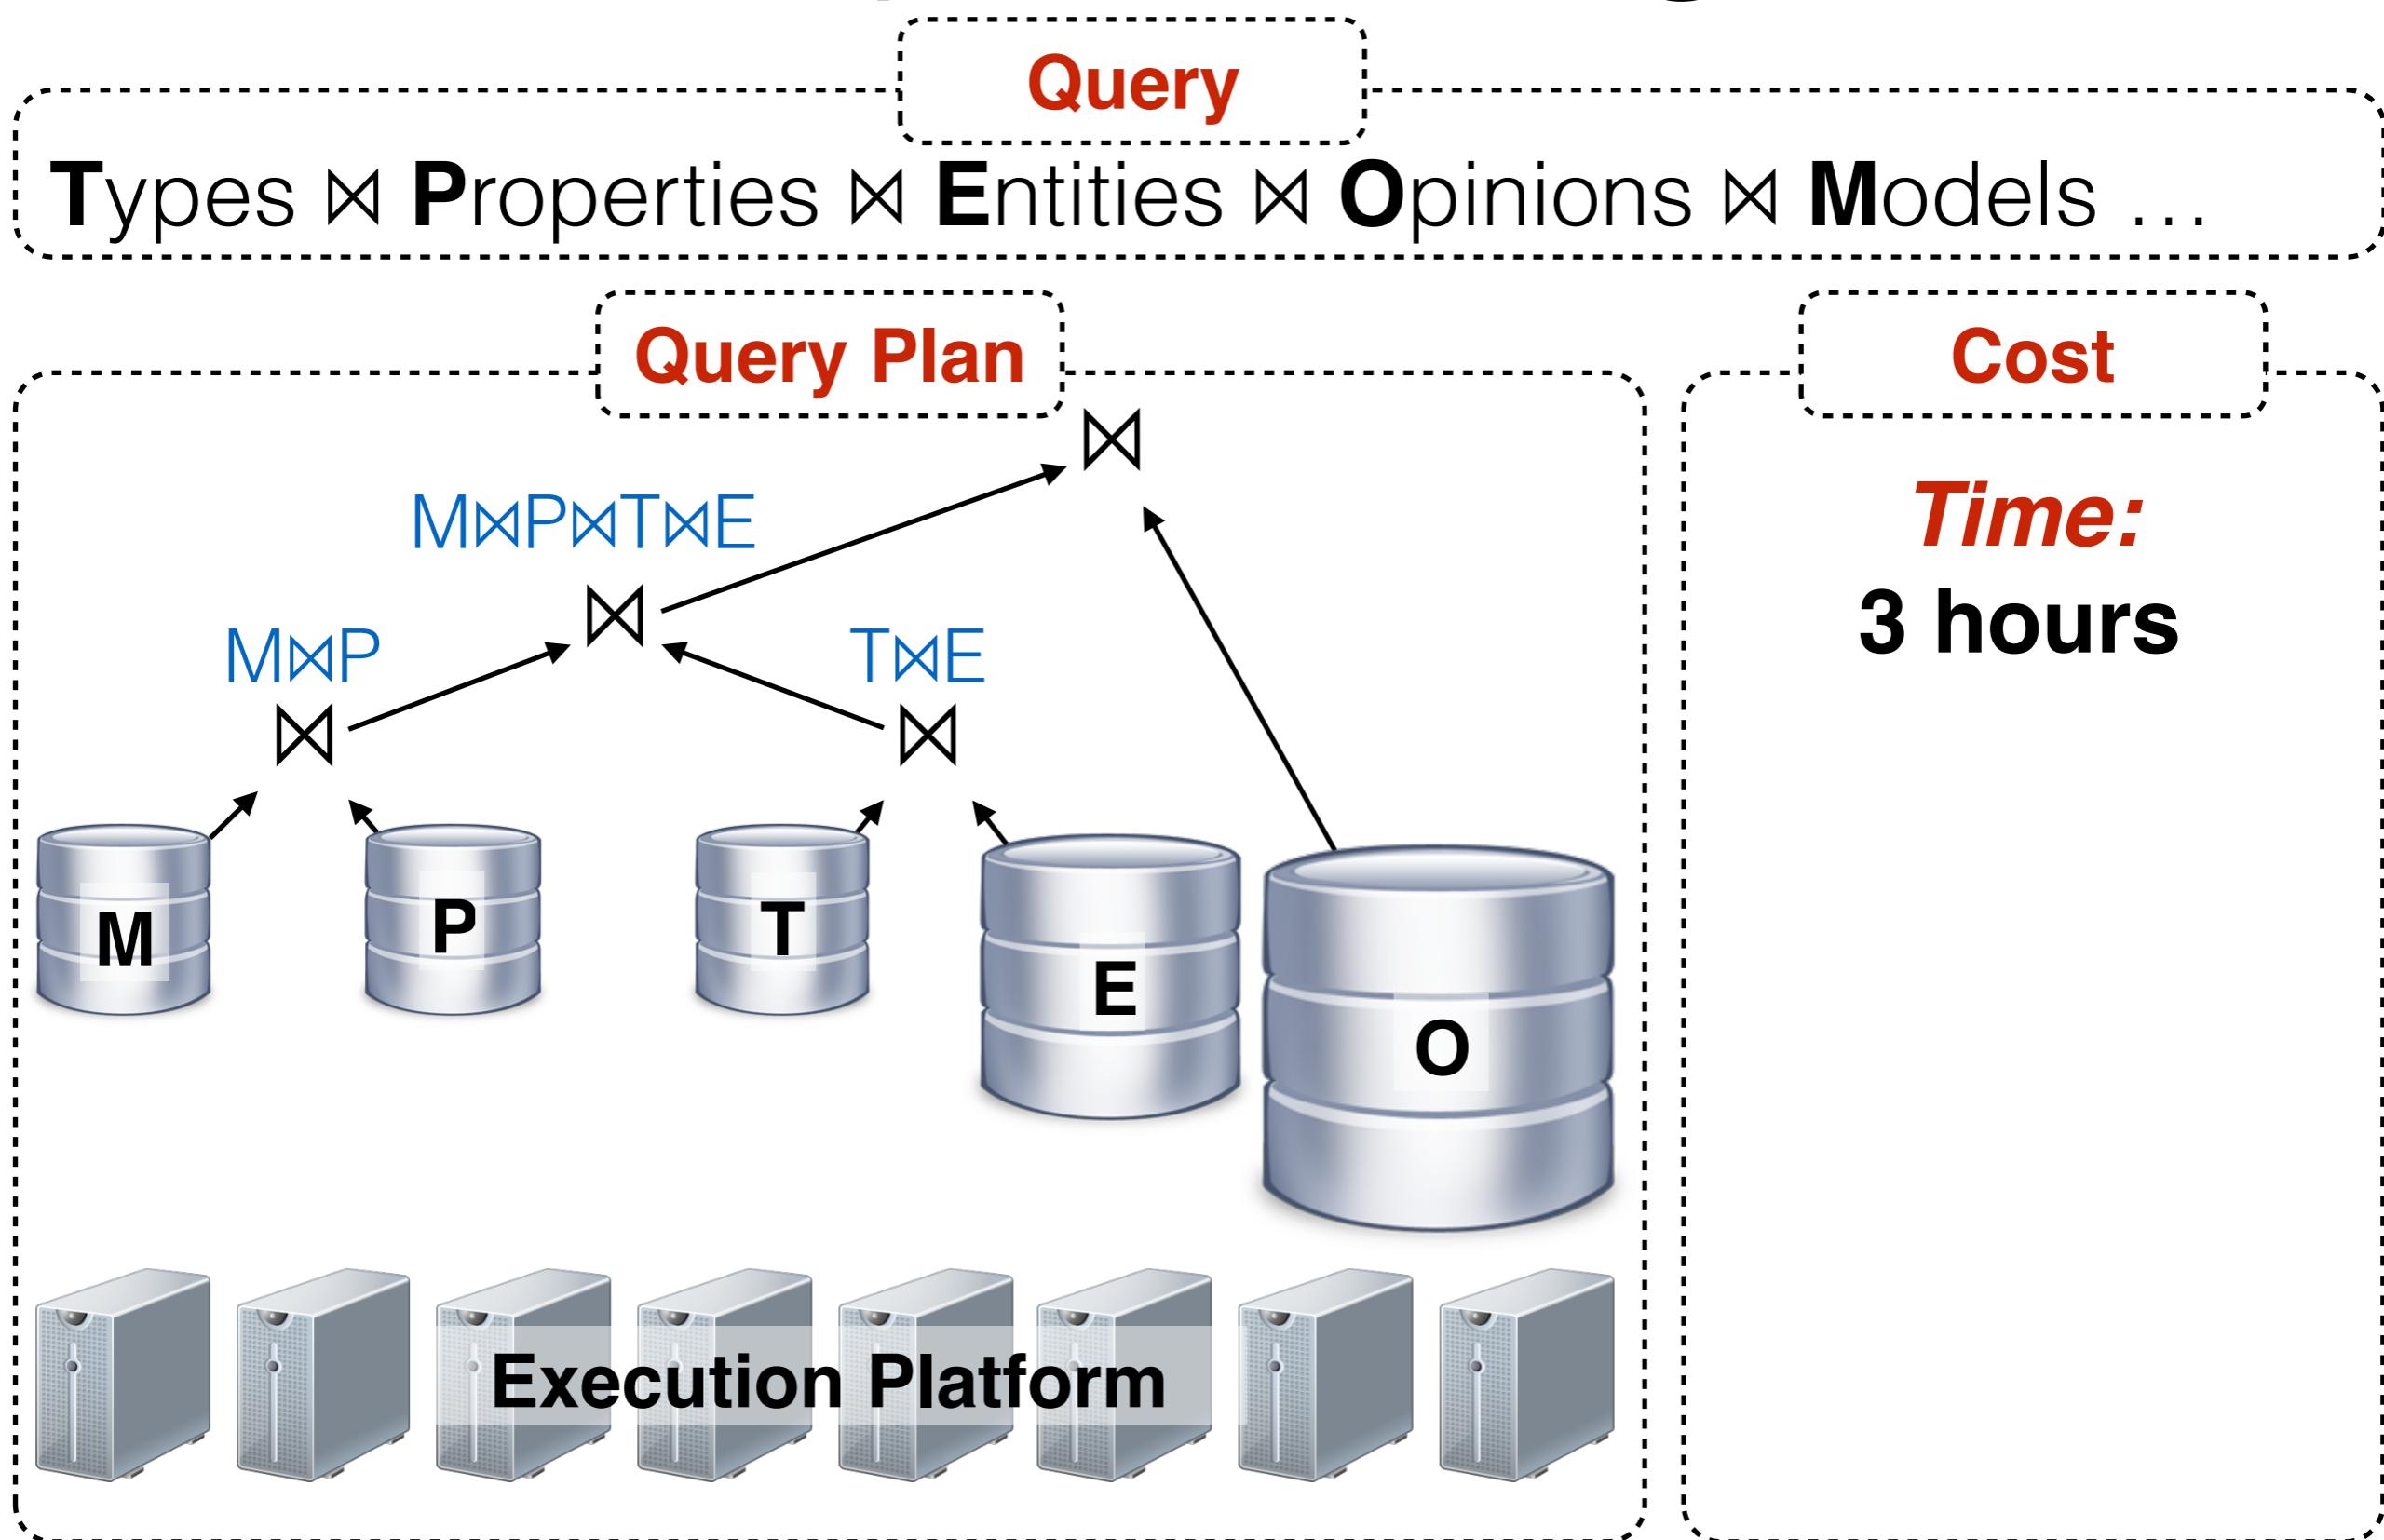
Example: Google



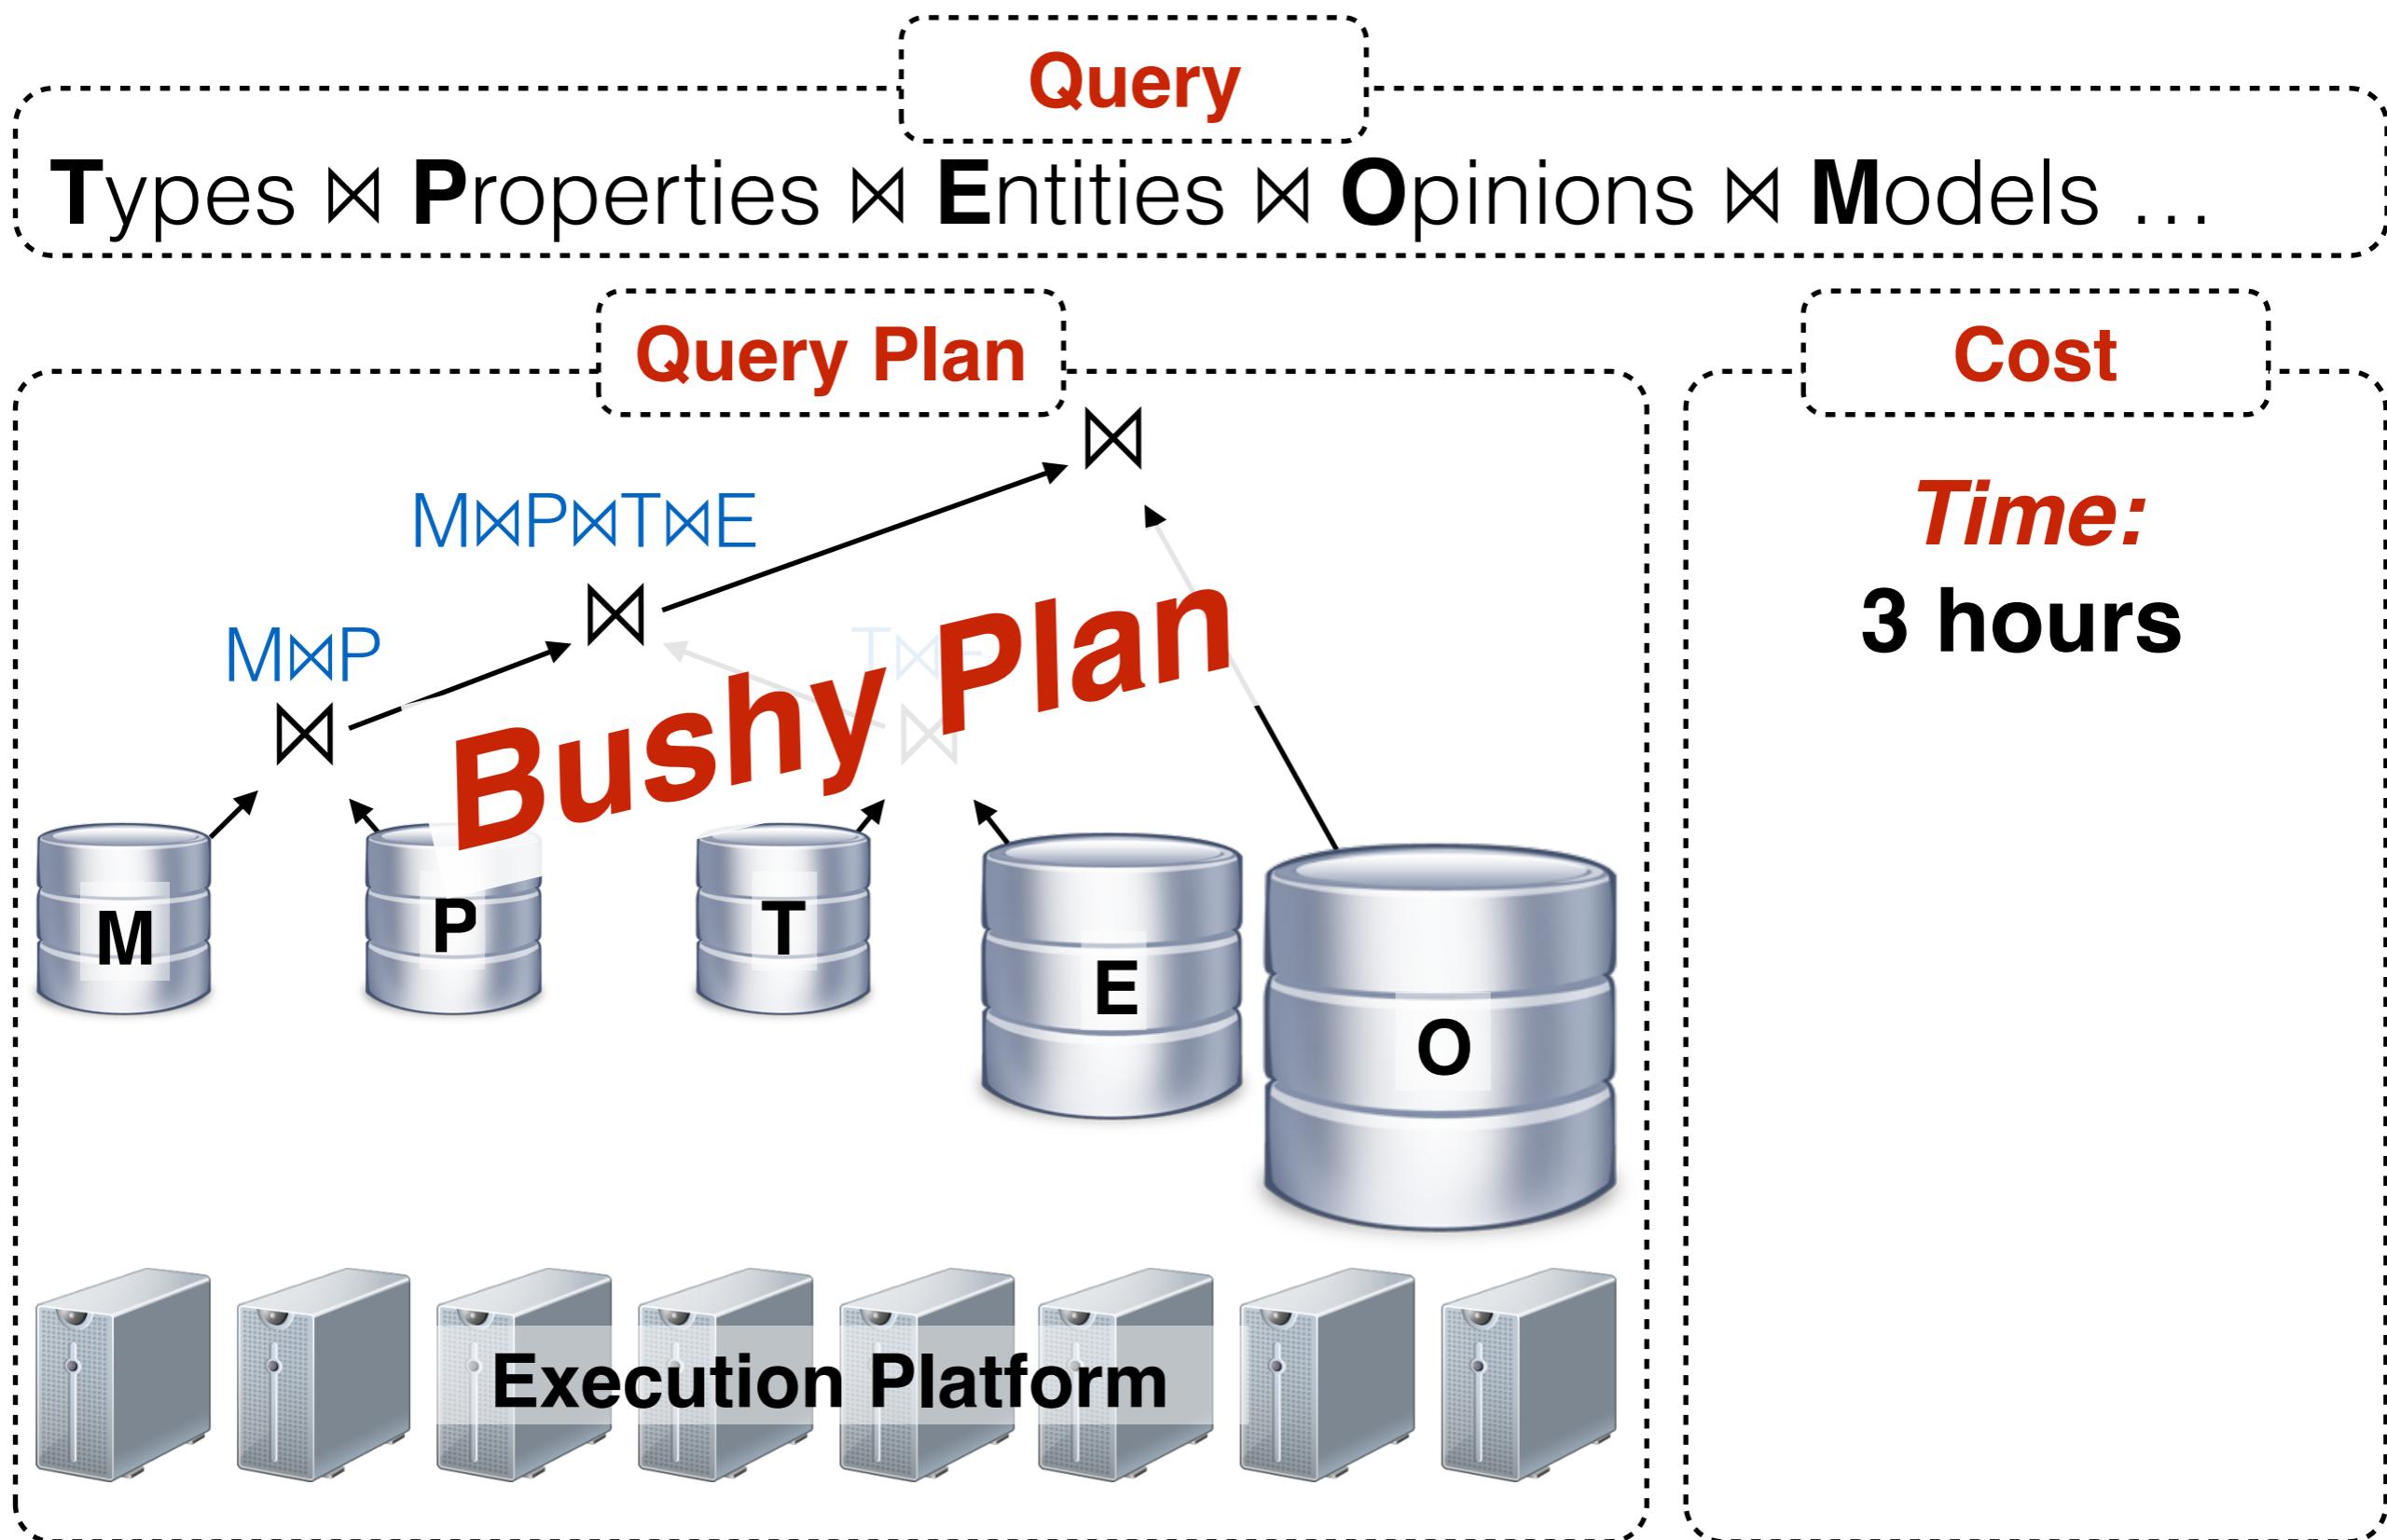
Example: Google



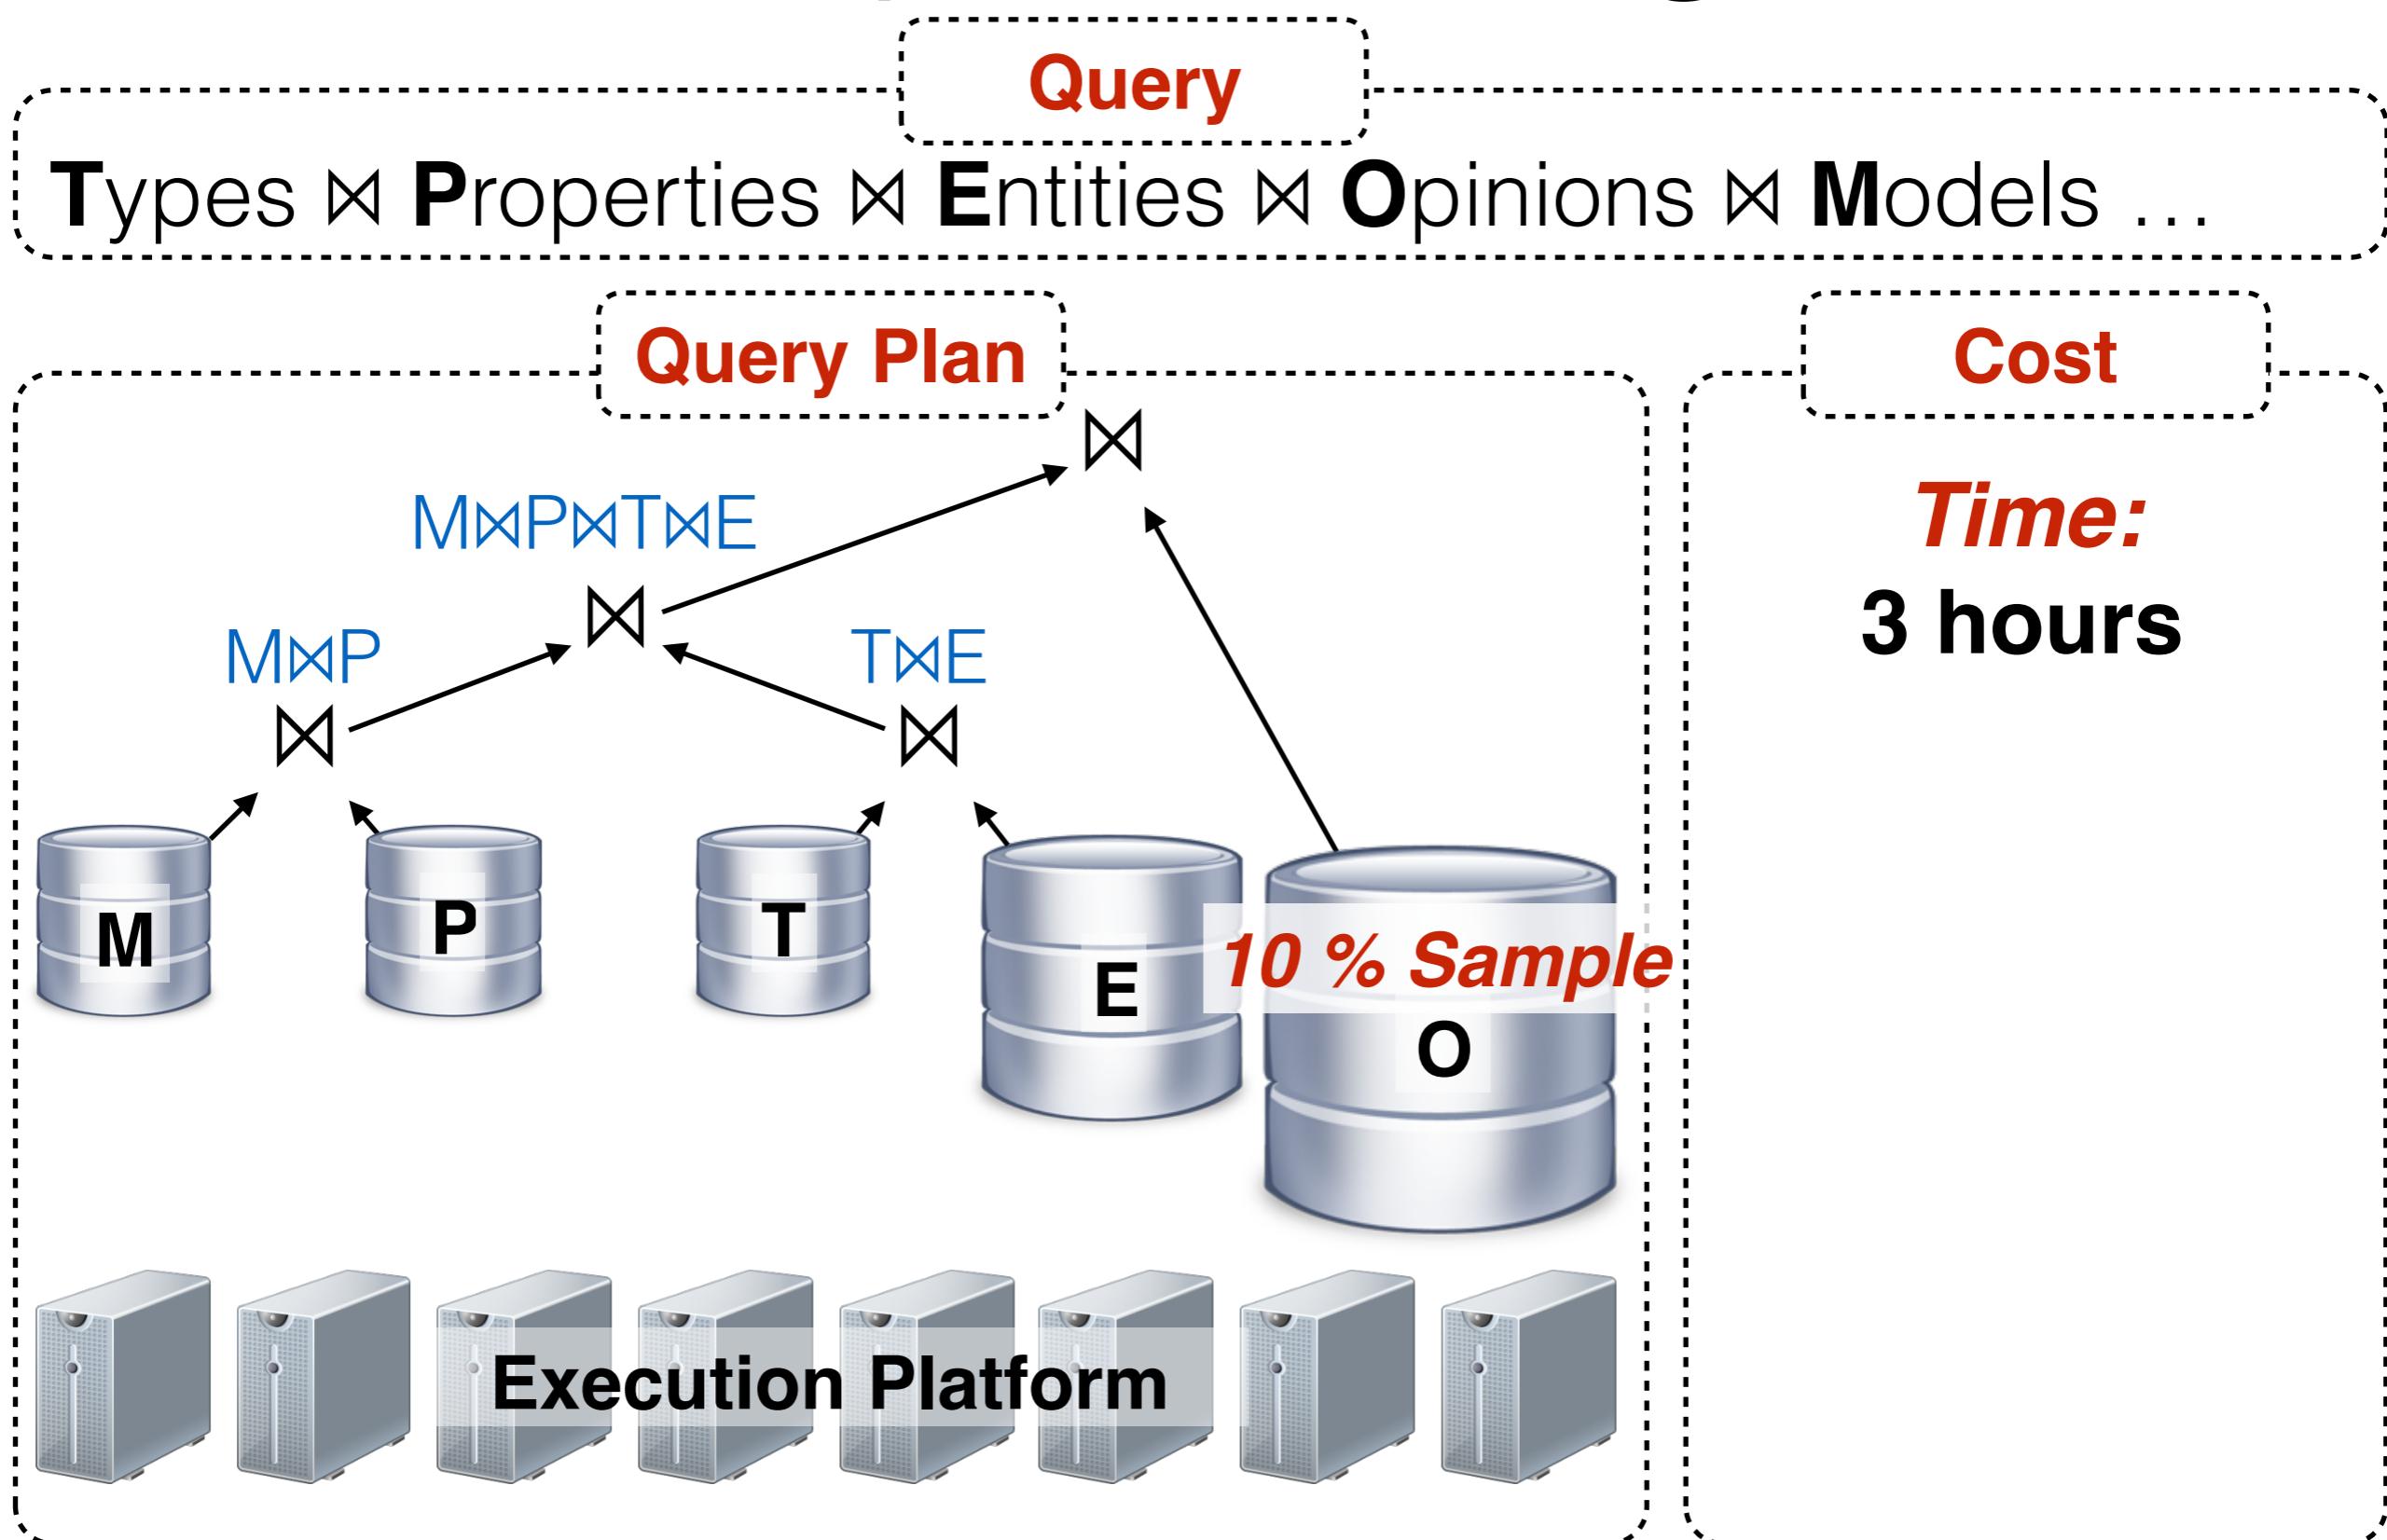
Example: Google



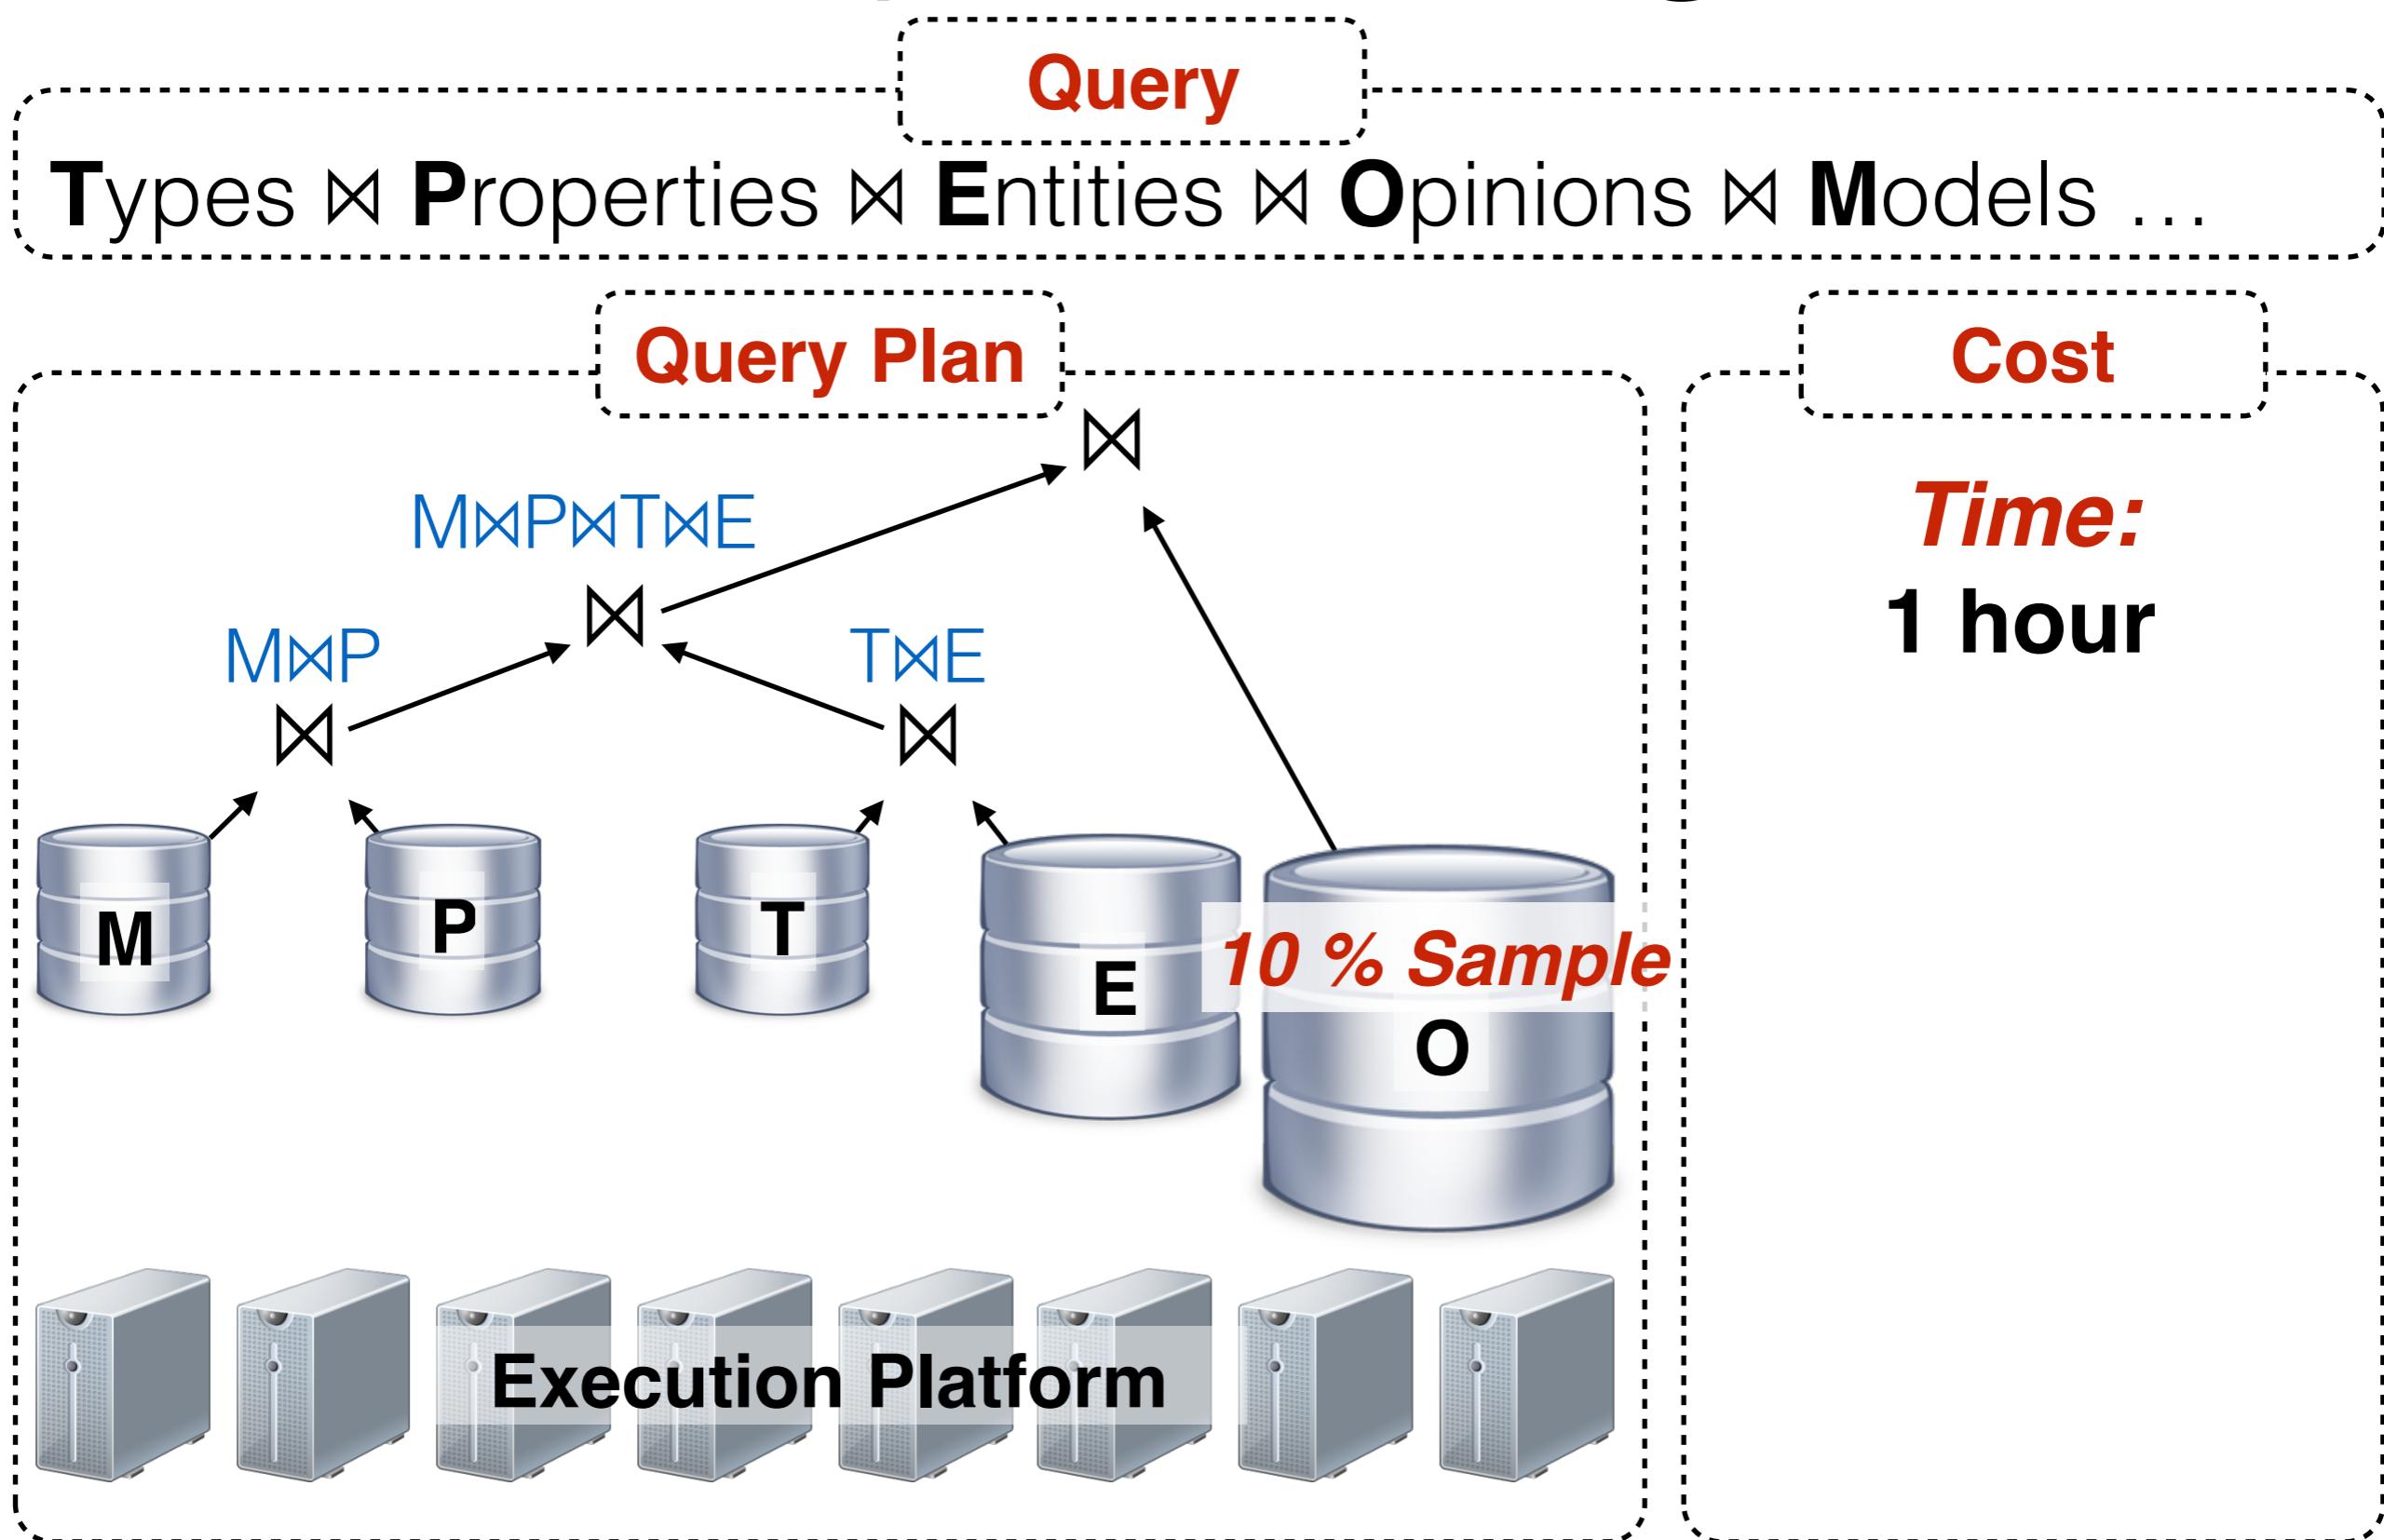
Example: Google



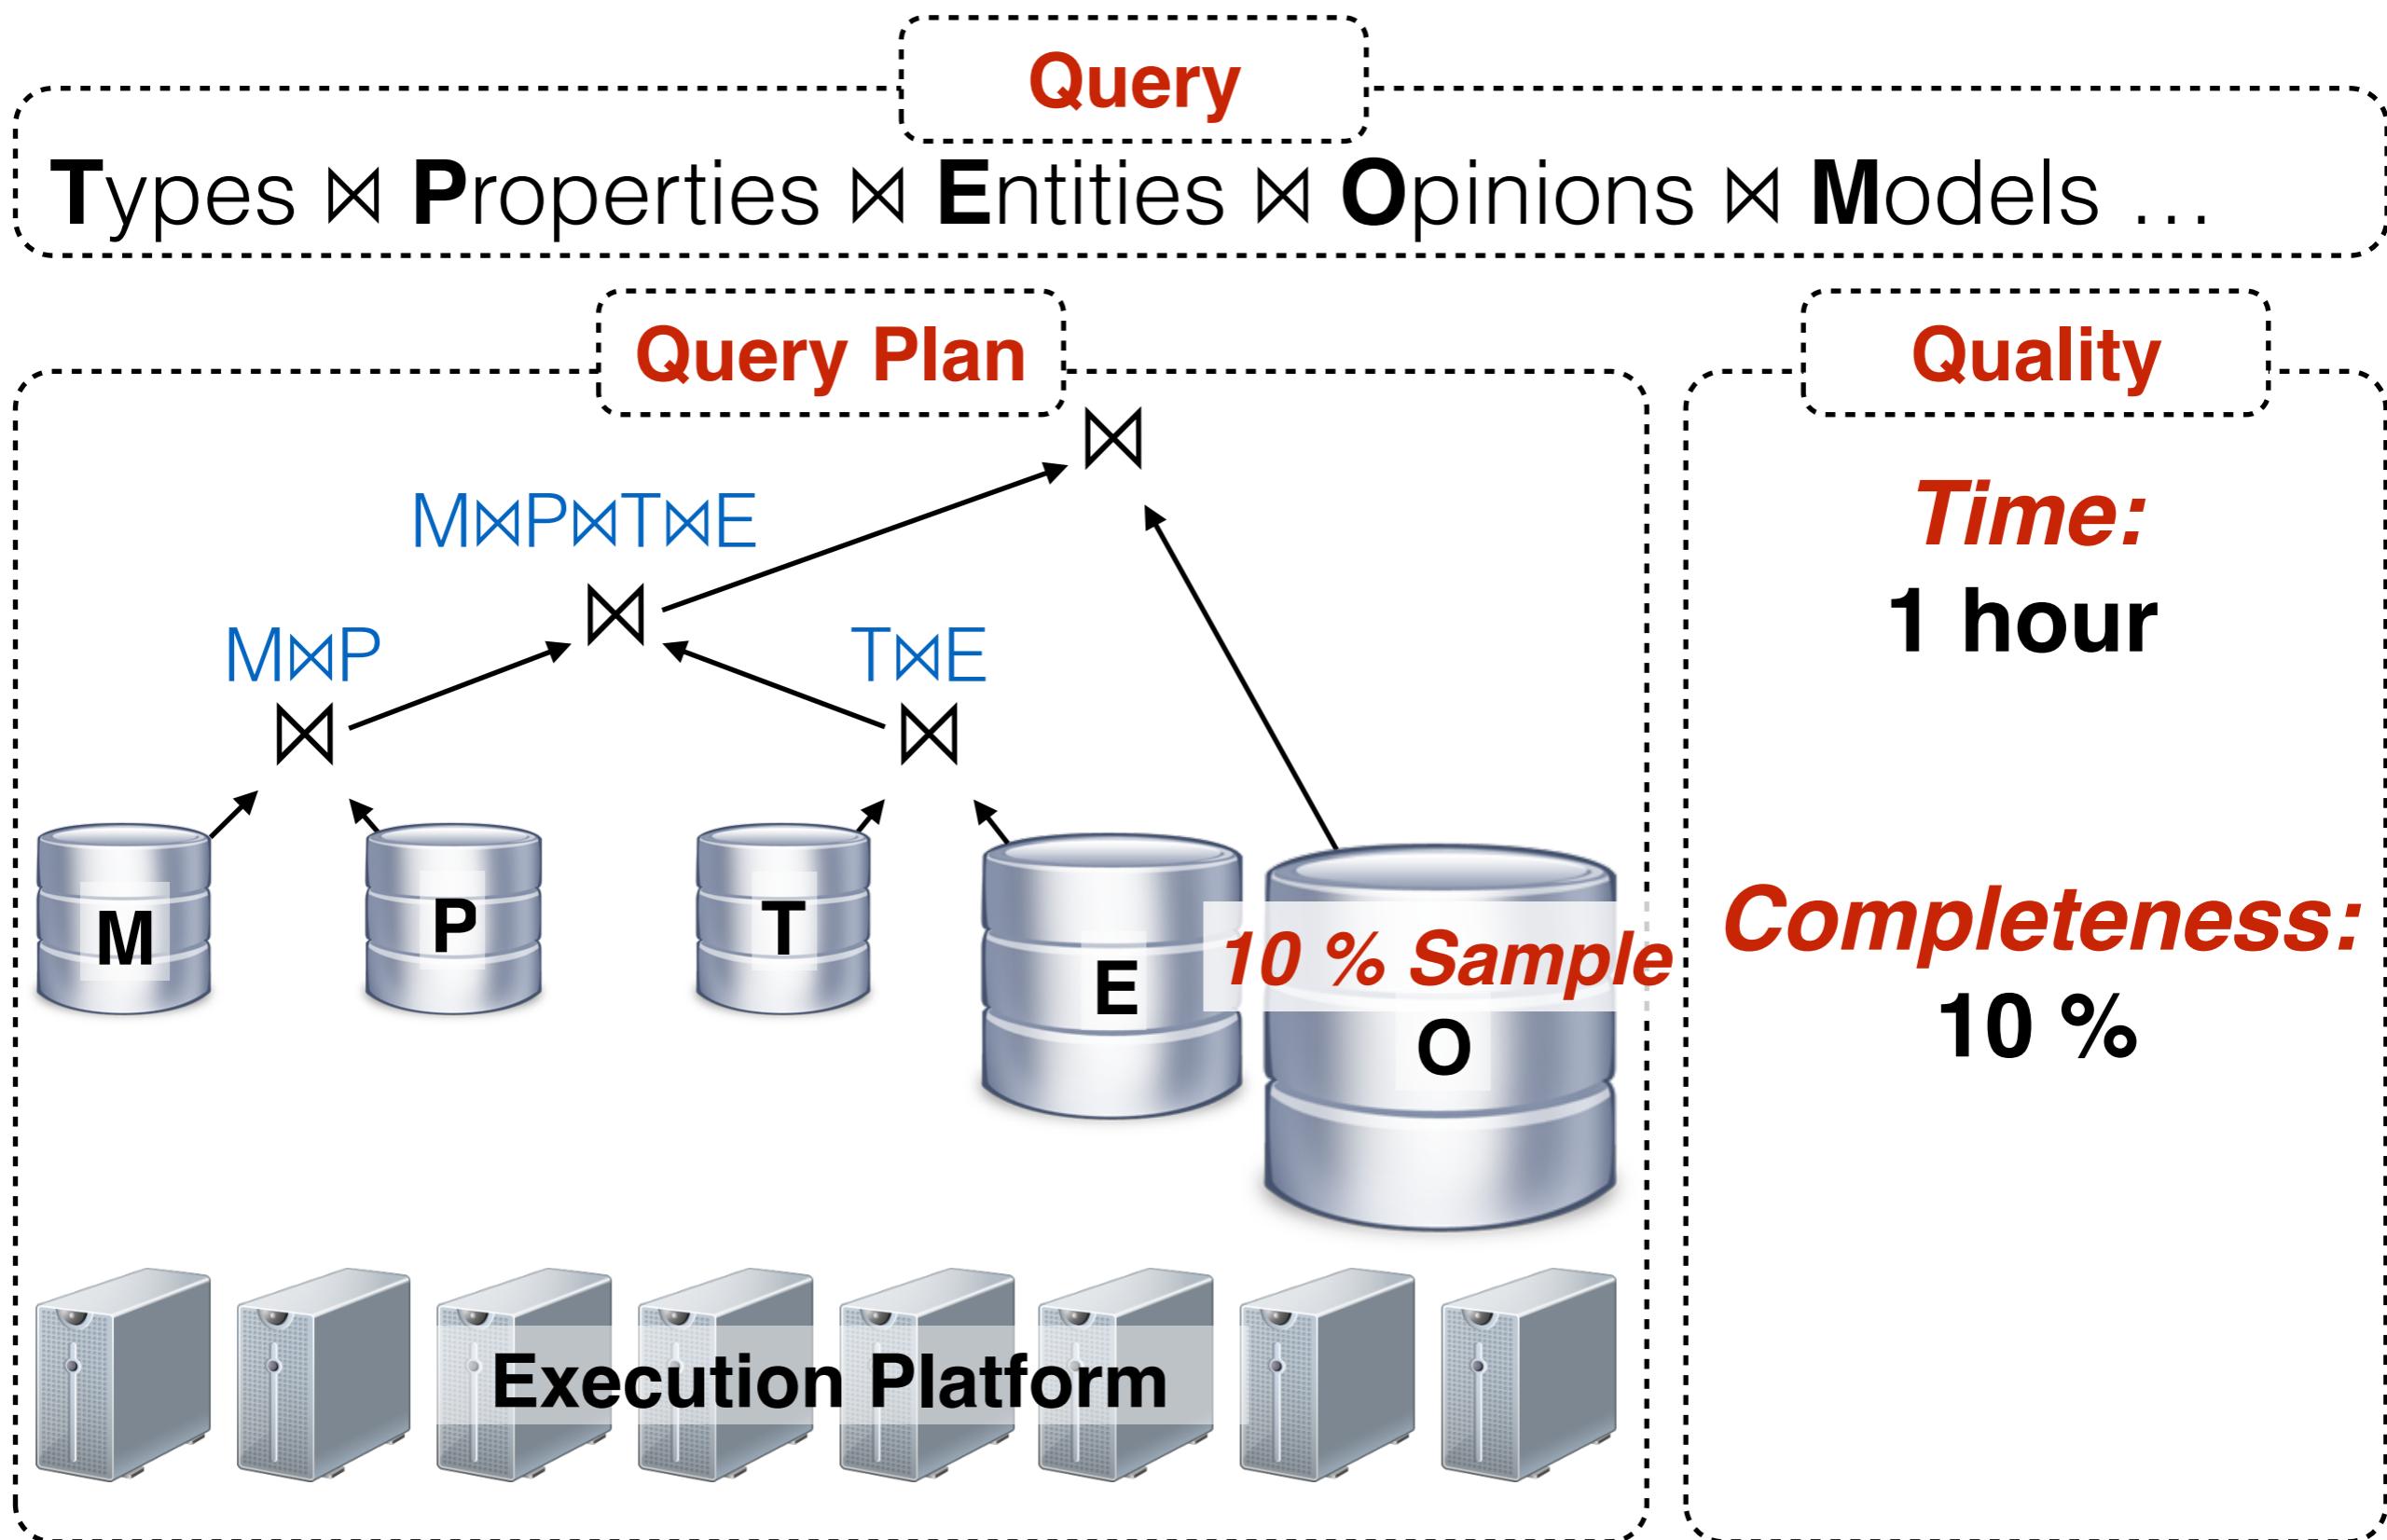
Example: Google



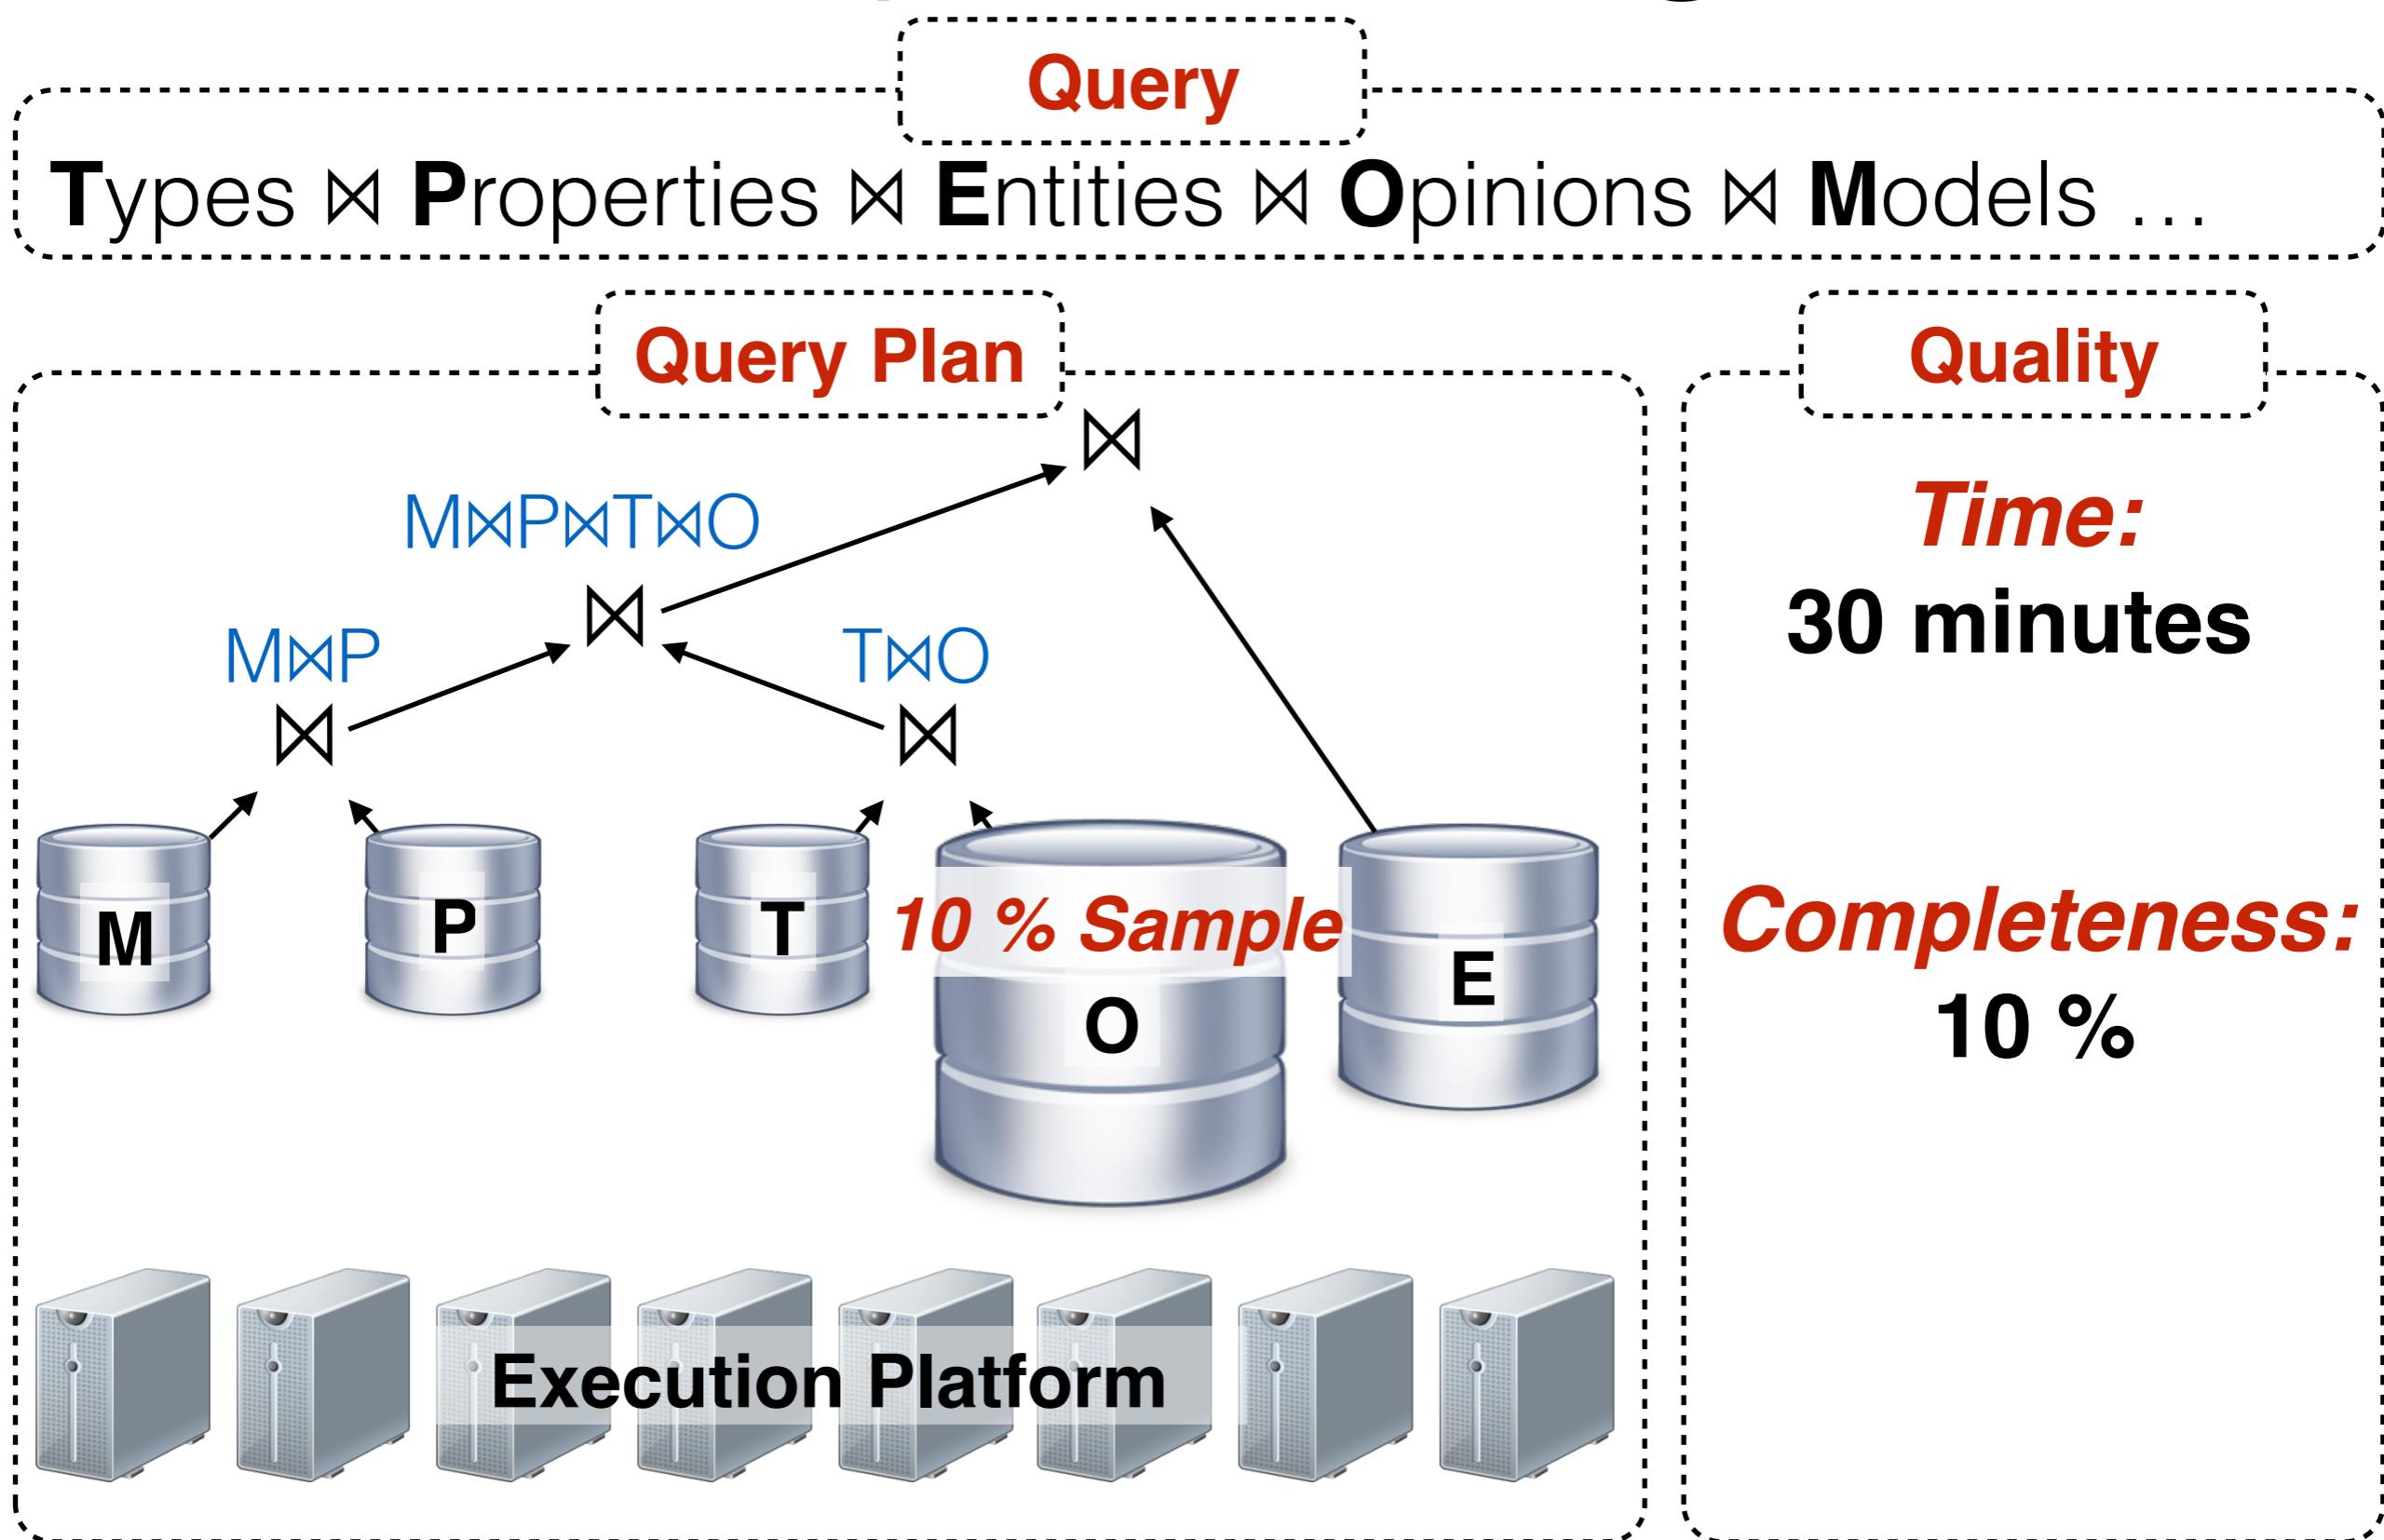
Example: Google



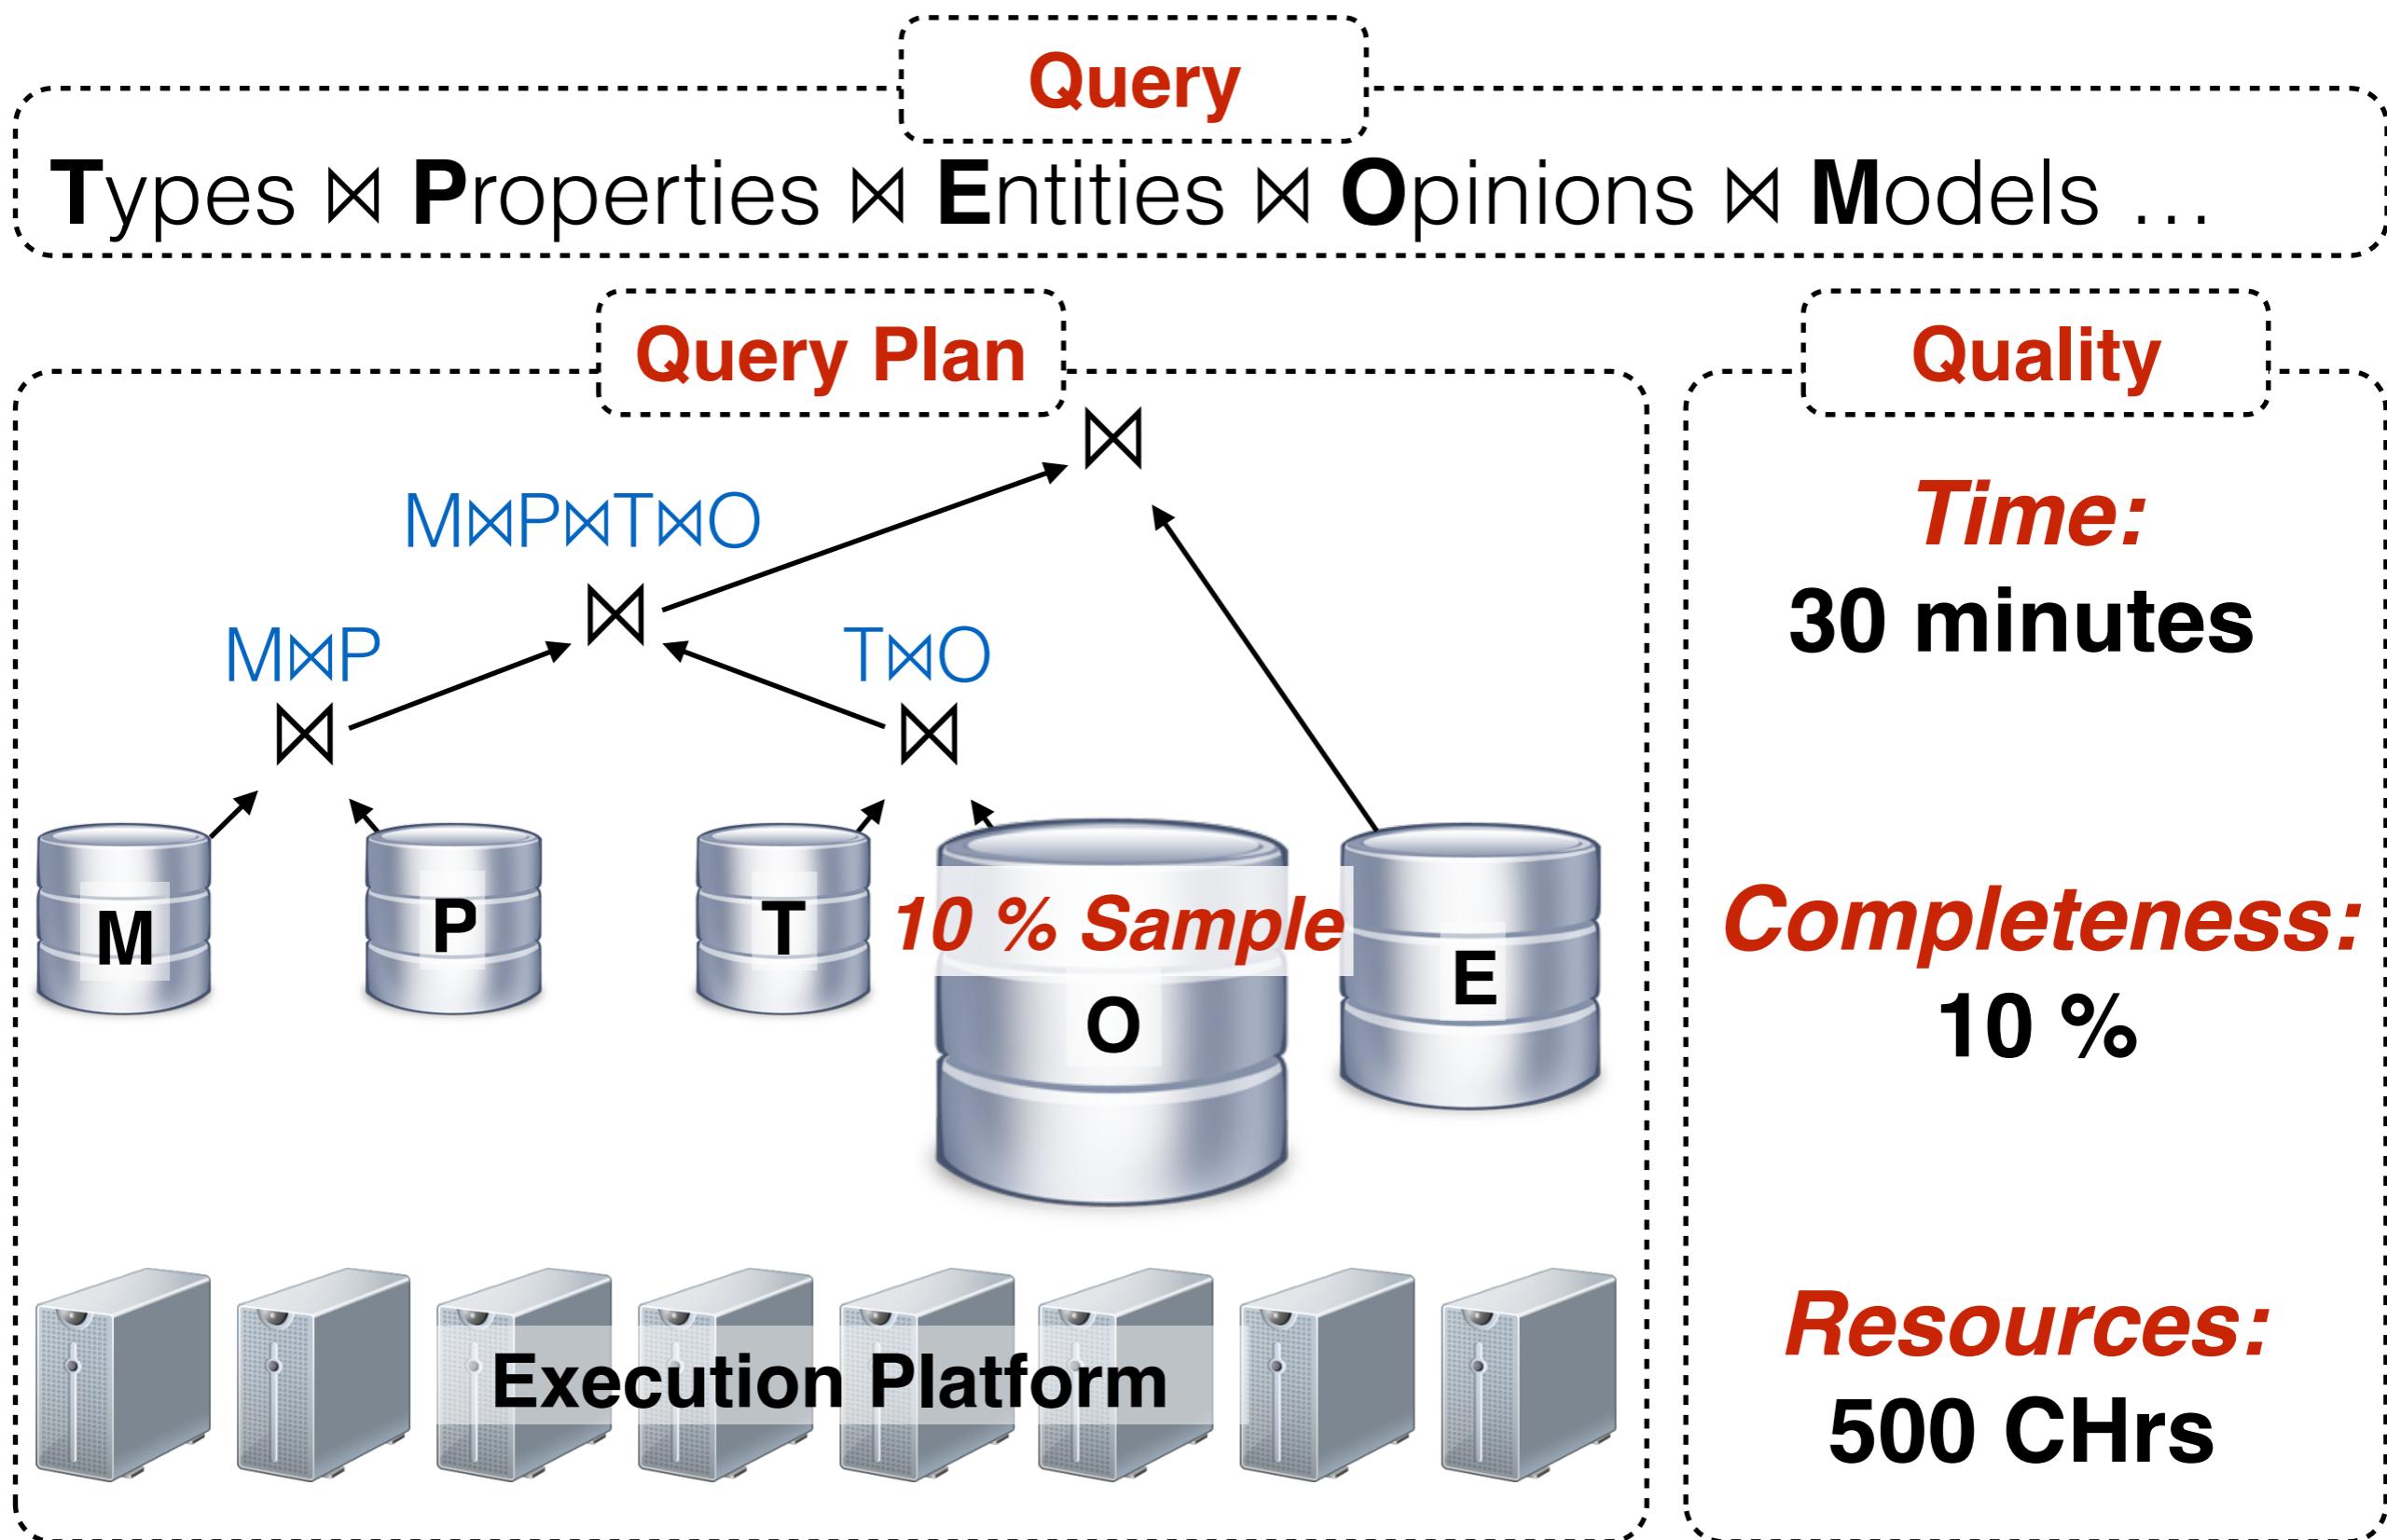
Example: Google



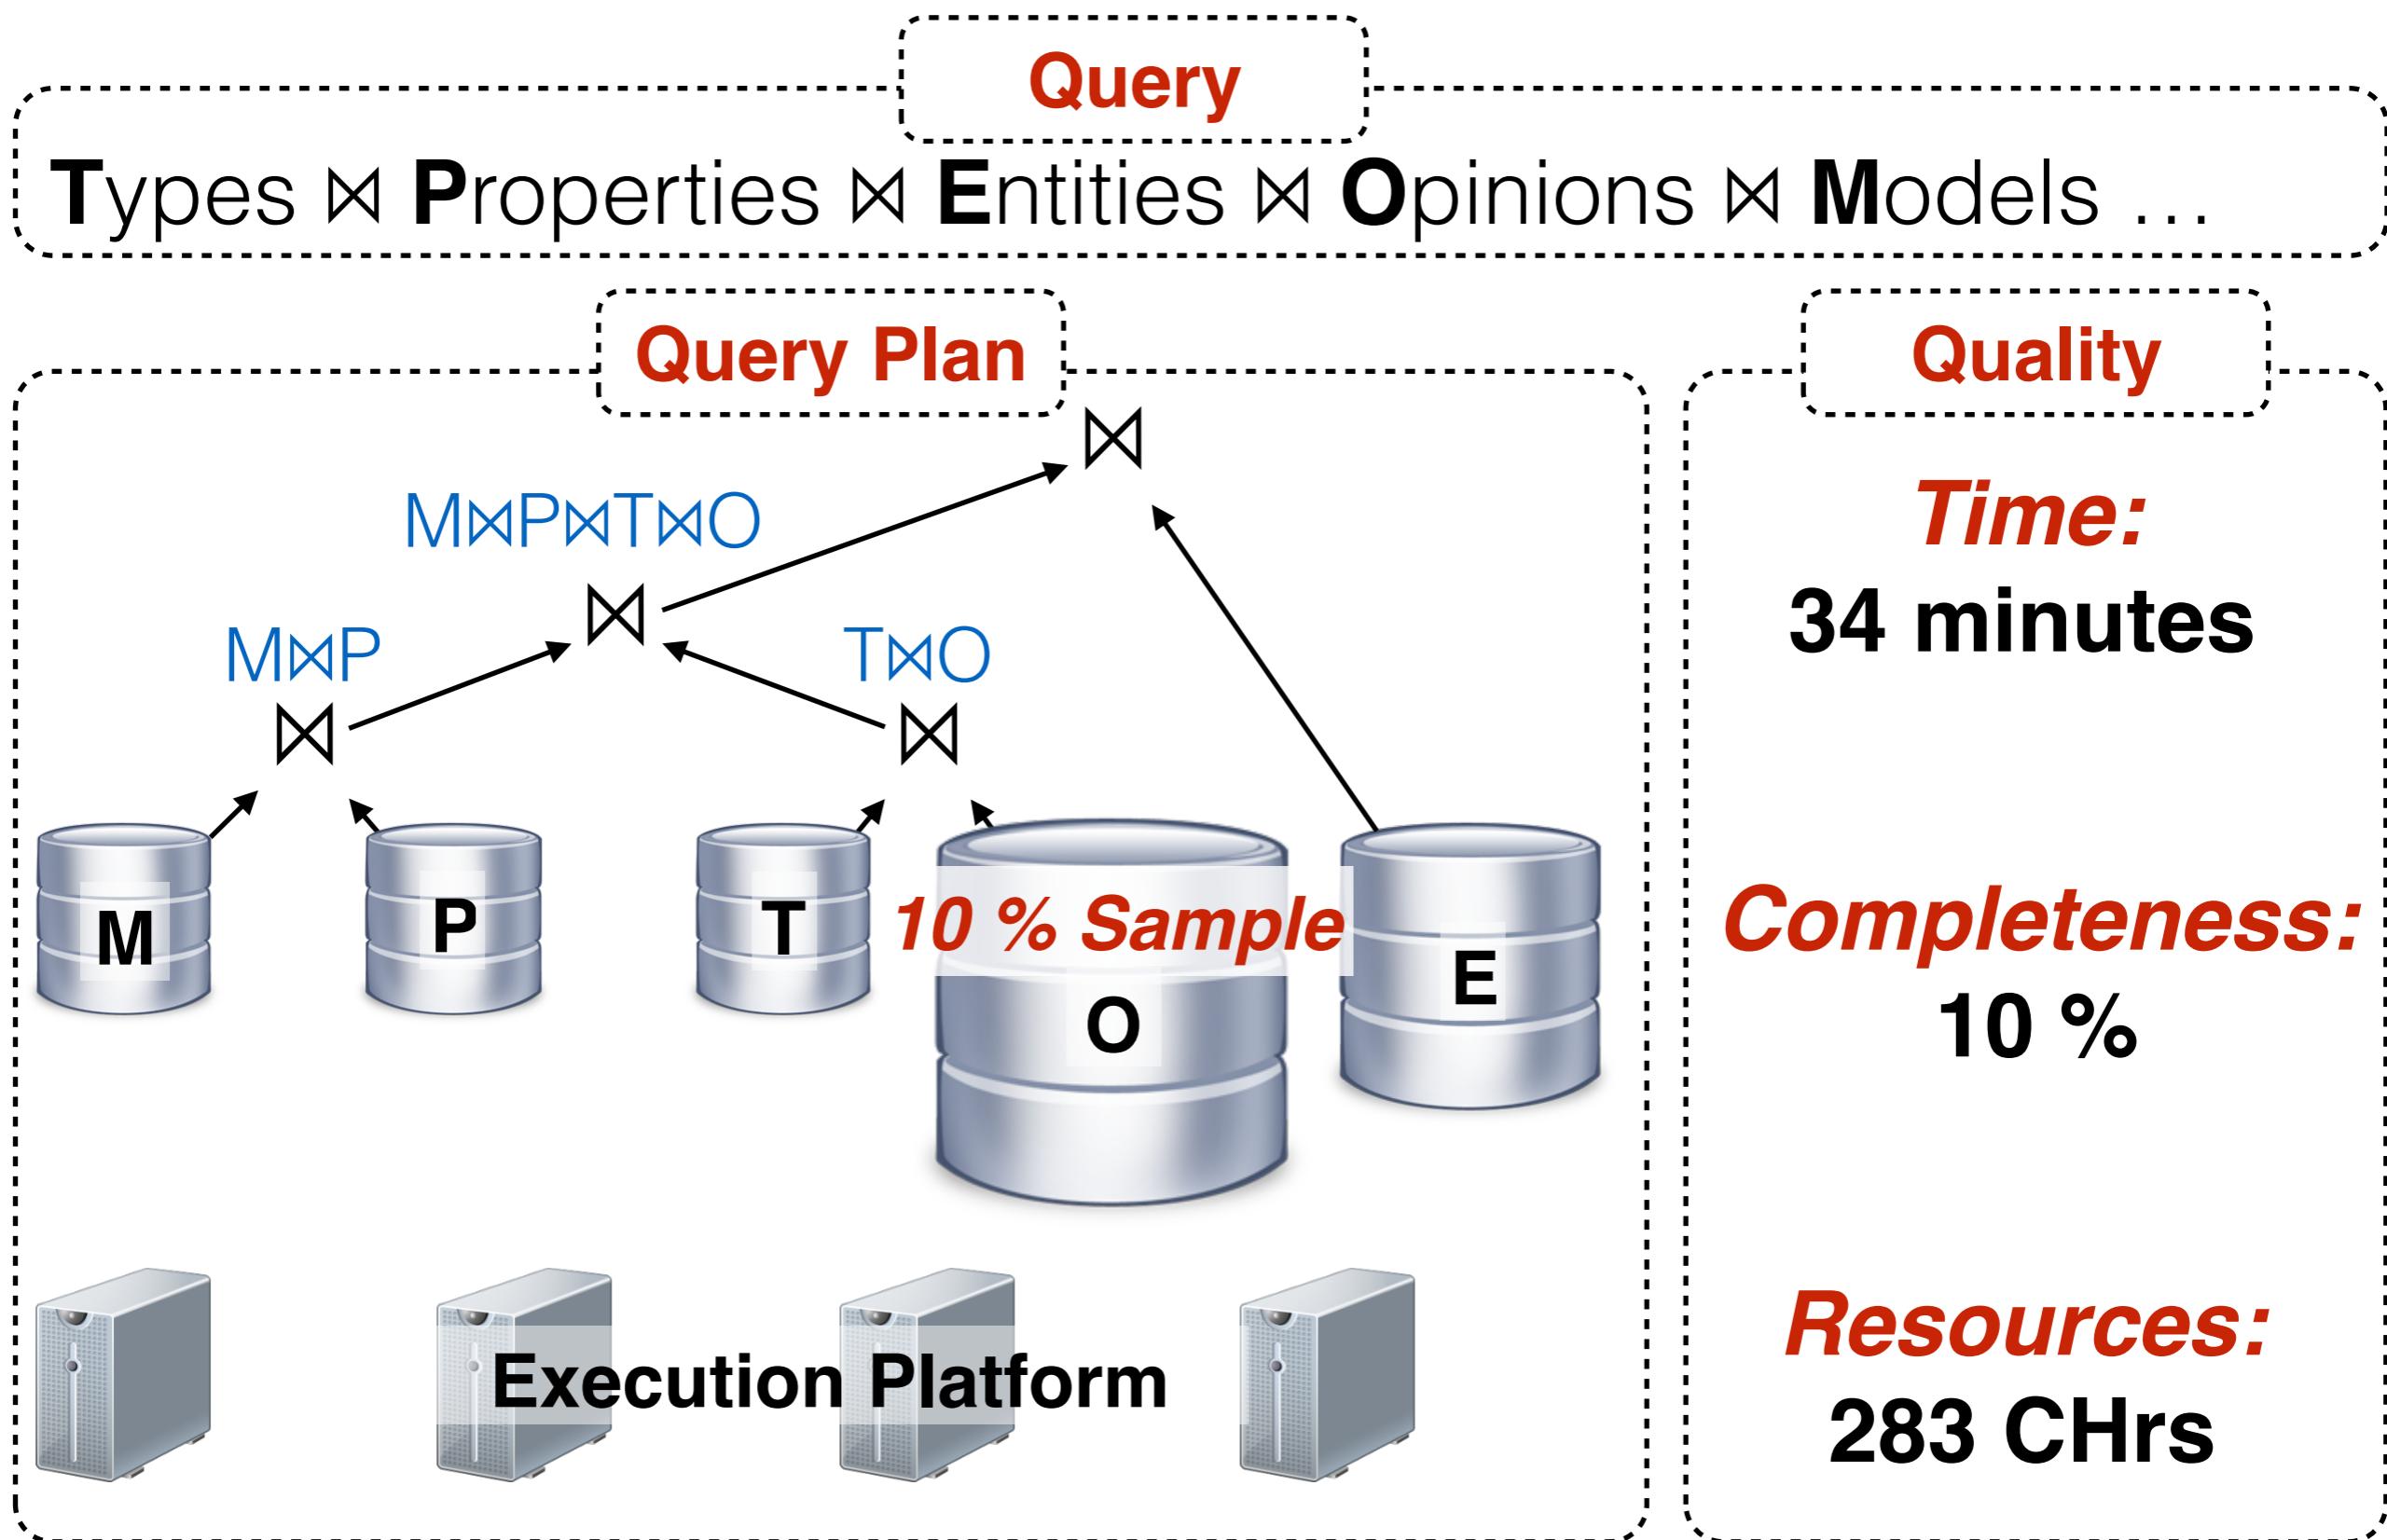
Example: Google



Example: Google



Example: Google

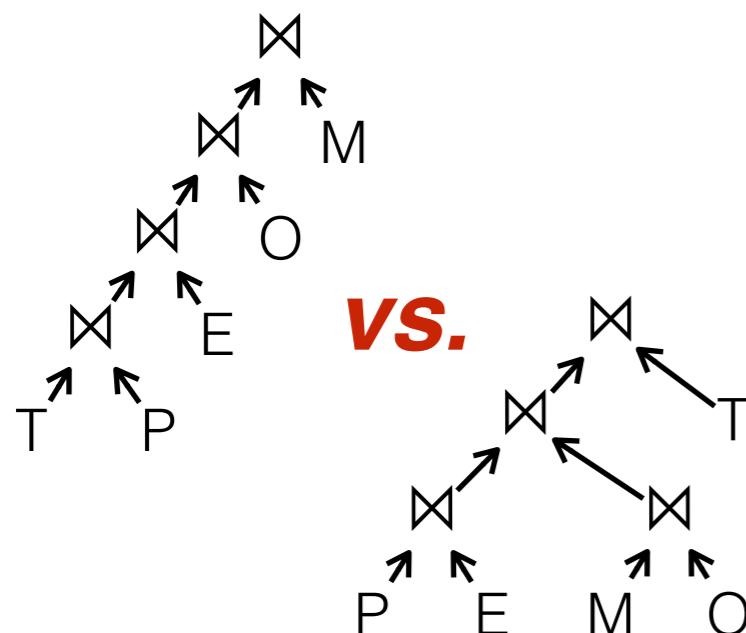


Example Summary

Query

Types \bowtie Properties \bowtie Entities \bowtie Opinions \bowtie Models ...

Join Order

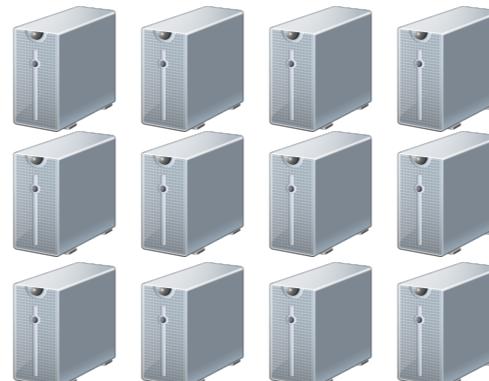


Decisions

Platform



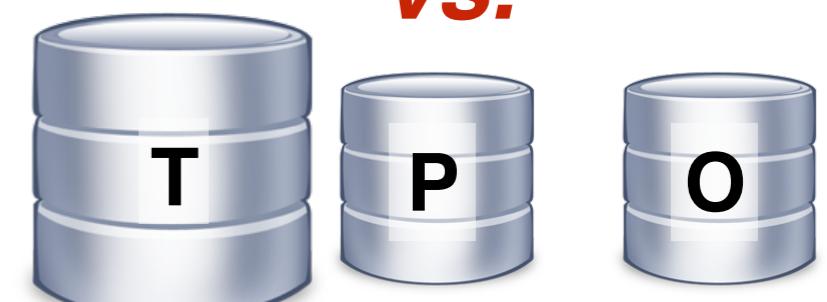
vs.



Sample Sizes



vs.



Execution
Time



Resource
Consumption

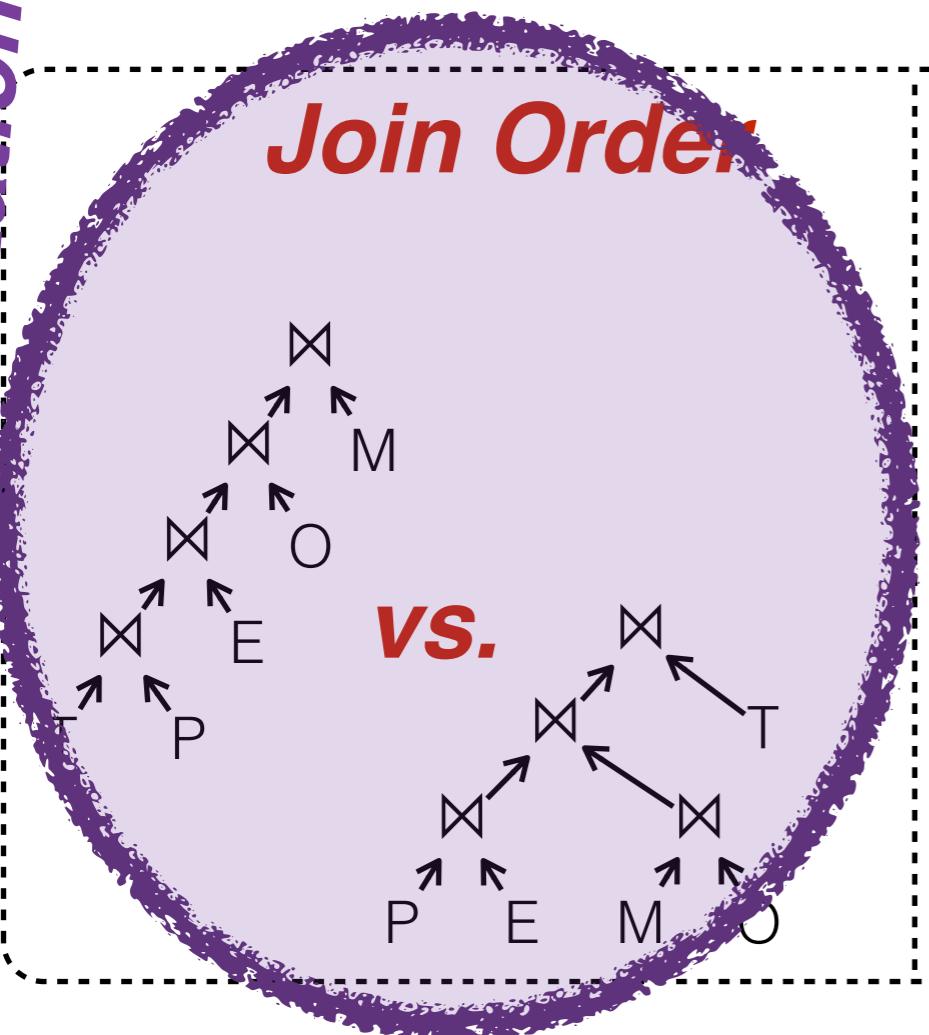


Result
Completeness

Metrics

Example Summary

Classical Query Optimization



Query

Types \bowtie Properties \bowtie Entities \bowtie Opinions \bowtie Models ...

Decisions

Platform



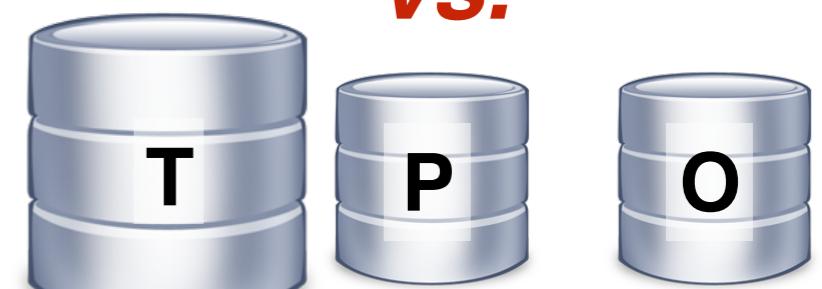
vs.



Sample Sizes



vs.



Metrics

Resource Consumption

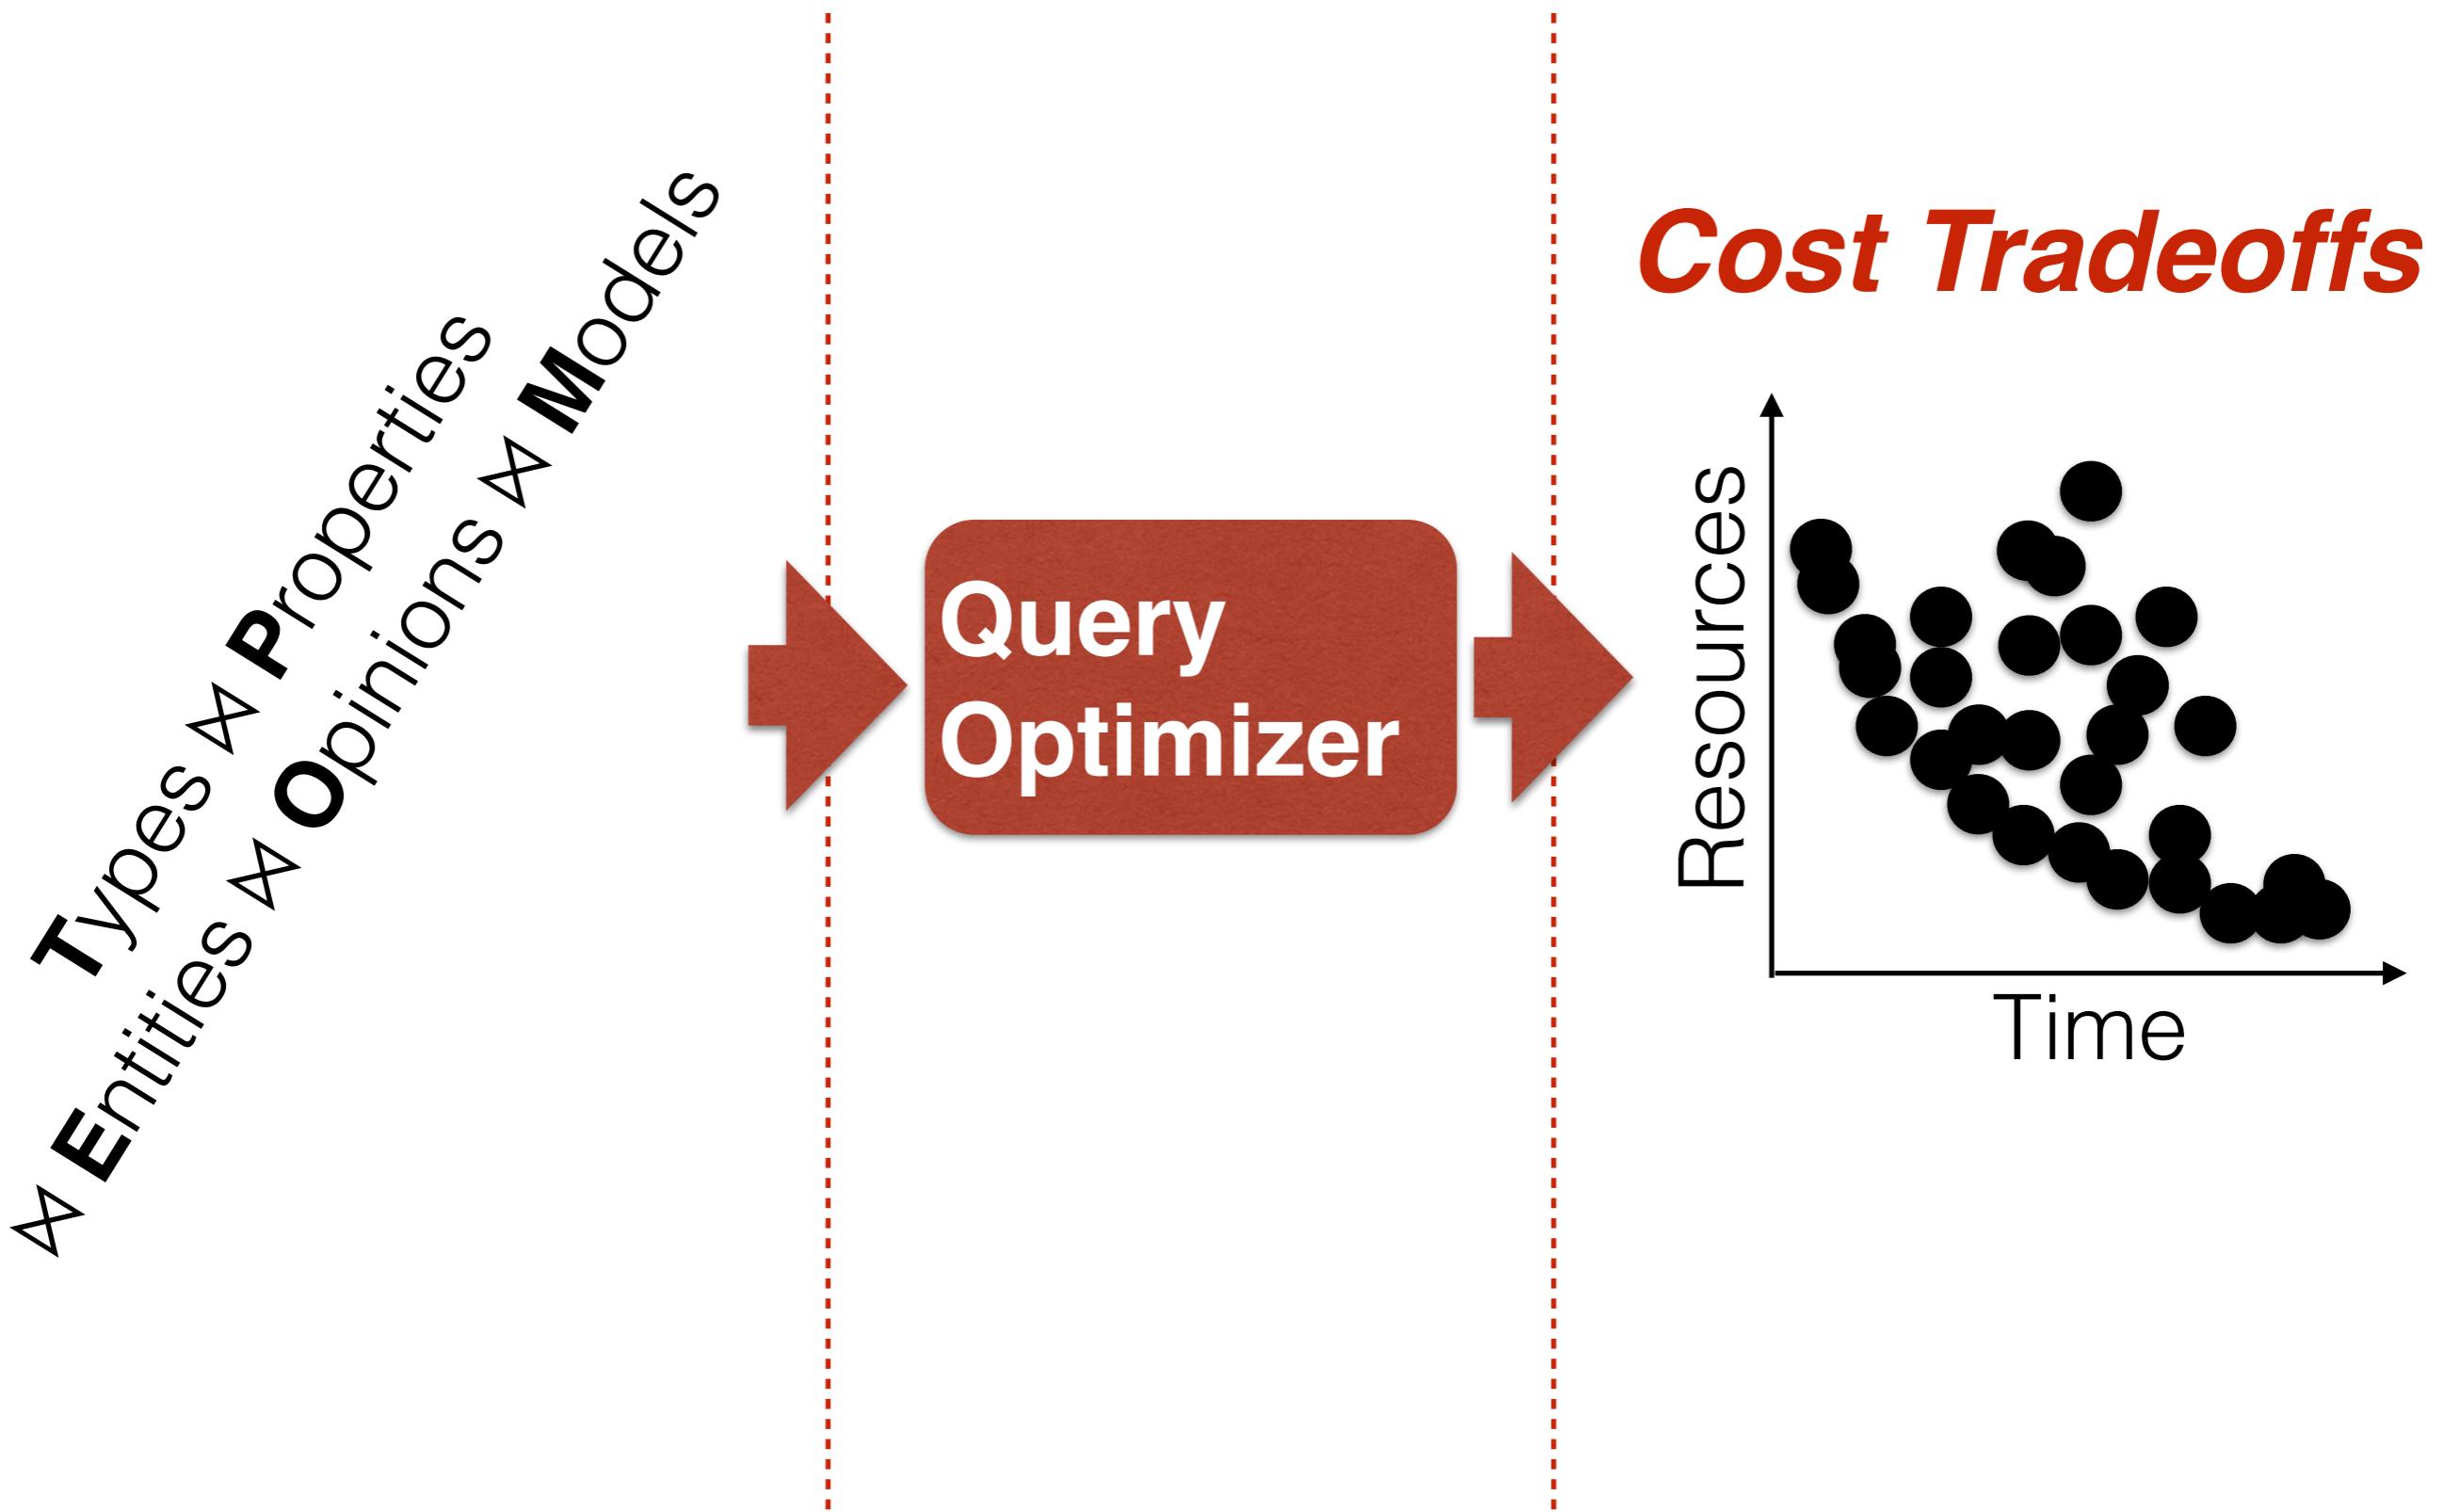
Execution Time



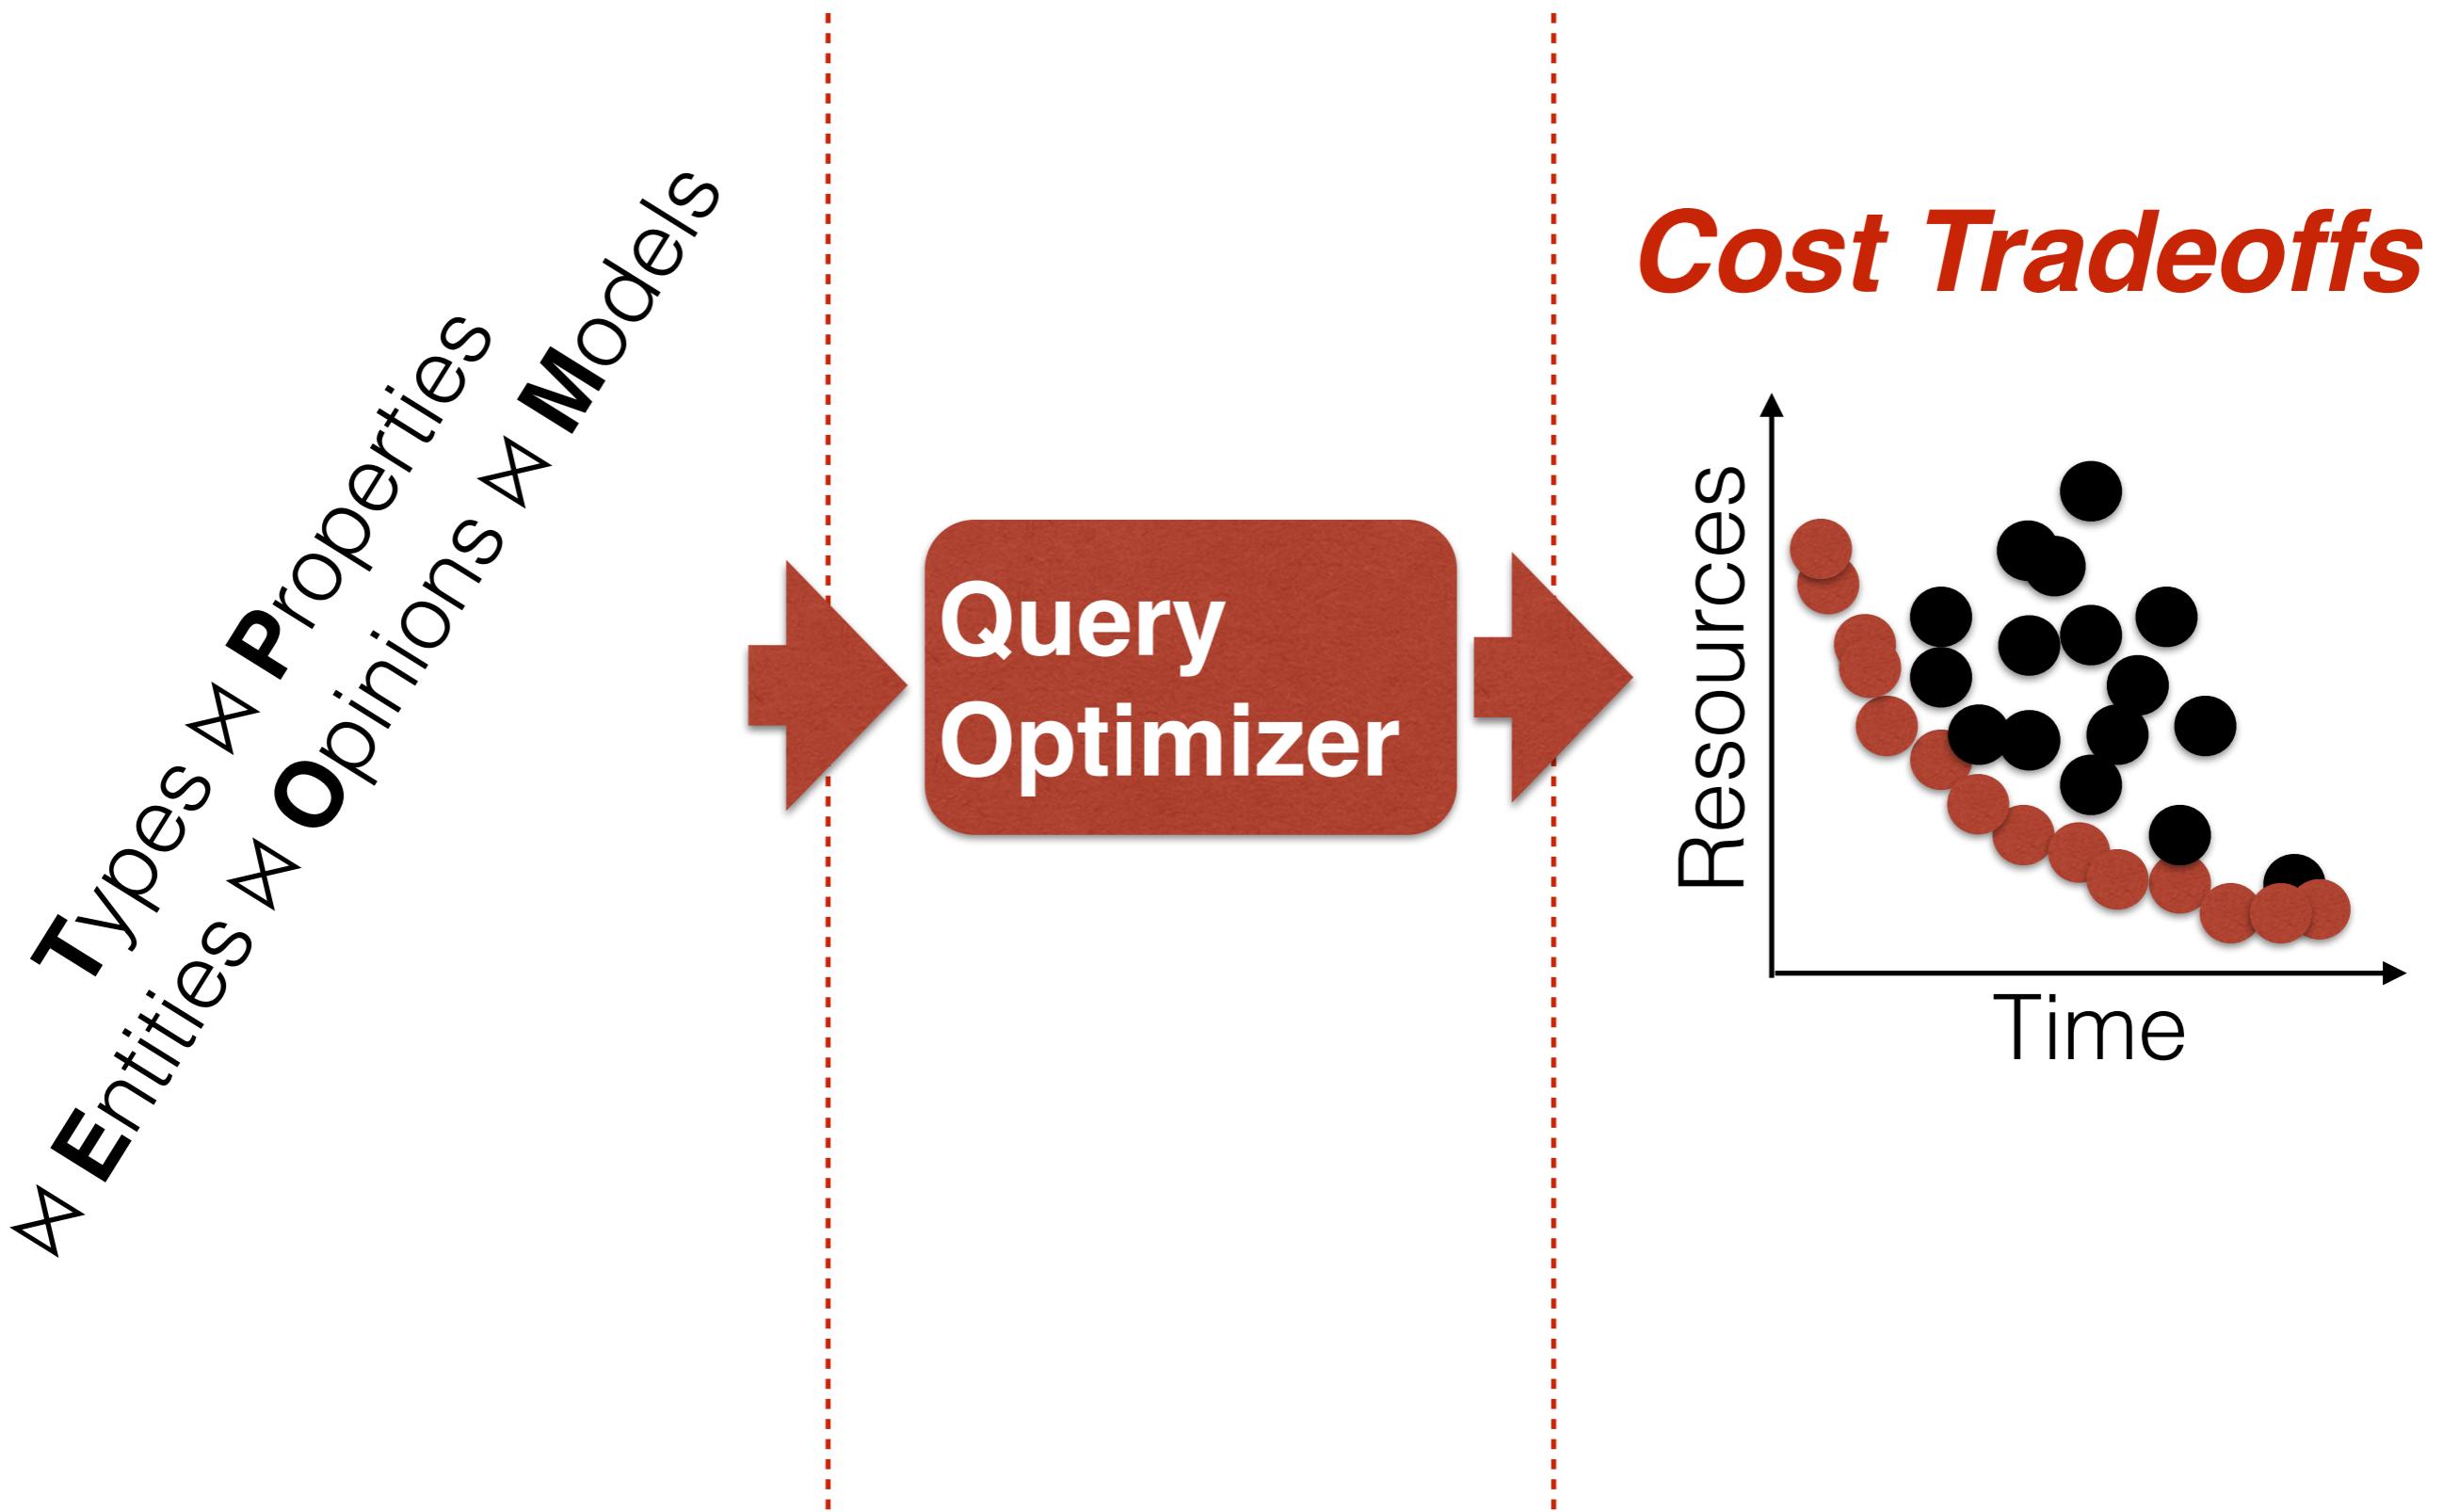
Result Completeness



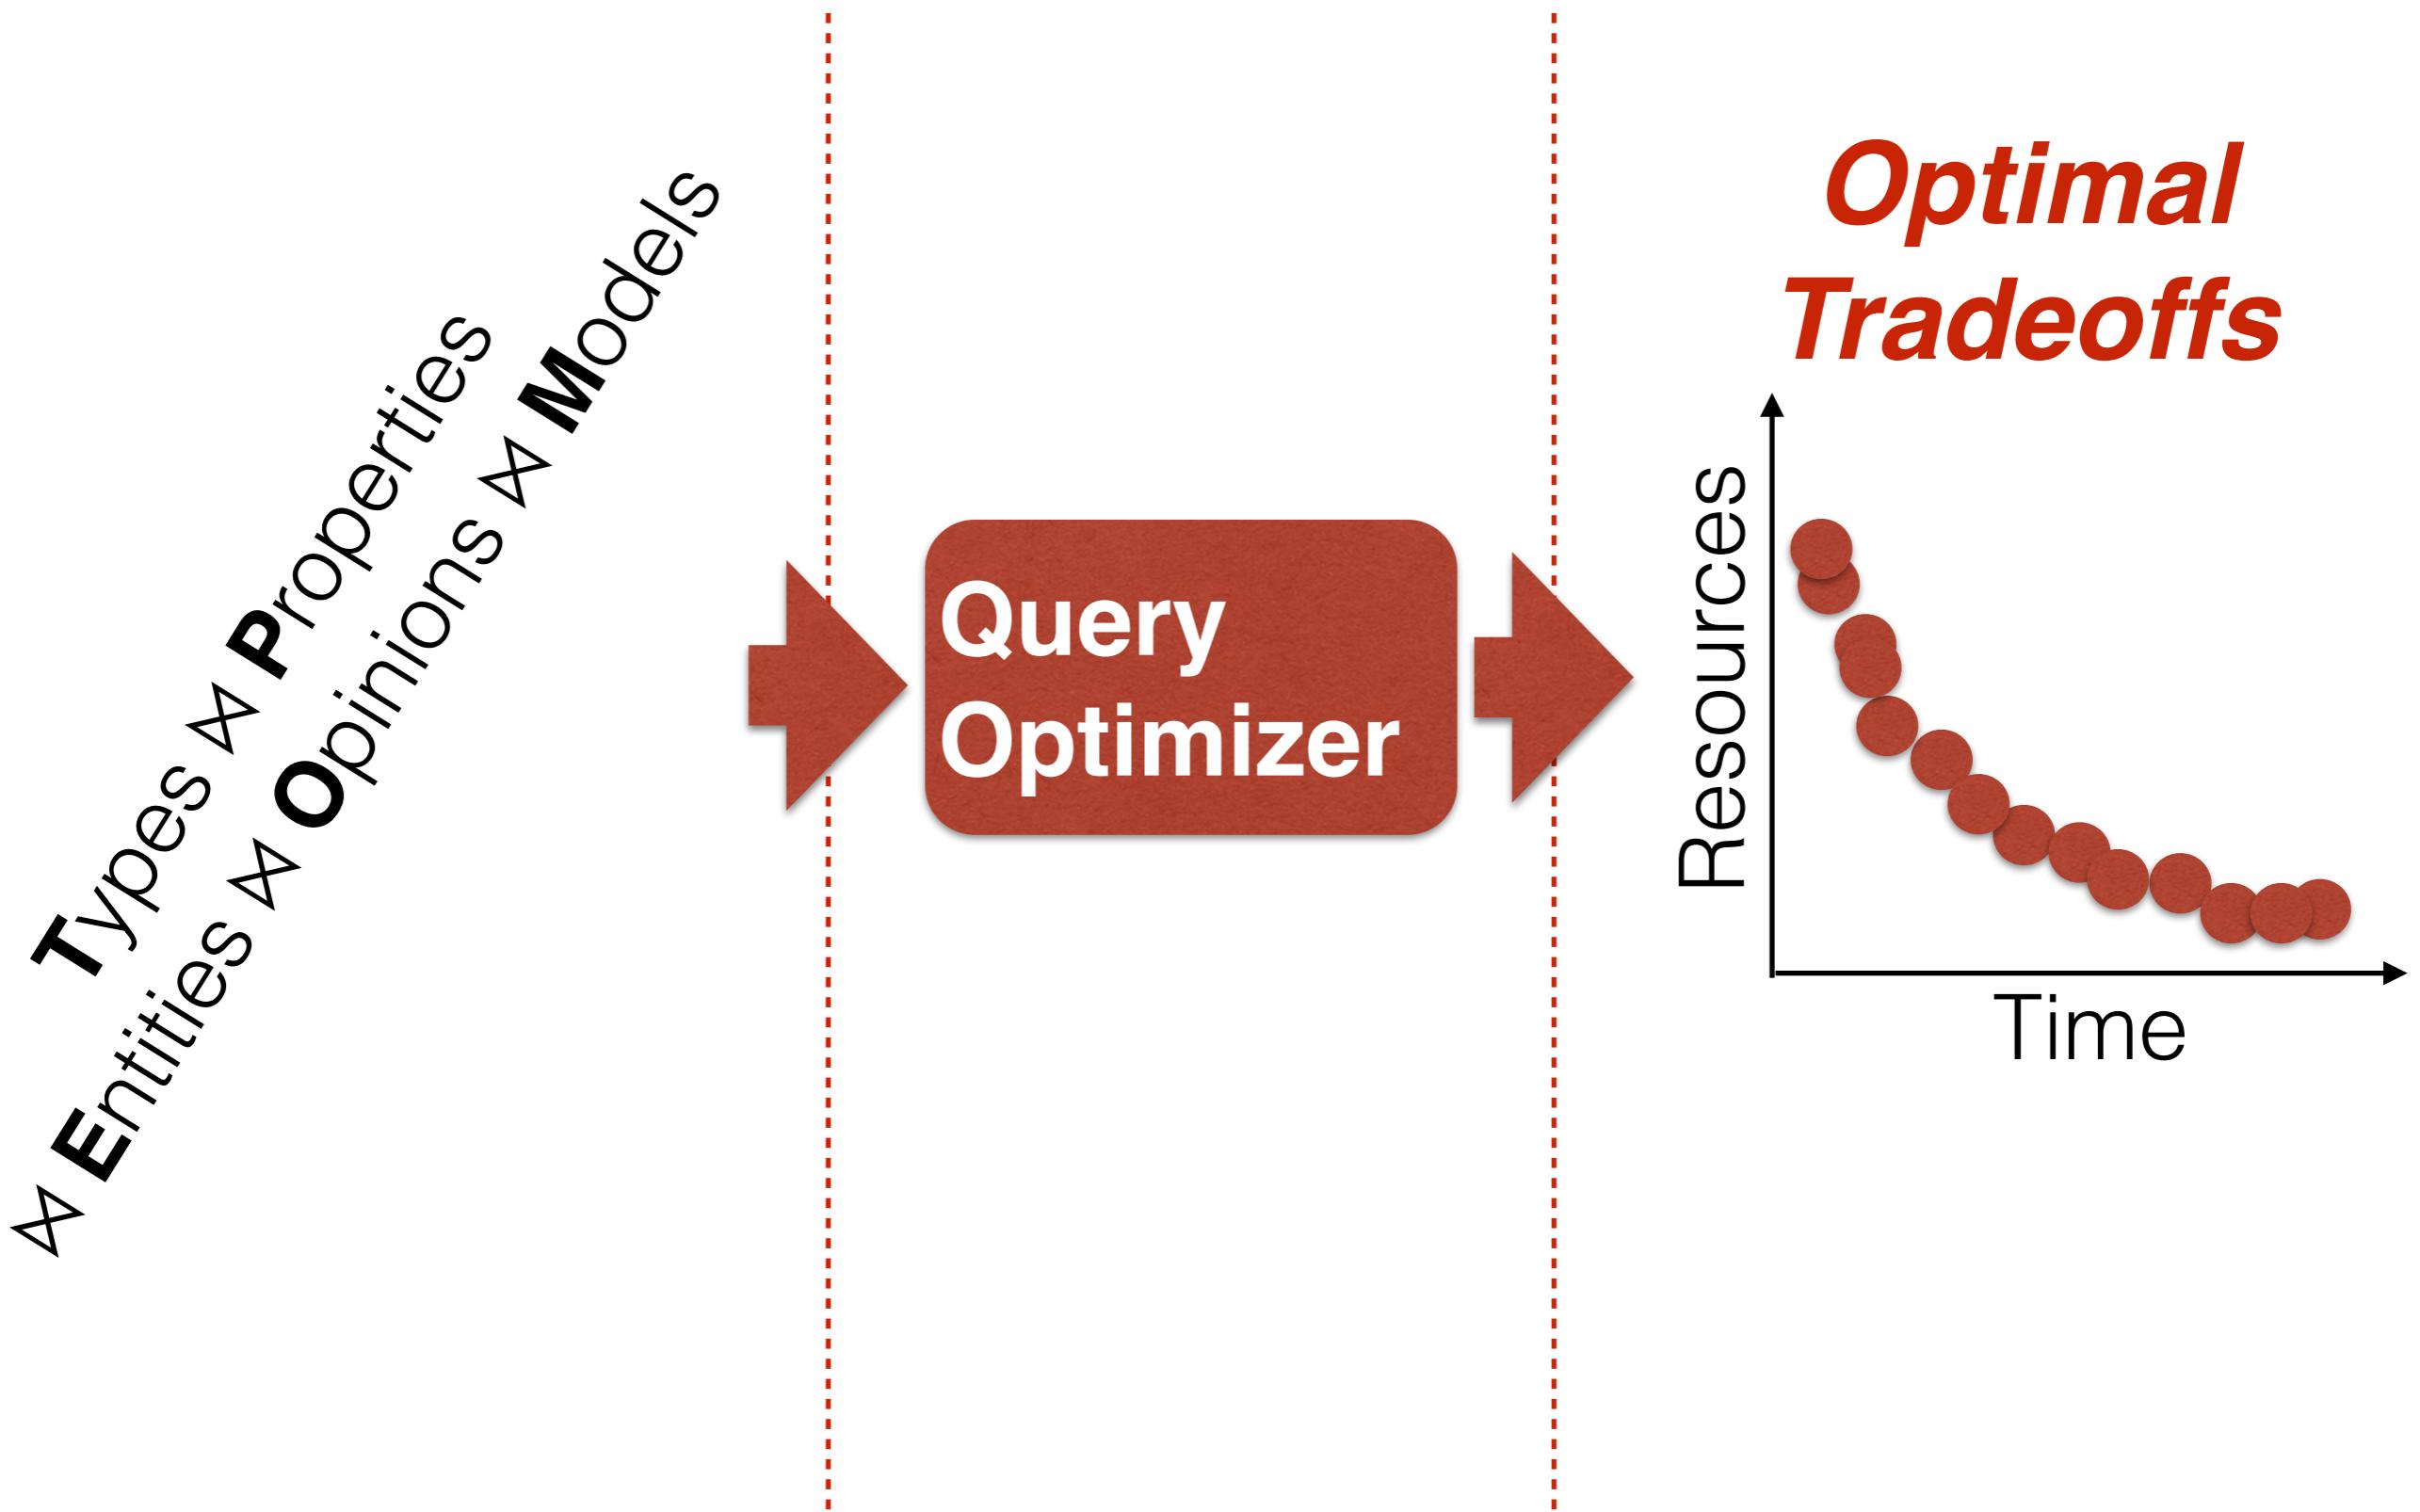
Example: Google



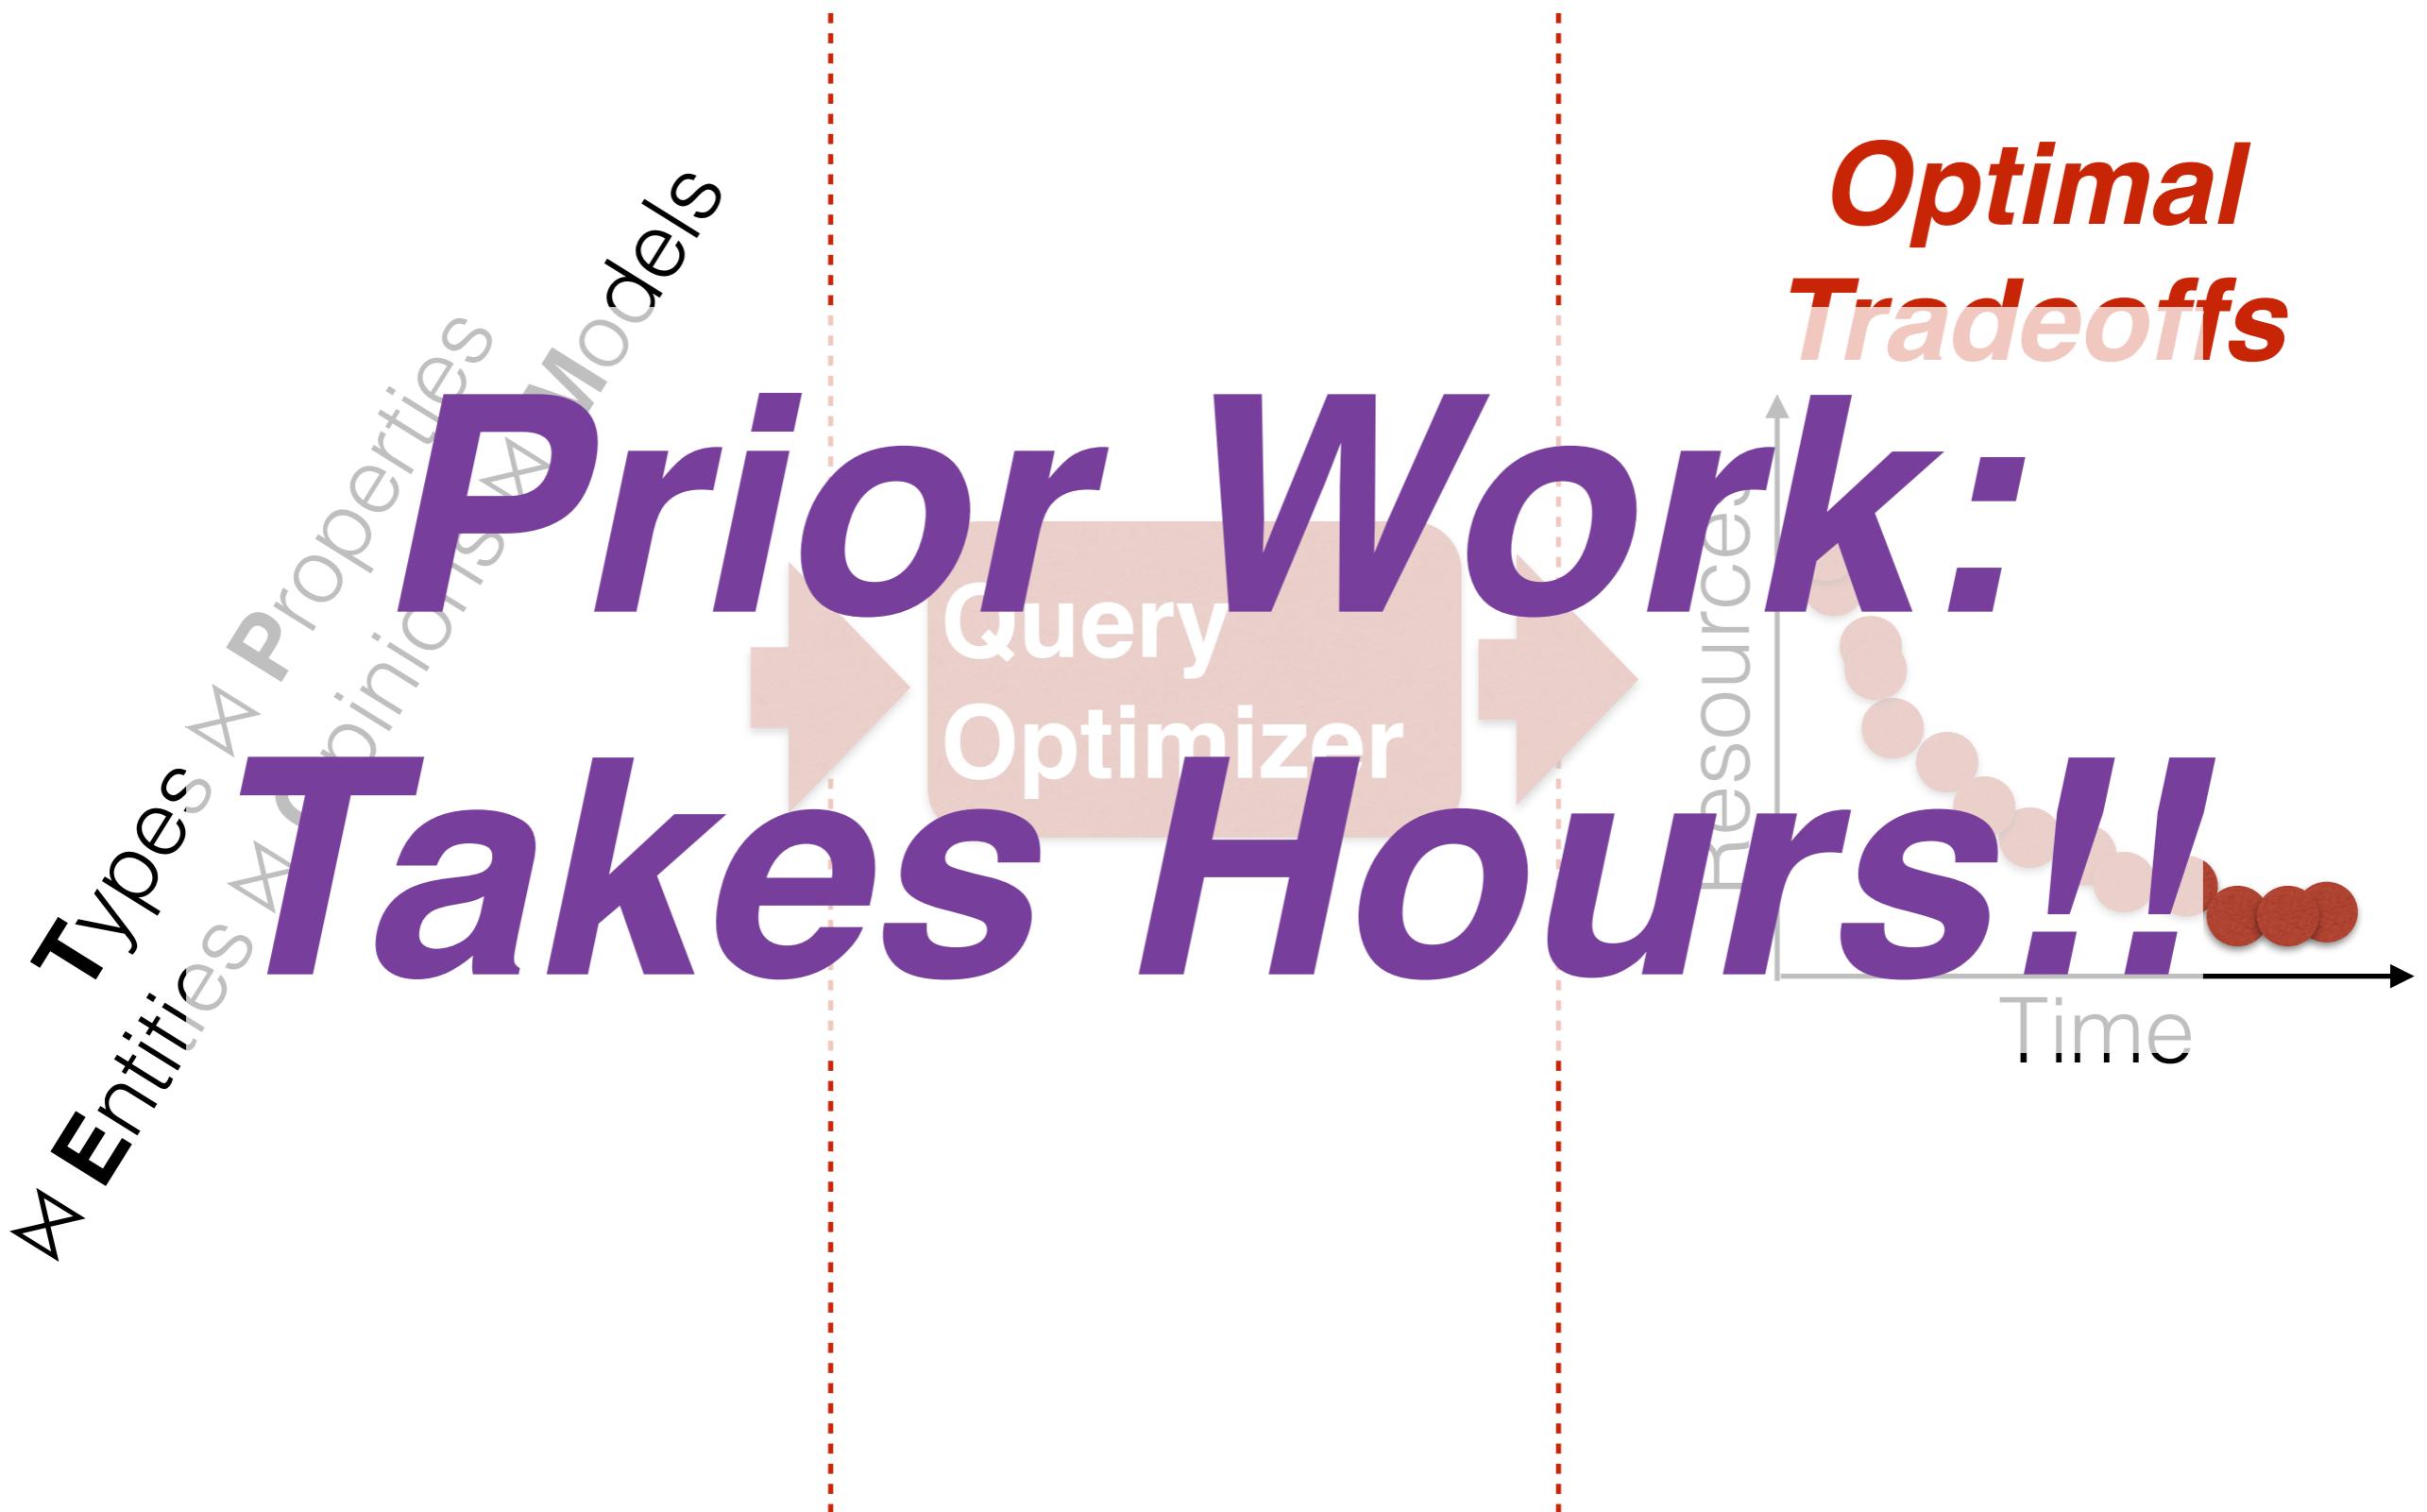
Example: Google



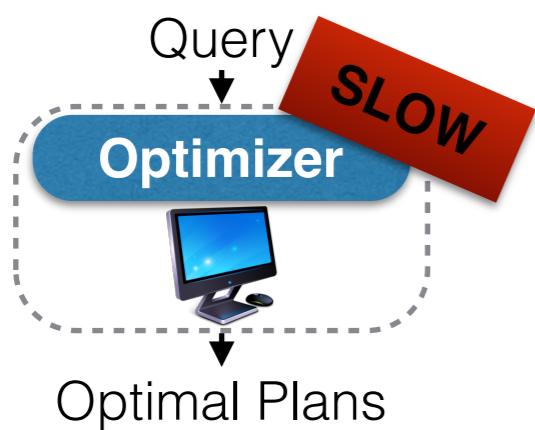
Example: Google



Example: Google



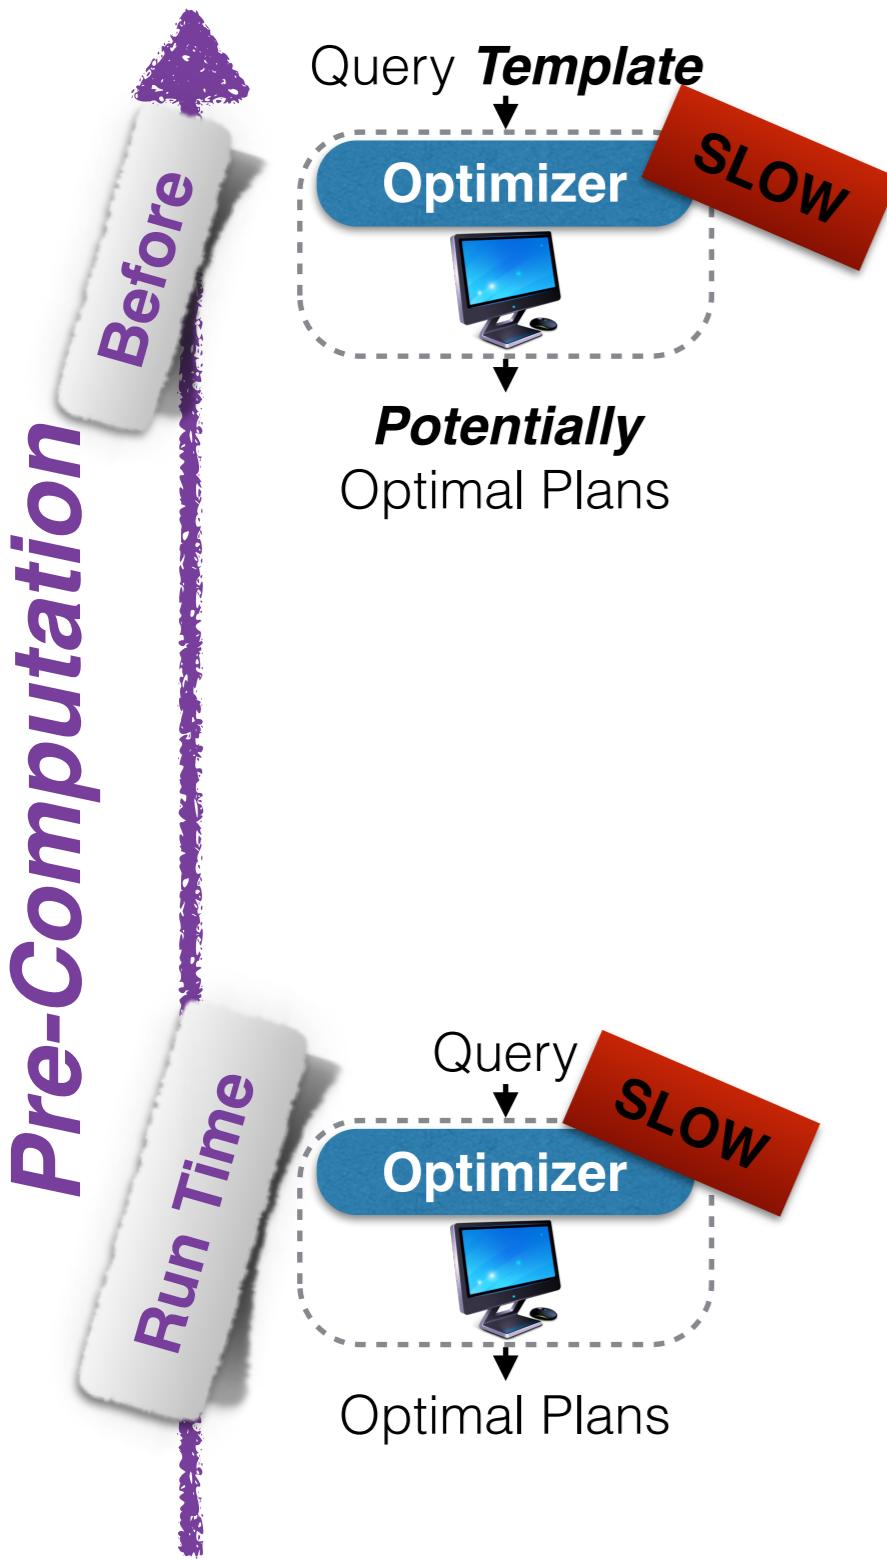
Dissertation Overview



Dissertation Overview



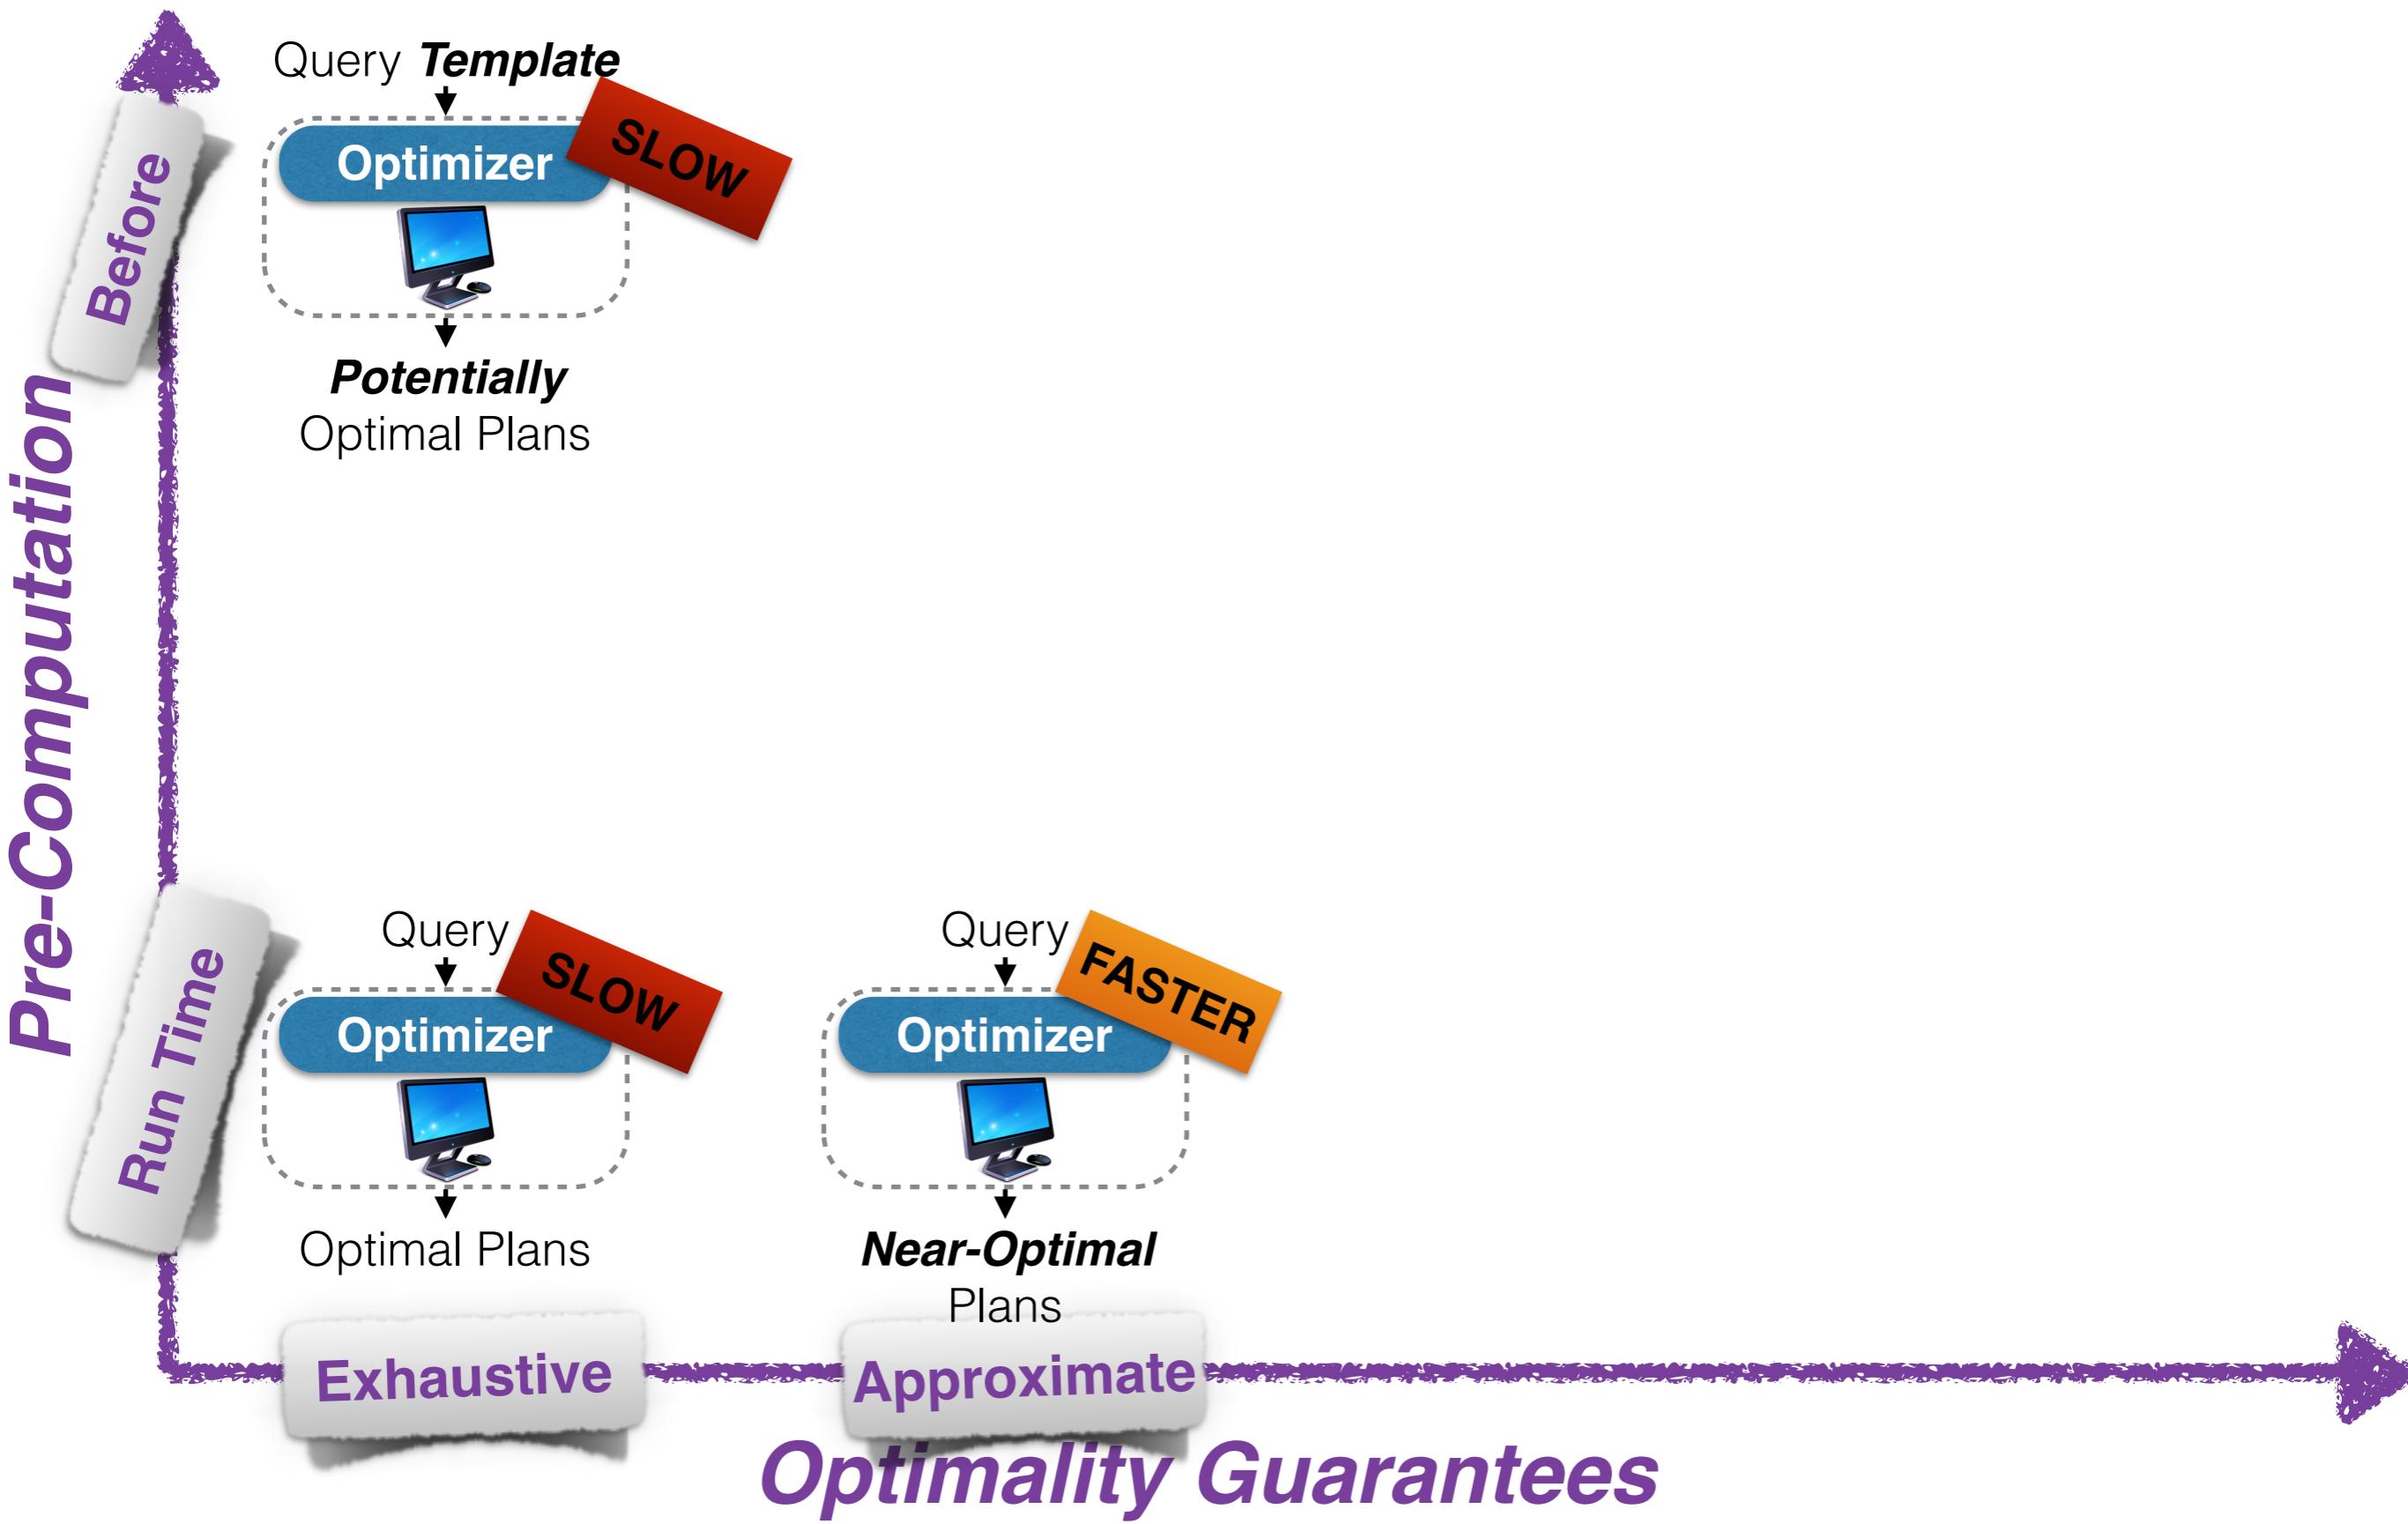
Dissertation Overview



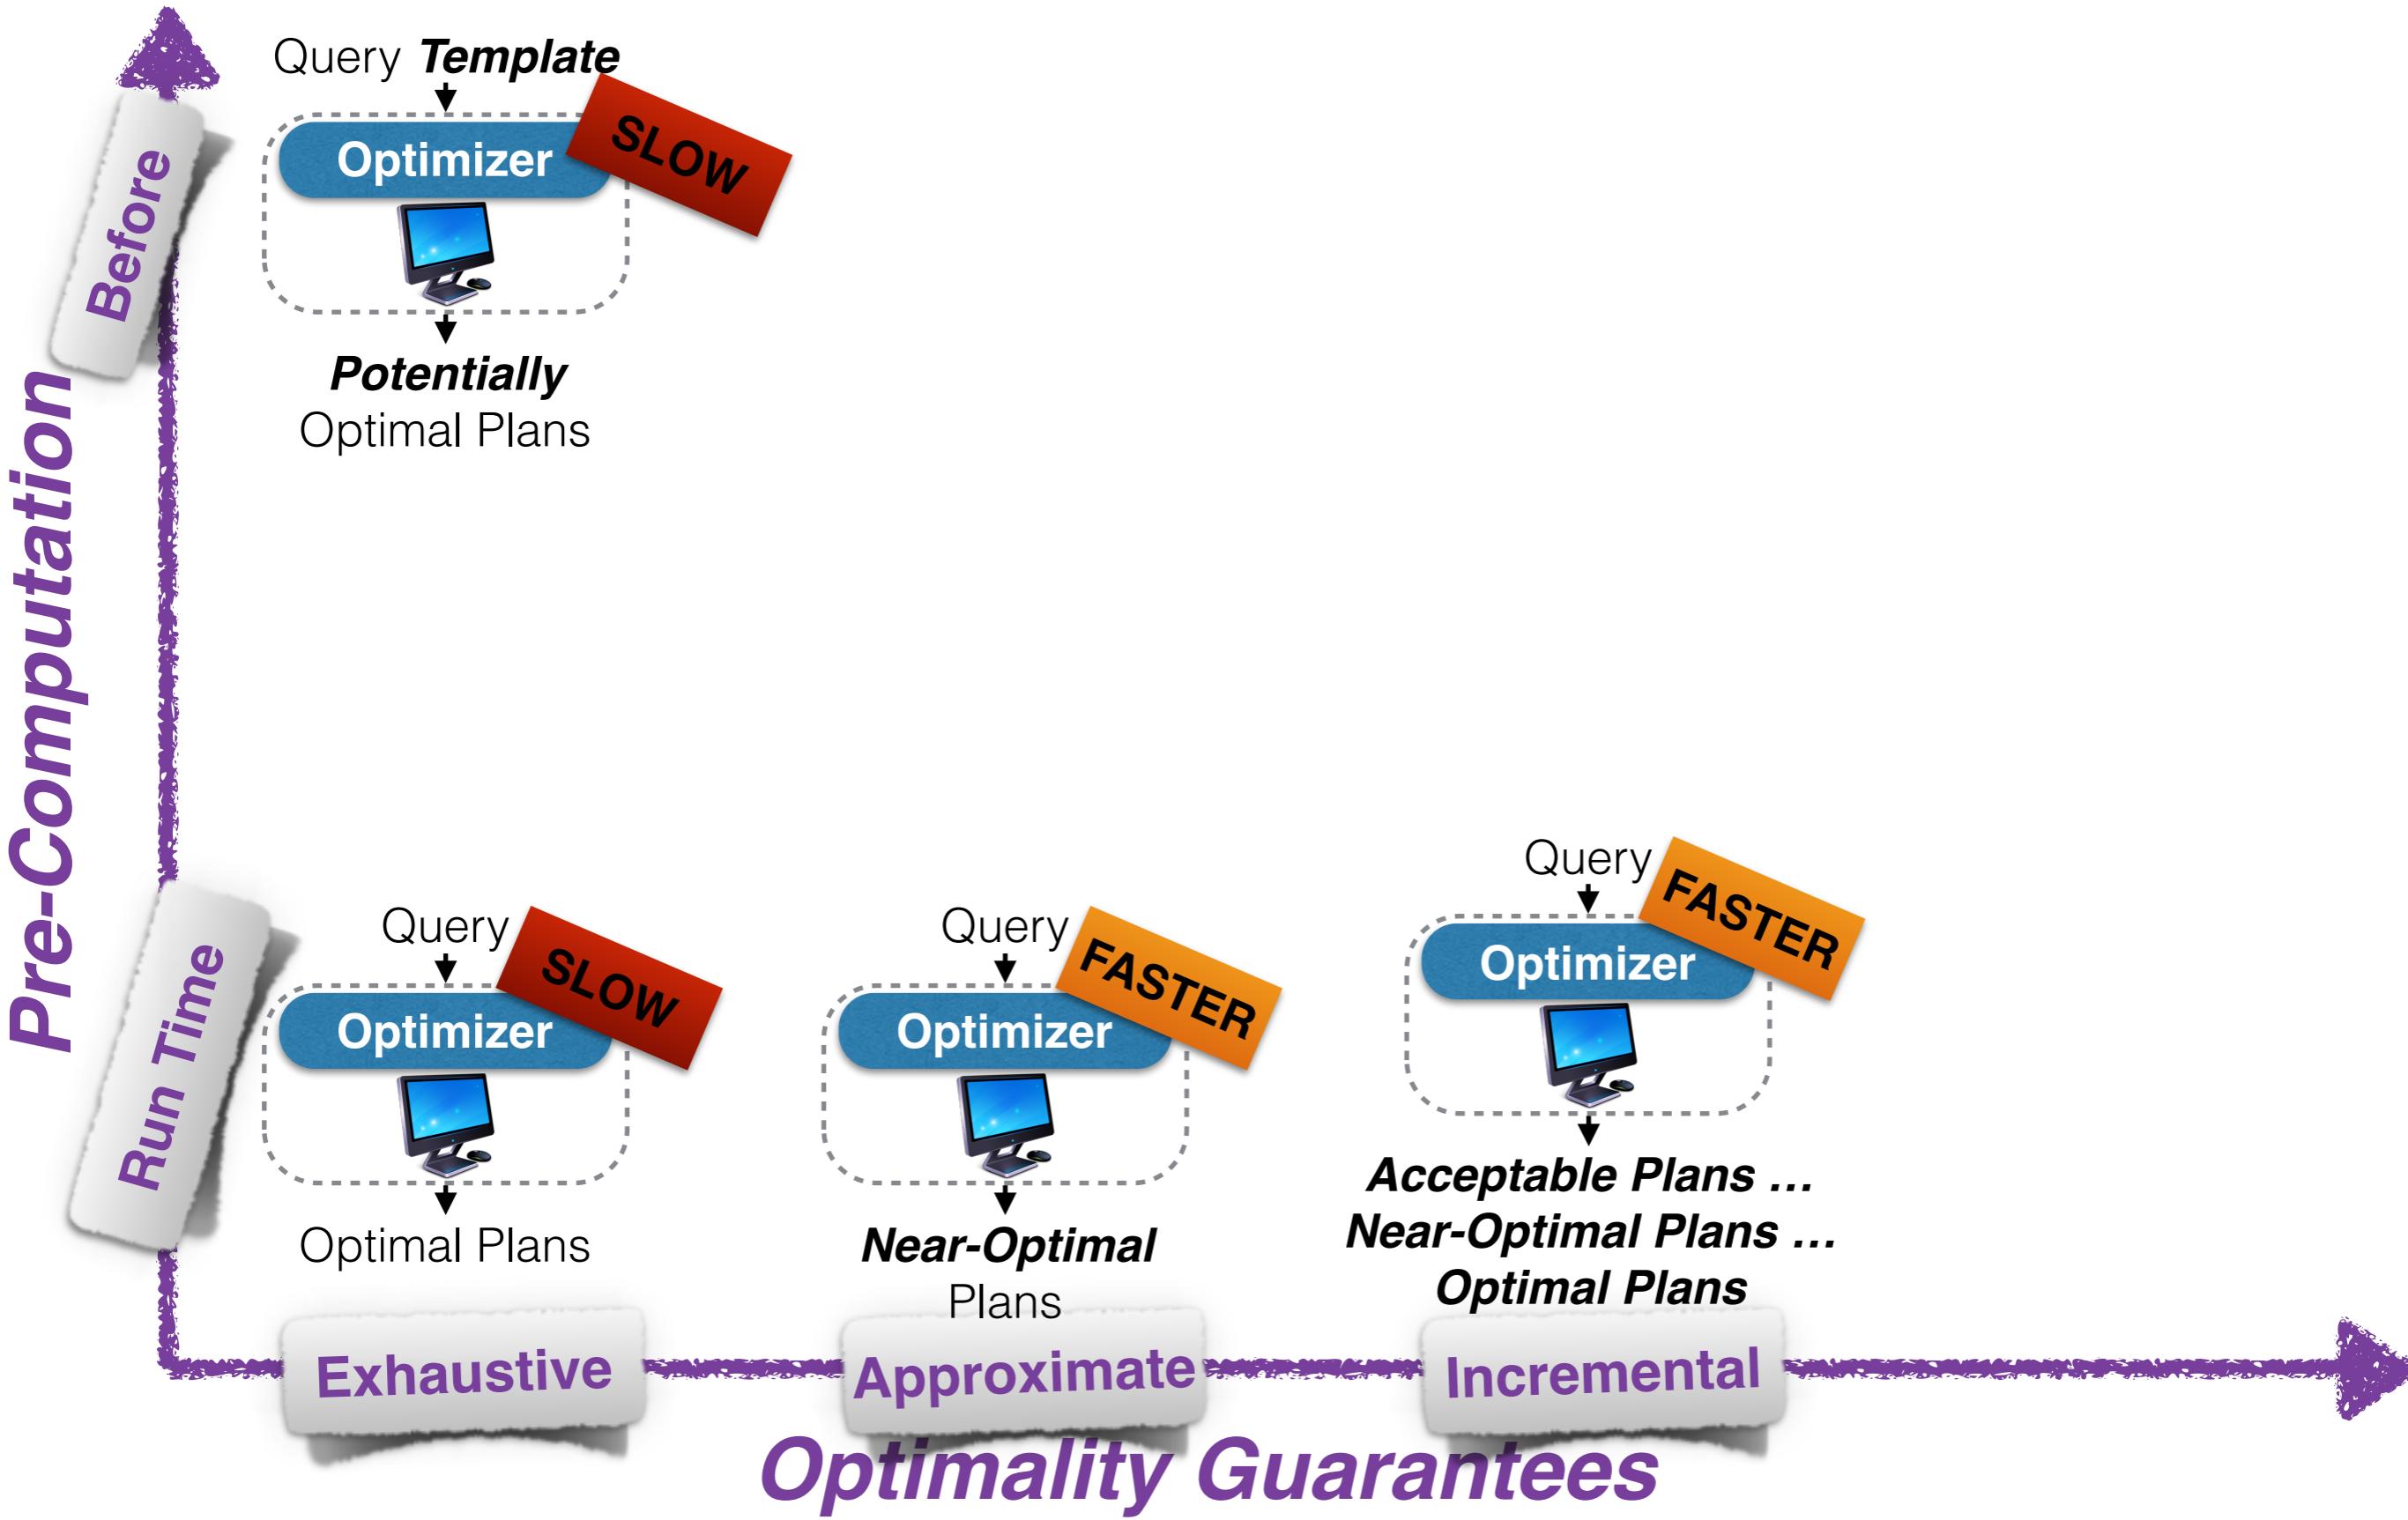
Dissertation Overview



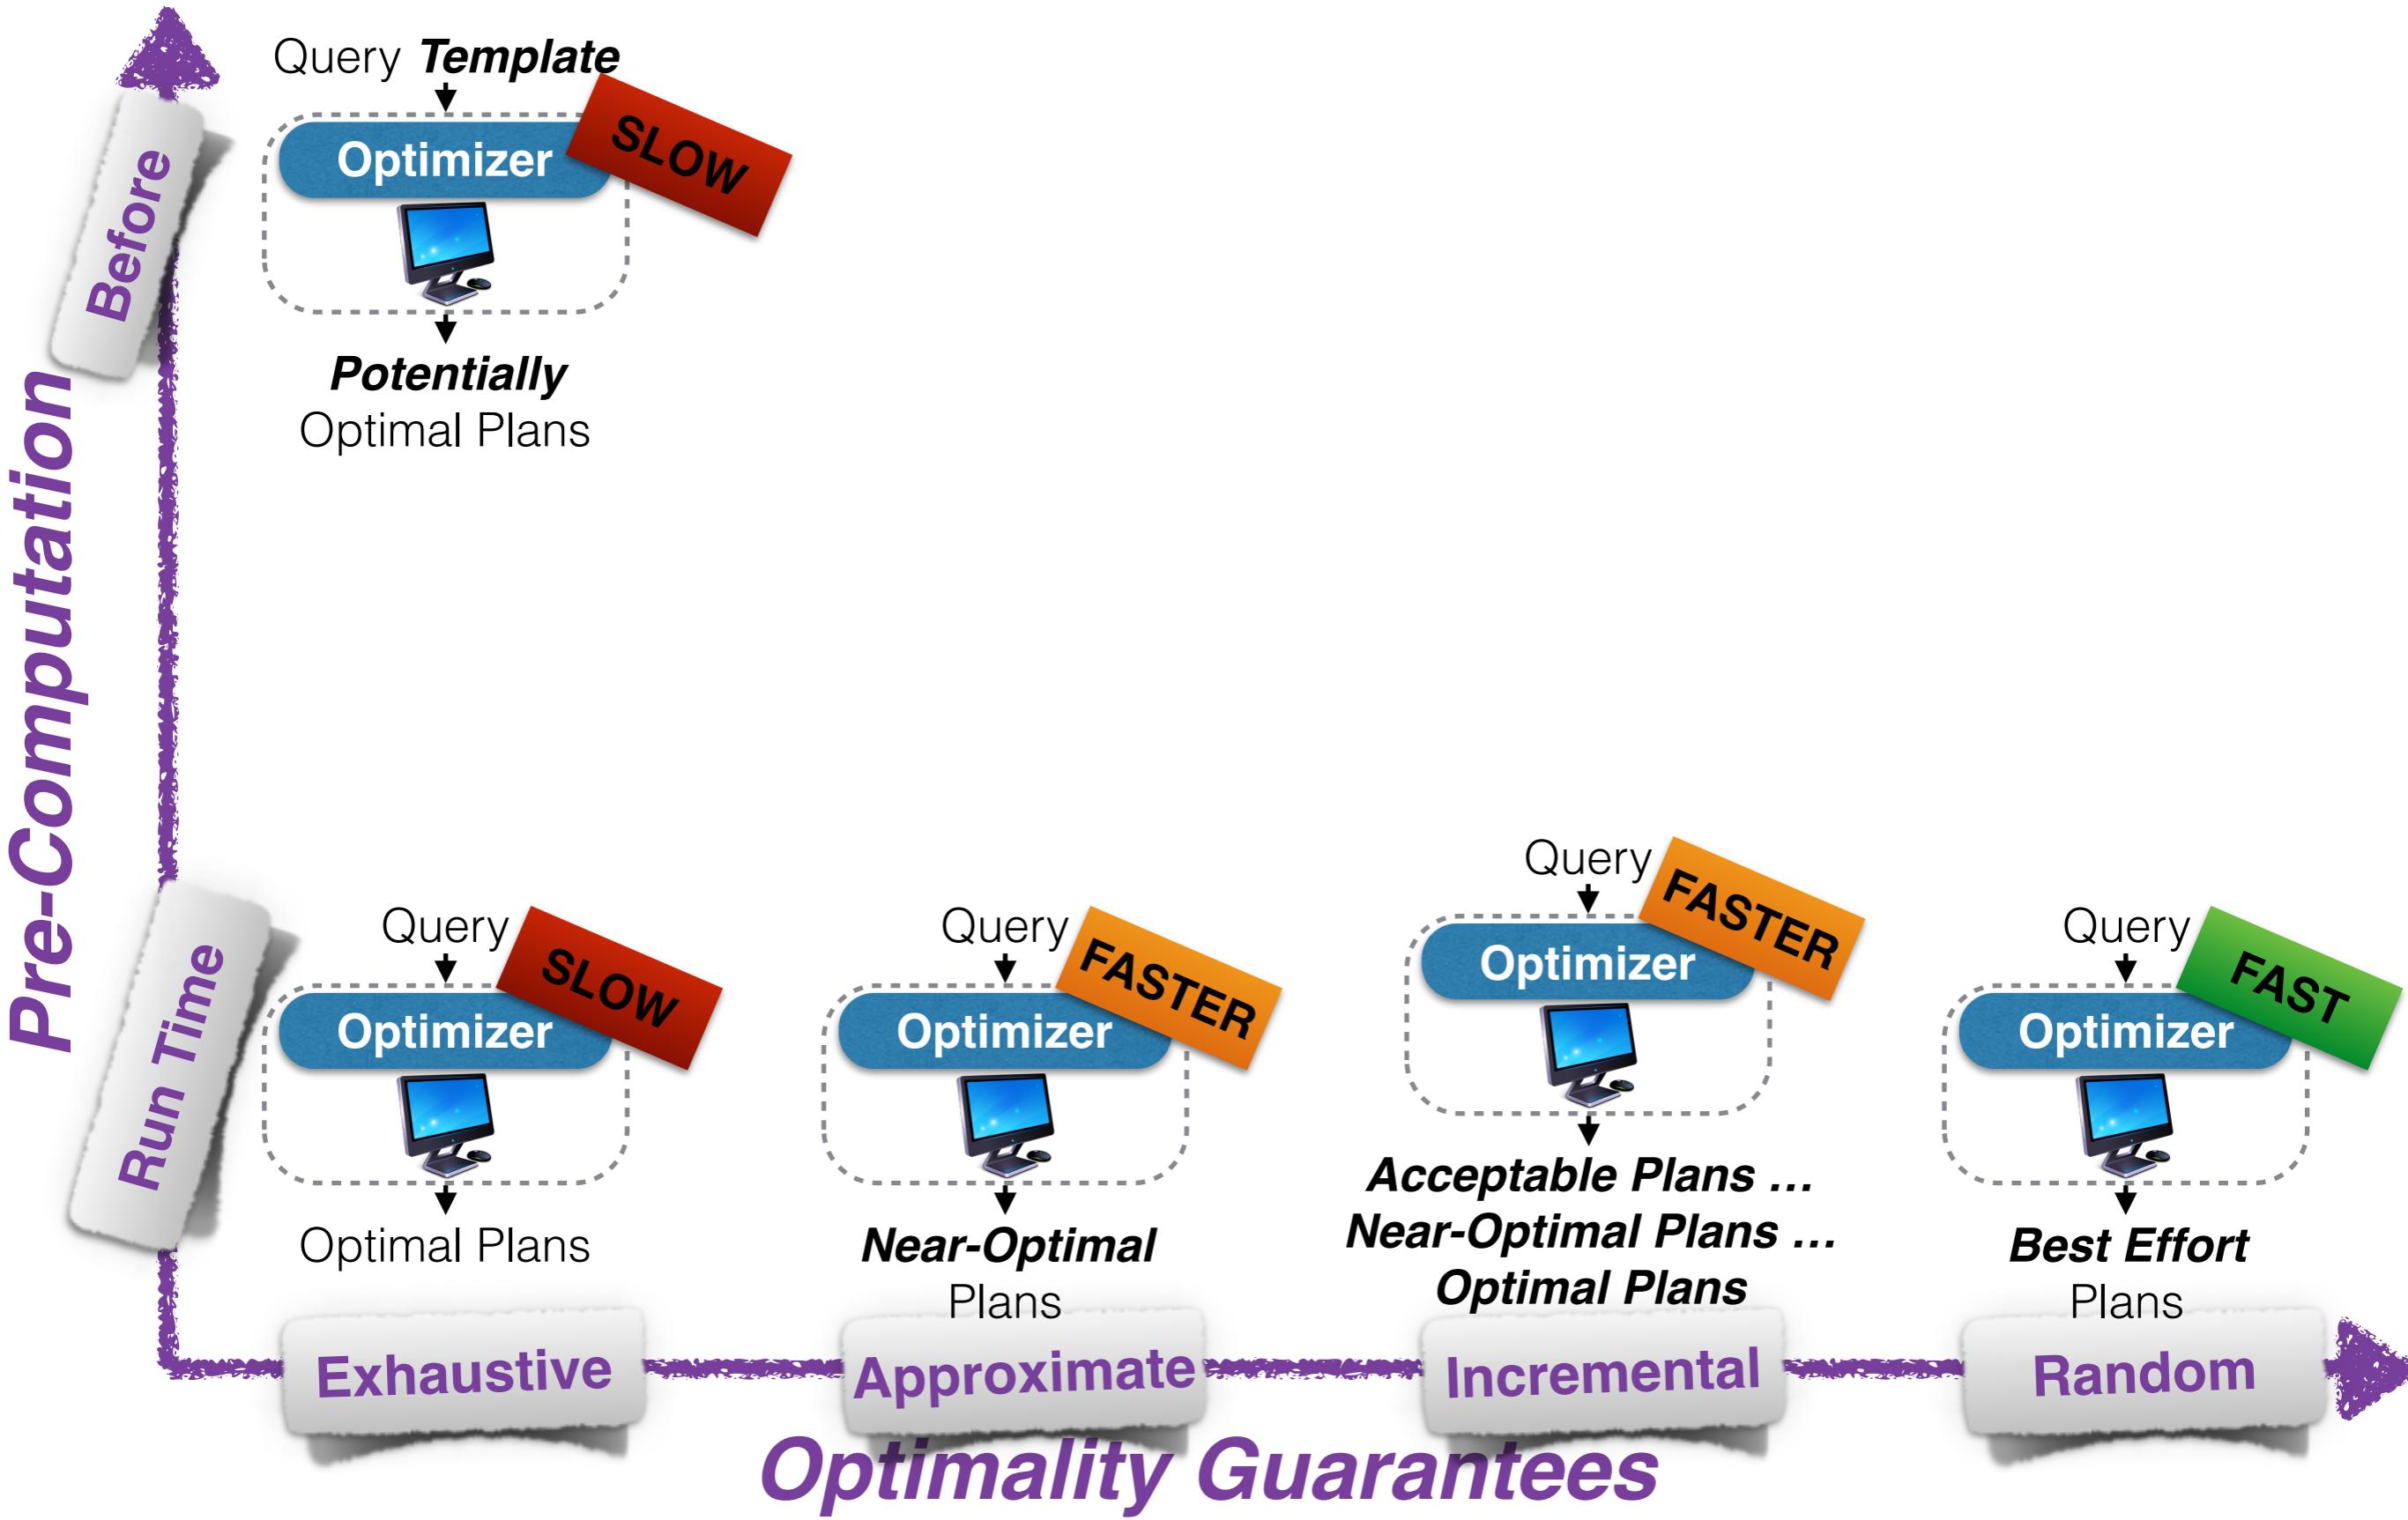
Dissertation Overview



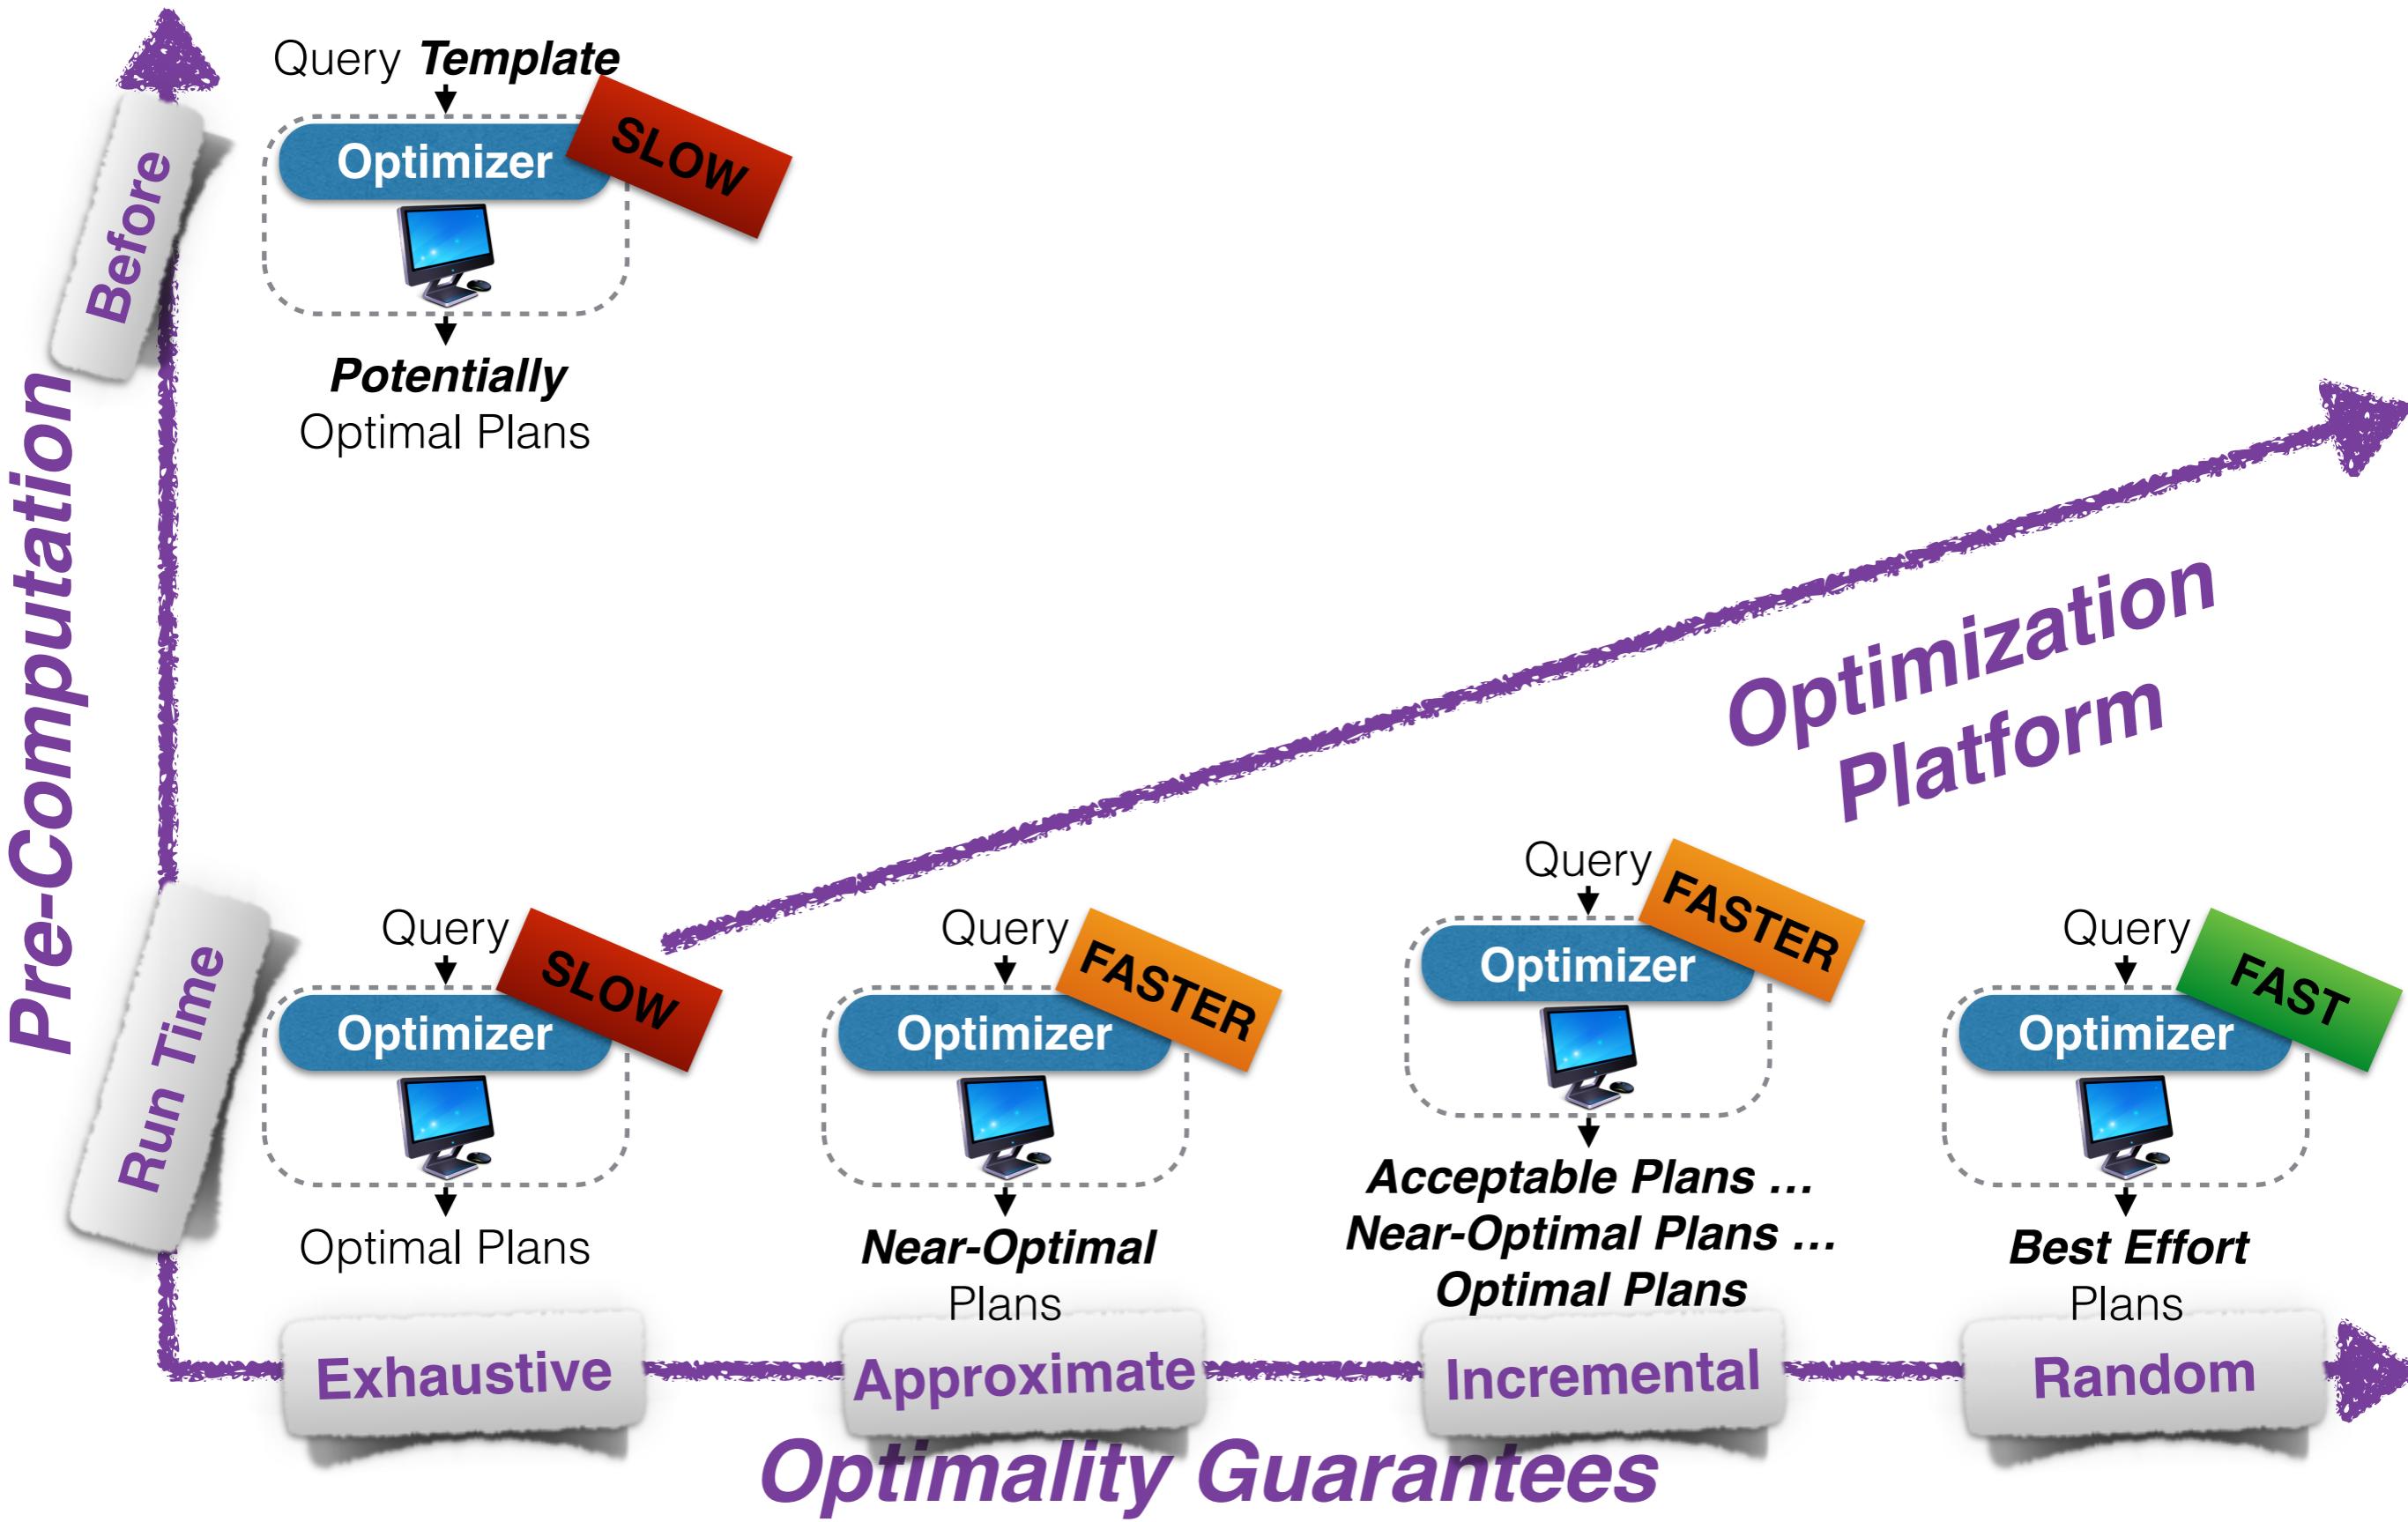
Dissertation Overview



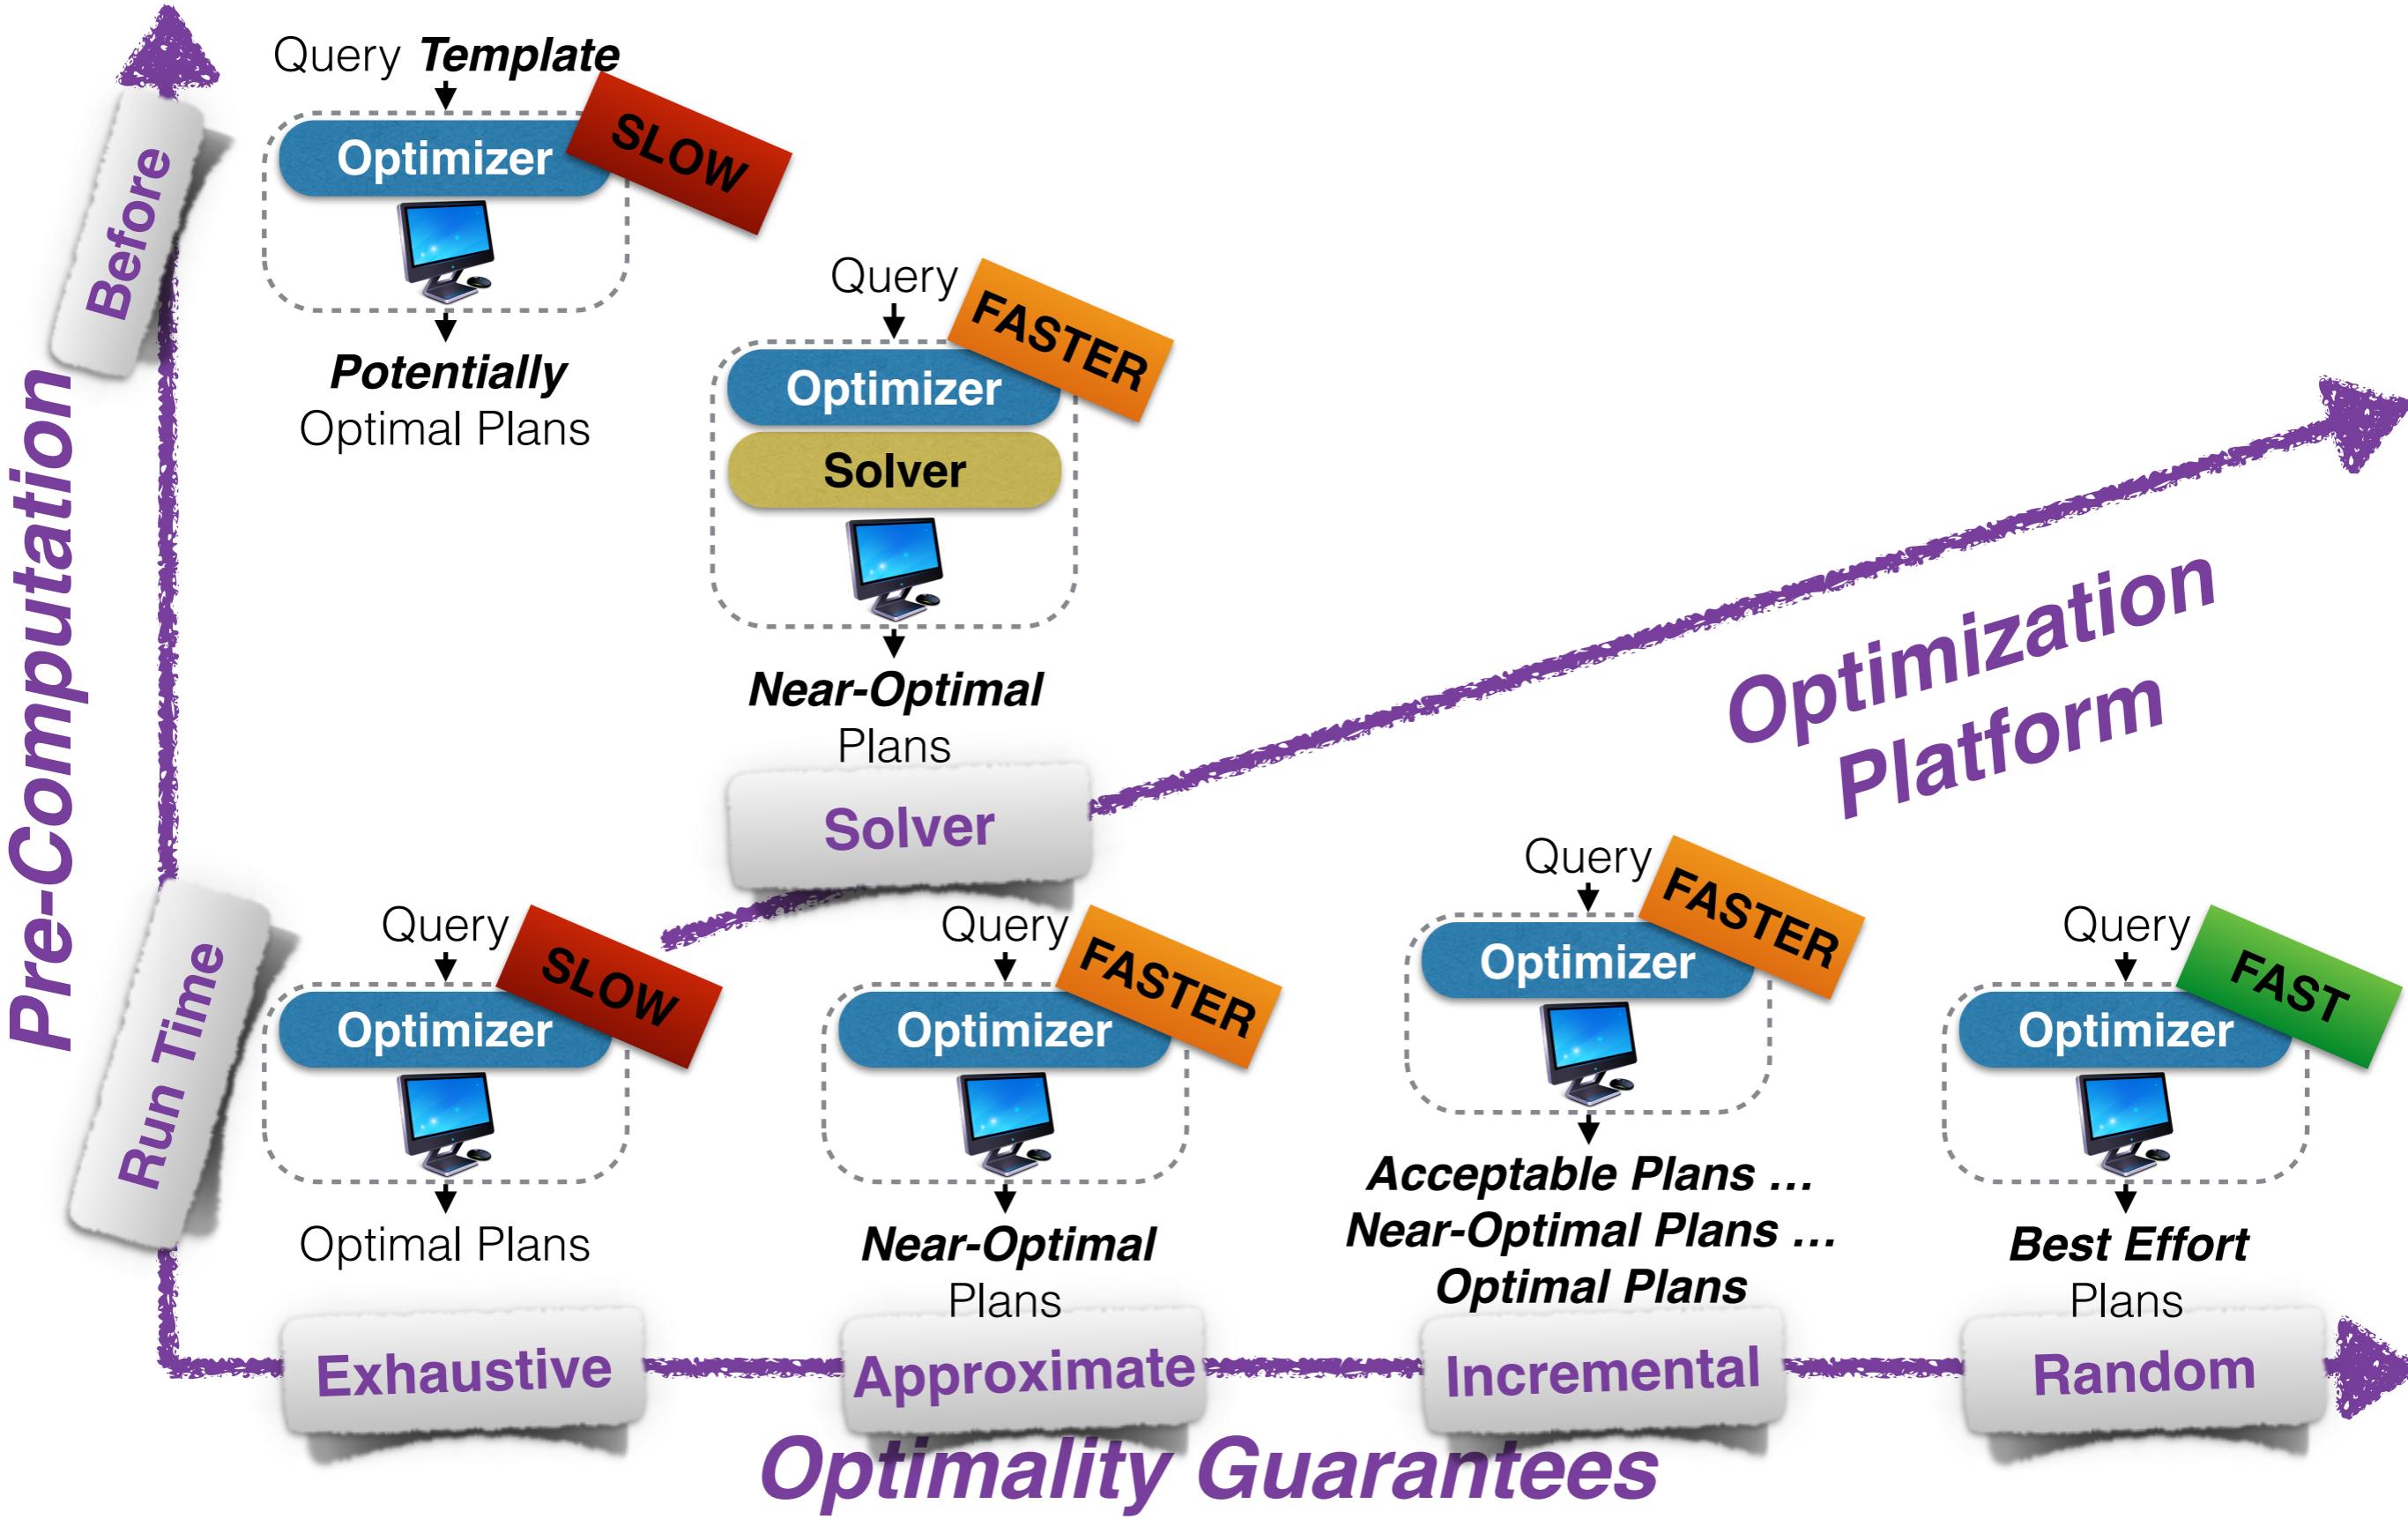
Dissertation Overview



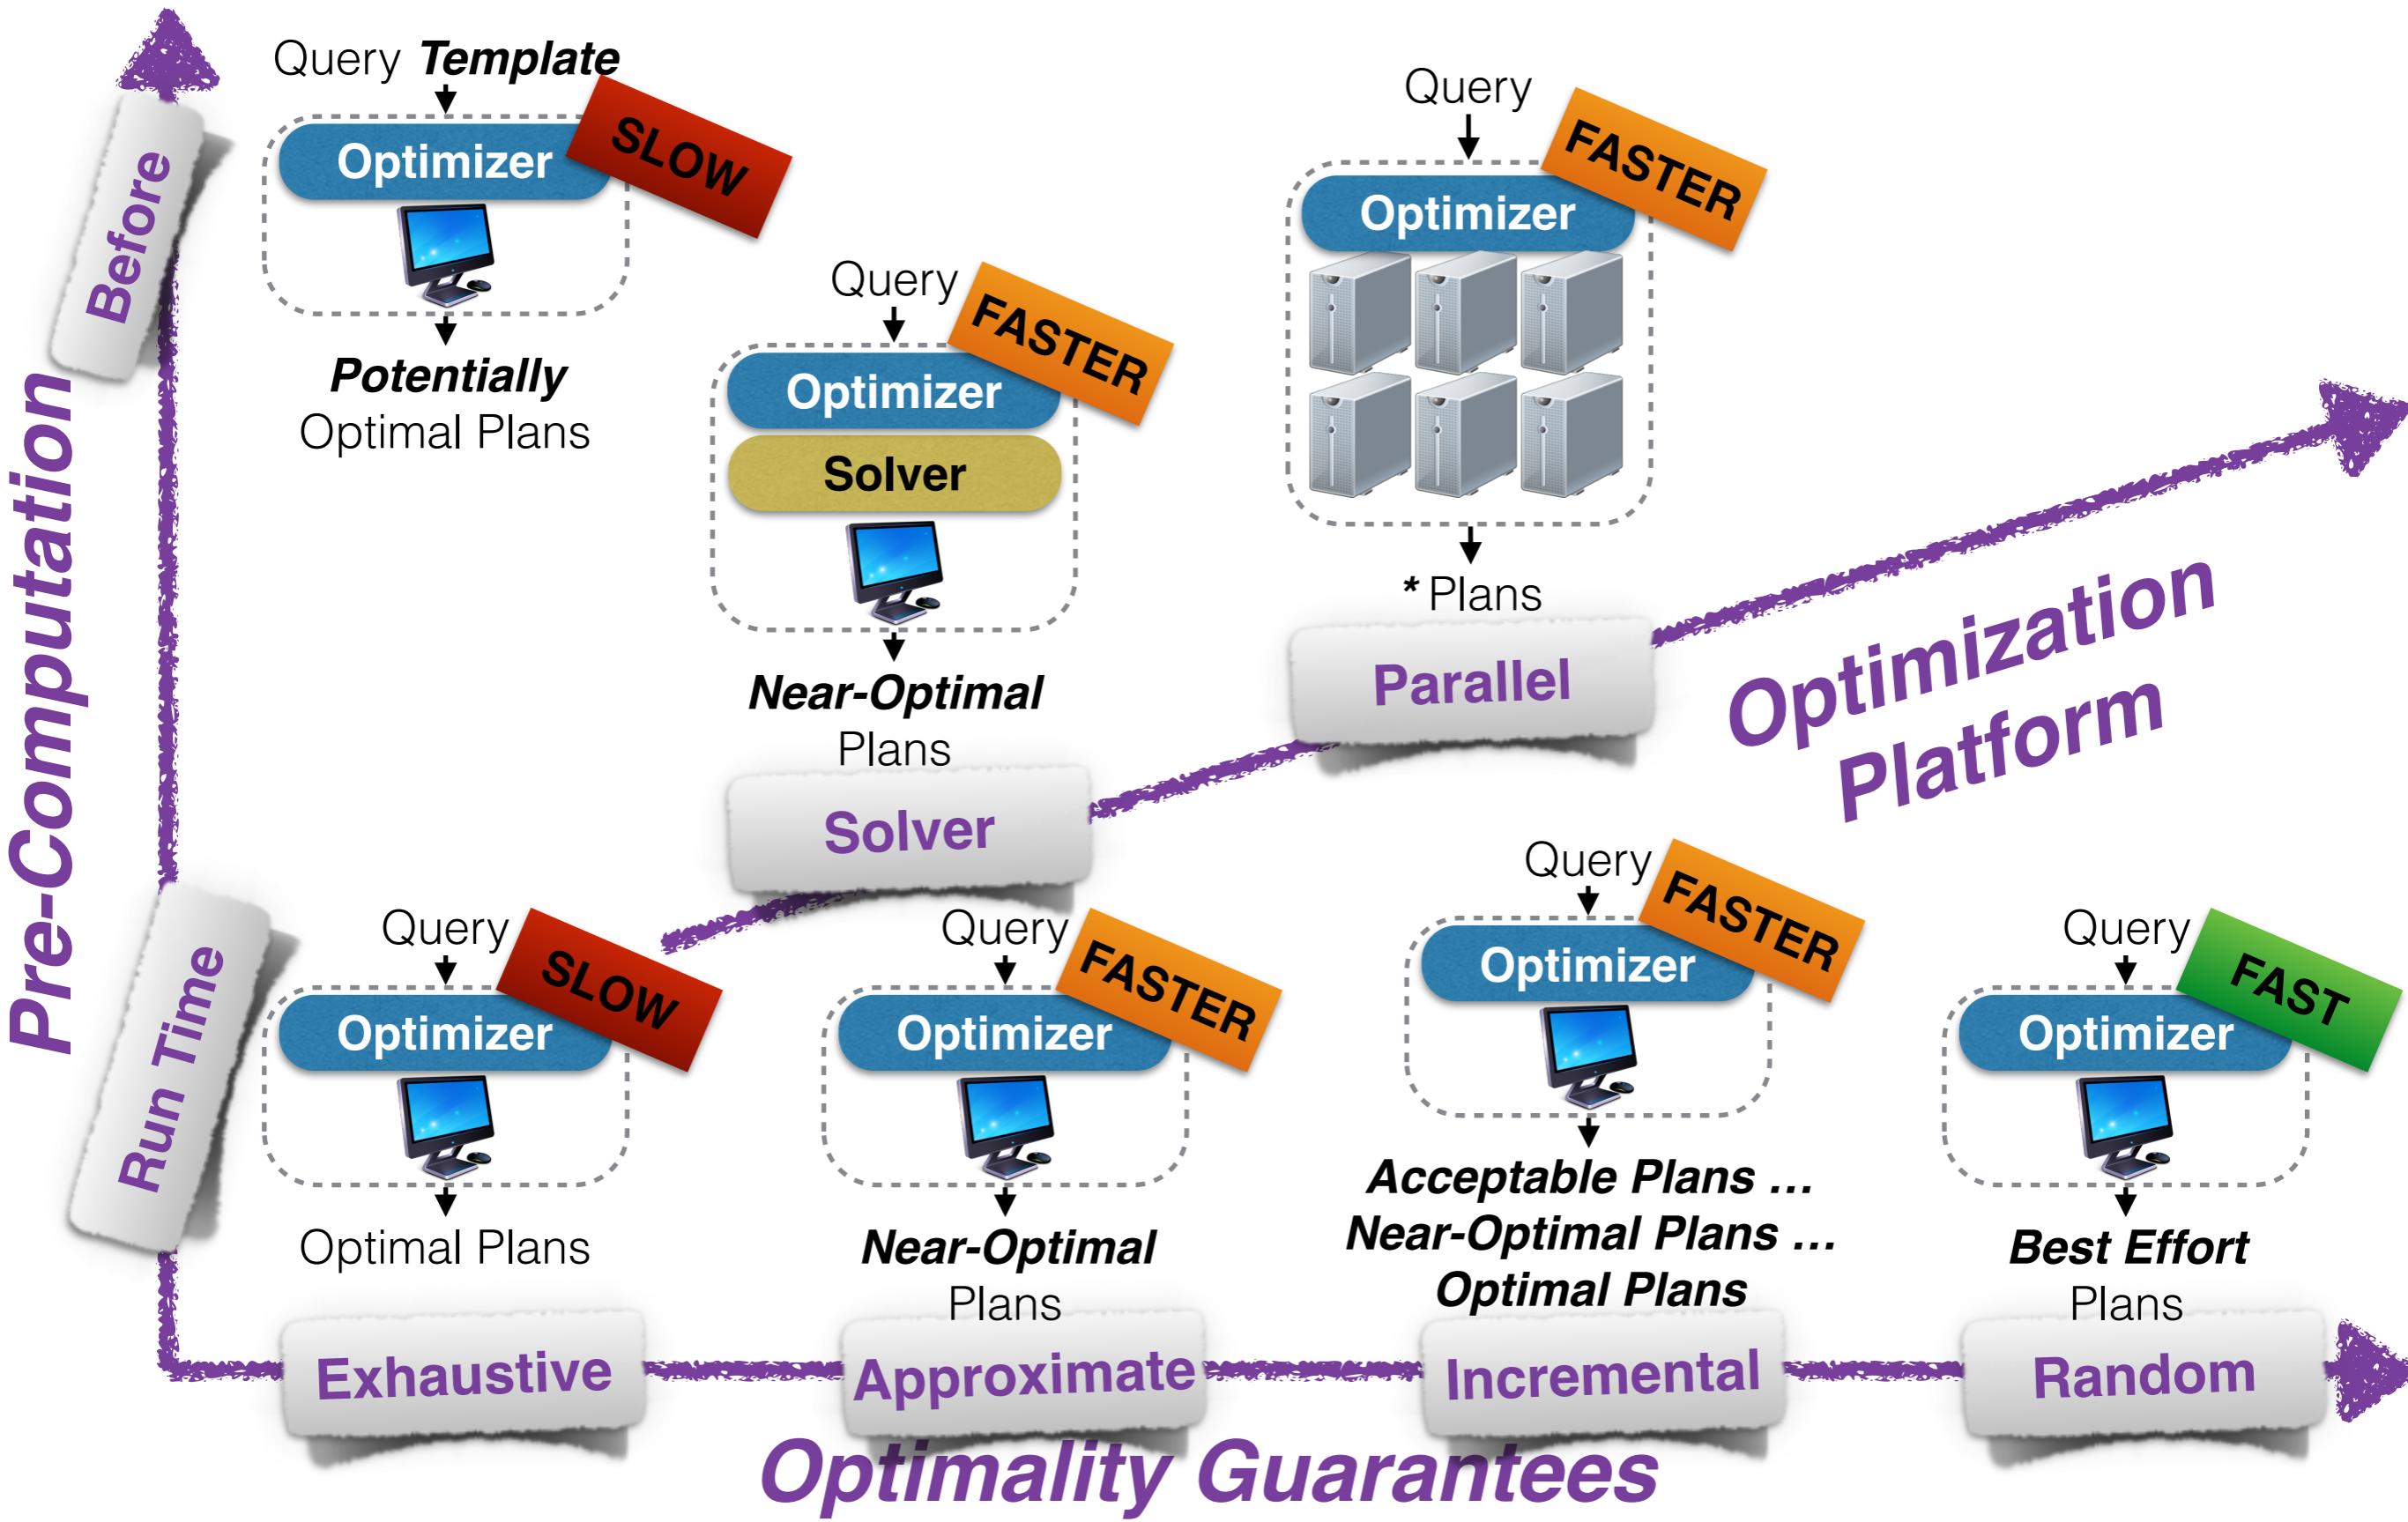
Dissertation Overview



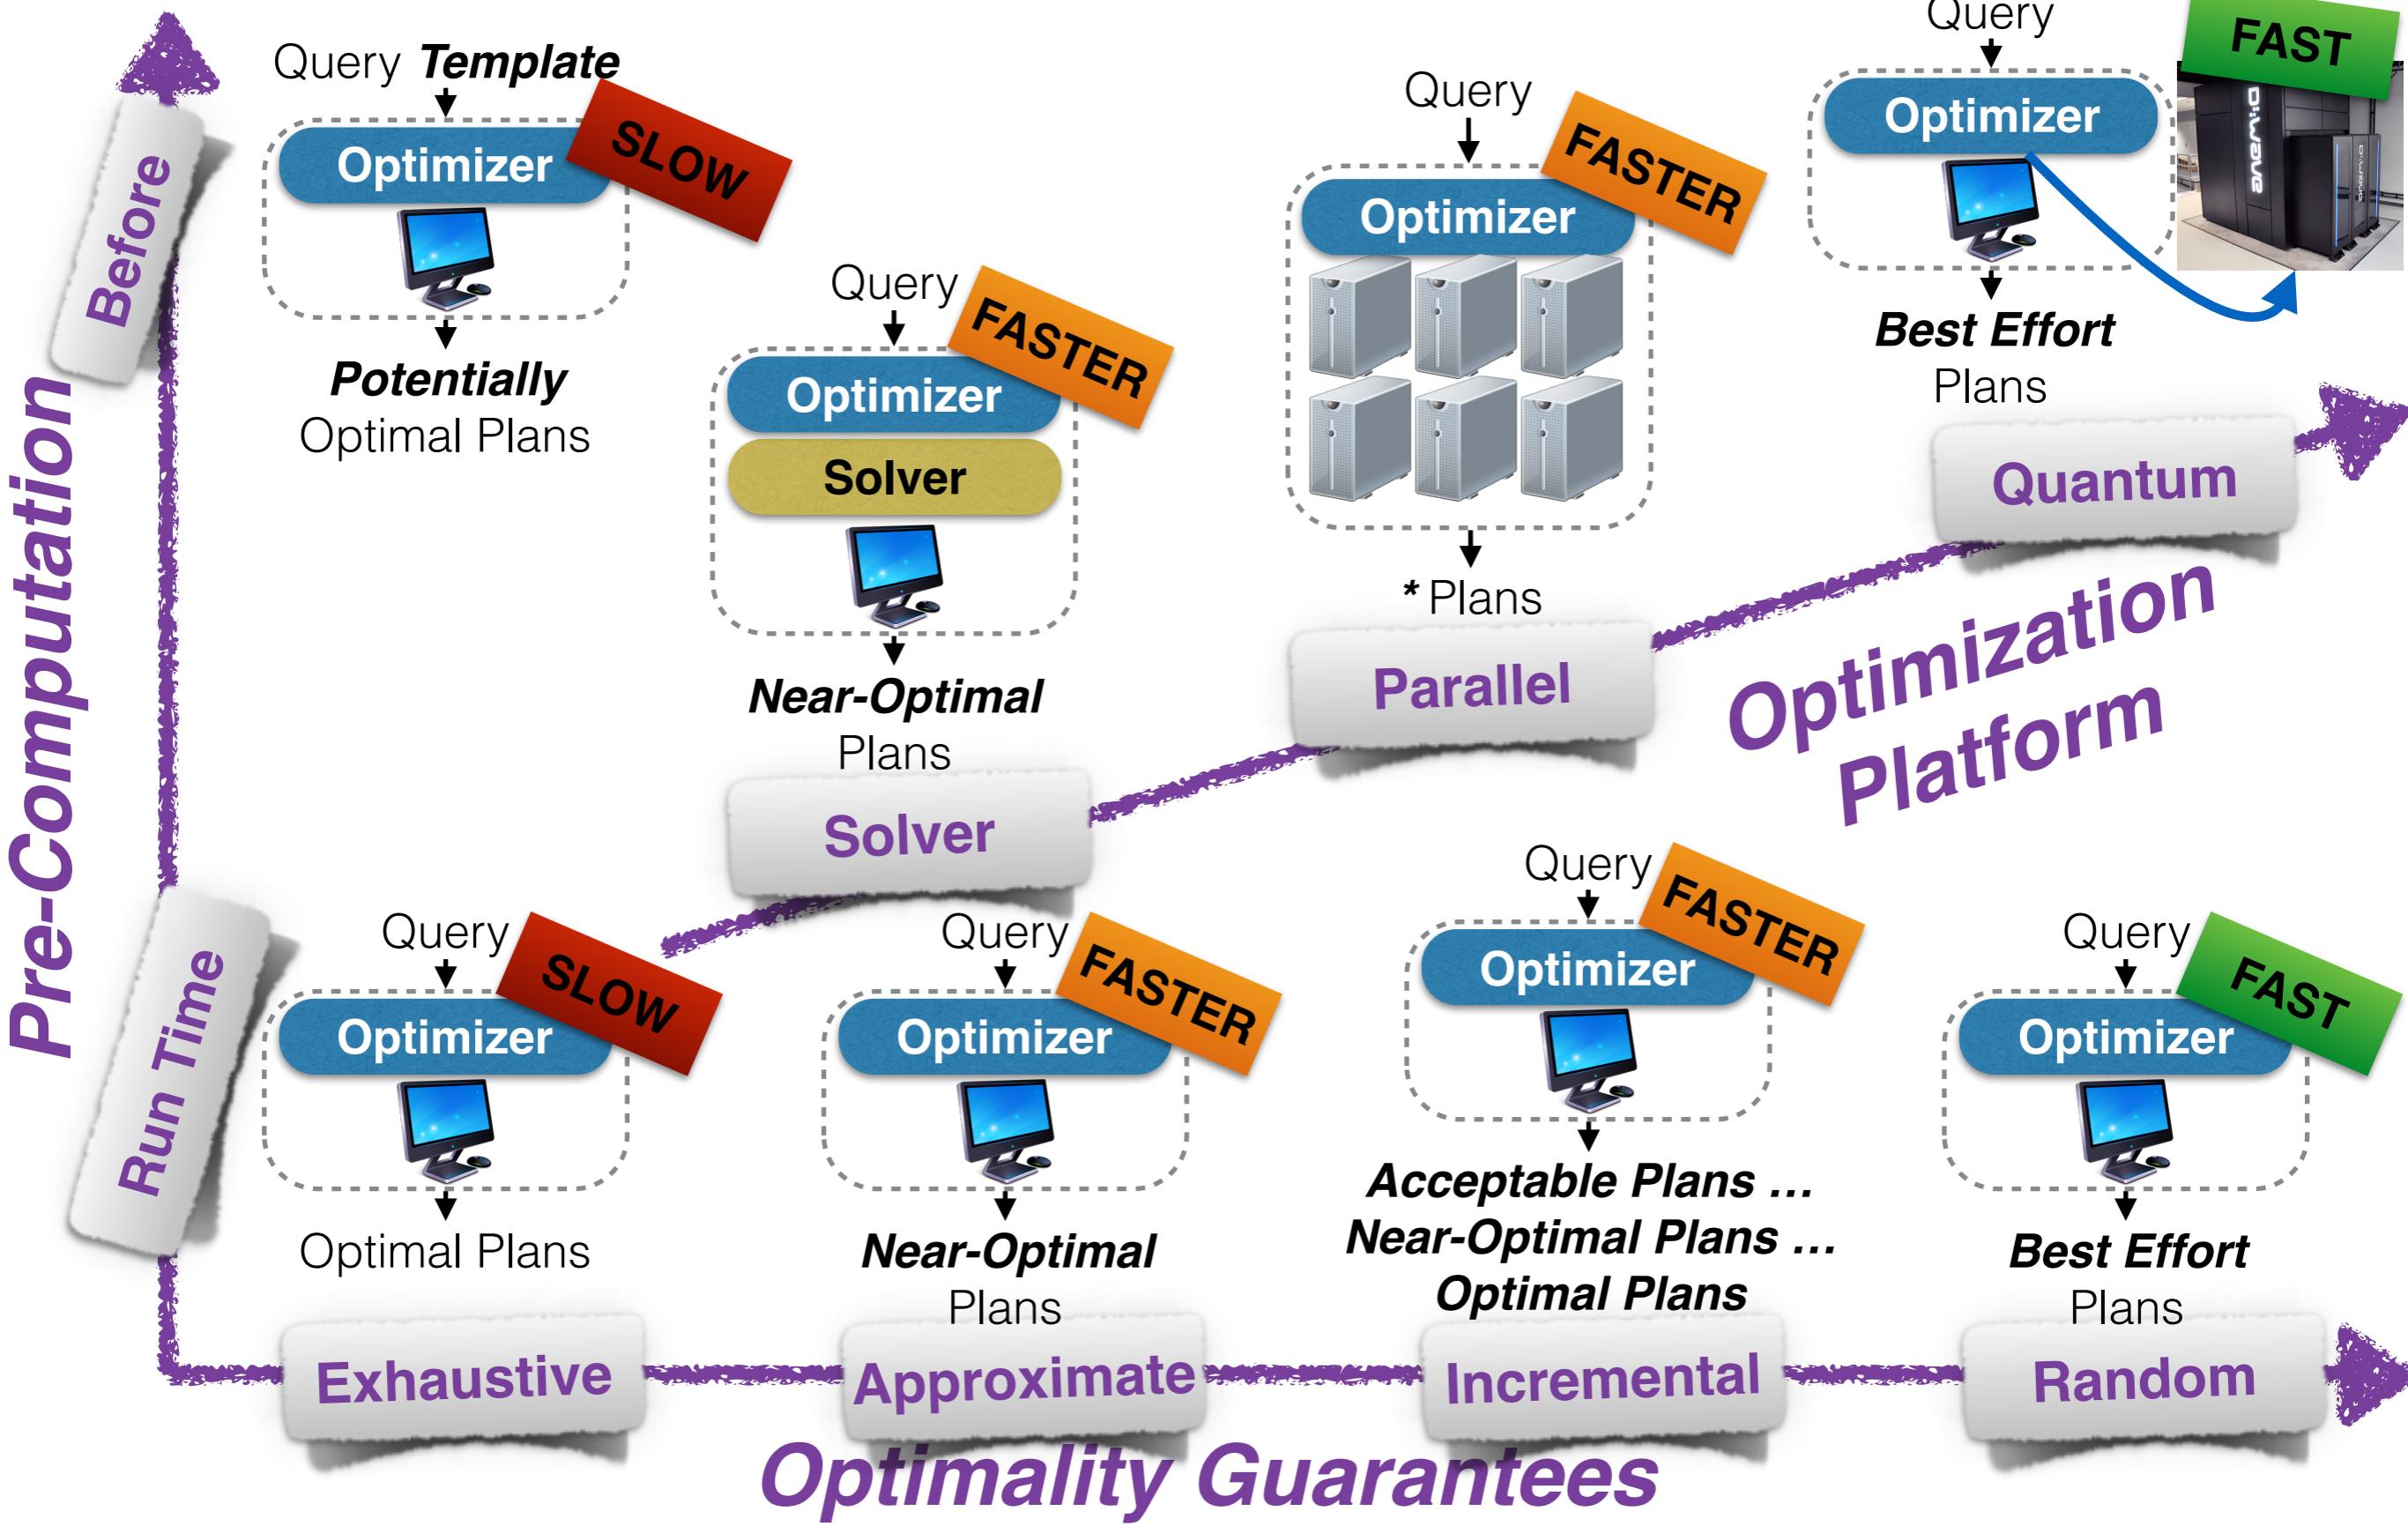
Dissertation Overview



Dissertation Overview

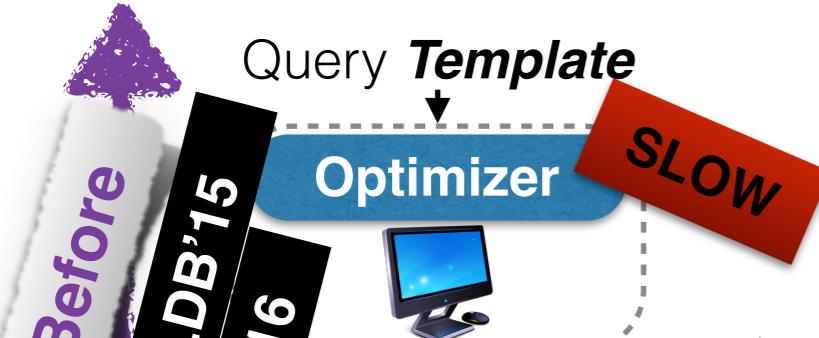


Dissertation Overview



Dissertation Overview

Pre-Computation



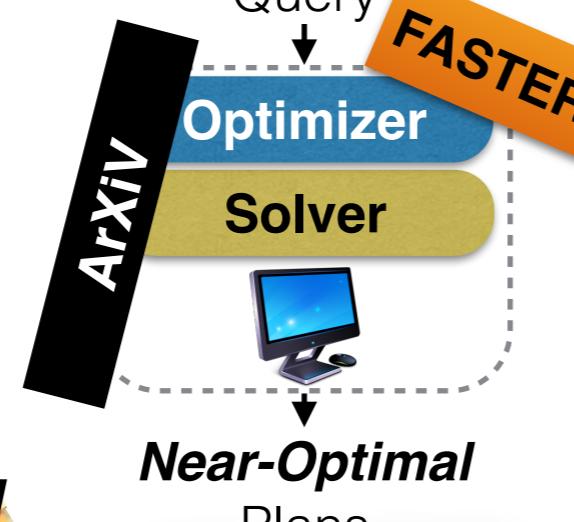
Potentially Optimal Plans



SIGMOD Research Highlight

CACM Research Highlight

Run Time

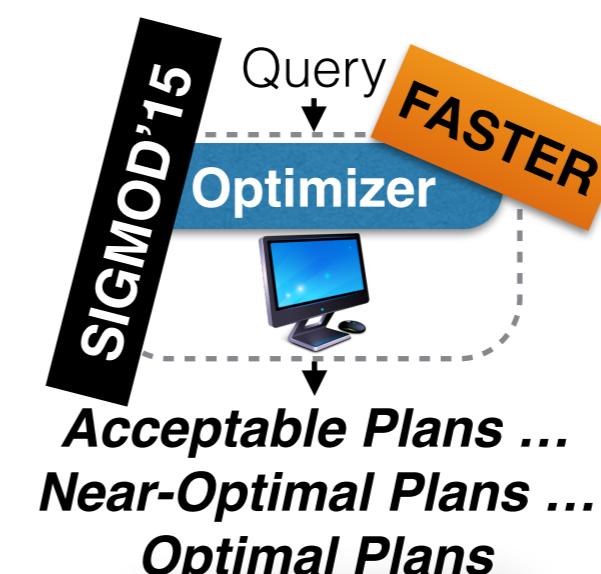


Solver



Exhaustive

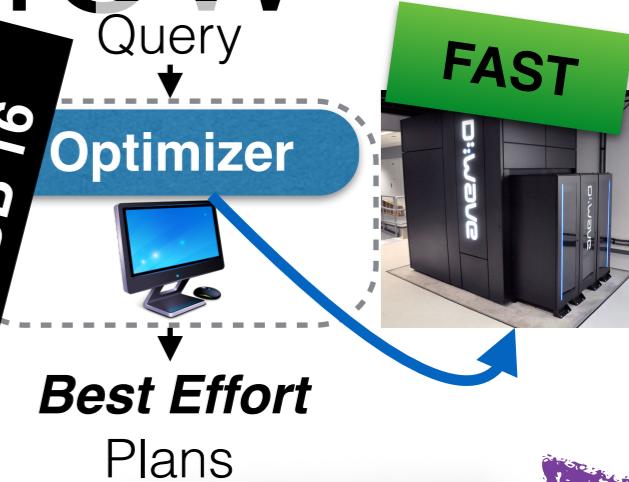
Approximate



Random

Optimality Guarantees

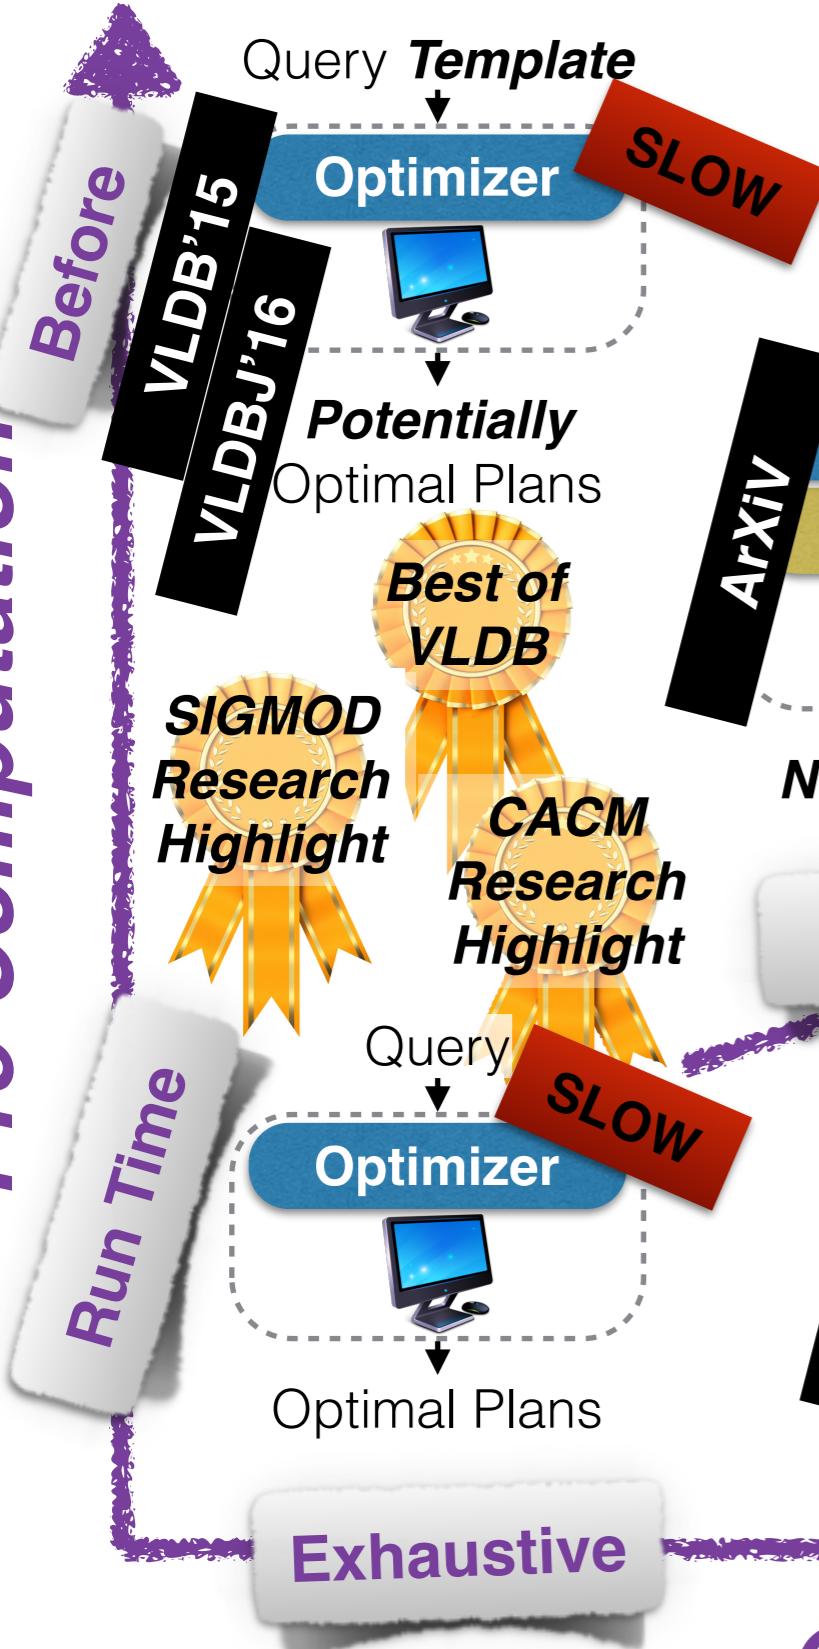
Optimization Platform



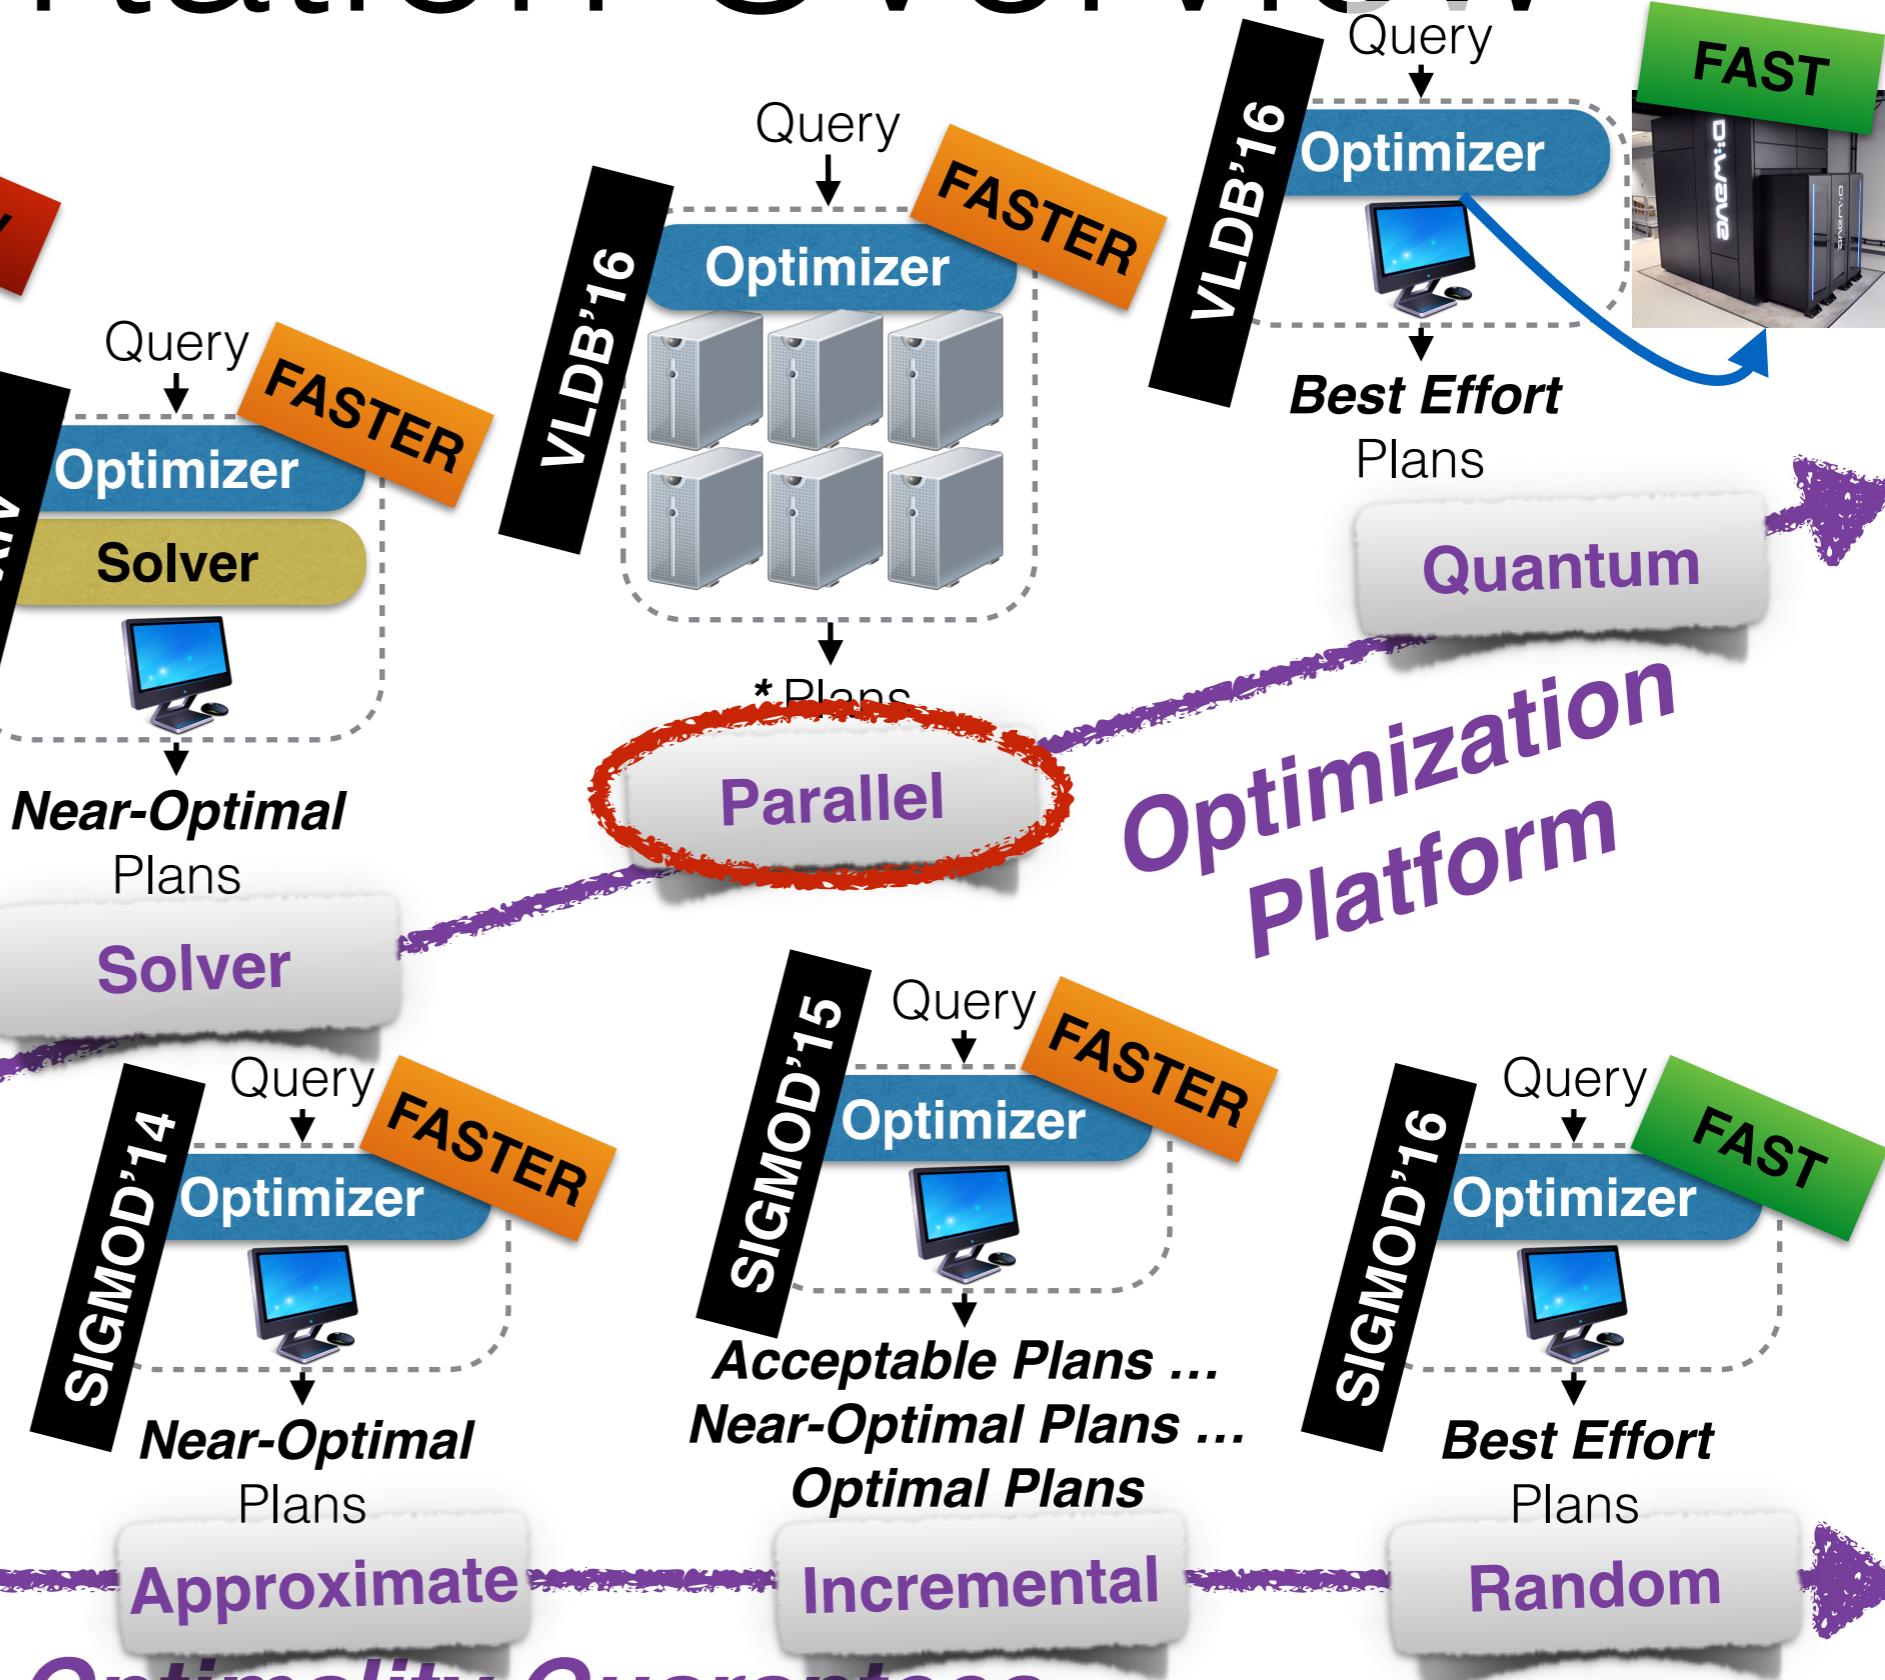
Plans

Dissertation Overview

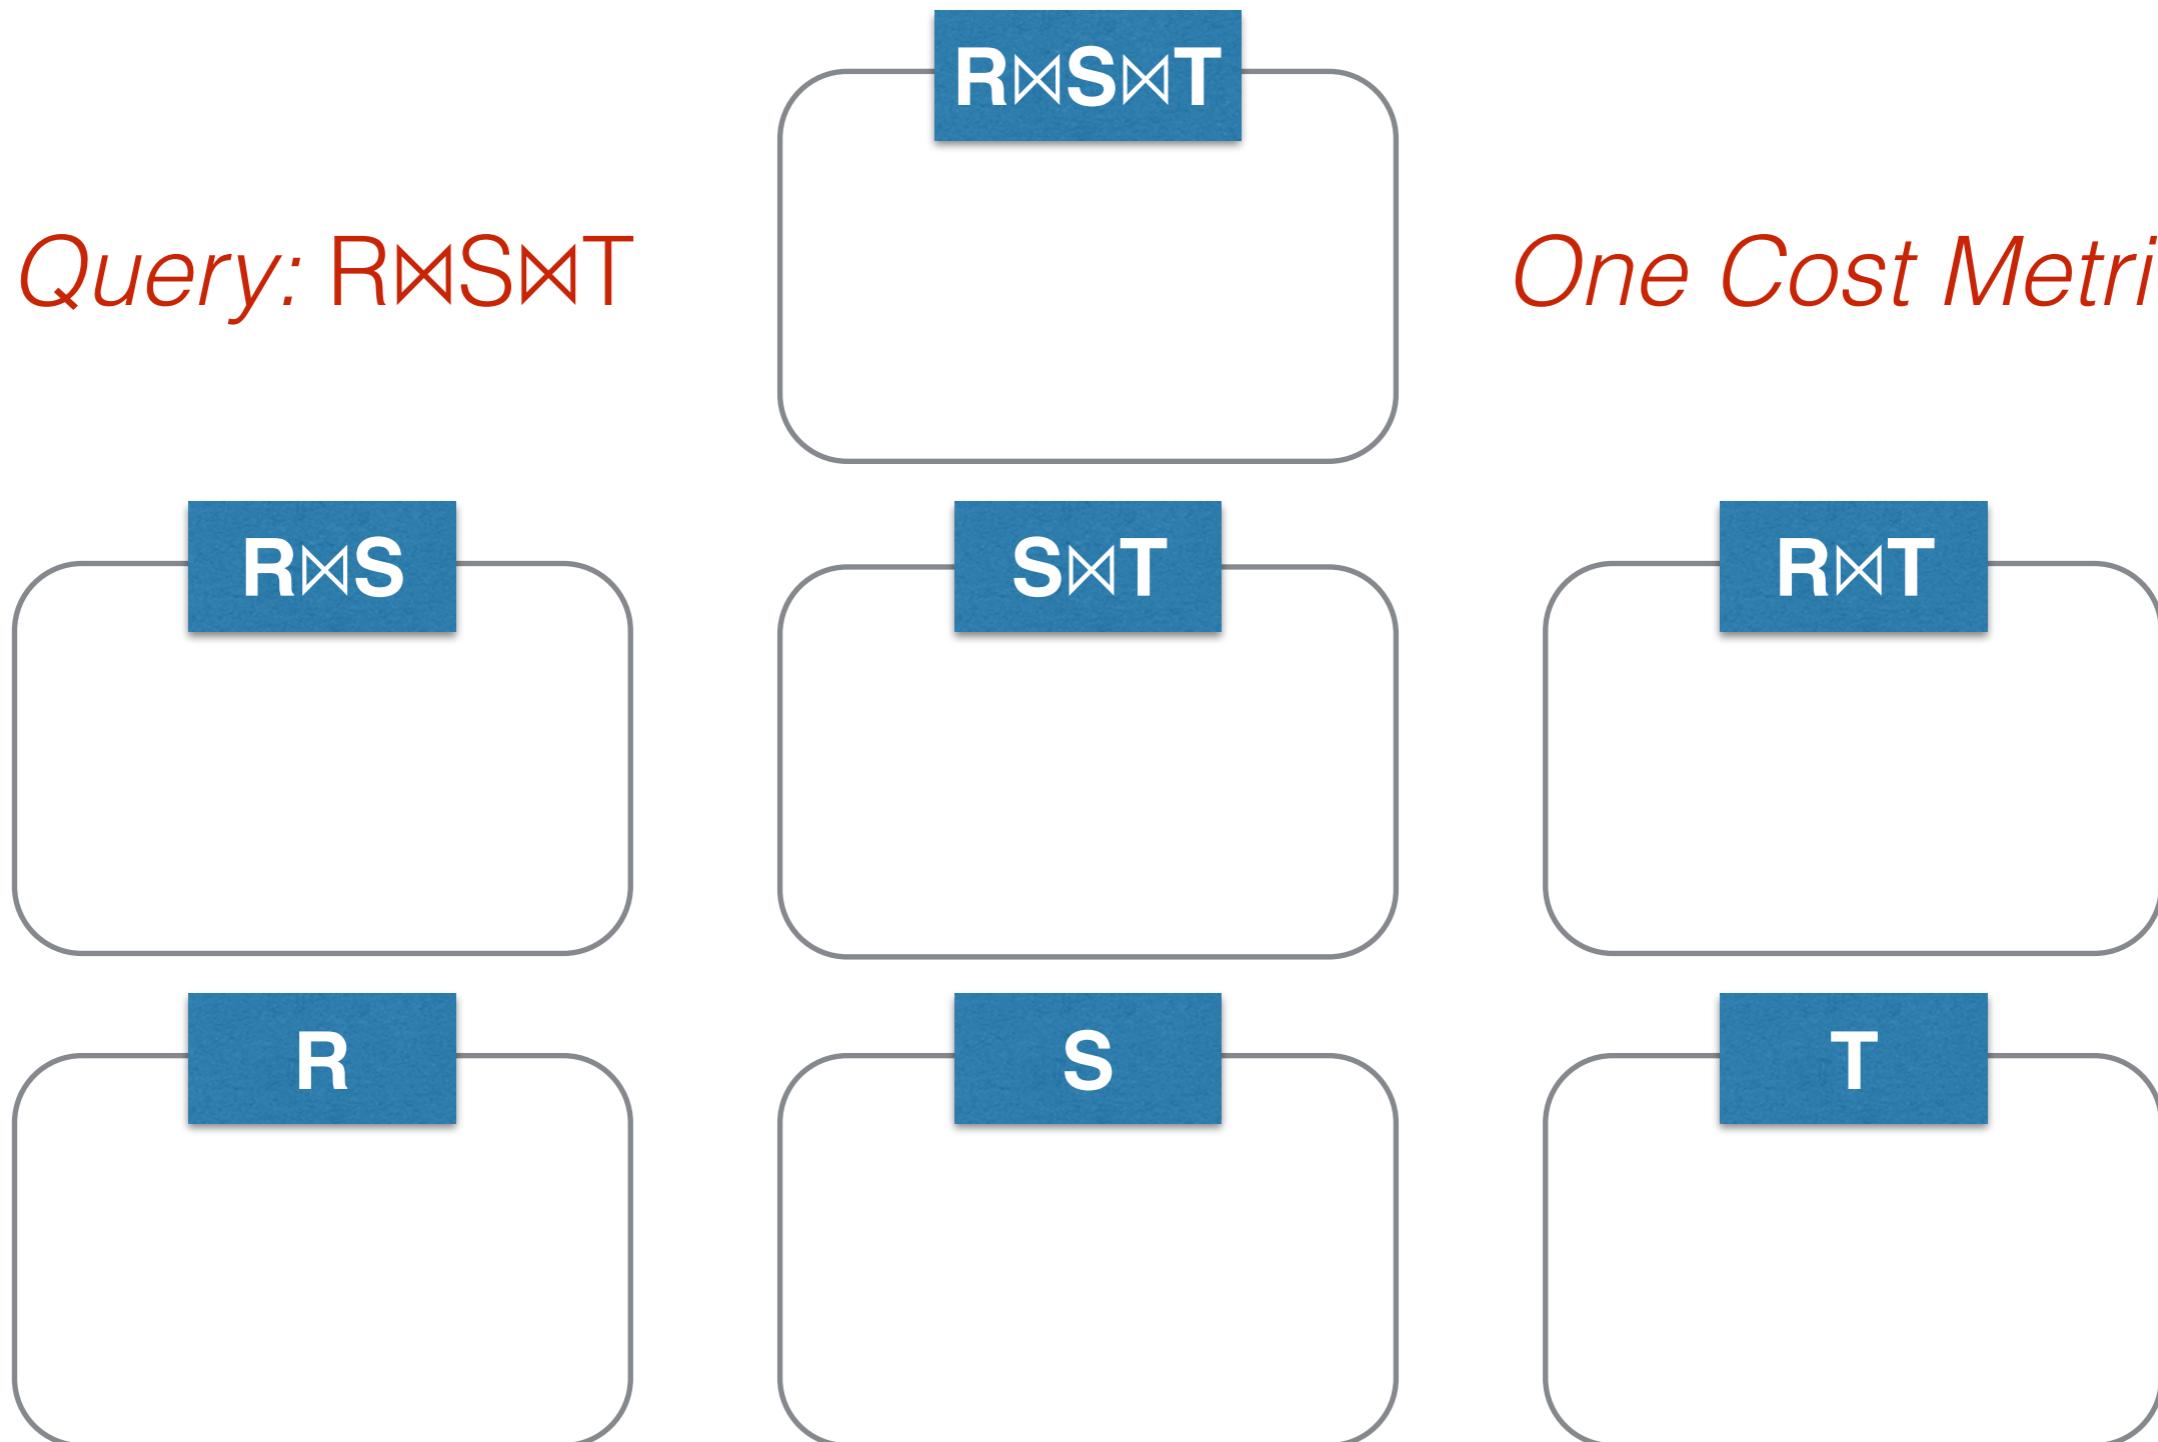
Pre-Computation



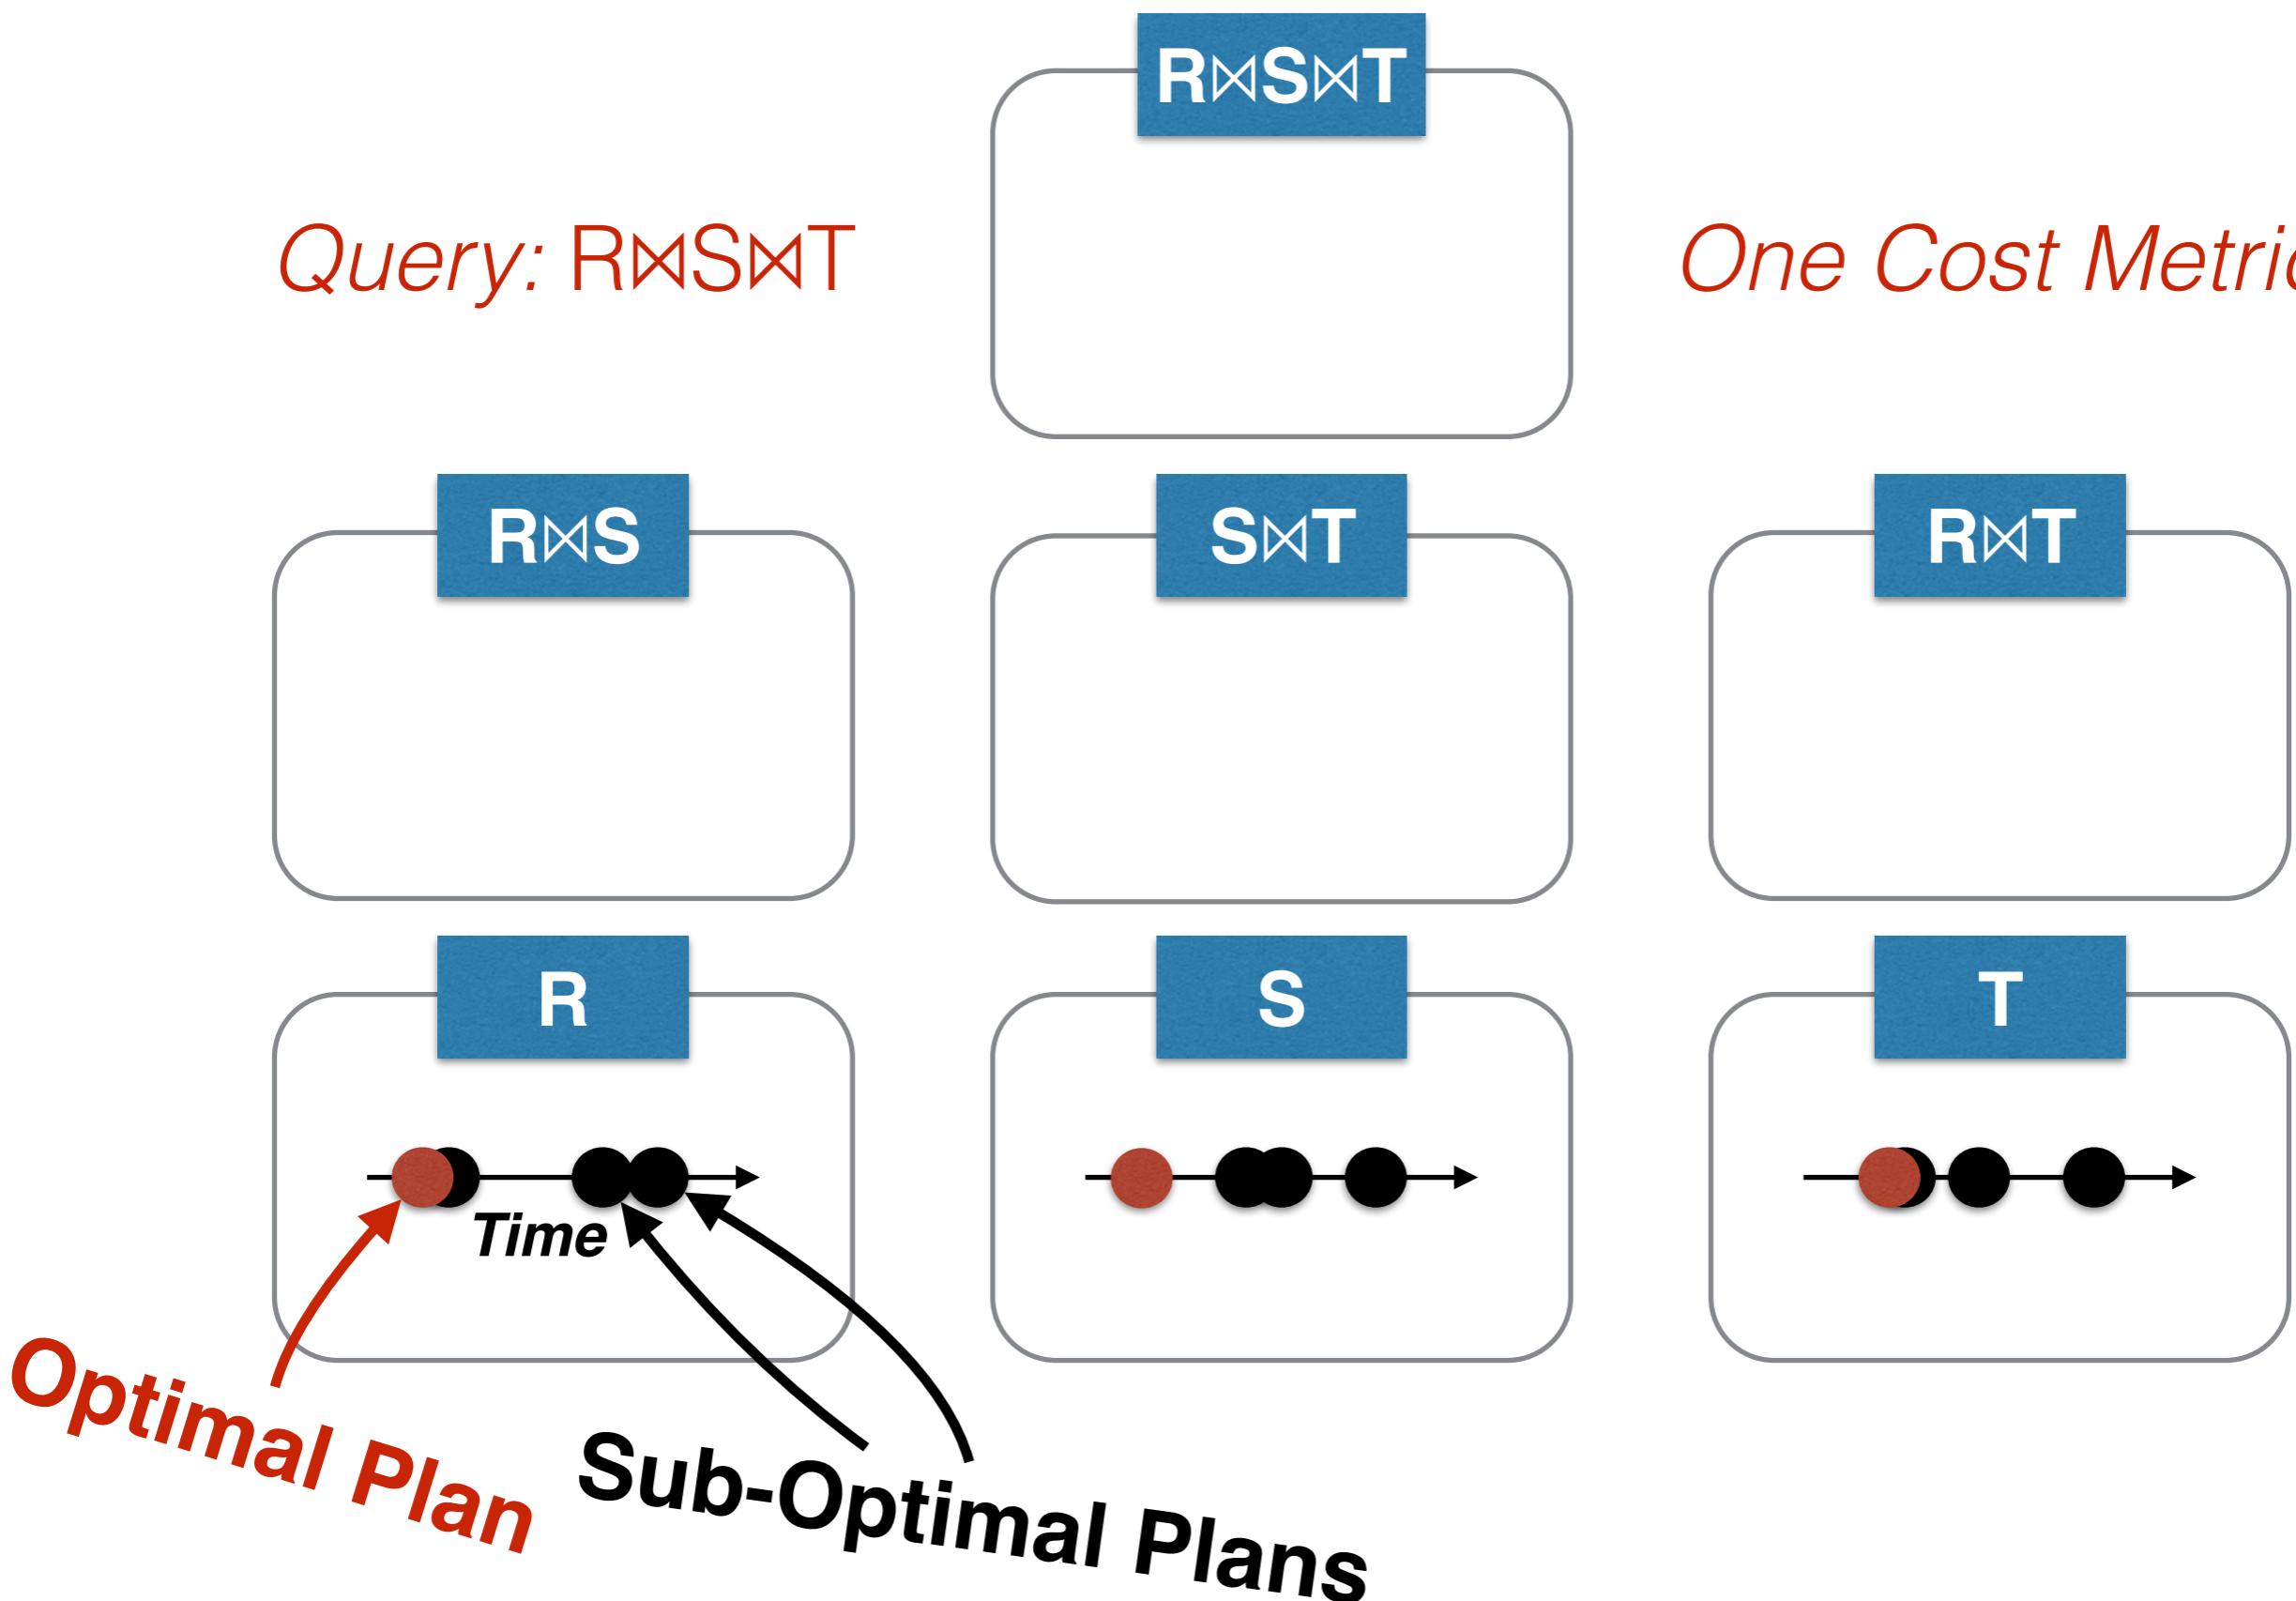
Optimality Guarantees



Classical Query Optimization



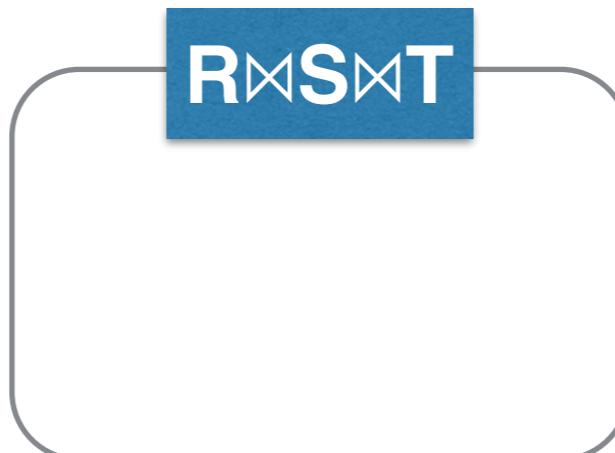
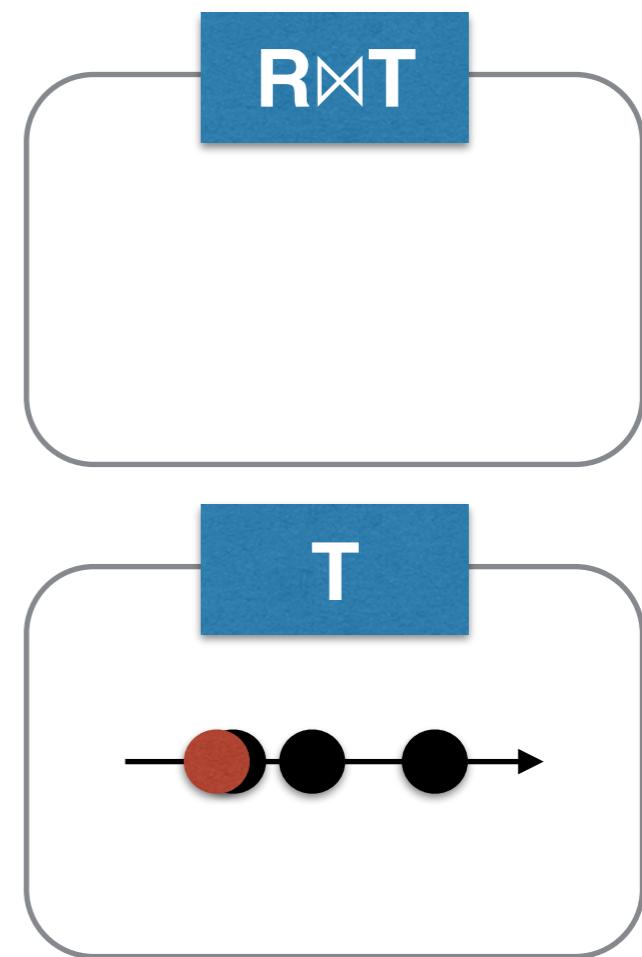
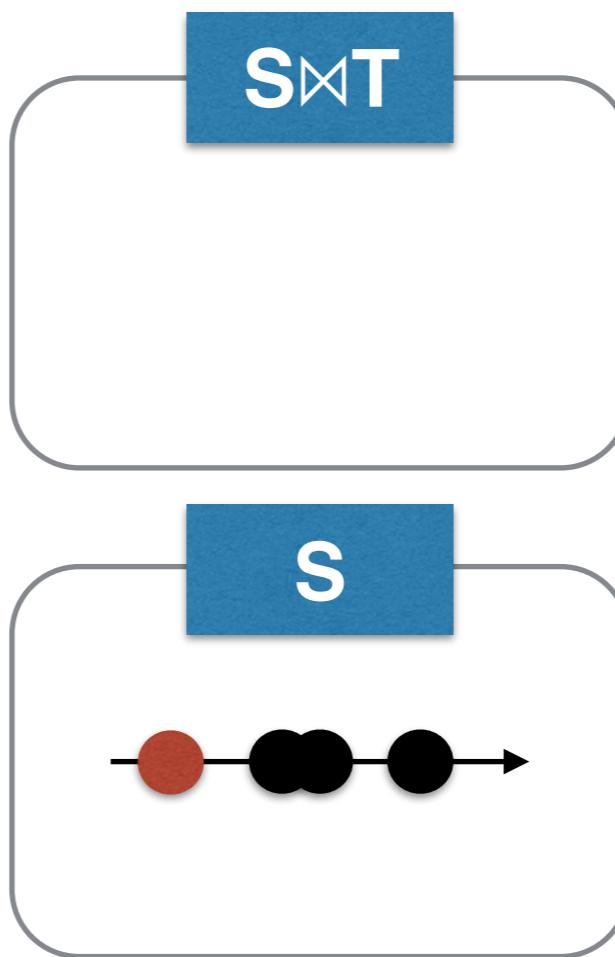
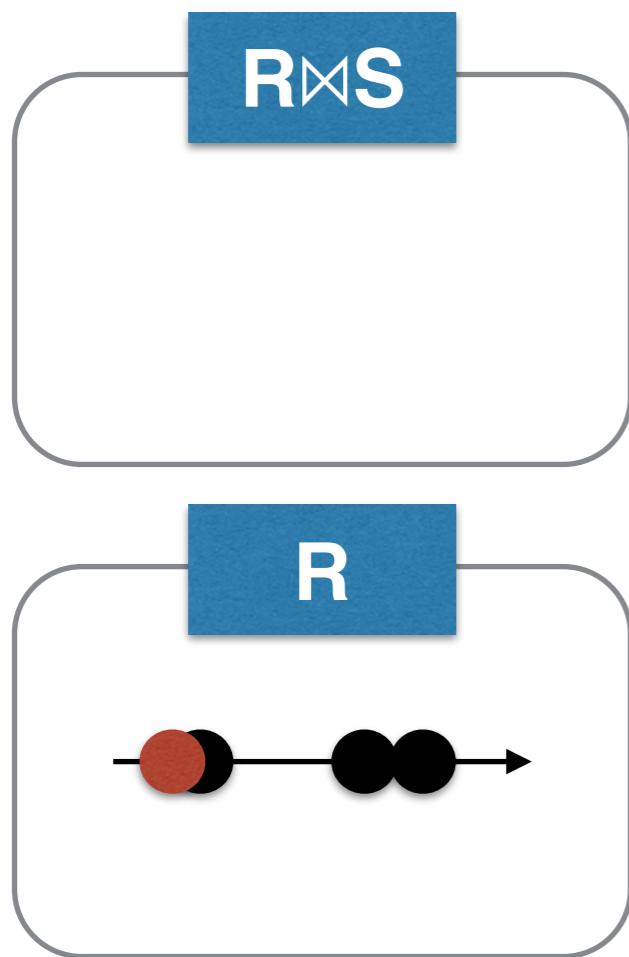
Classical Query Optimization



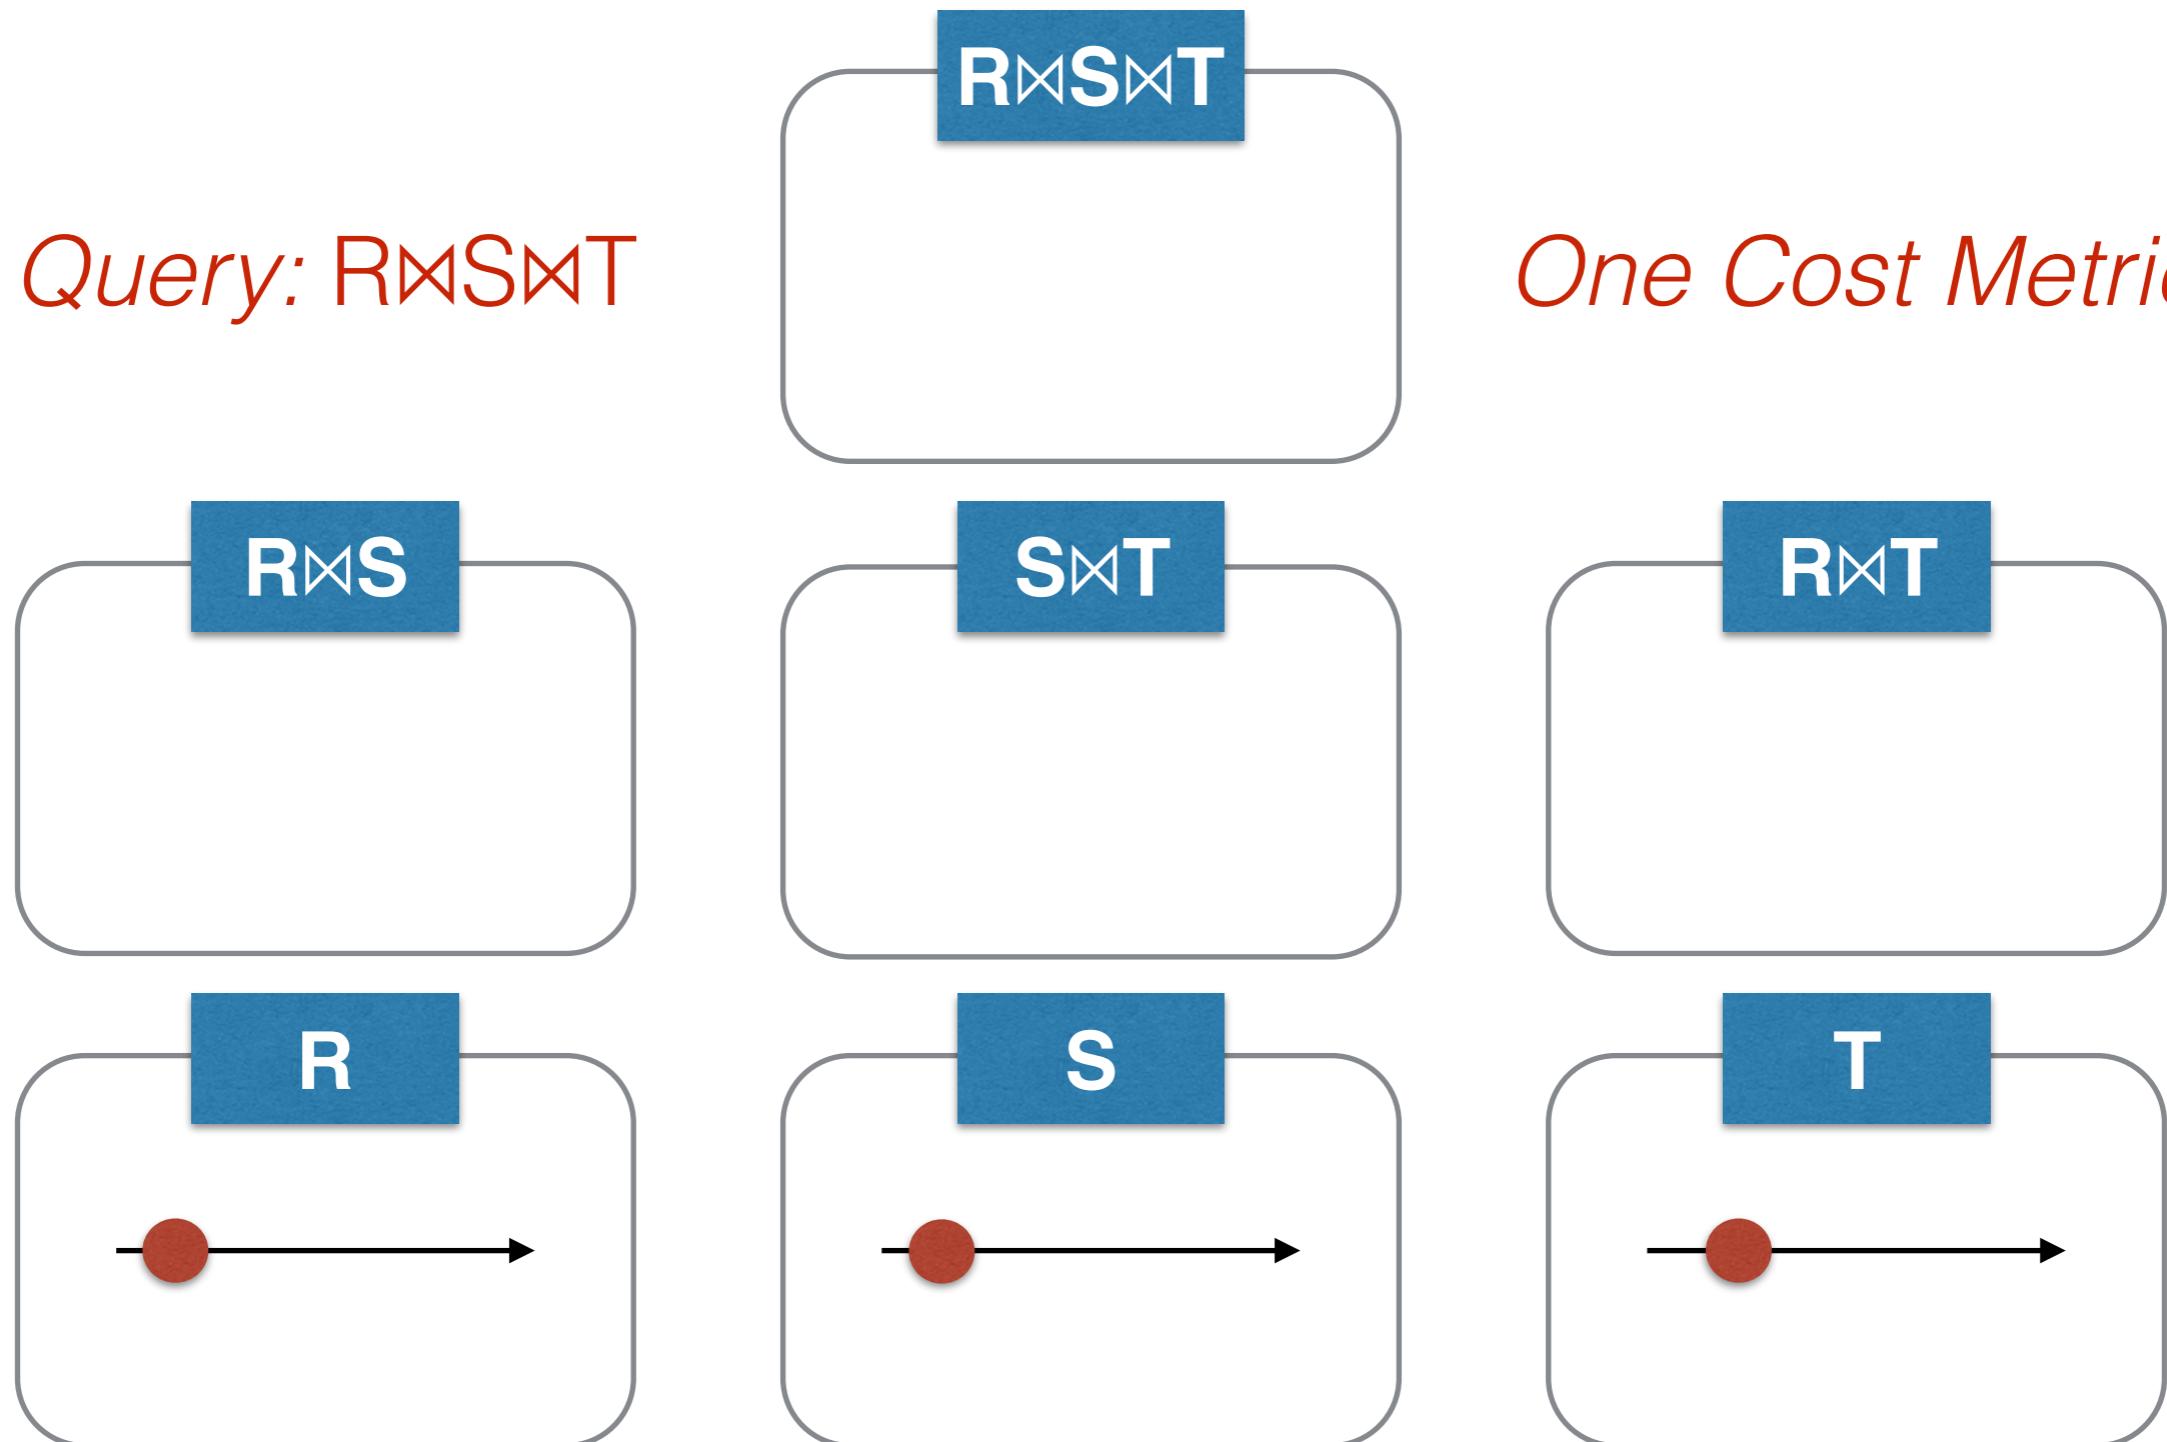
Classical Query Optimization

Query: R⊗S⊗T

One Cost Metric



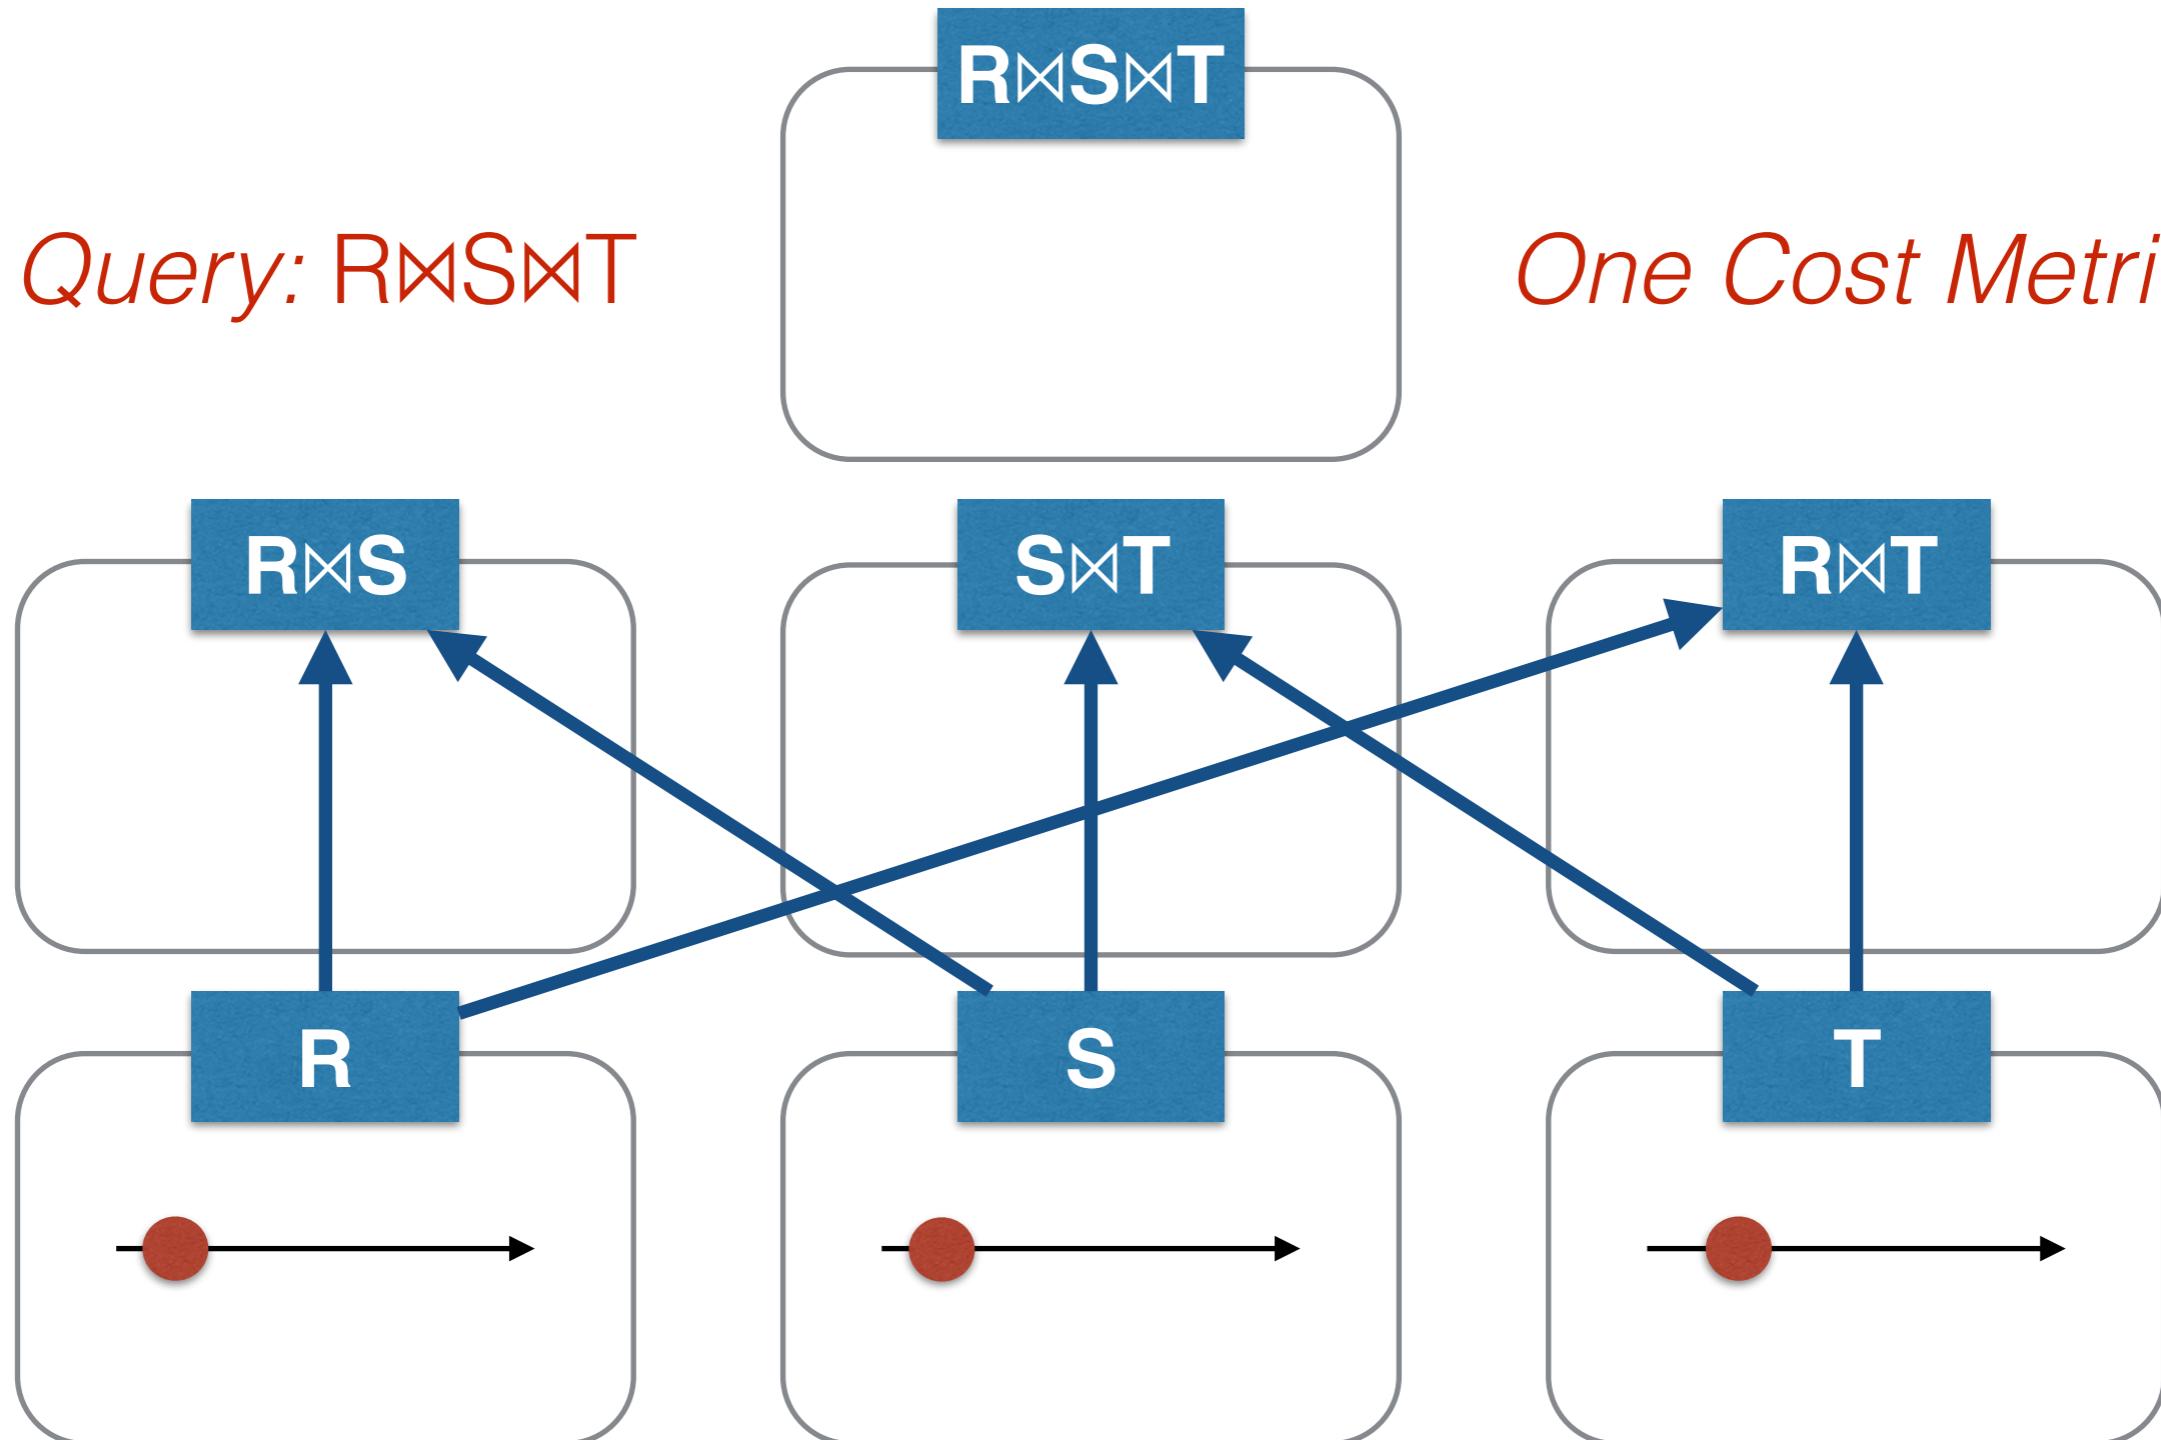
Classical Query Optimization



Classical Query Optimization

Query: R \bowtie S \bowtie T

One Cost Metric



Classical Query Optimization

Query: R \bowtie S \bowtie T

R \bowtie S \bowtie T

One Cost Metric

R \bowtie S

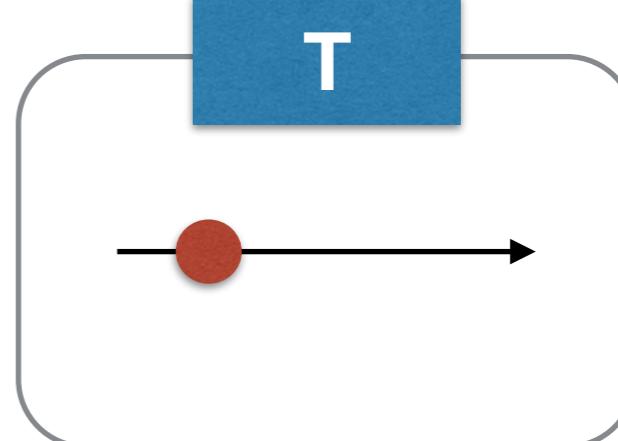
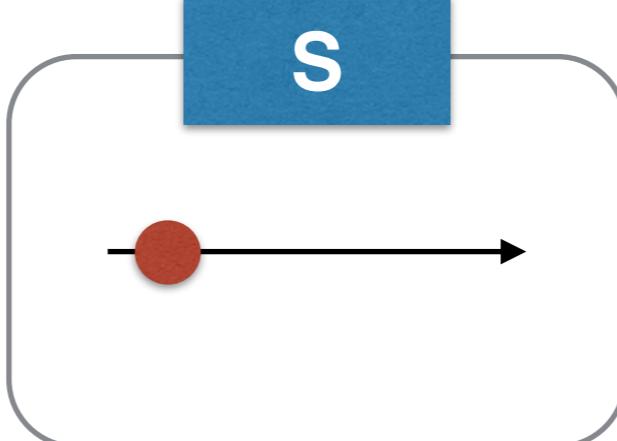
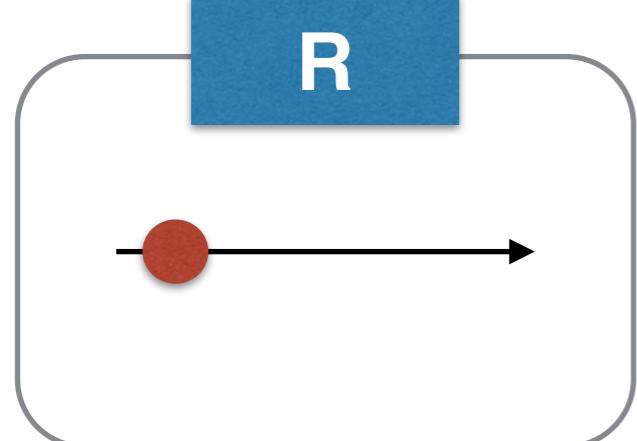
S \bowtie T

R \bowtie T

R

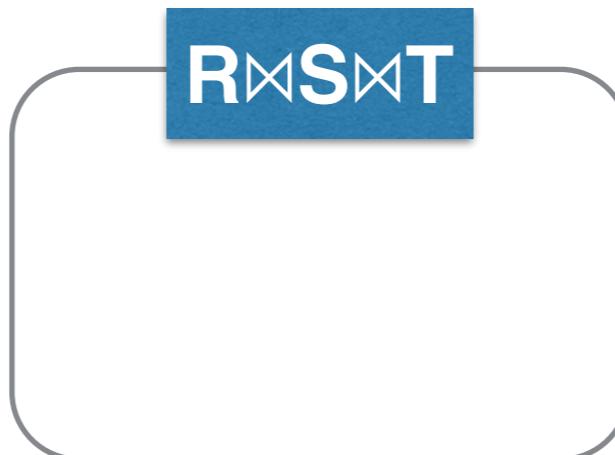
S

T

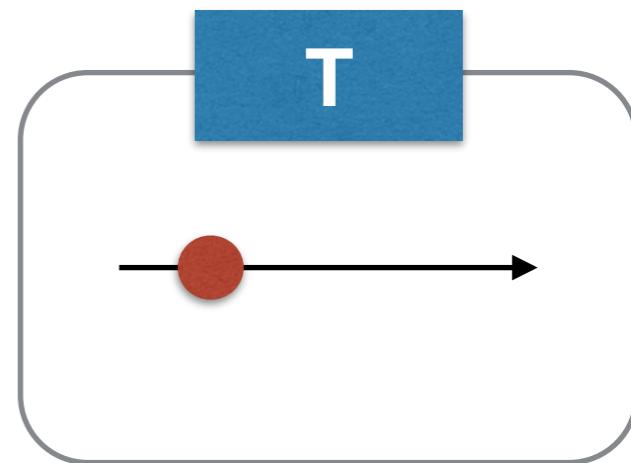
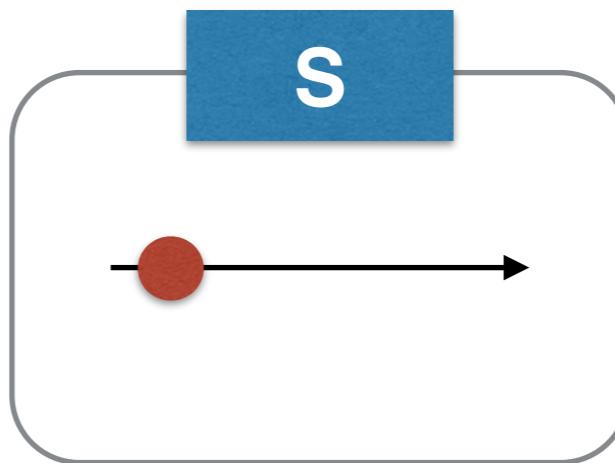
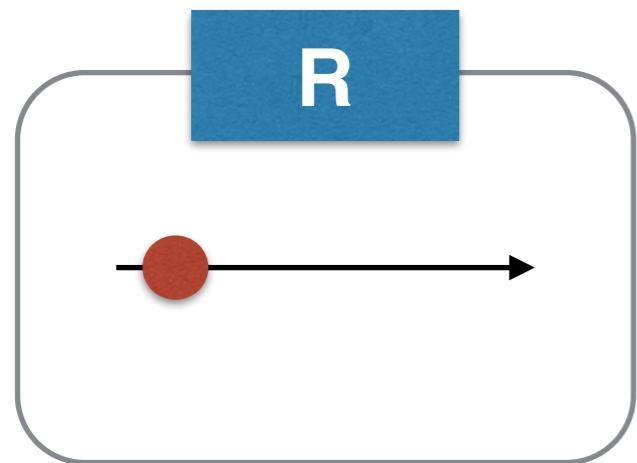
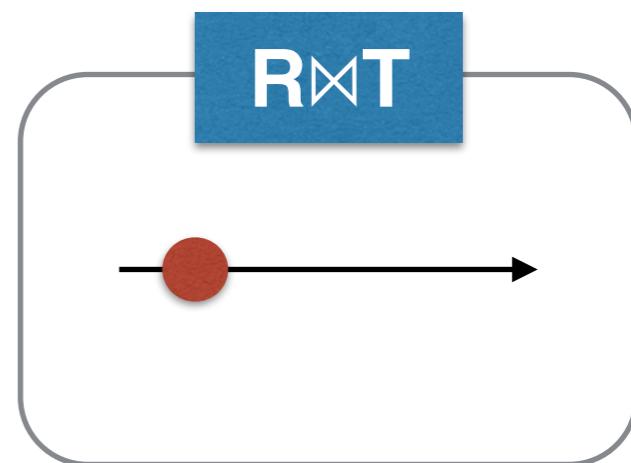
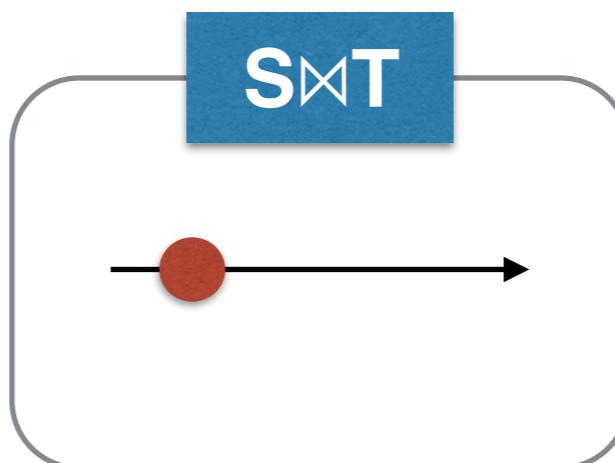
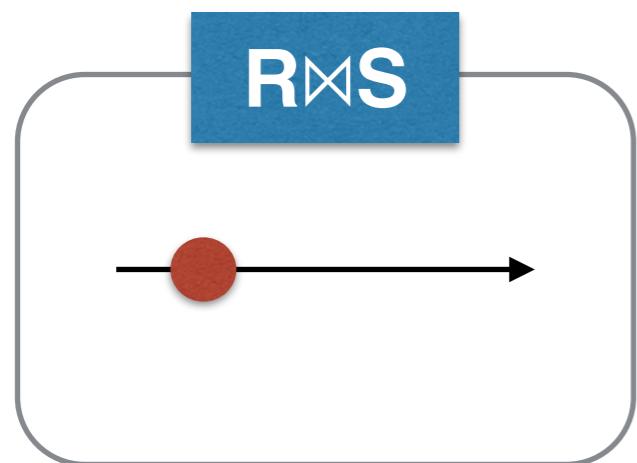


Classical Query Optimization

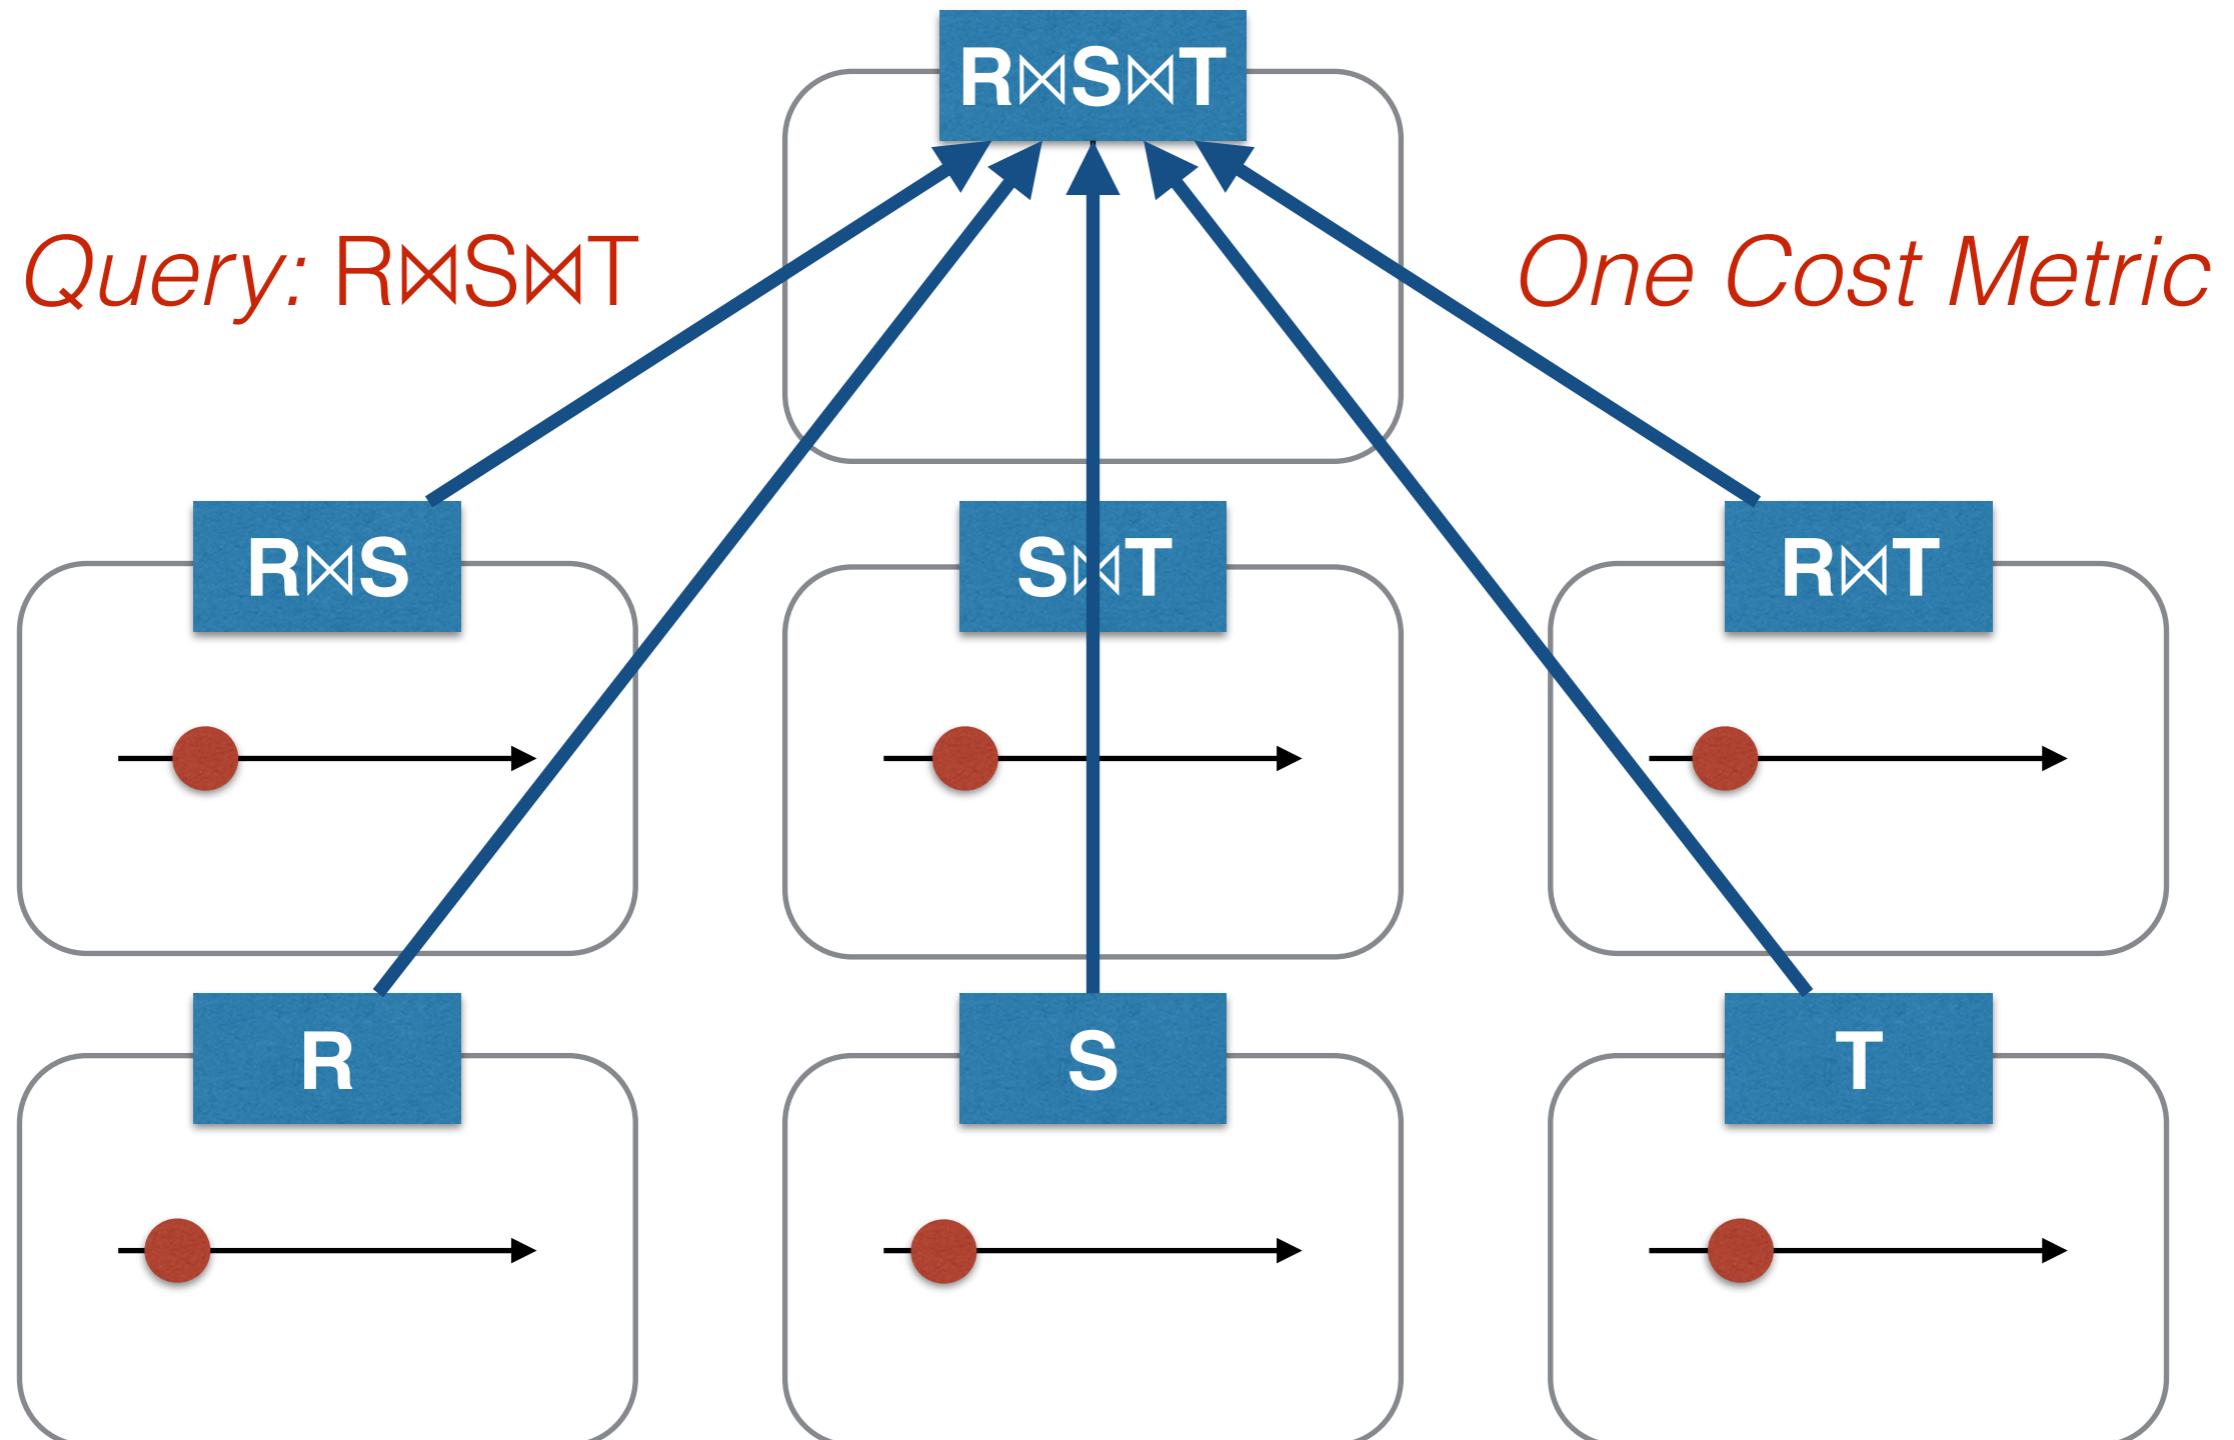
Query: R \bowtie S \bowtie T



One Cost Metric

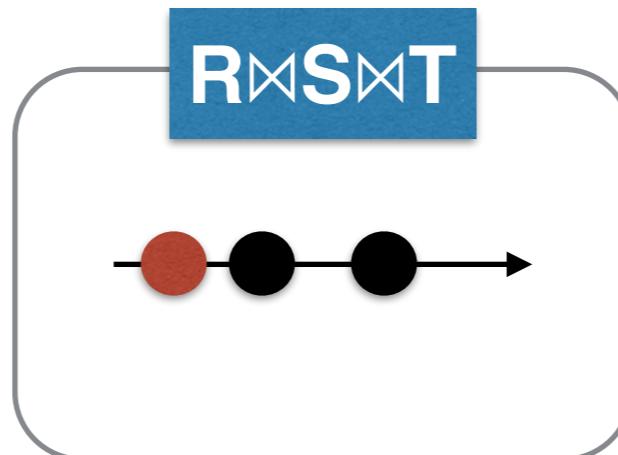


Classical Query Optimization

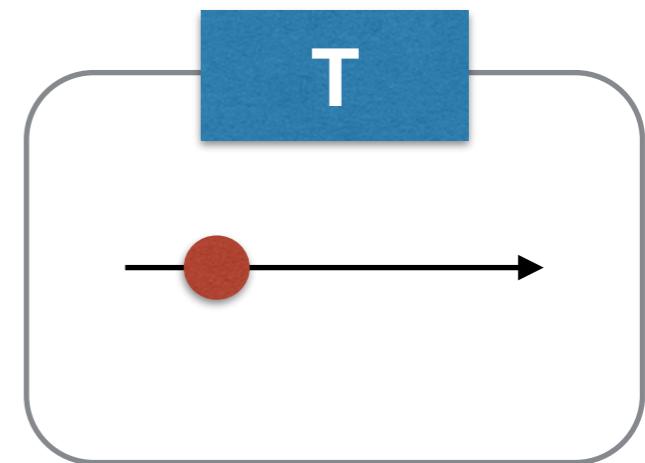
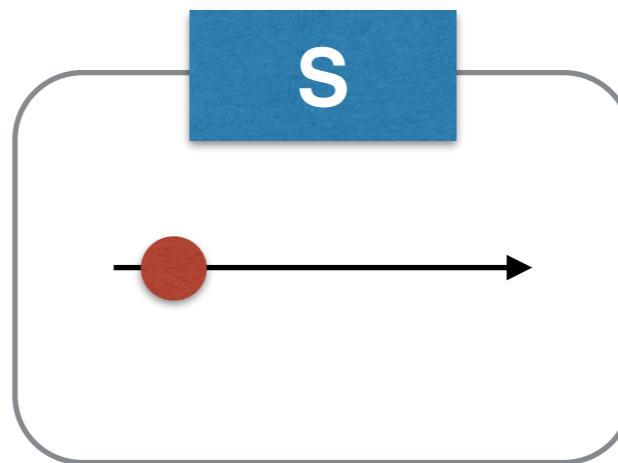
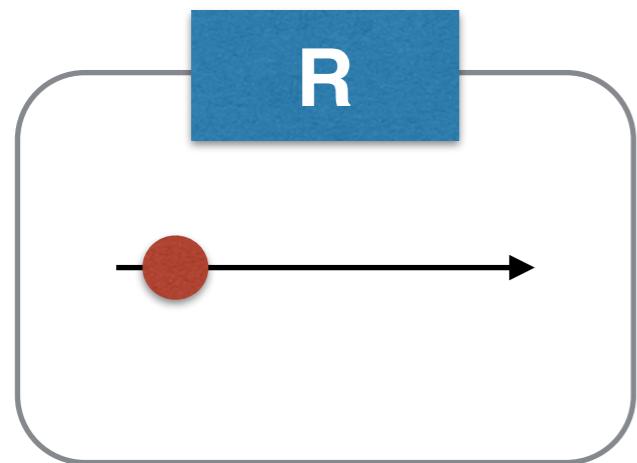
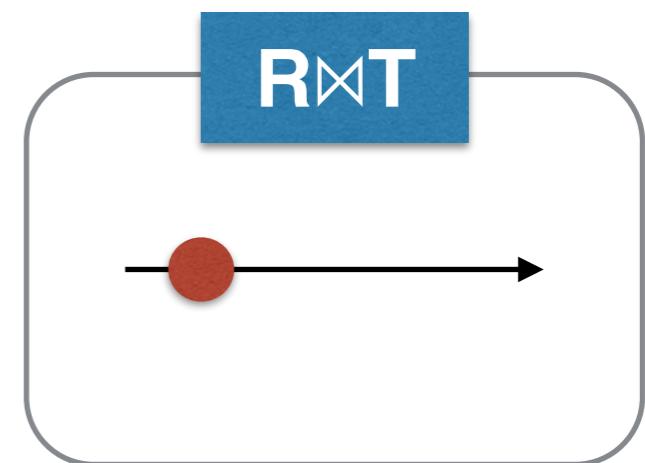
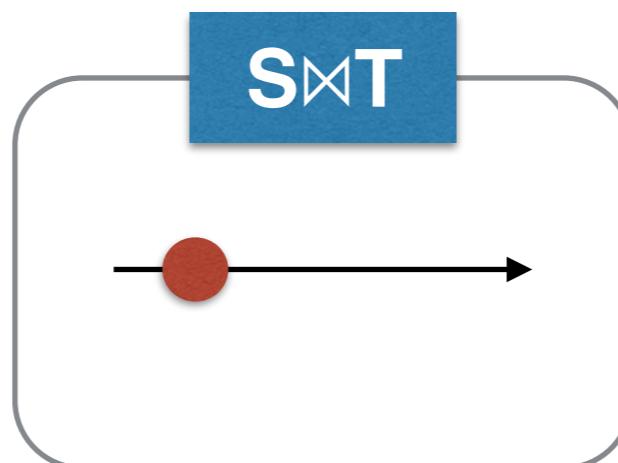
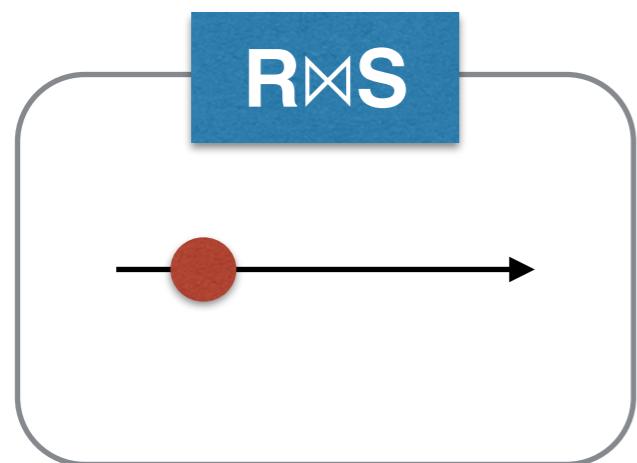


Classical Query Optimization

Query: R \bowtie S \bowtie T

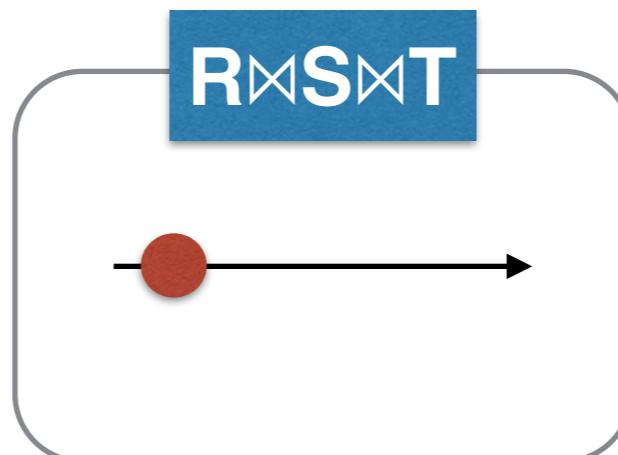


One Cost Metric

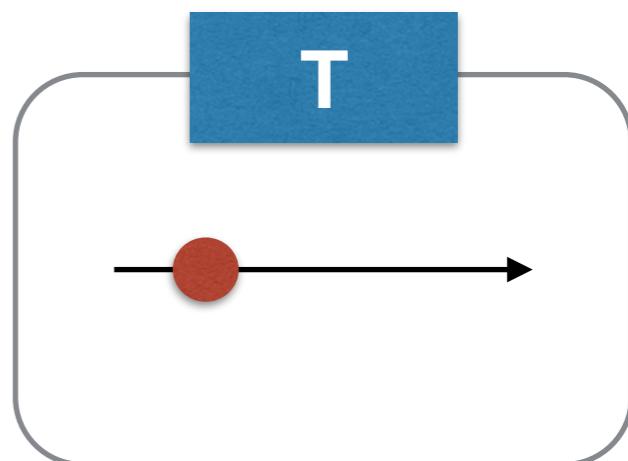
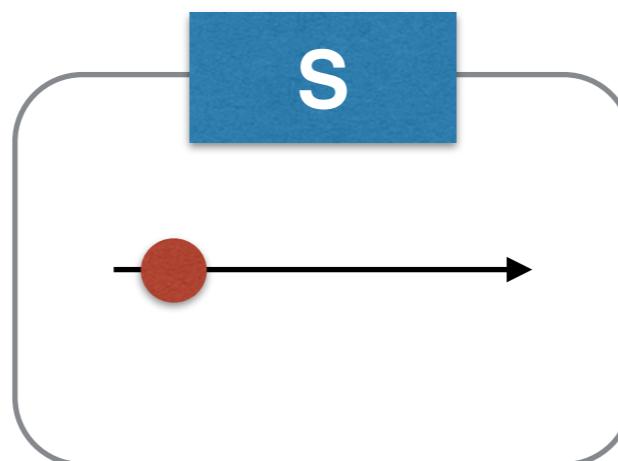
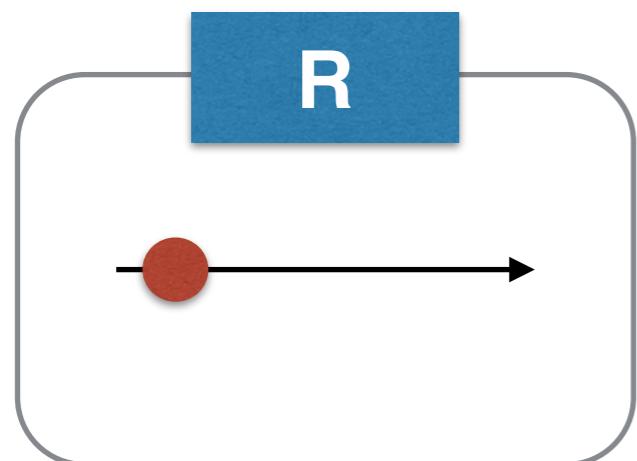
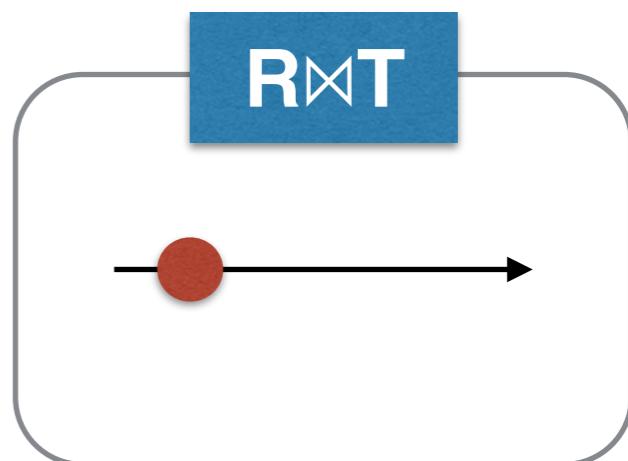
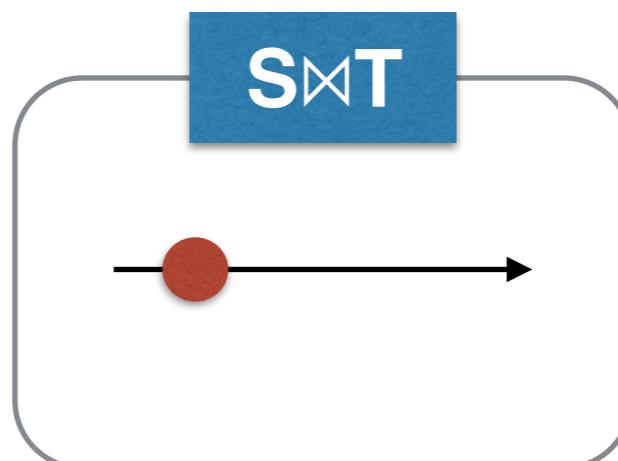
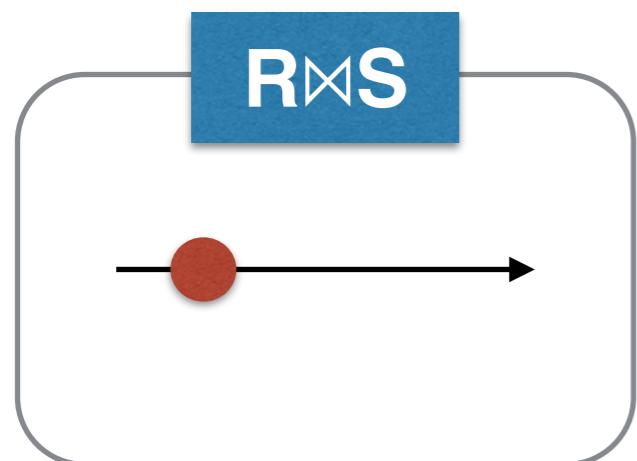


Classical Query Optimization

Query: R⊗S⊗T

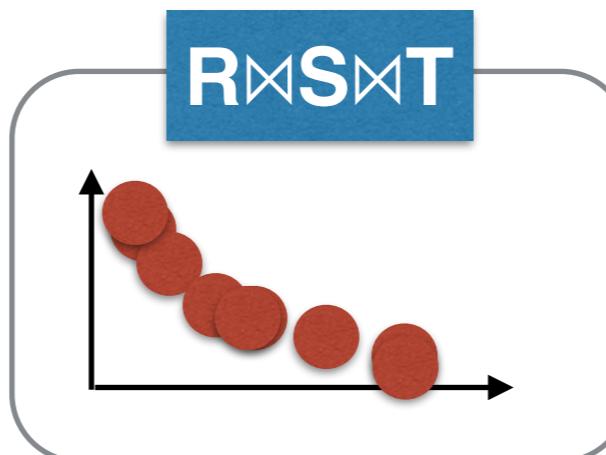


One Cost Metric

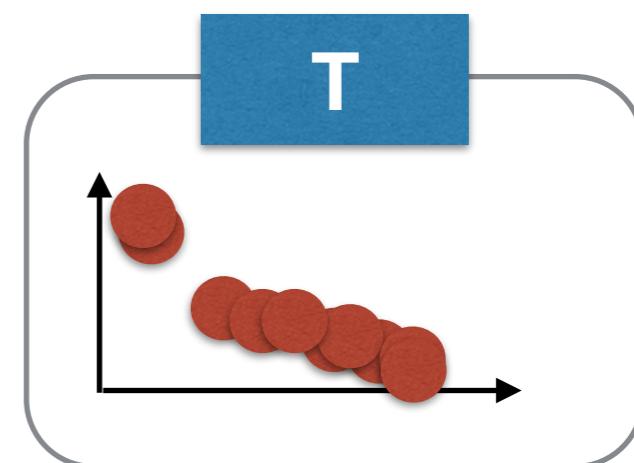
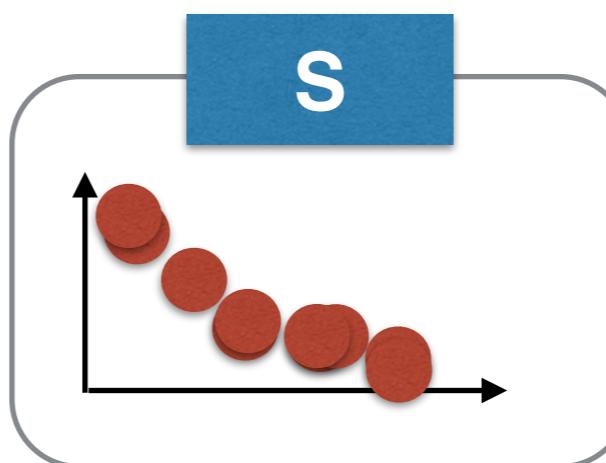
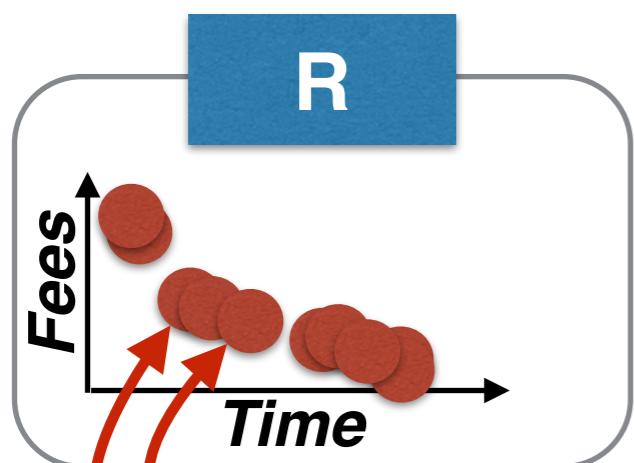
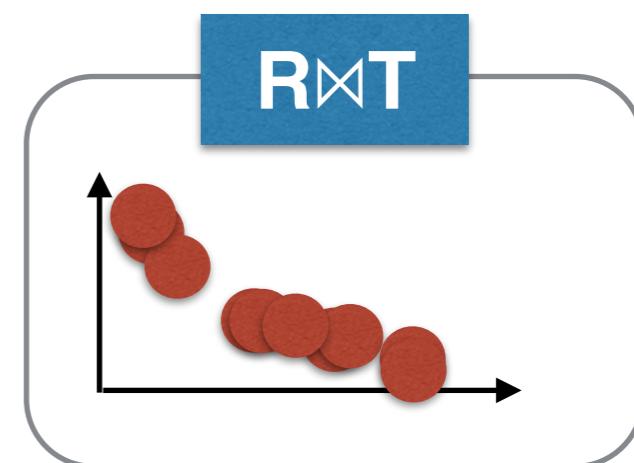
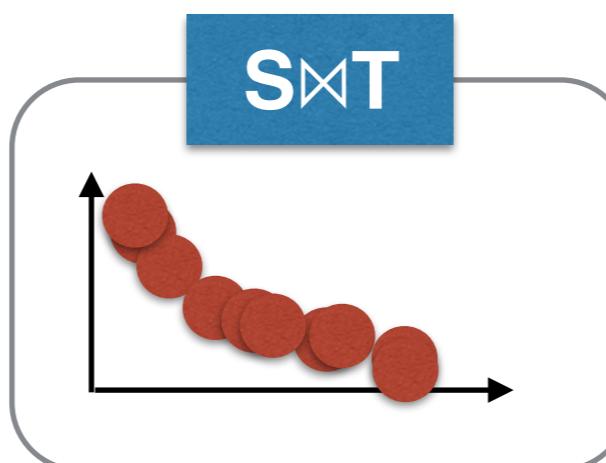
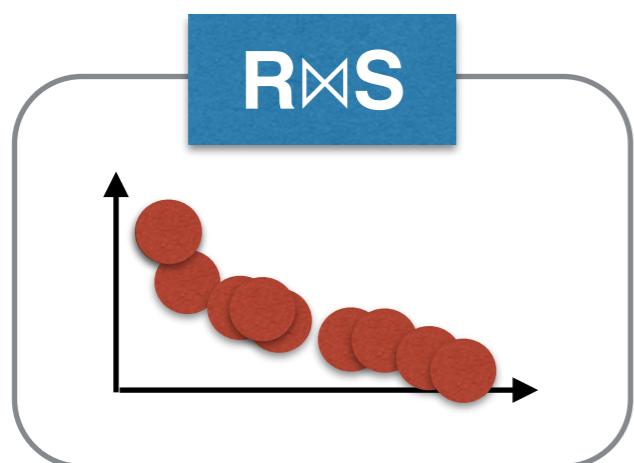


Multi-Objective Query Optimization

Query: R⊗S⊗T



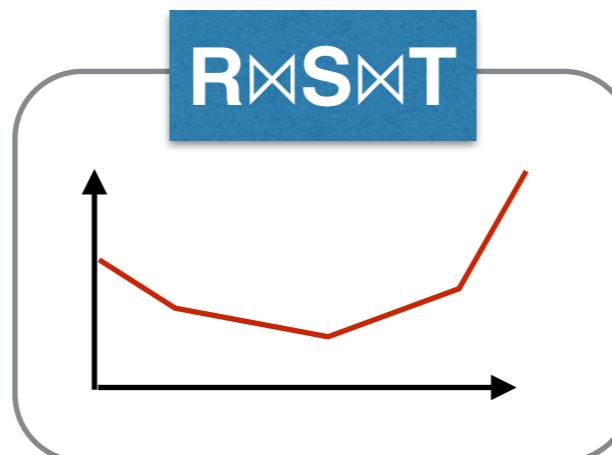
Two Cost Metrics



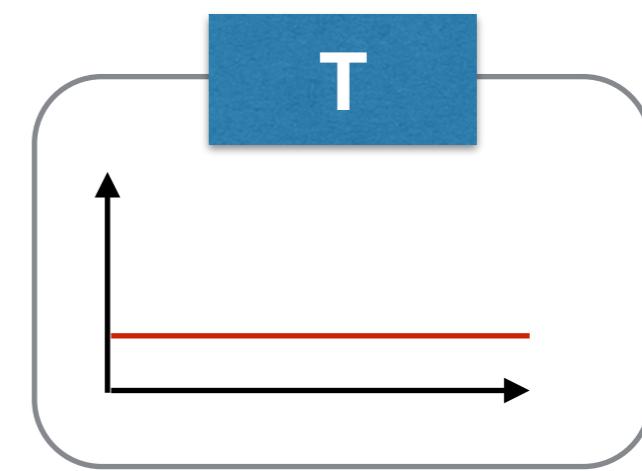
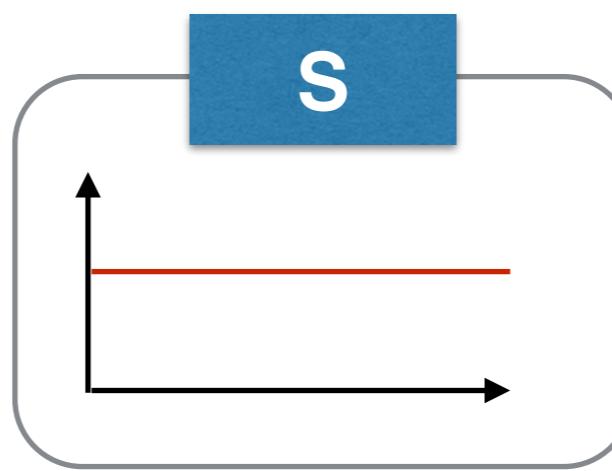
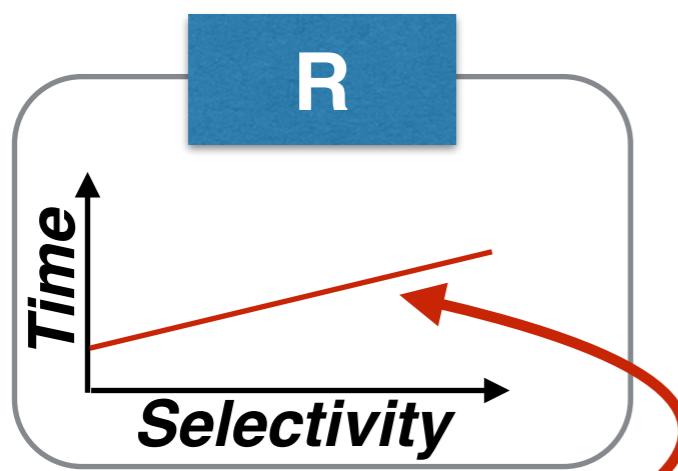
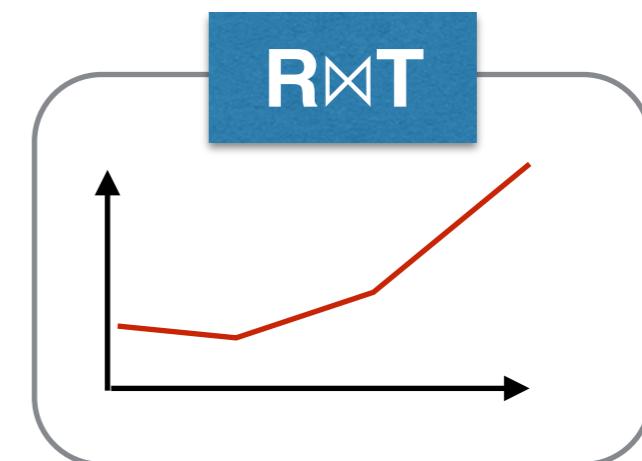
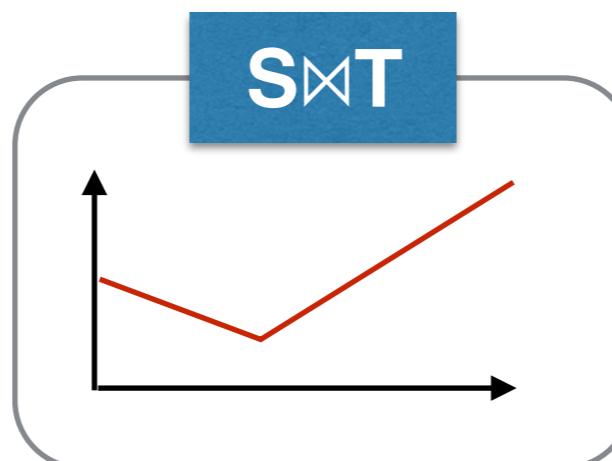
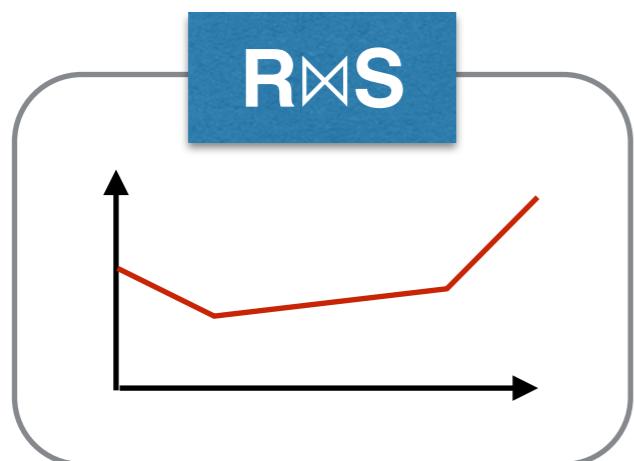
*Optimal
Plans*

Parametric Query Optimization

Query: $R \bowtie S \bowtie T$



*One Cost Metric
One Unknown*



Cost Function

Parallelization Challenges

Sub-Queries

R,S,T,U,V

R,T,U,V

S,T,U,V

R,S,T,V

R,S,U,V

R,S,T,U

R,U,V

S,T,U

S,T,V

S,U,V

T,U,V

R,S,T

R,S,U

R,S,V

R,T,U

R,T,V

S,U

S,V

T,U

T,V

U,V

R,S

R,T

R,U

R,V

S,T

R

S

T

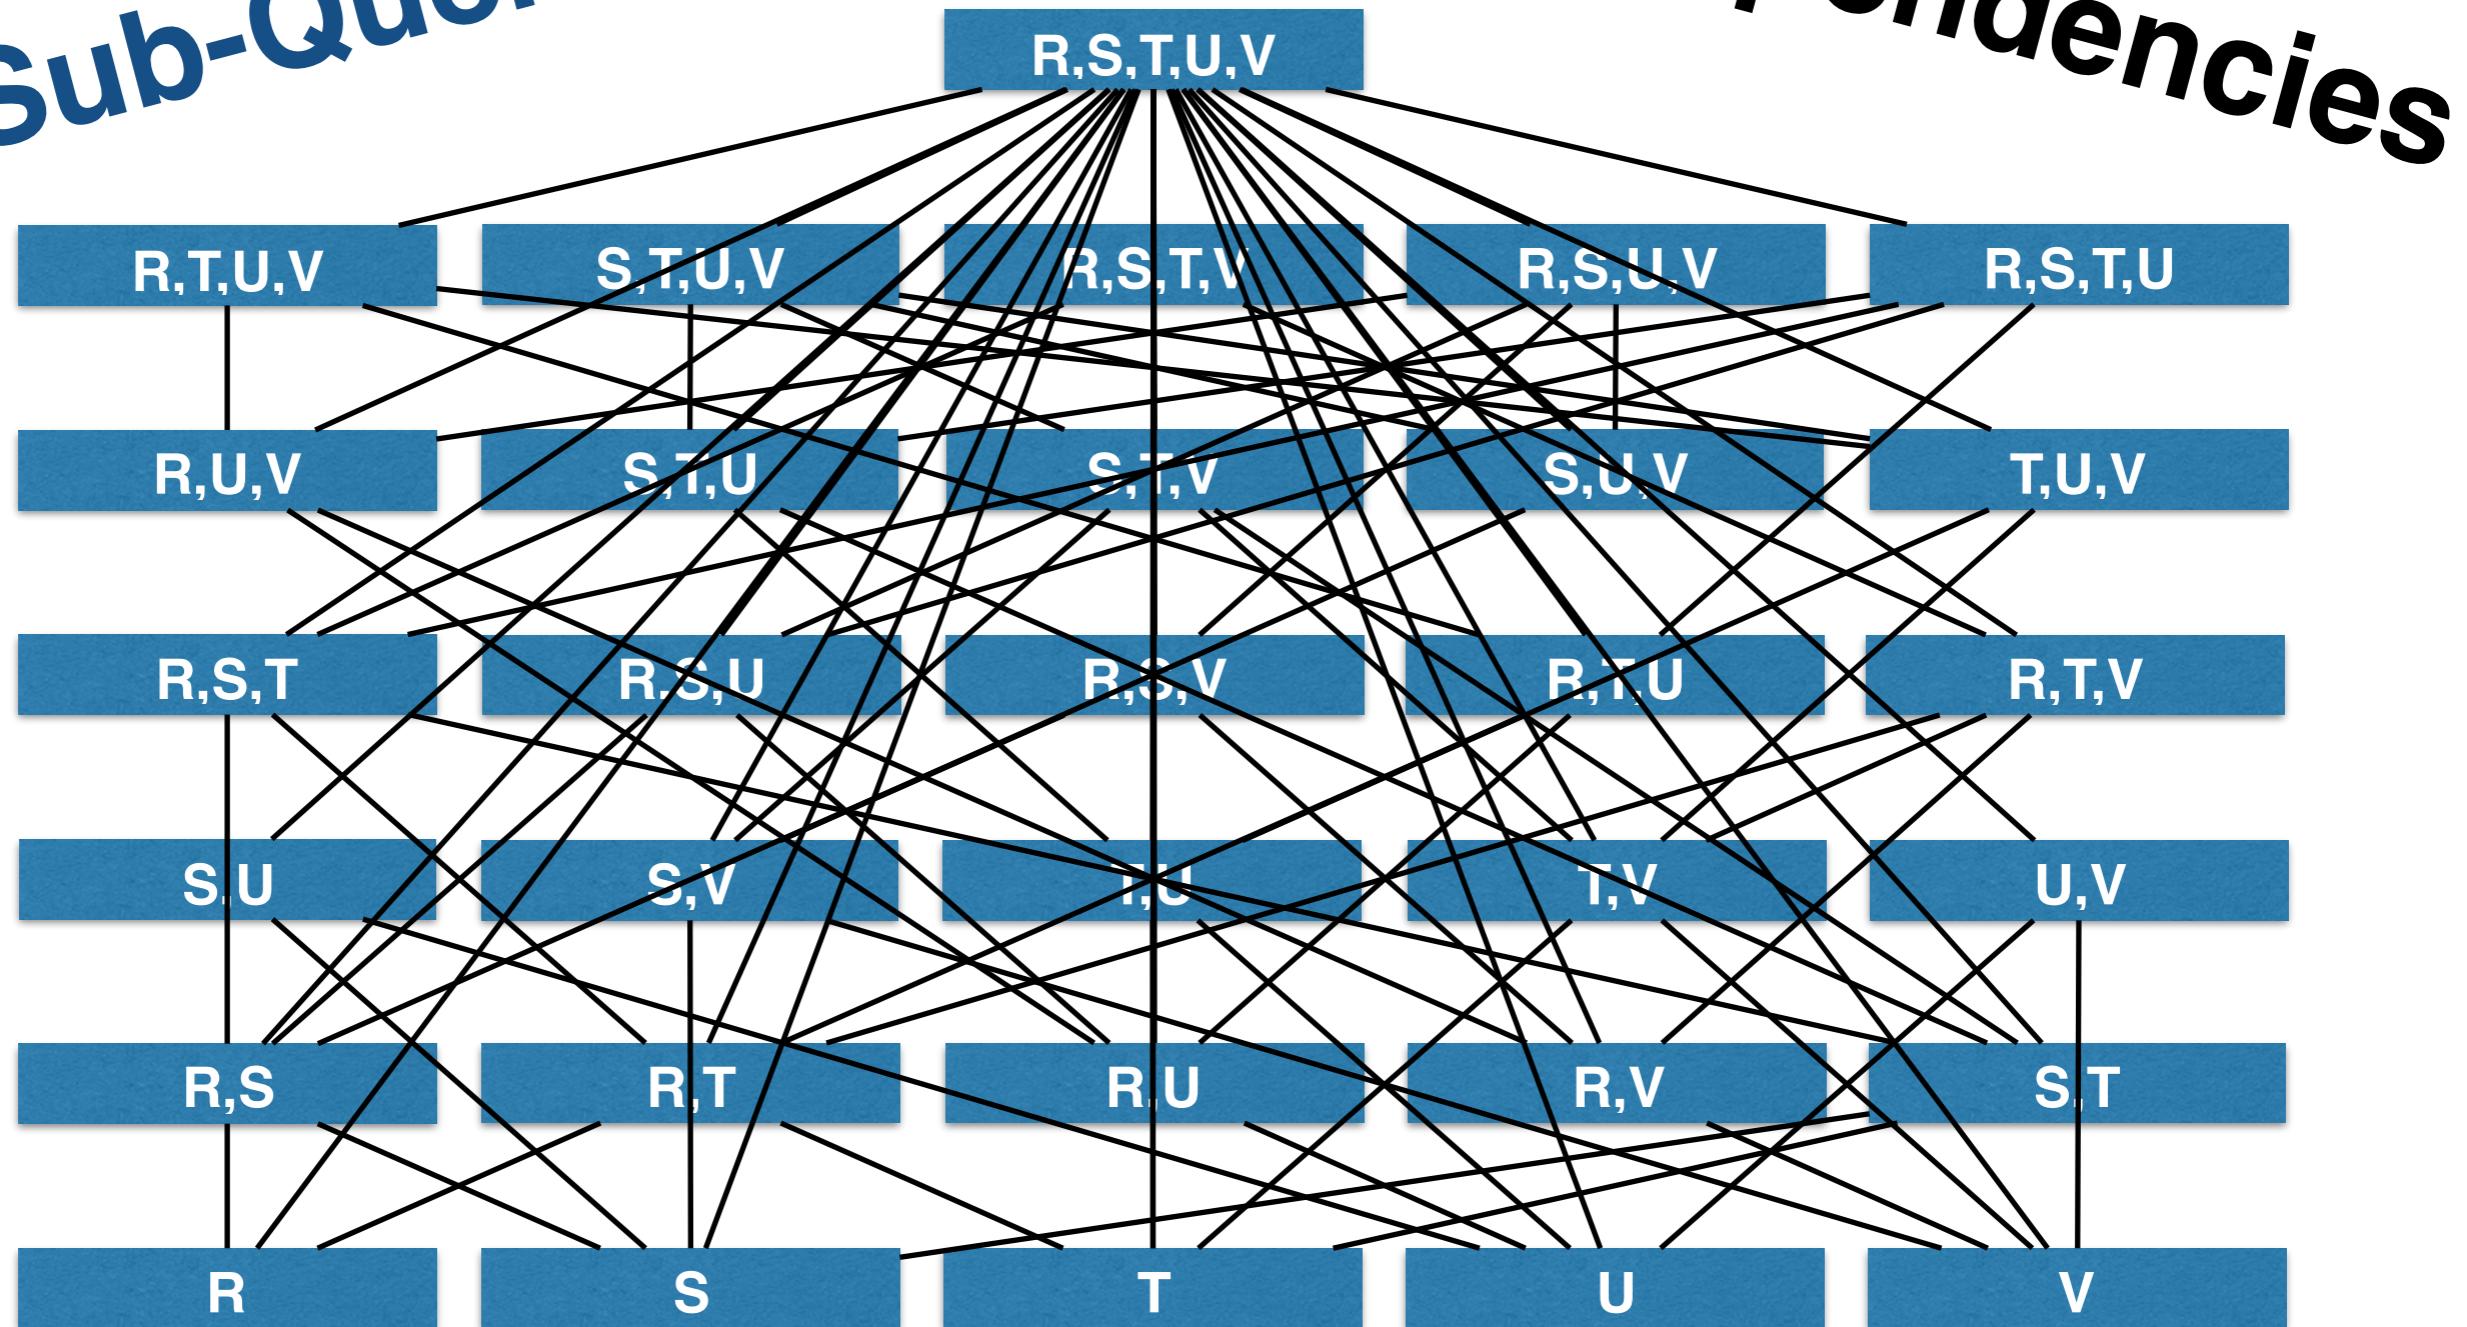
U

V

Parallelization Challenges

Sub-Queries

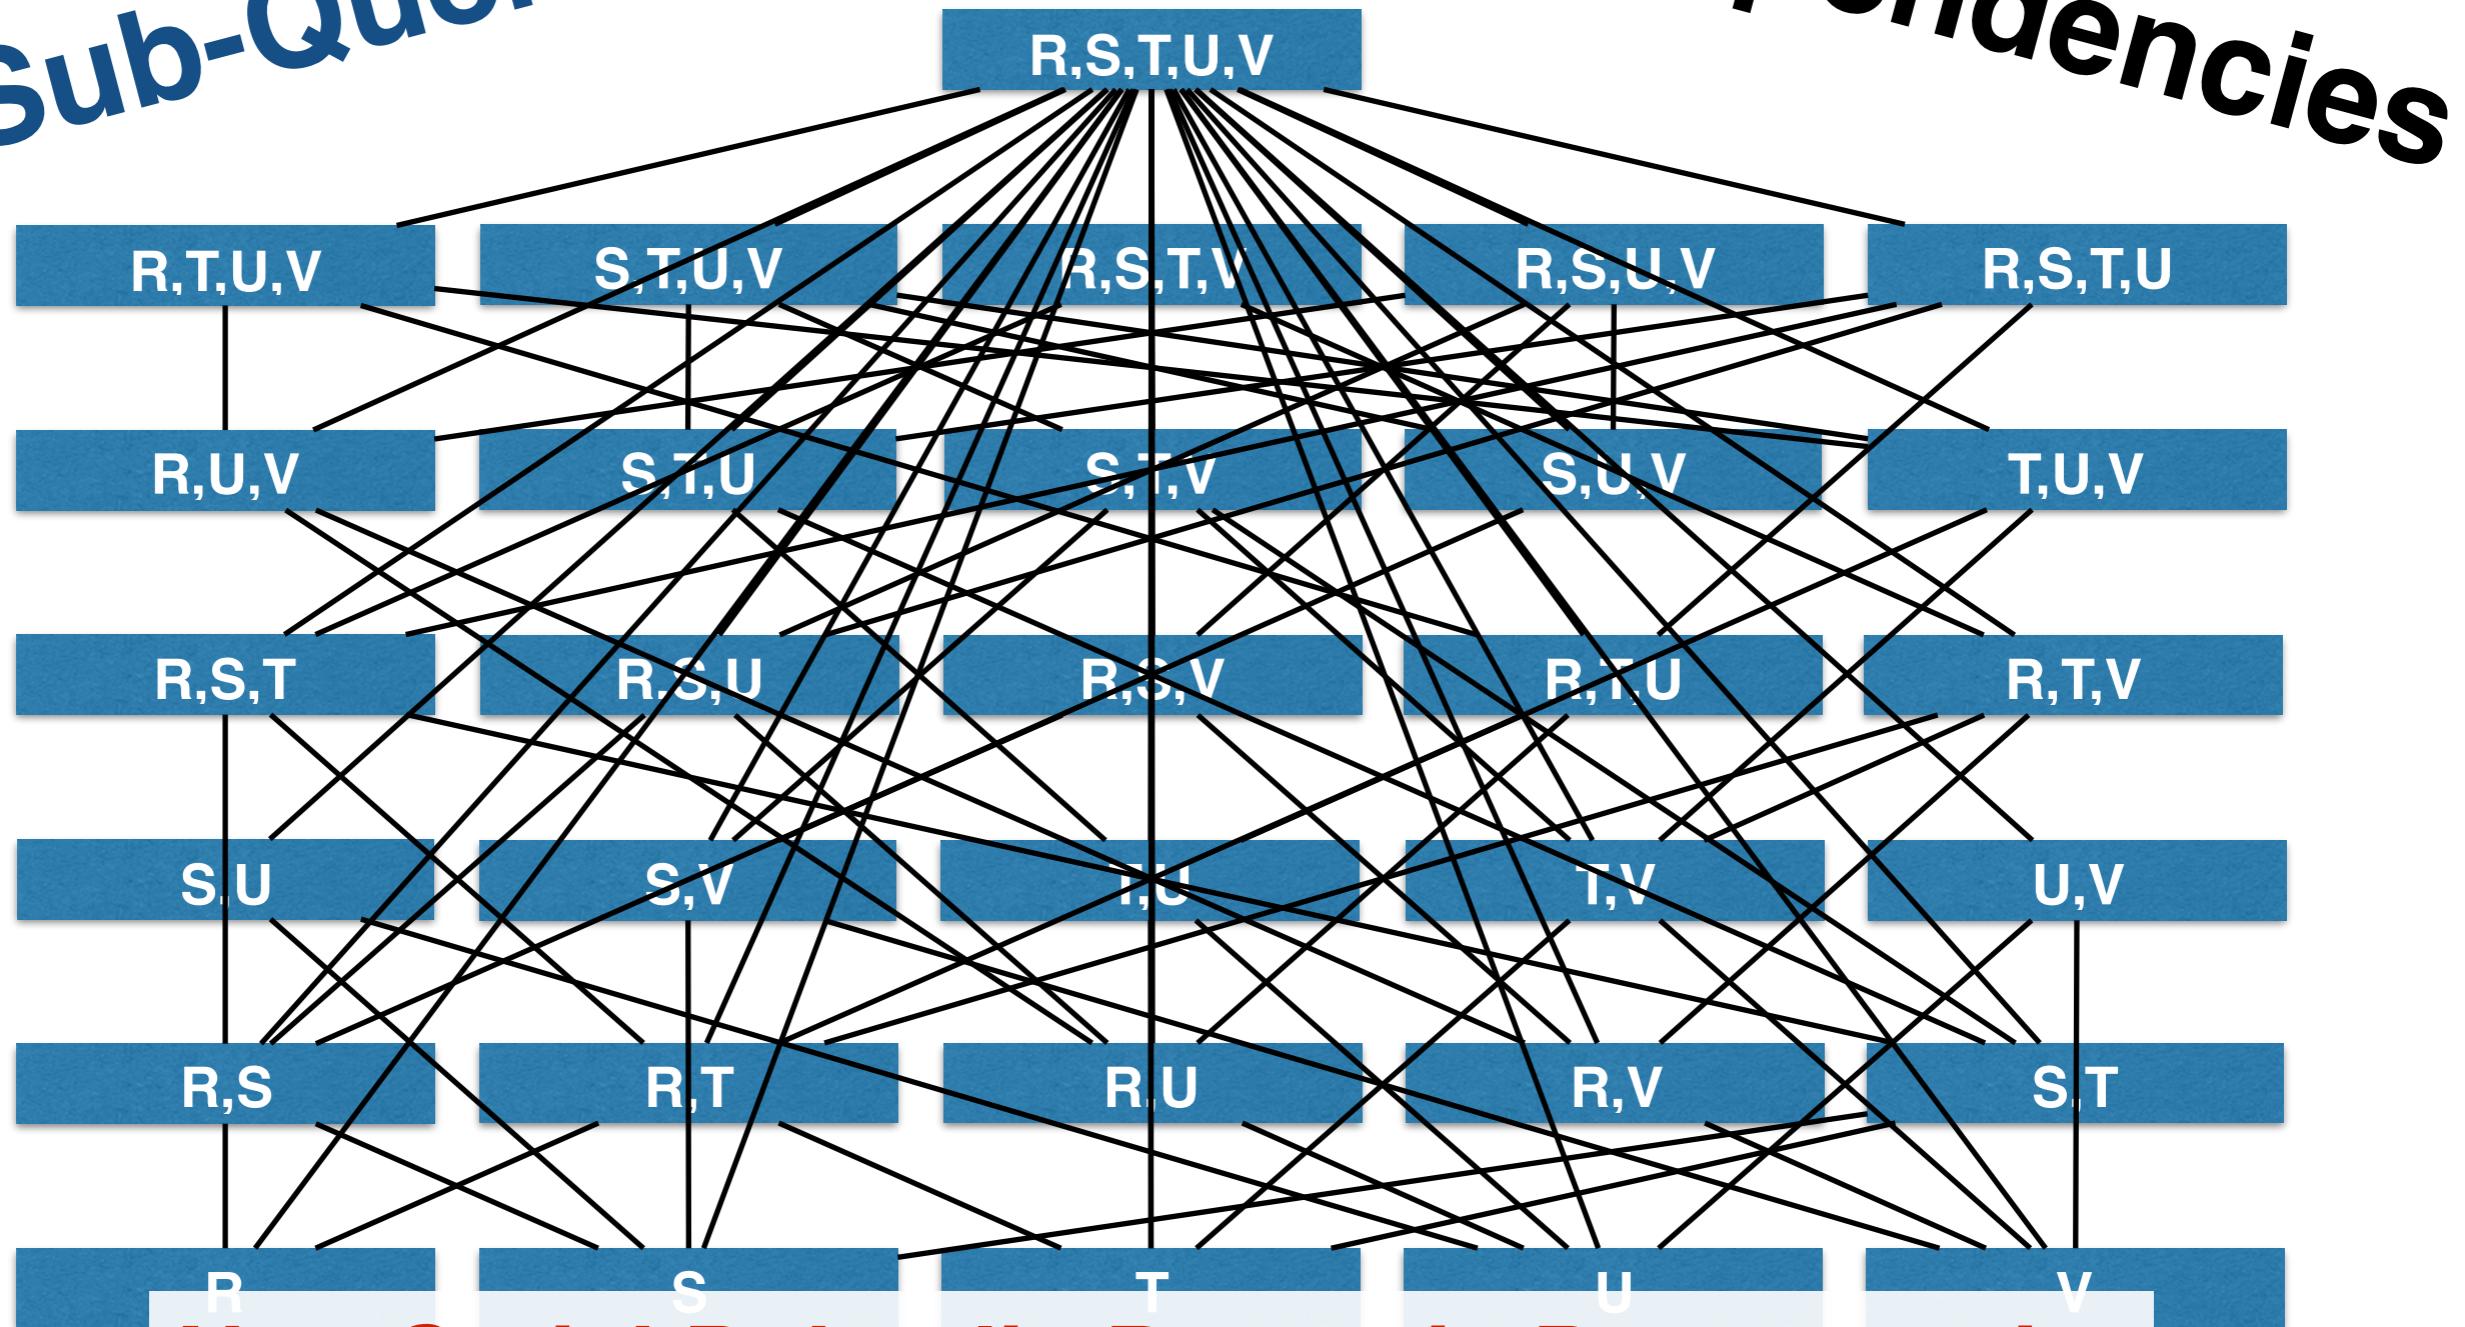
Dependencies



Parallelization Challenges

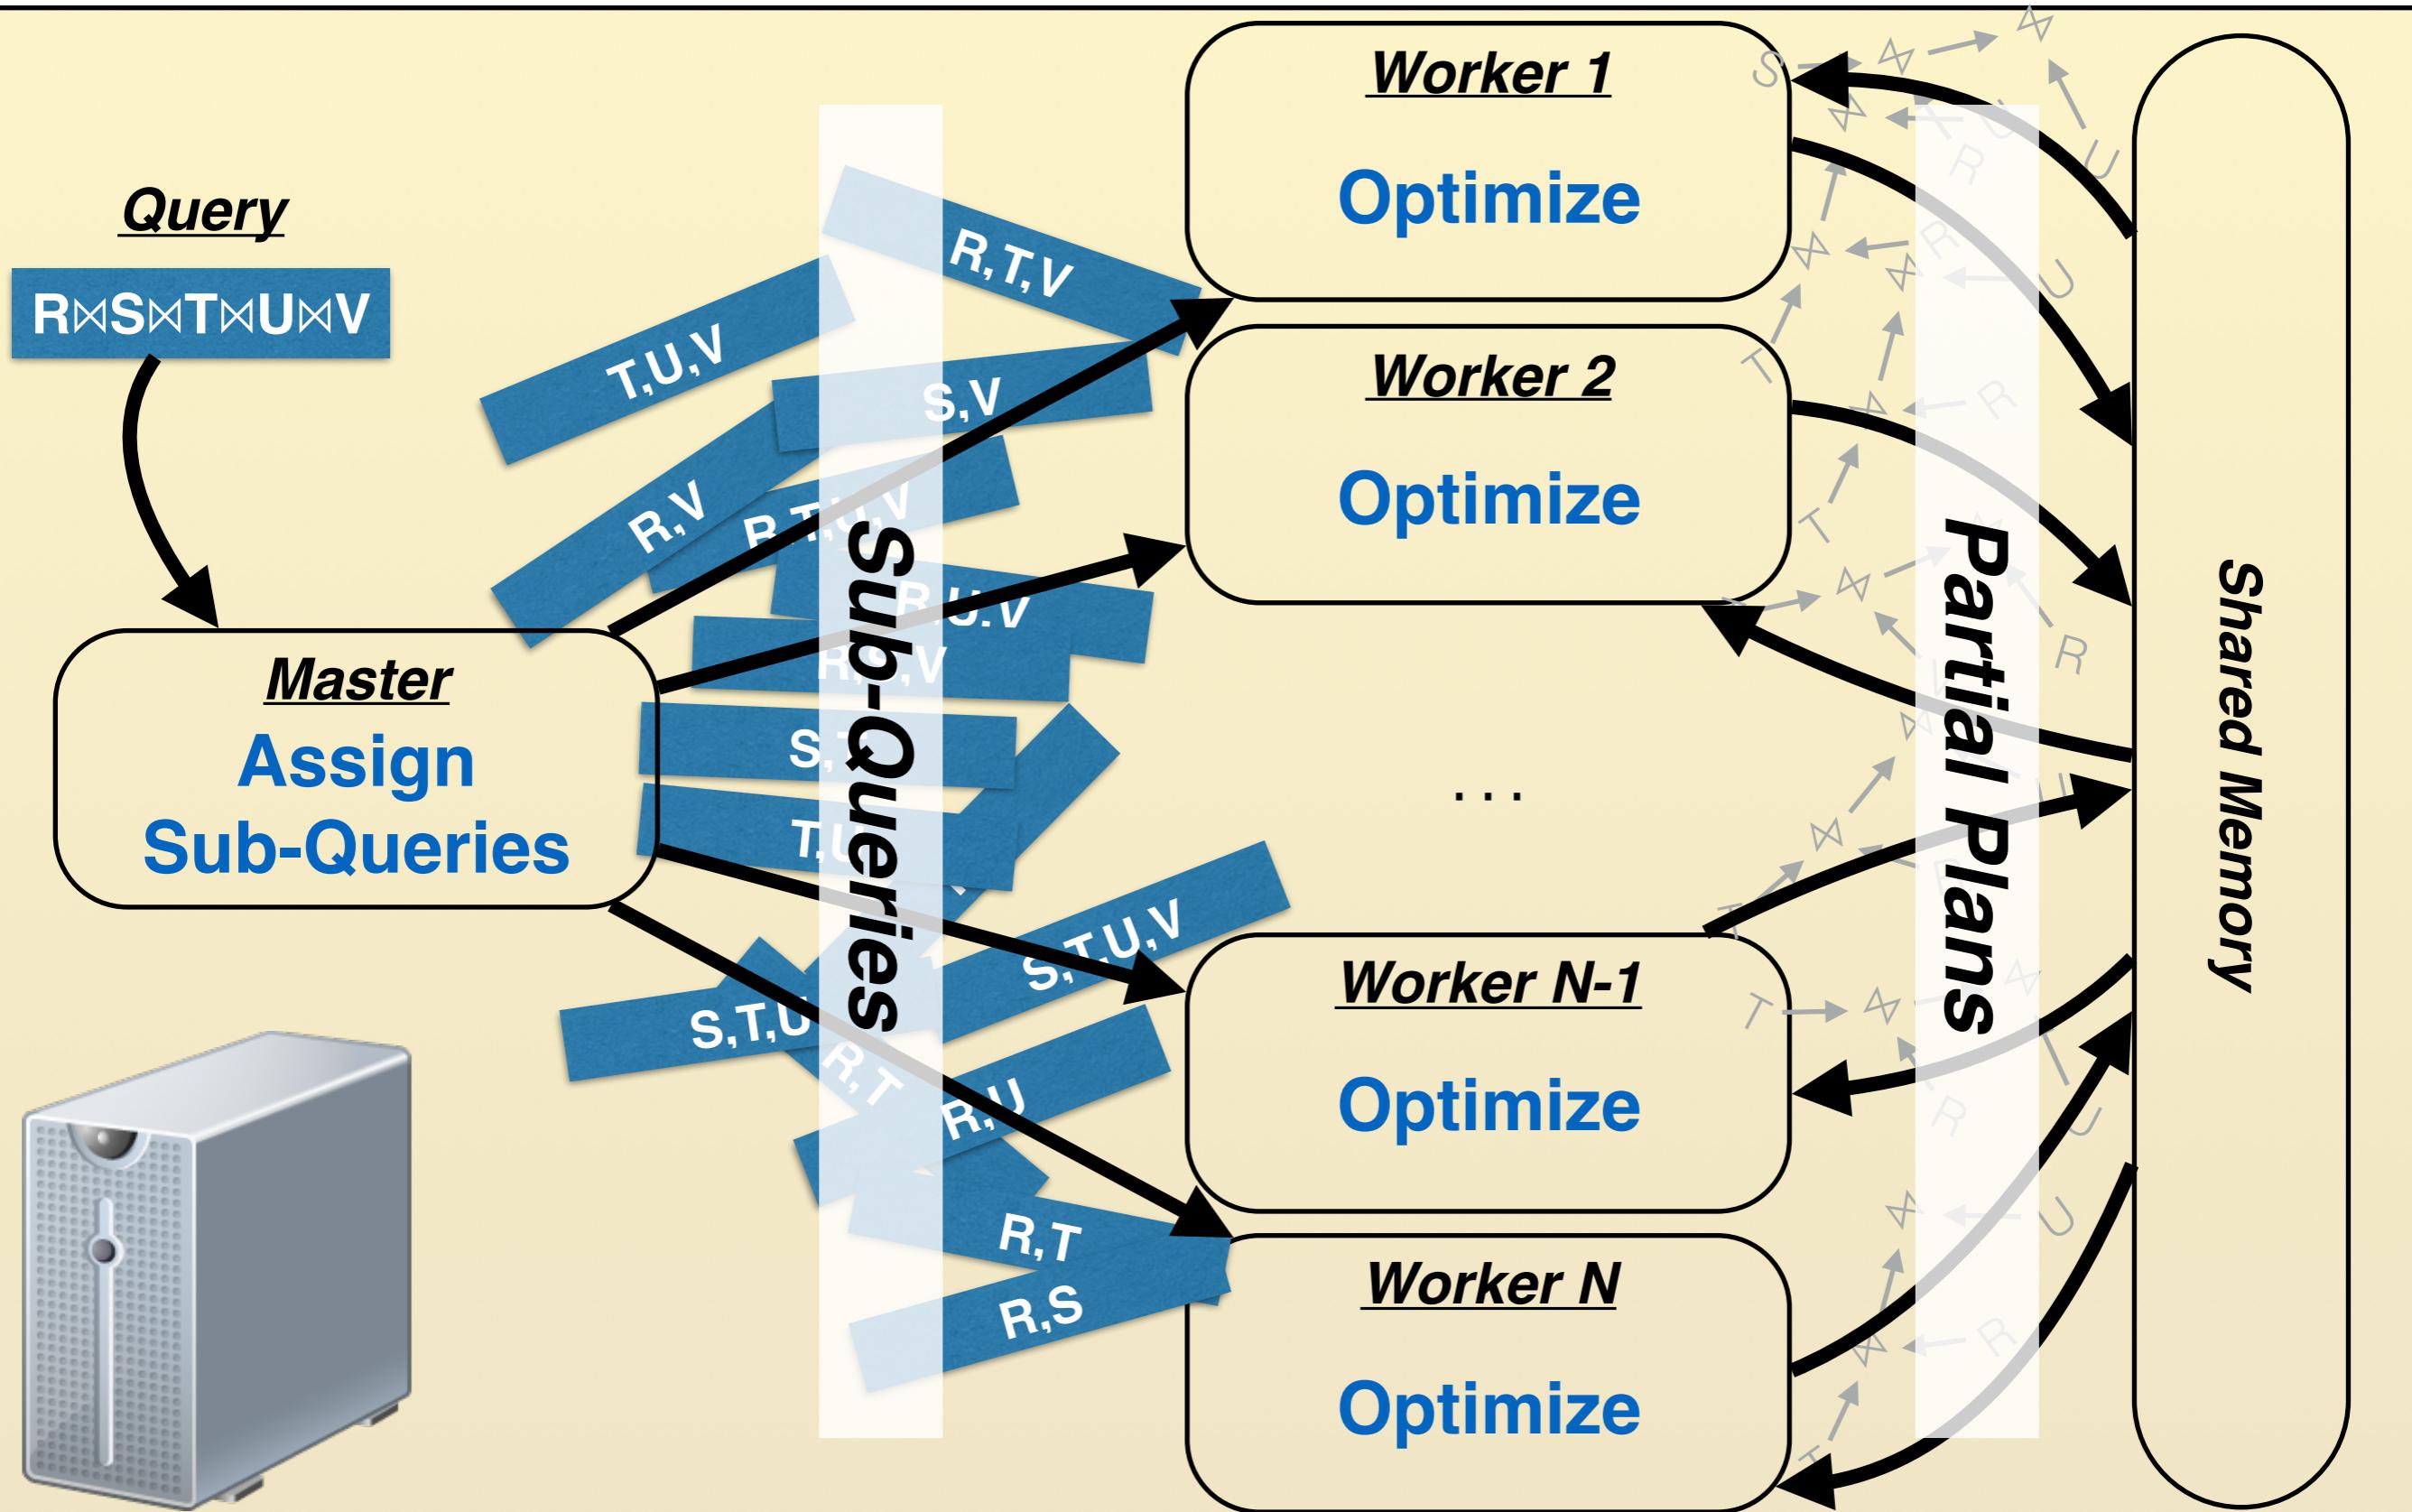
Sub-Queries

Dependencies

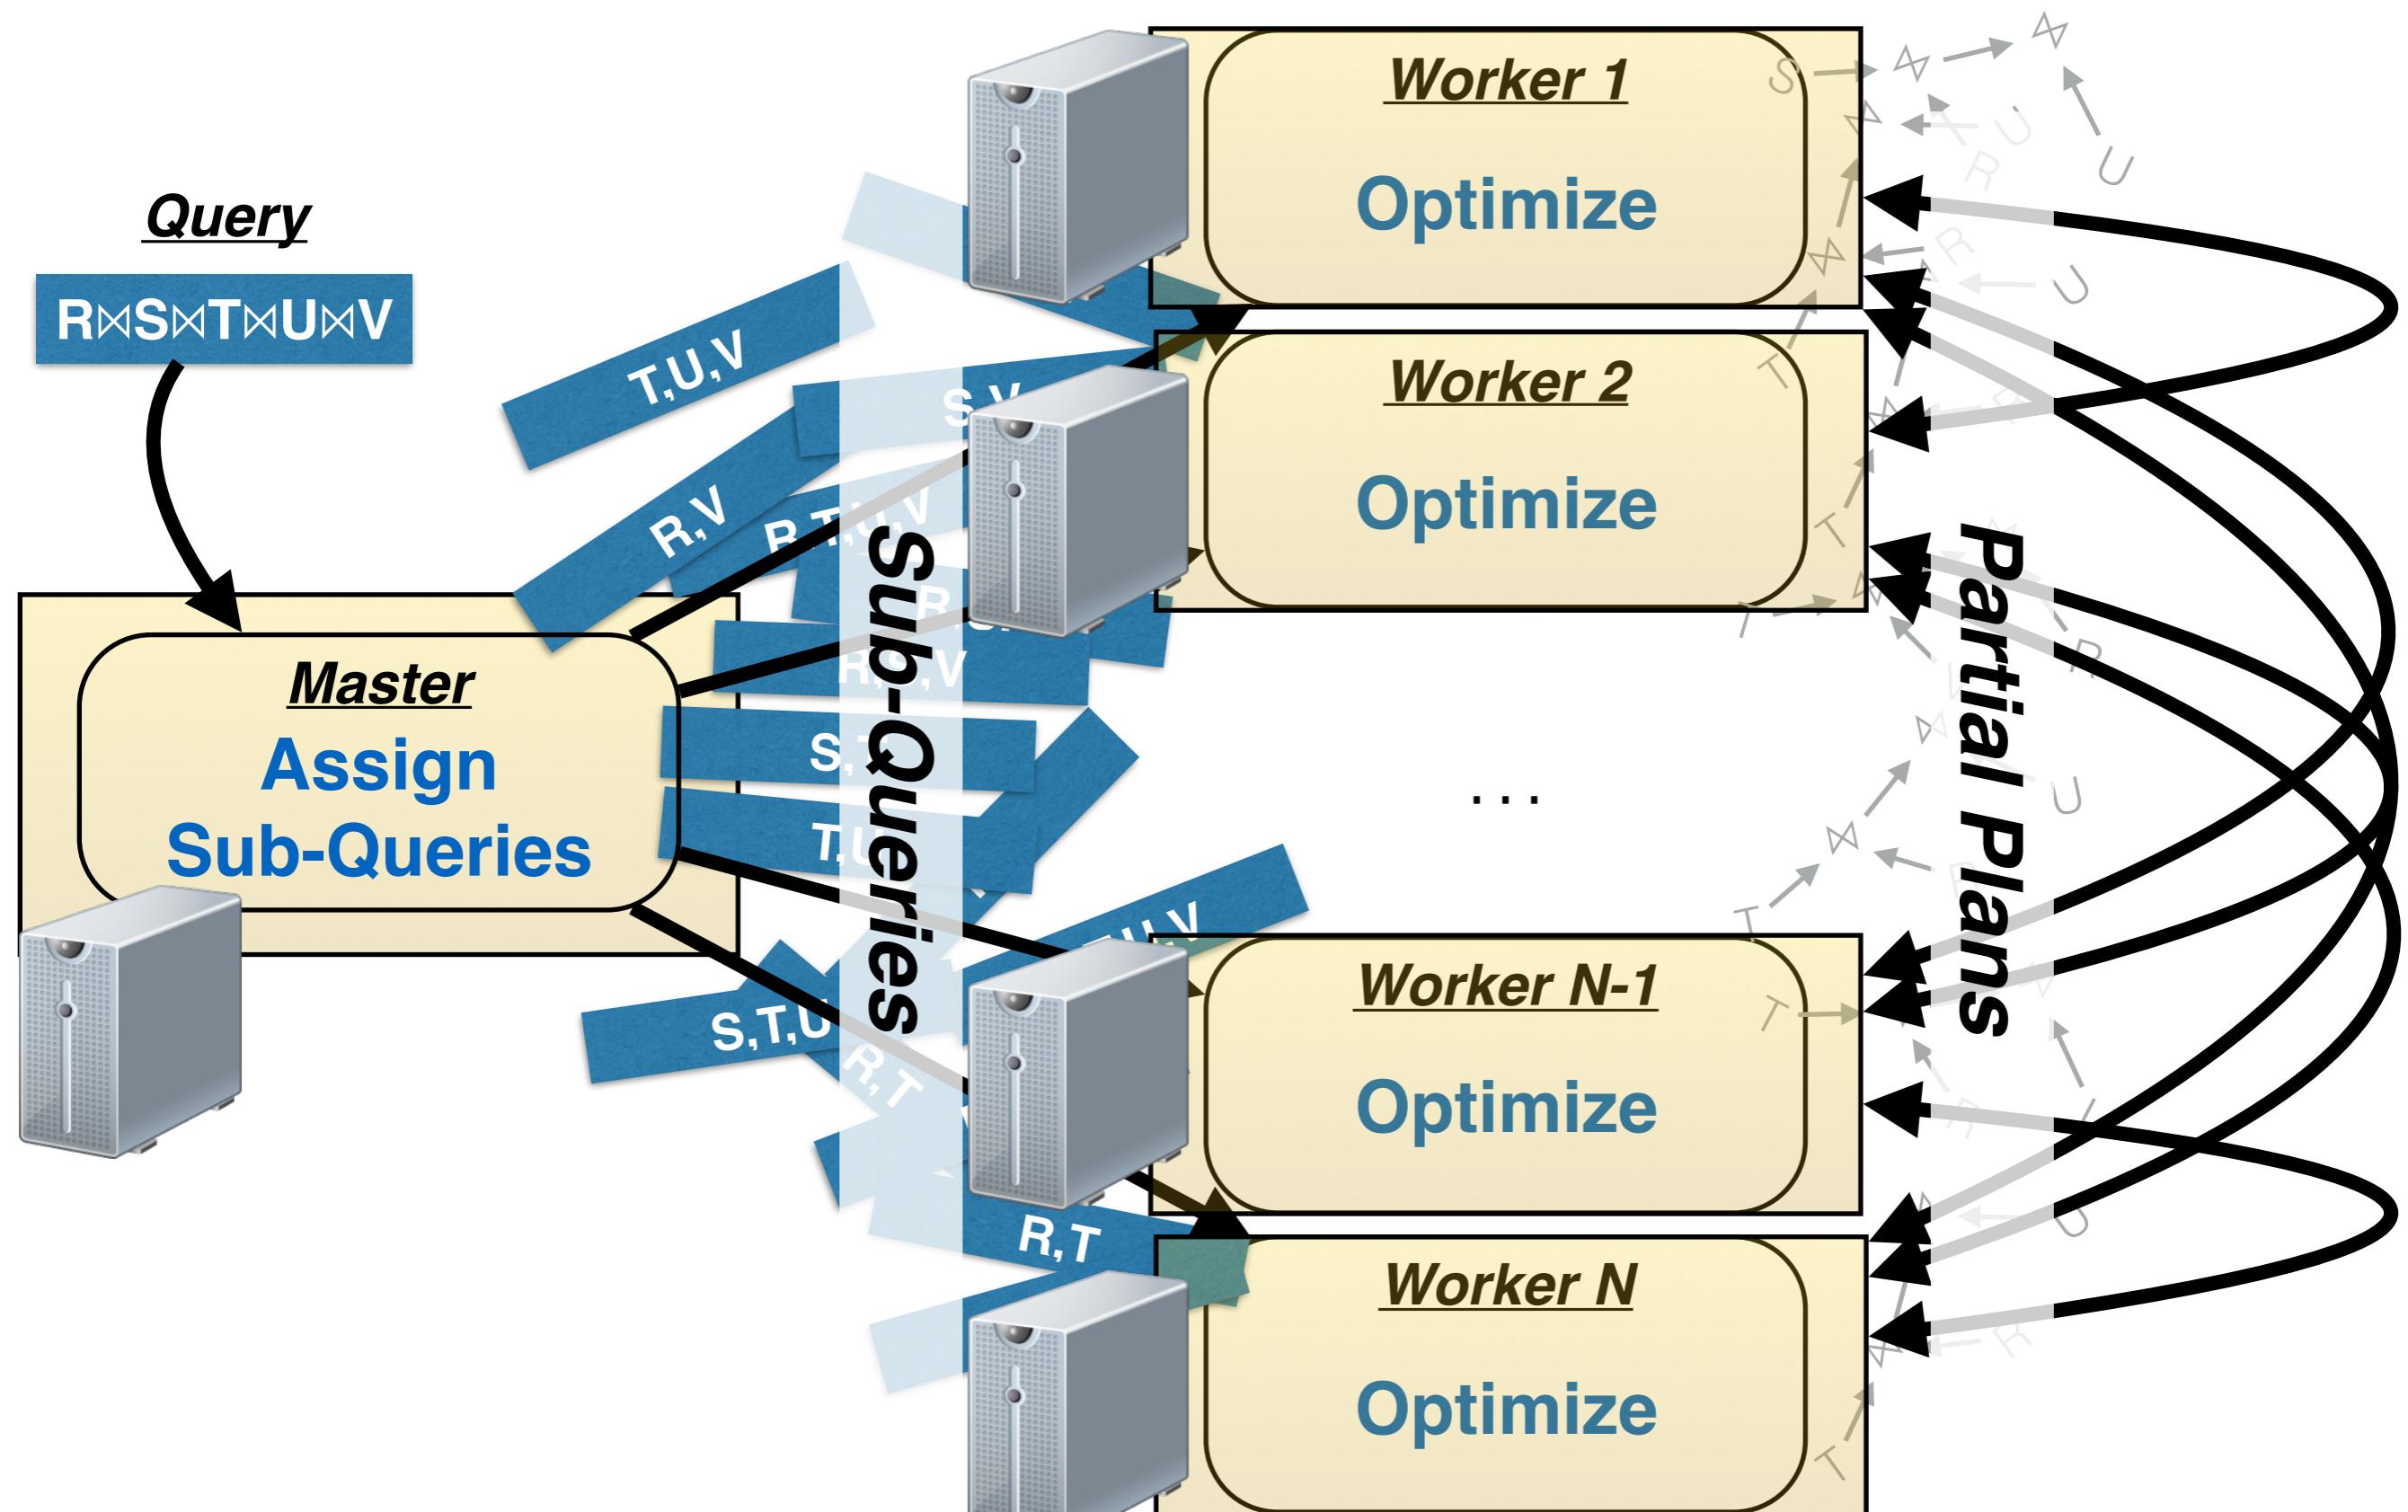


Non-Serial Polyadic Dynamic Programming

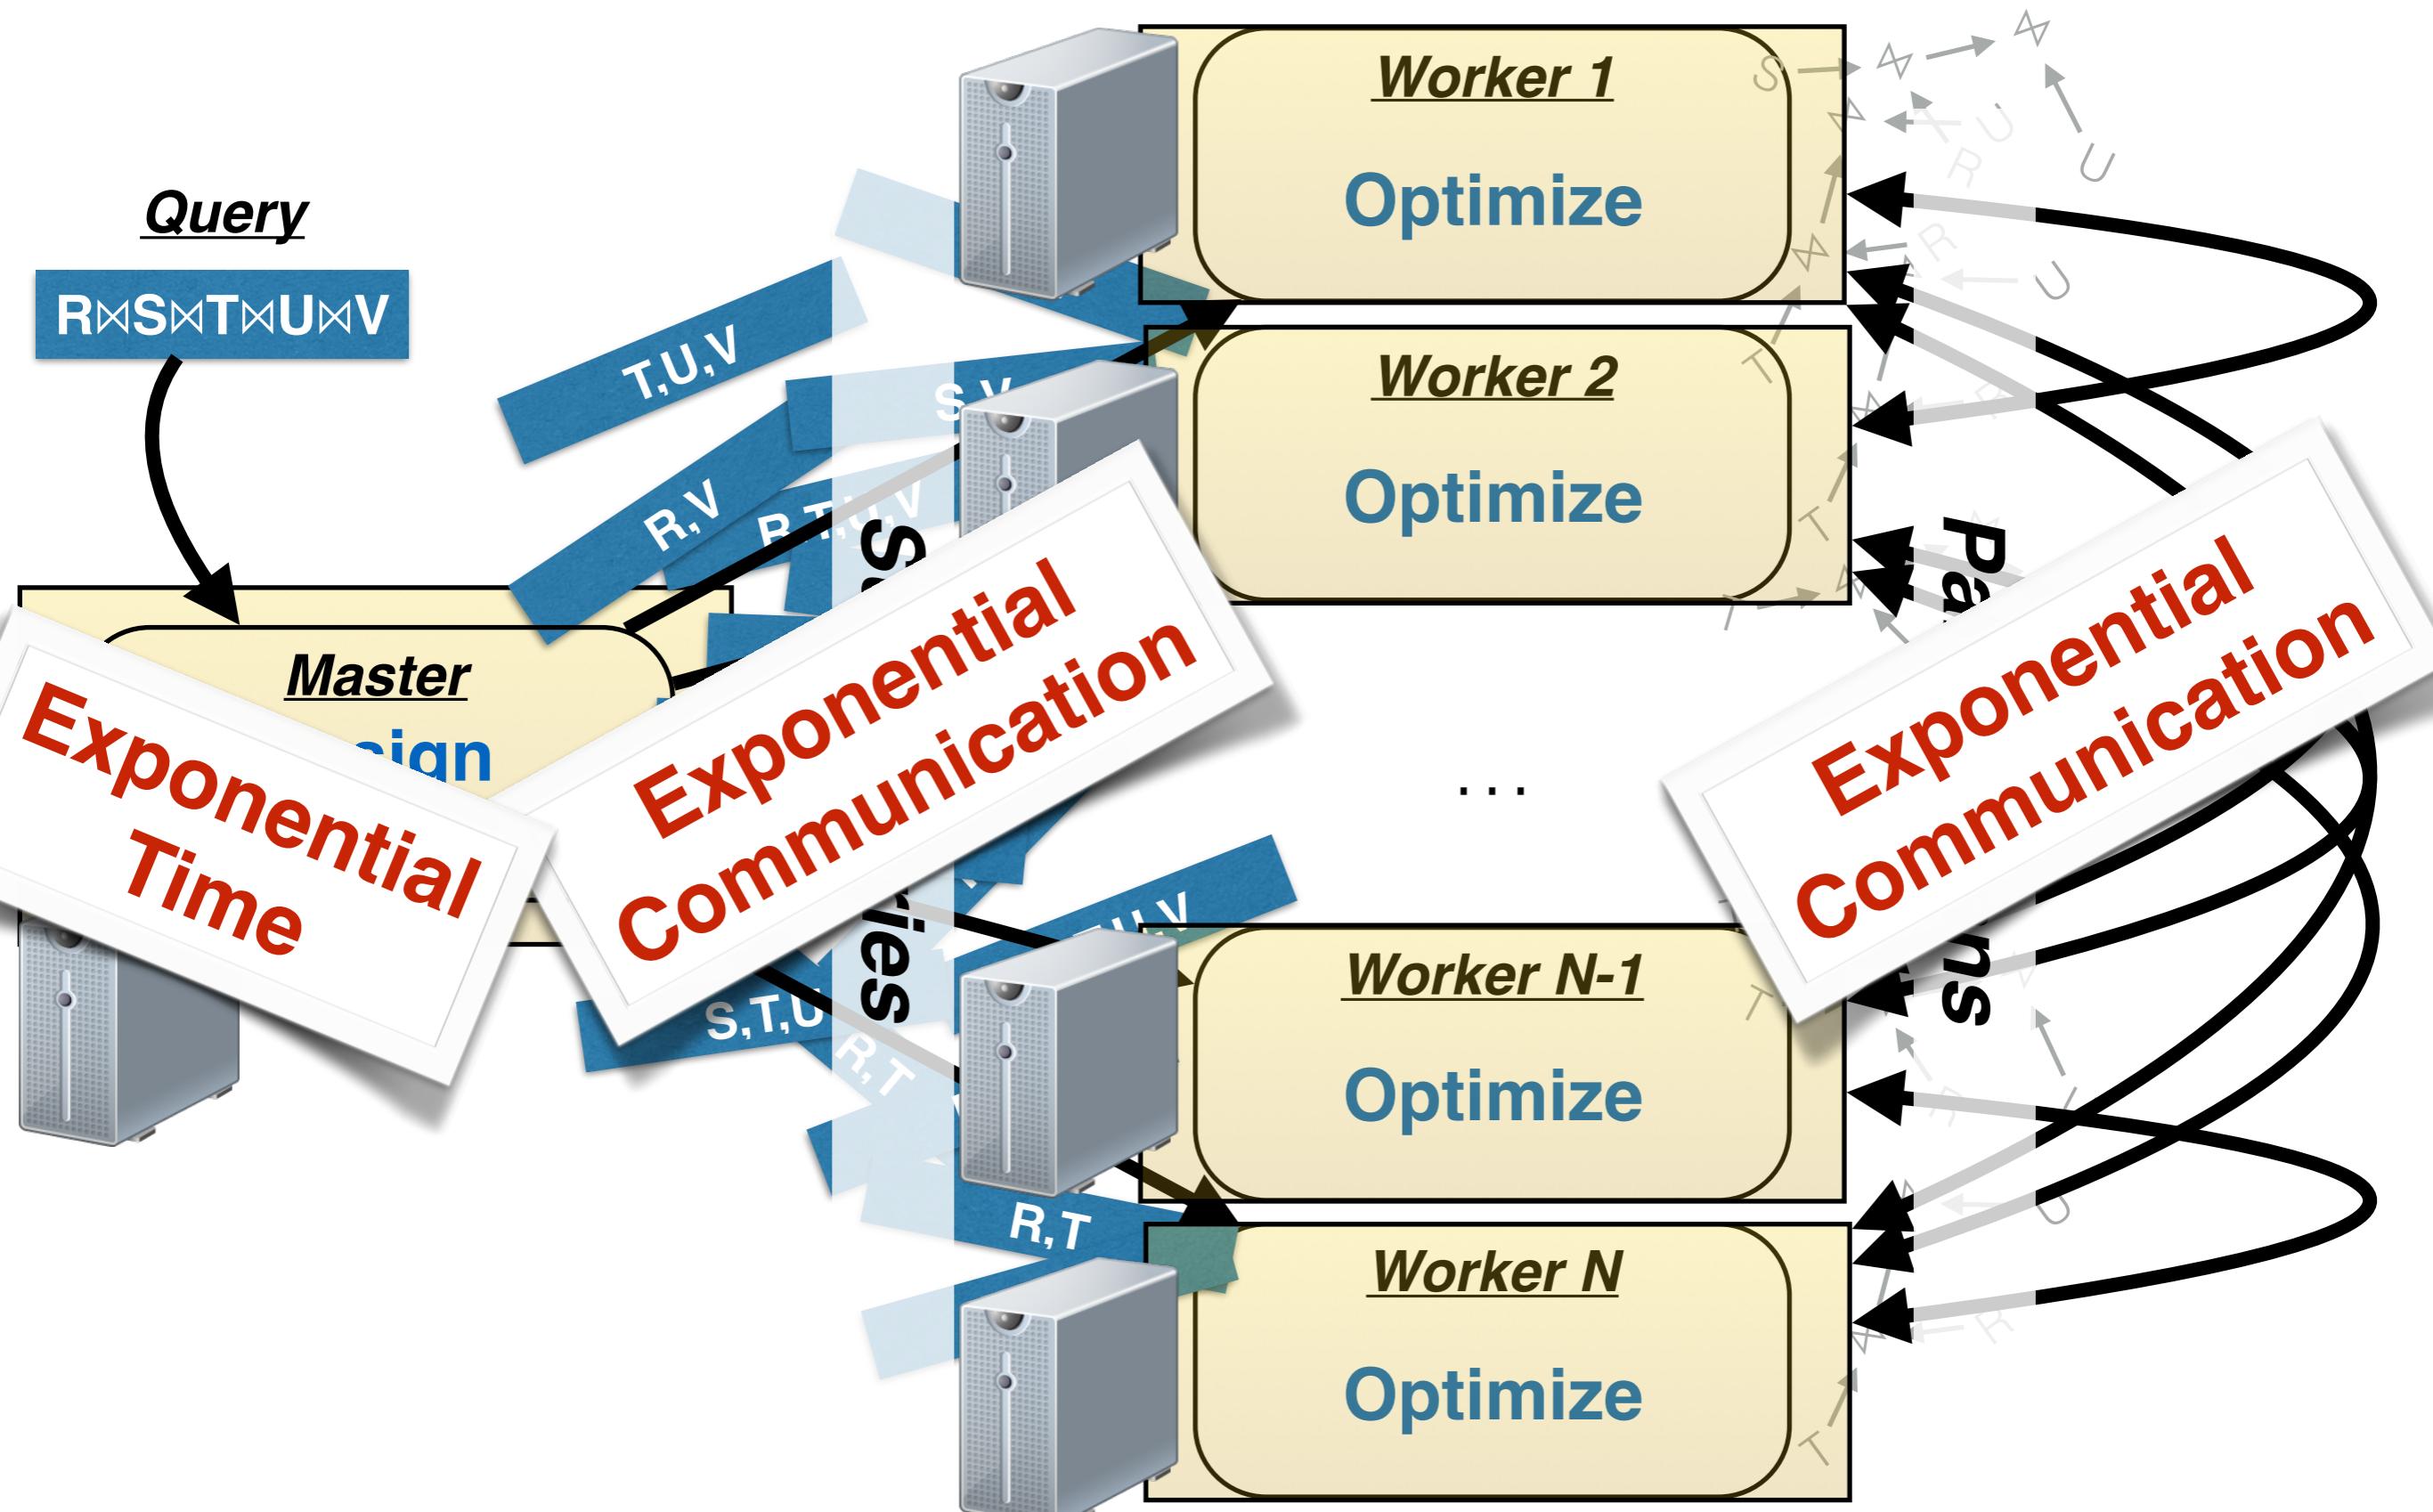
Prior Art



Prior Art



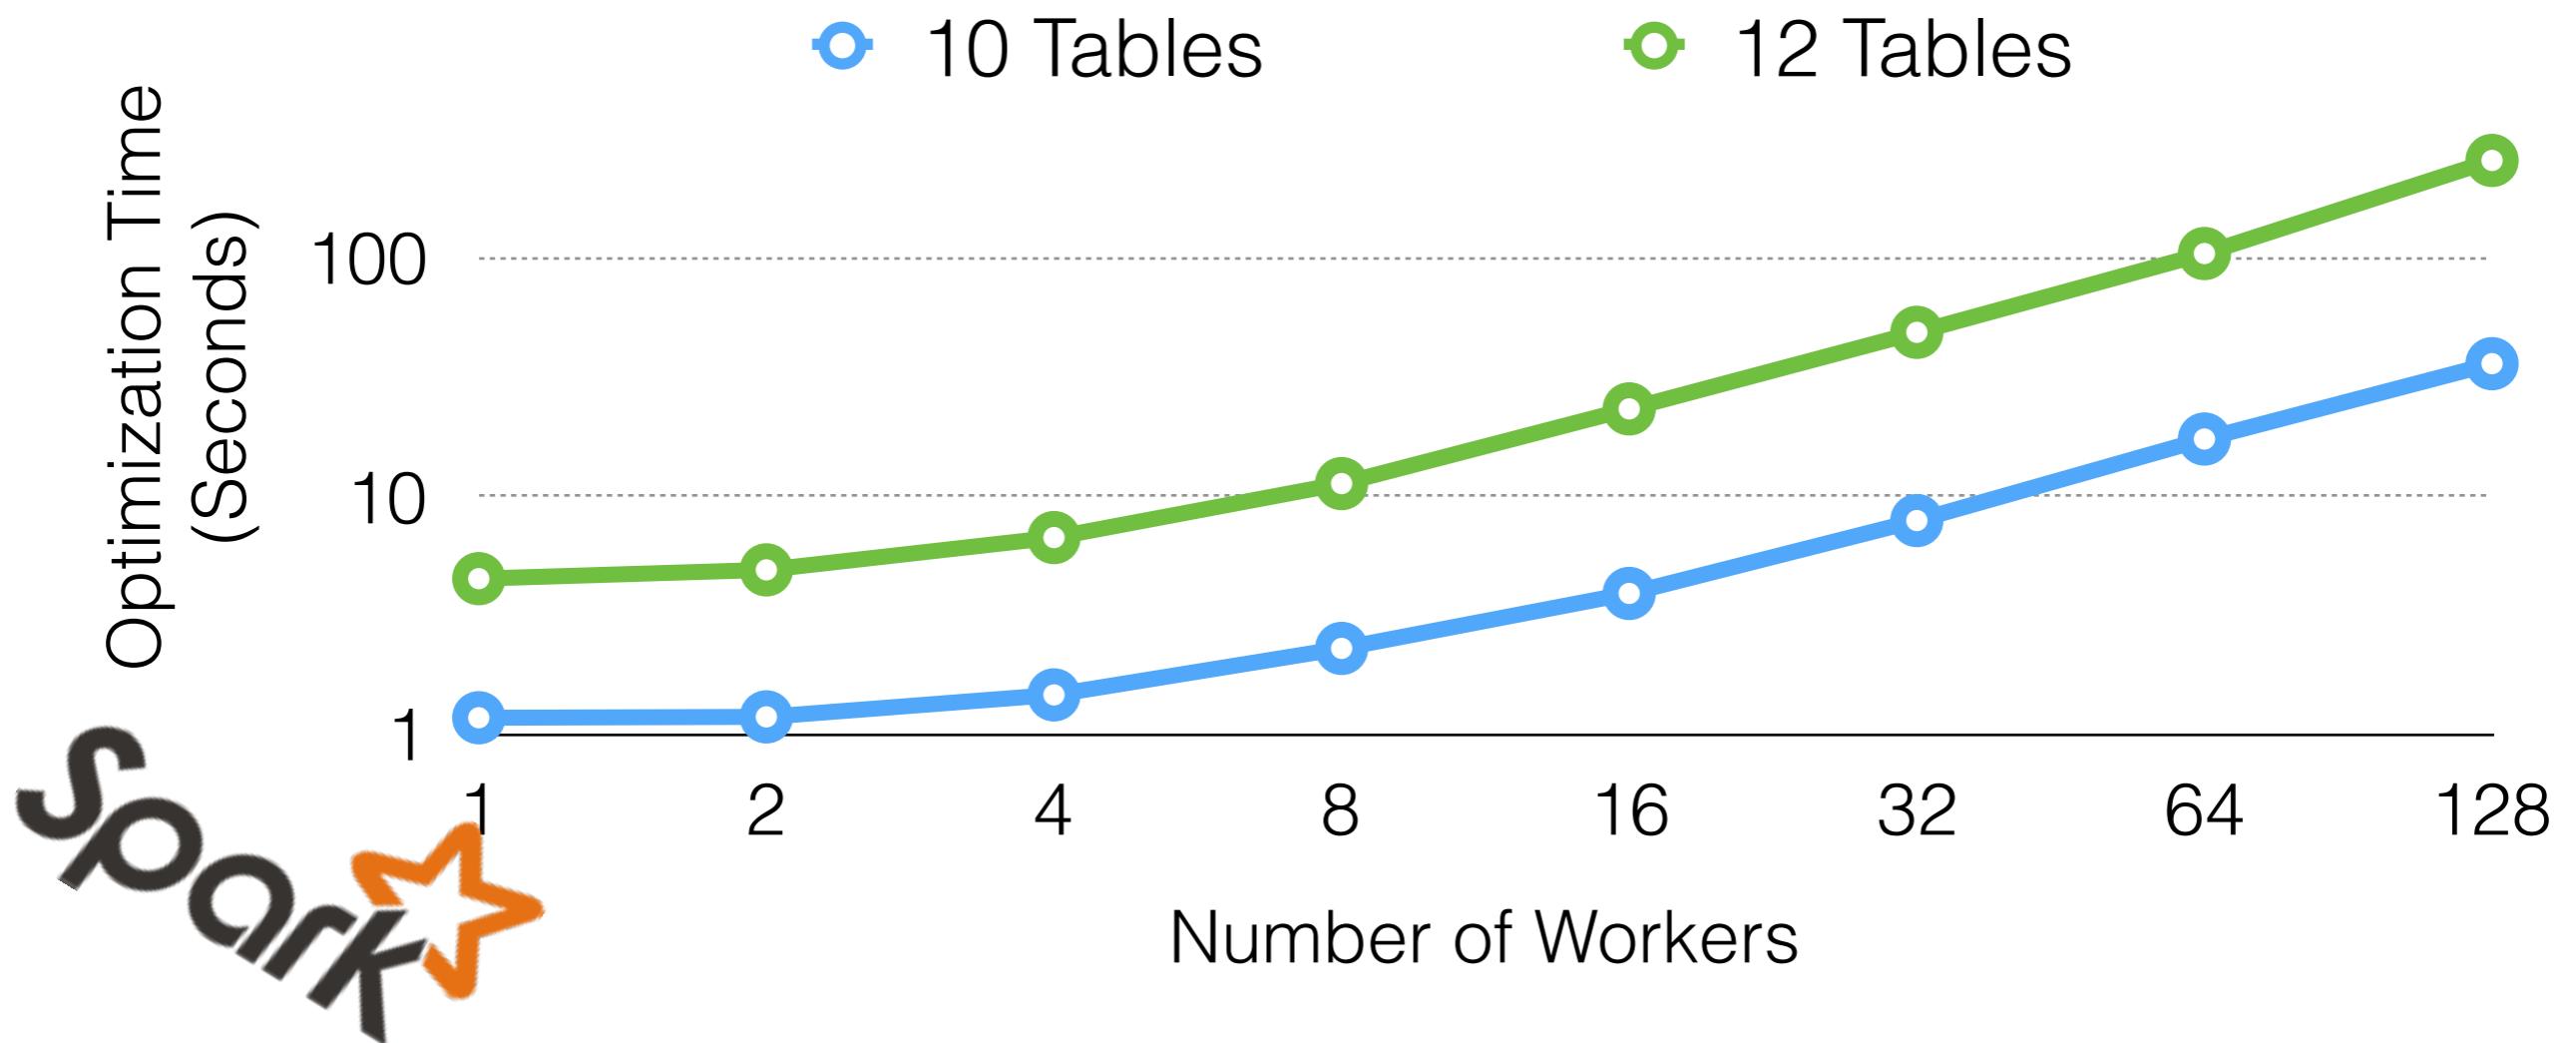
Prior Art



Experimental Results

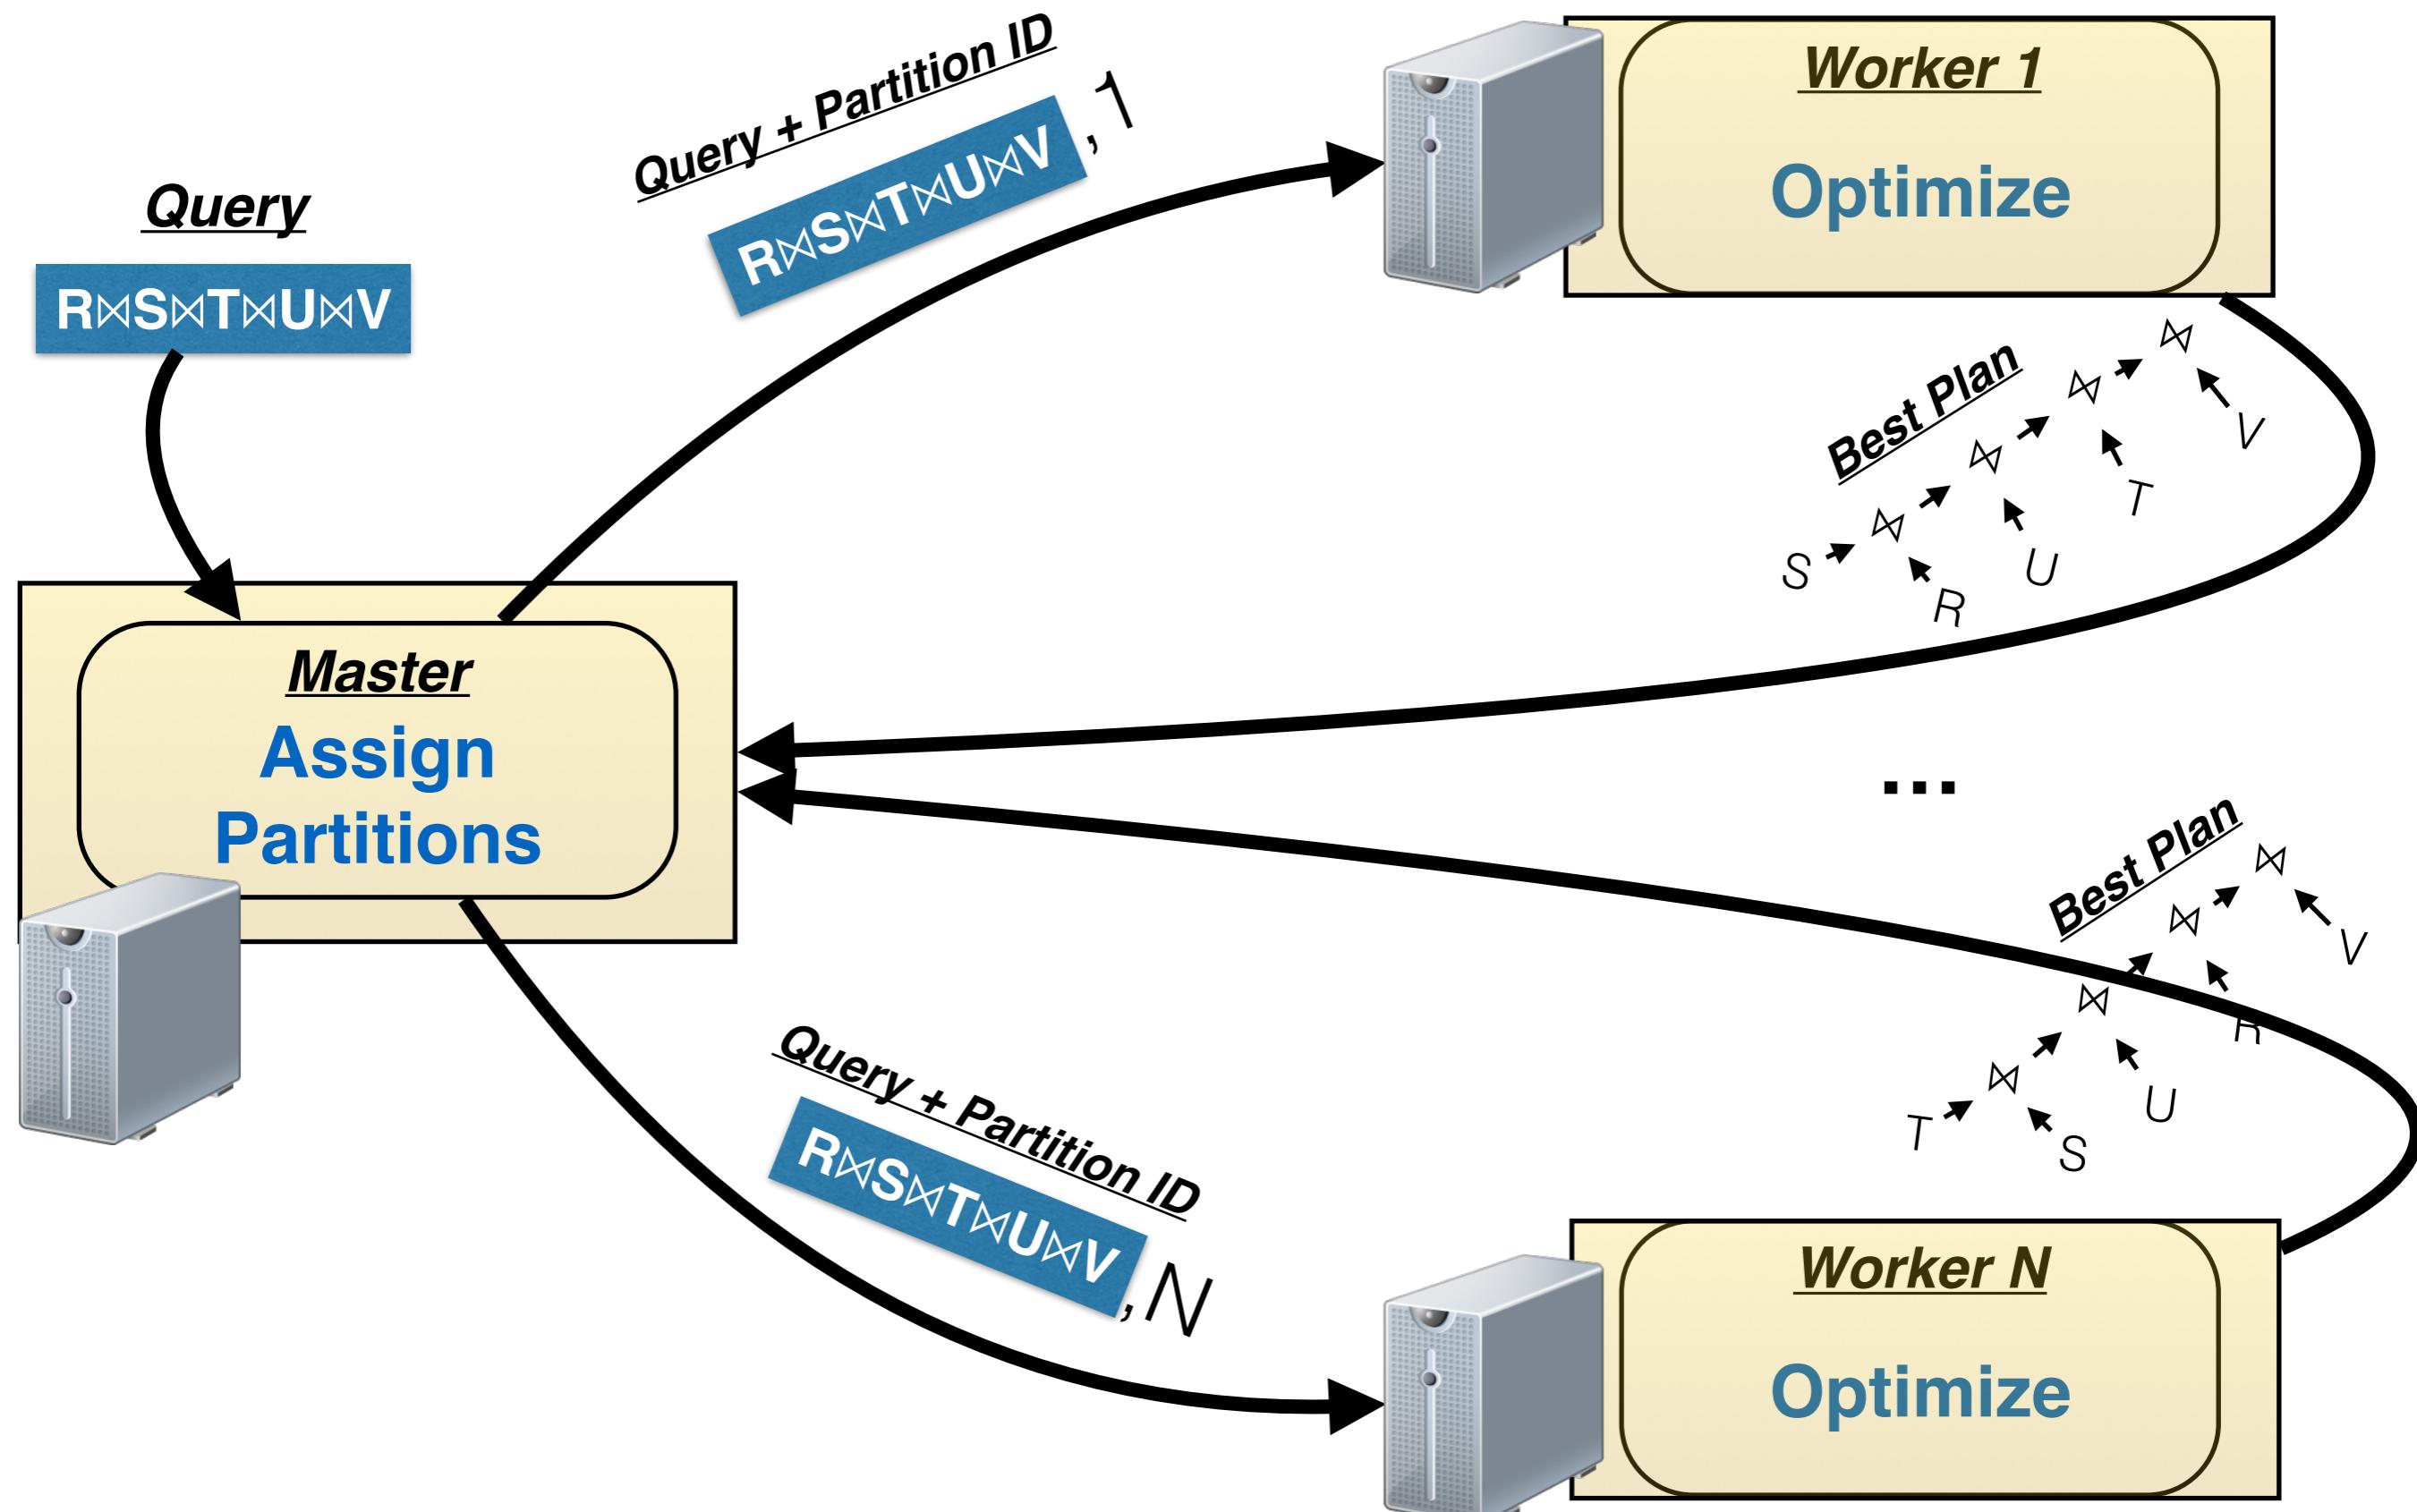
Experimental Setup

Randomly generated queries; two execution cost metrics; approximated optimization
Spark 1.5 on Yarn 2.7; 100 nodes with 2x Intel Xeon 2.6 GHz, 128 GB memory

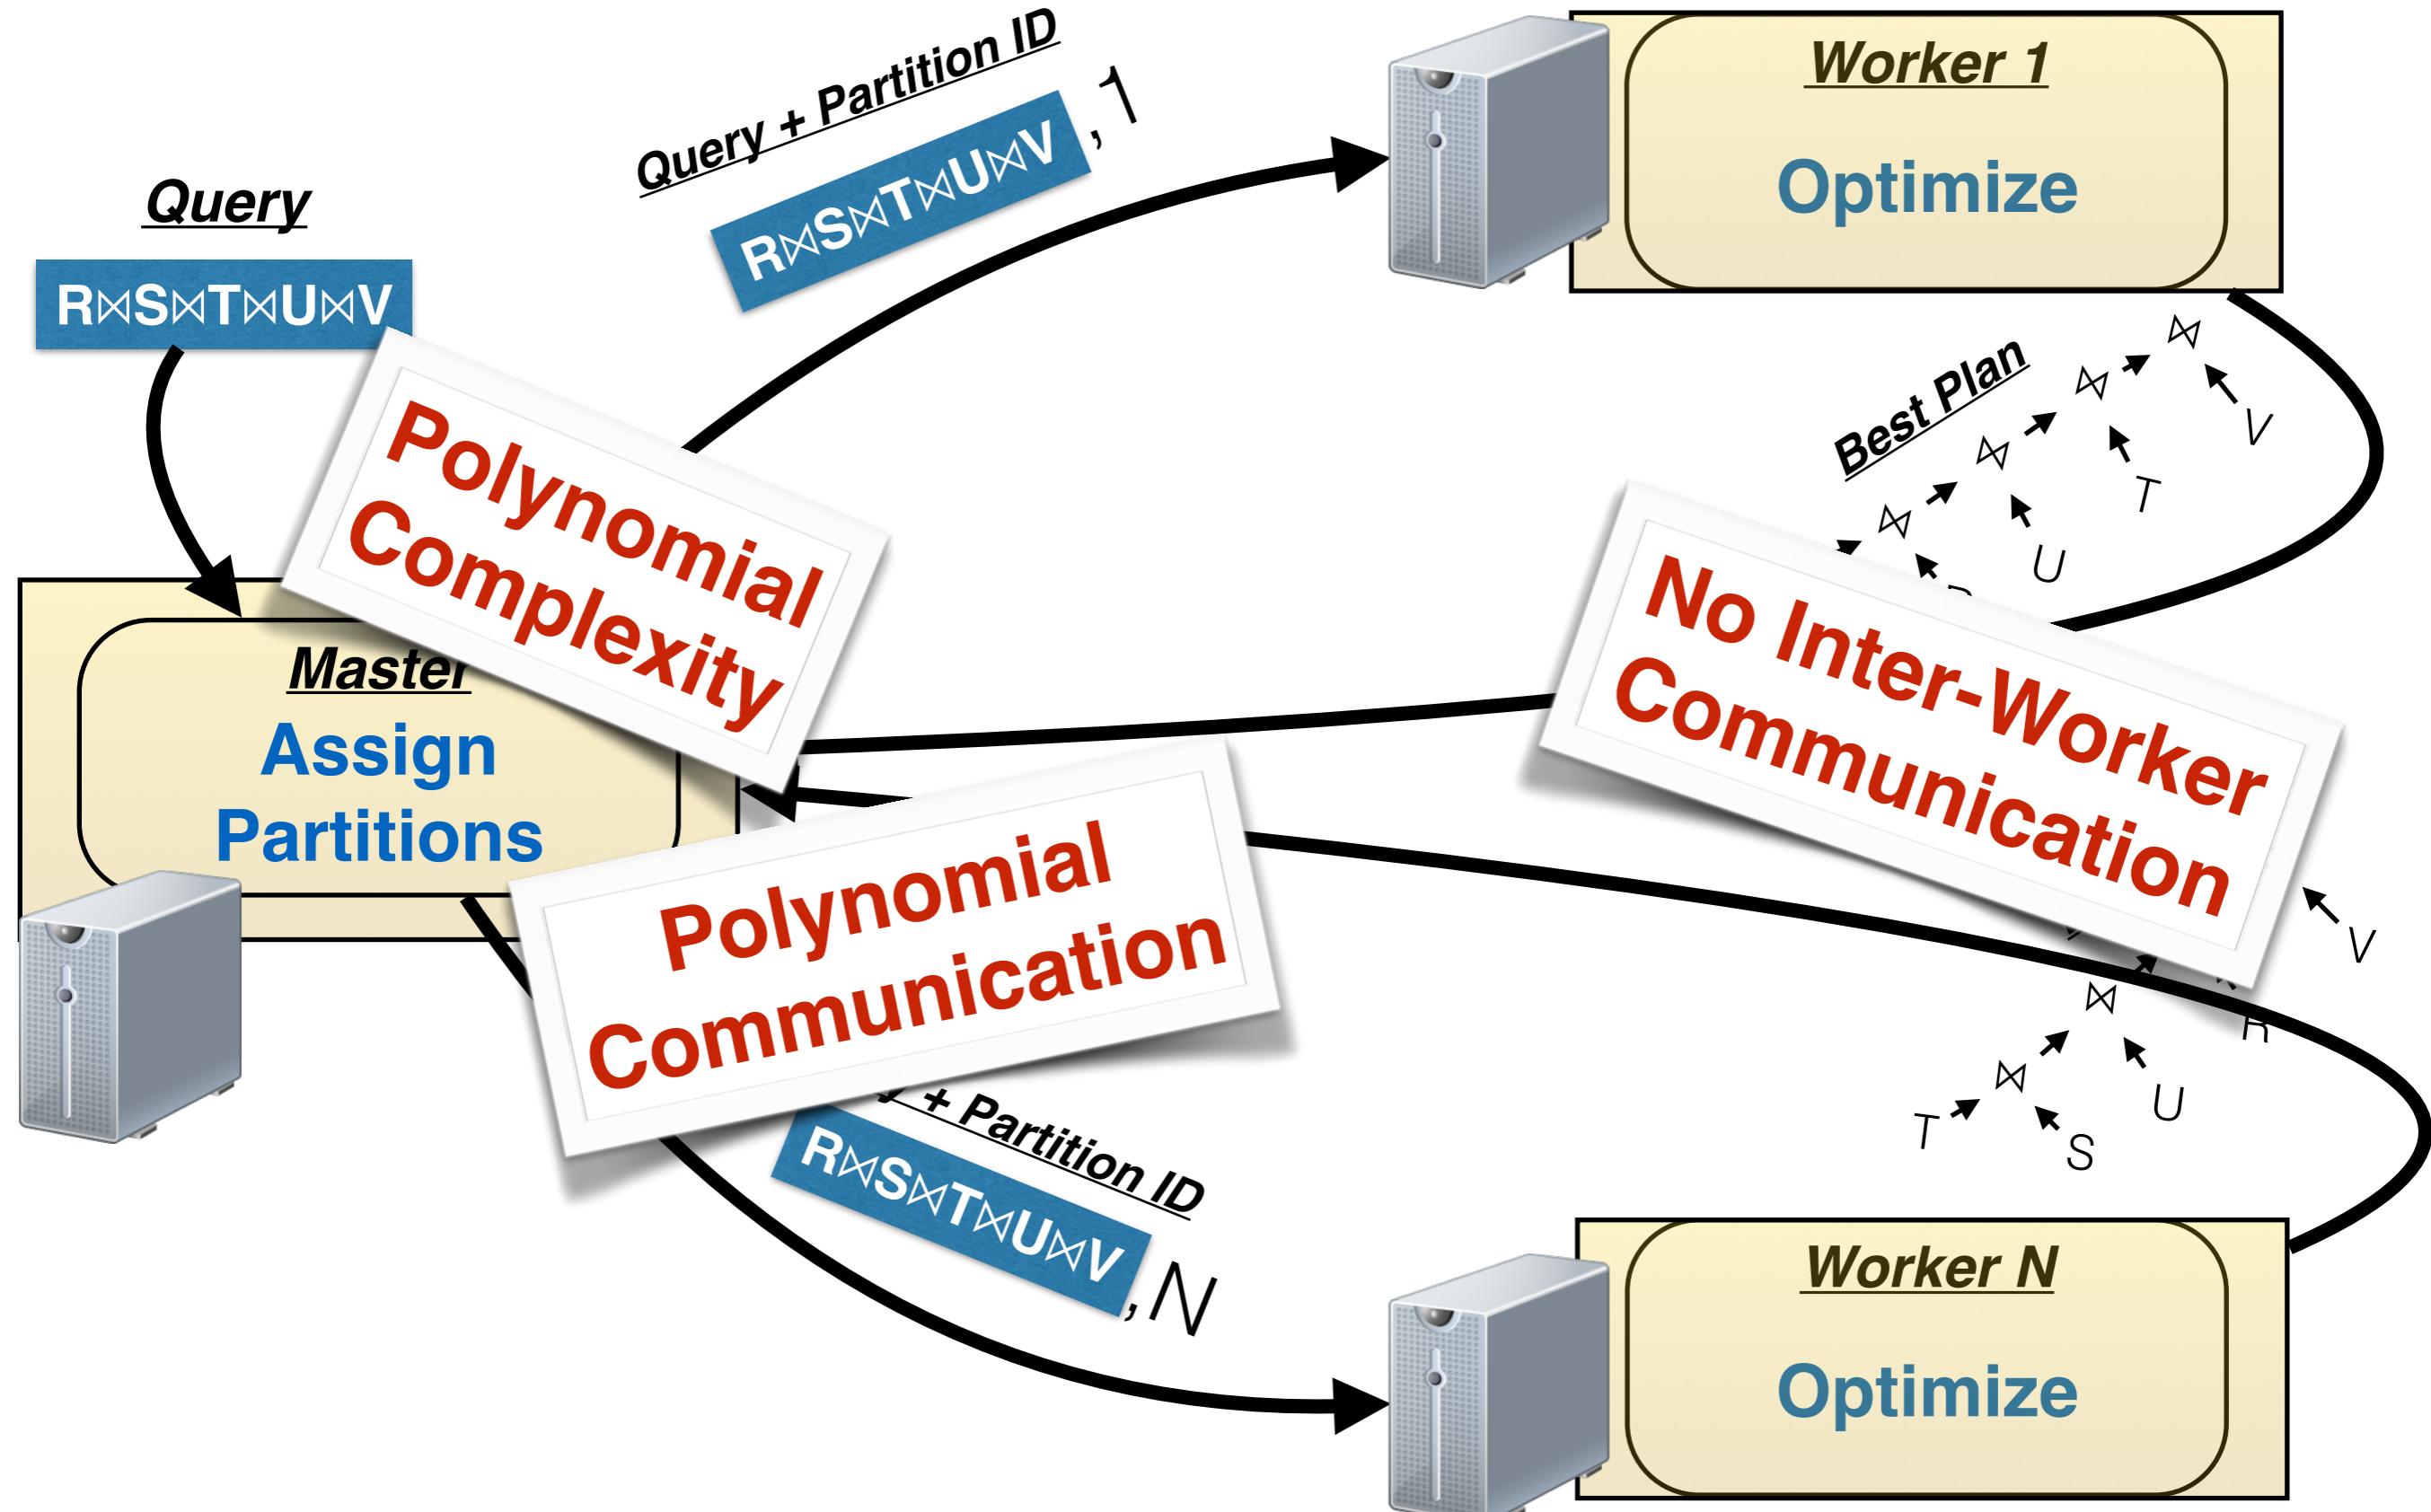


Scalability of previous parallel query optimization algorithms

New Approach

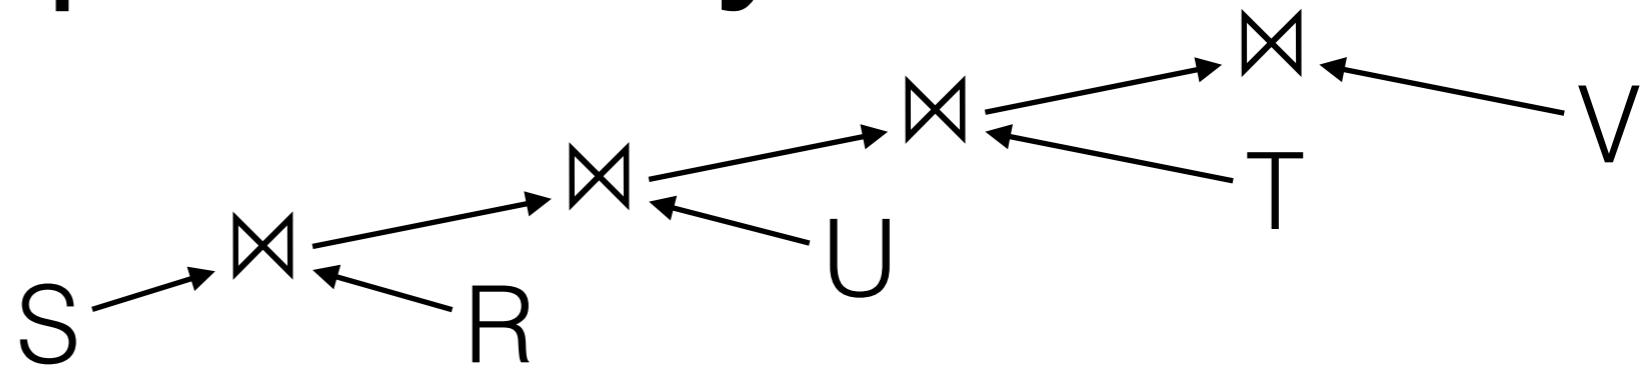


New Approach



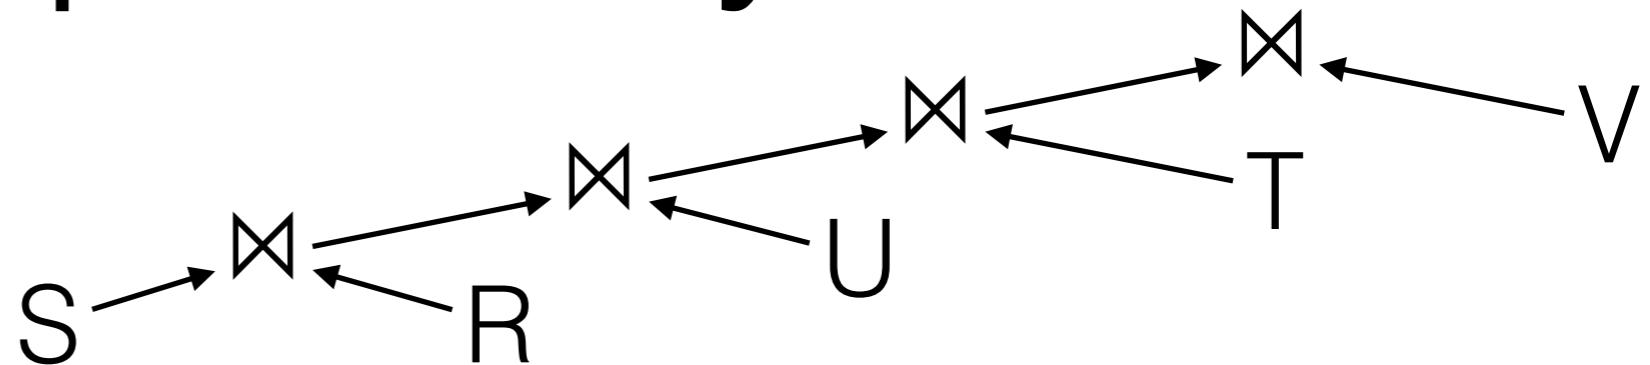
Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Unary Constraints

First: **R**

First: **S**

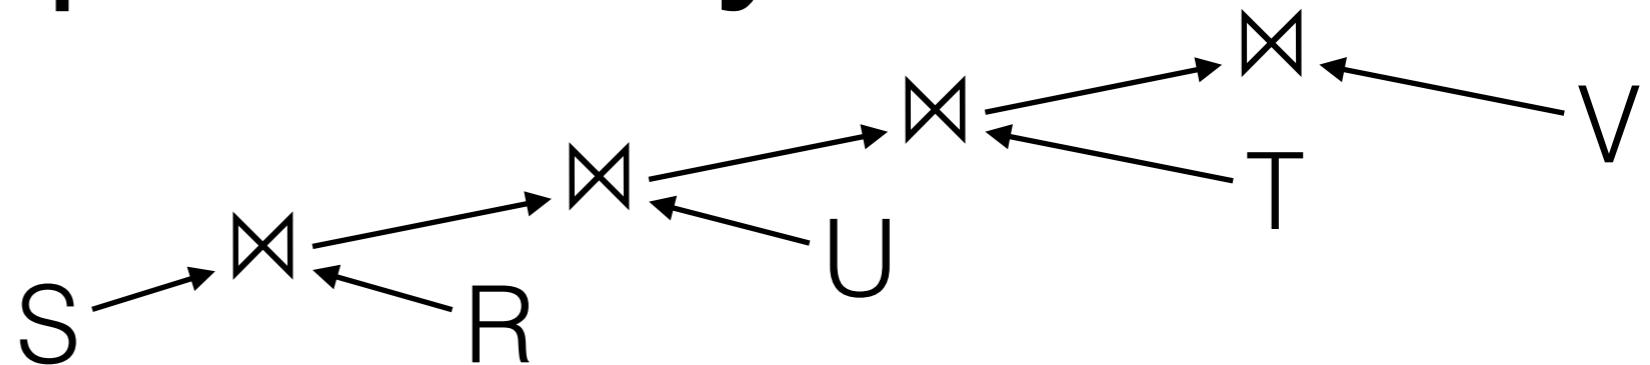
First: **T**

First: **U**

First: **V**

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Unary Constraints

First: **R**

First: **S**

First: **T**

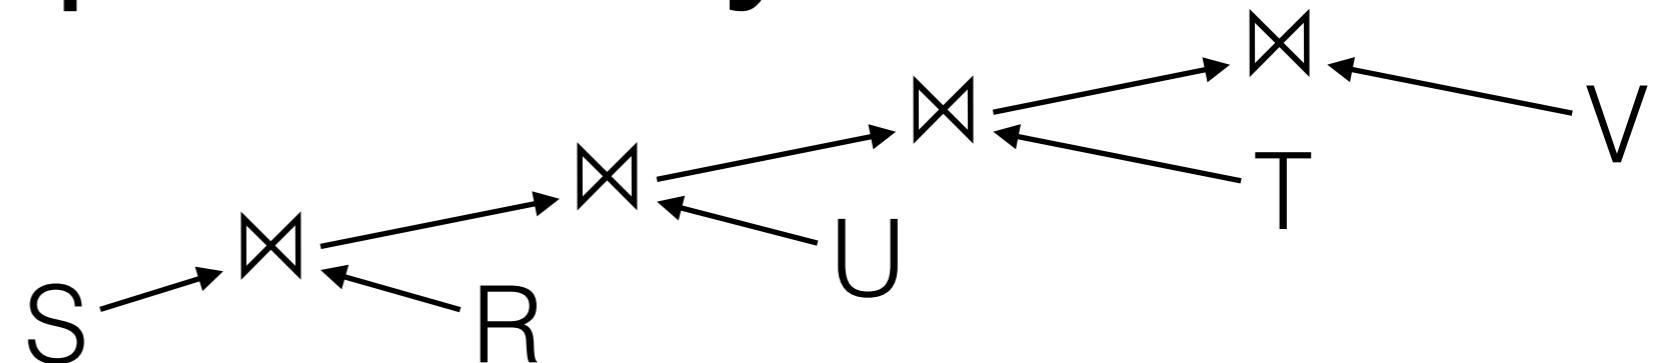
First: **U**

First: **V**



Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Unary Constraints

First: **R**

First: **S**

First: **T**

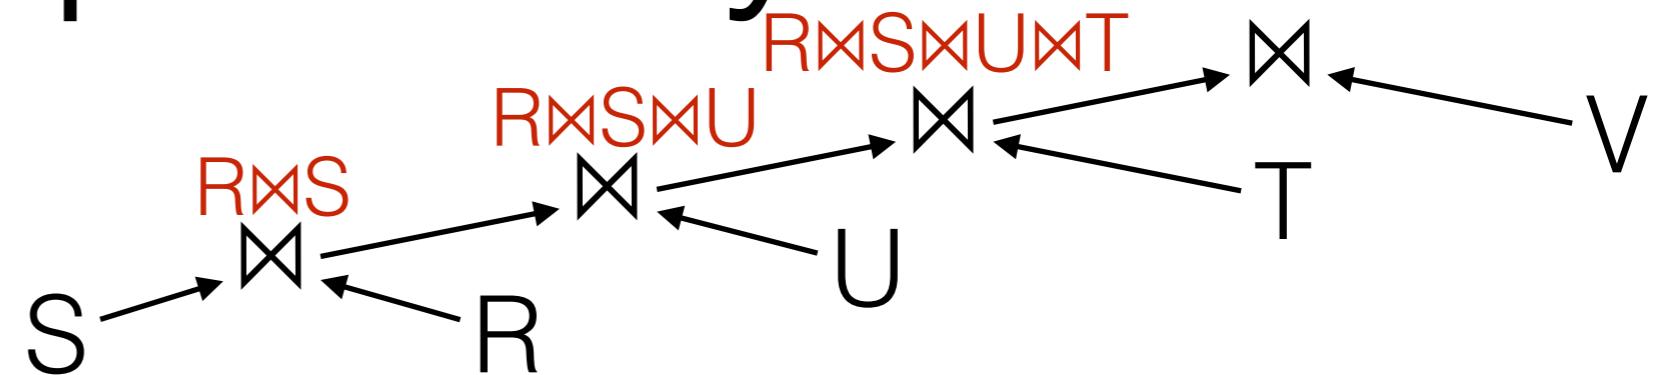
First: **U**

First: **V**

| | RS | RS | RS | RS |
|------------|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Unary Constraints

First: **R**

First: **S**

First: **T**

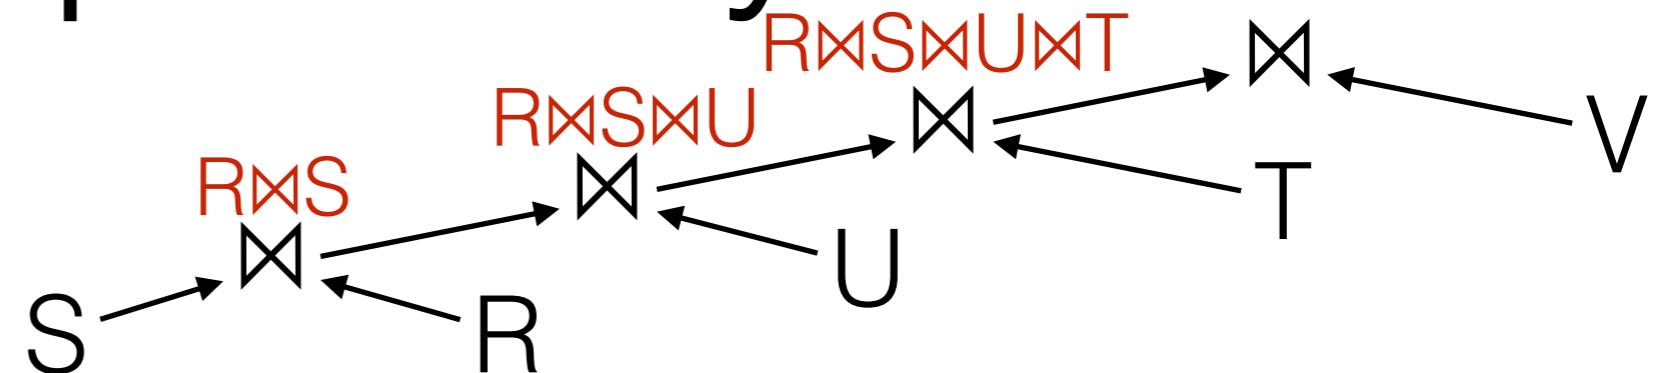
First: **U**

First: **V**

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Unary Constraints

First: **R**

First: **S**

First: **T**

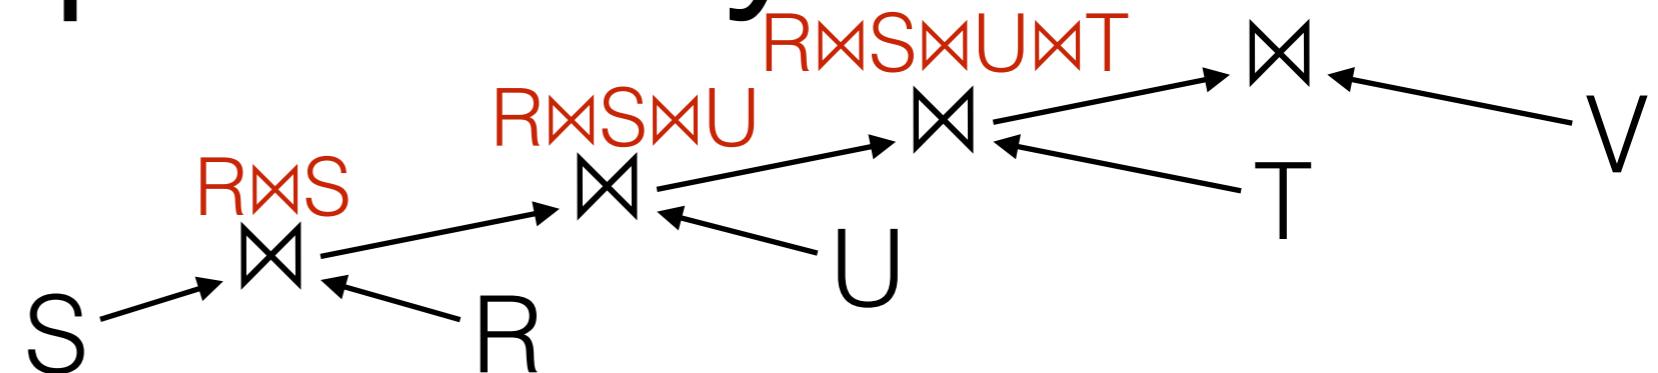
First: **U**

First: **V**

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

Left-Deep Plan ~
Table Order



Partitioning via Unary Constraints

Parallelism ~ Nr. Tables

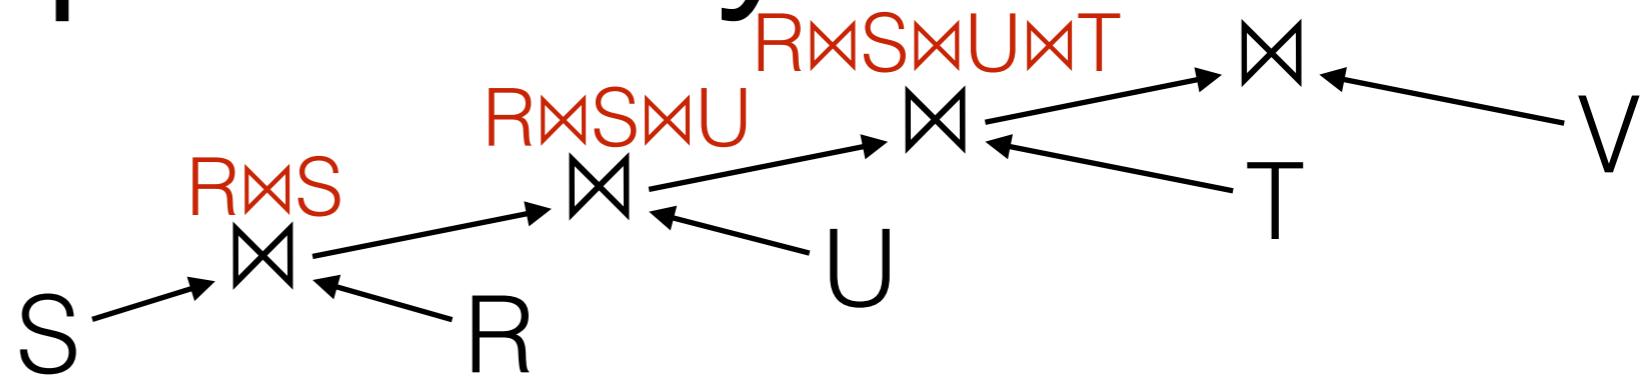
Time / 2

Sub-Query

| TUV | T,U | R,U | S,T,U | R,S,T,U |
|-----|-----|-----|-------|---------|
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

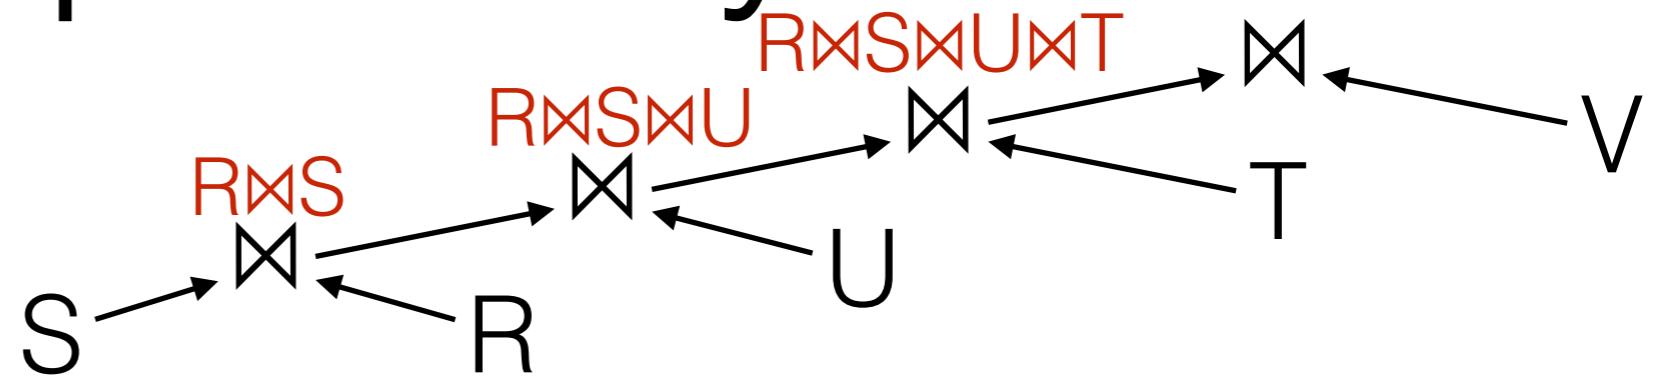
Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



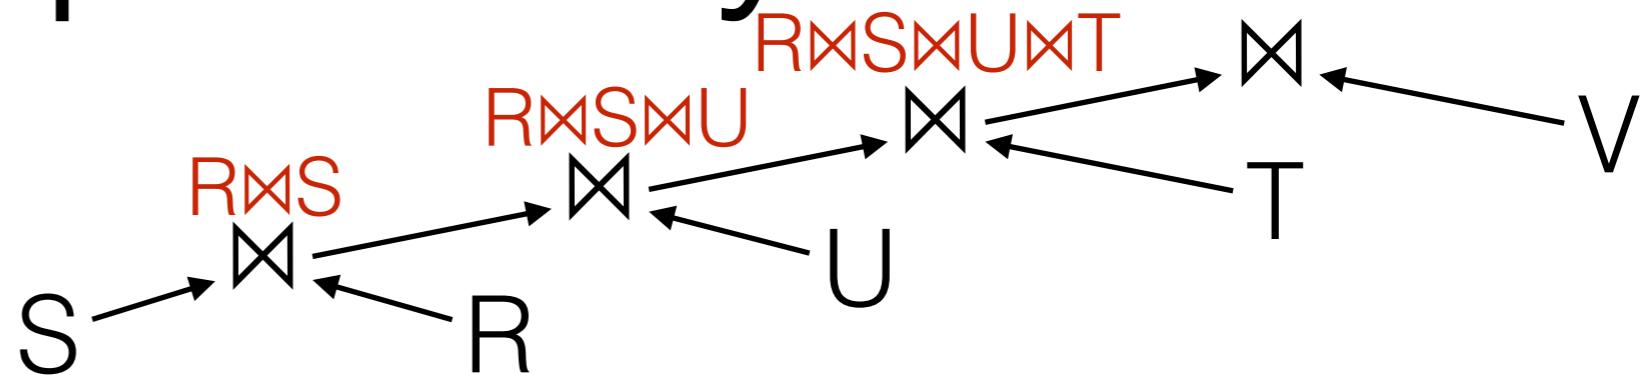
Partitioning via Binary Constraints

R Before **S**

S Before **R**

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Binary Constraints

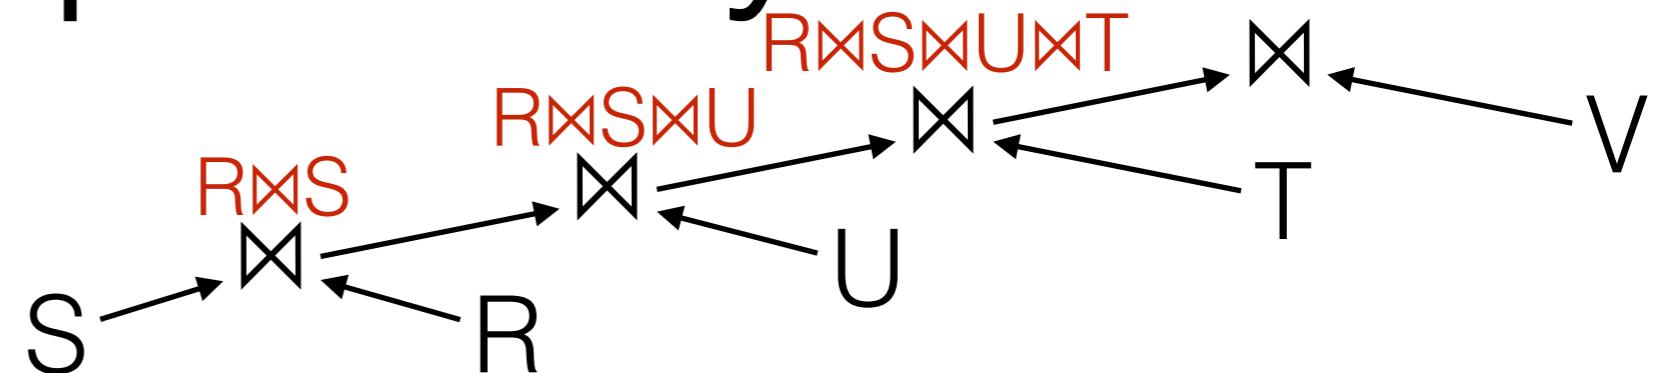
R Before S

S Before R

Left-Deep Query Plans

Plan Space

*Left-Deep Plan ~
Table Order*



Partitioning

R Before S

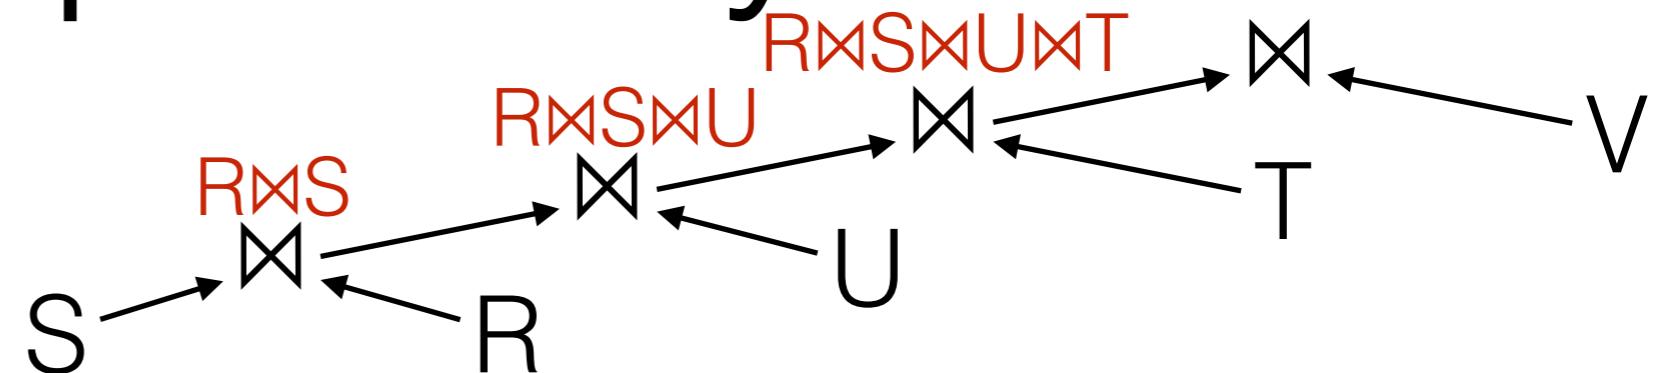
S Before R

Sub-Queries

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Binary Constraints

R Before S

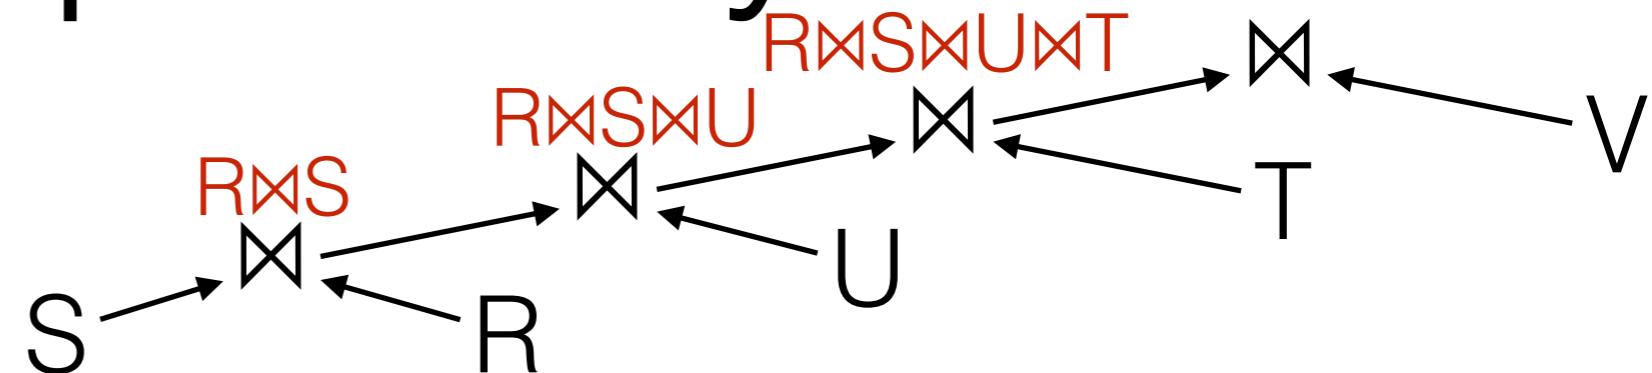
S Before R

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

Plan Space

*Left-Deep Plan ~
Table Order*



Partitioning

Partitioning via Binary Constraints

R Before S
T Before U

U Before T

S Before R

T Before U

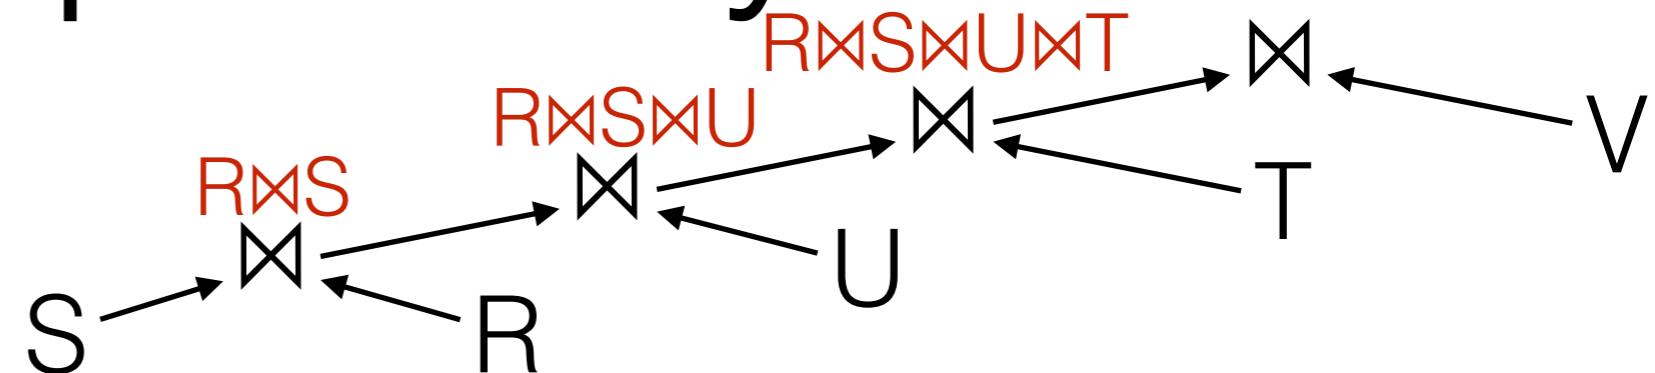
U Before T

Sub-Queries

| | RS | RS | RS | RS |
|------------|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Binary Constraints

R Before S
T Before U

U Before T

S Before R
U Before T

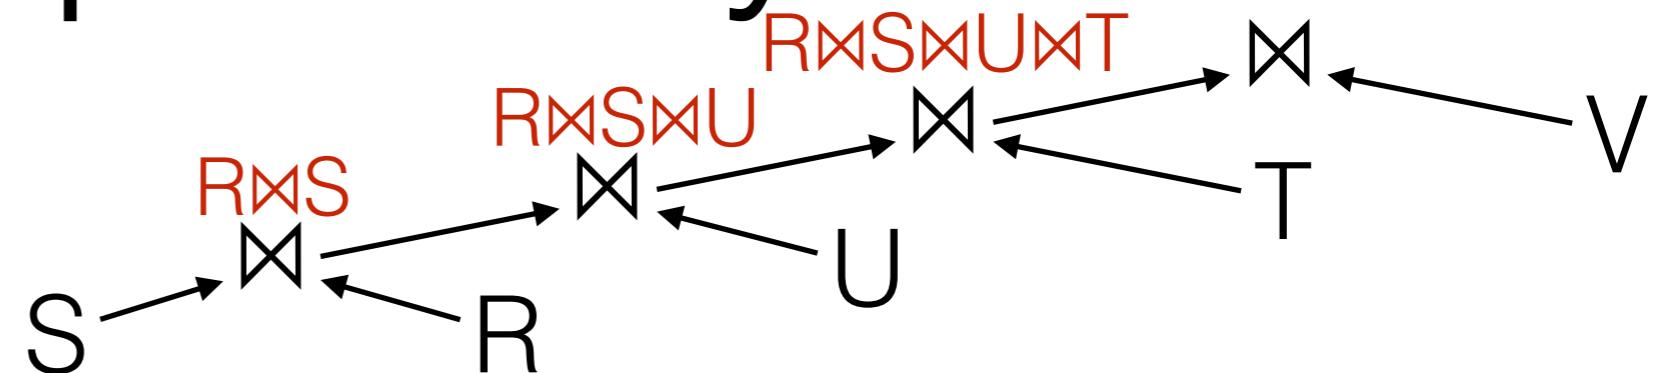
Partitioning

Sub-Queries

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

*Left-Deep Plan ~
Table Order*



Partitioning via Binary Constraints

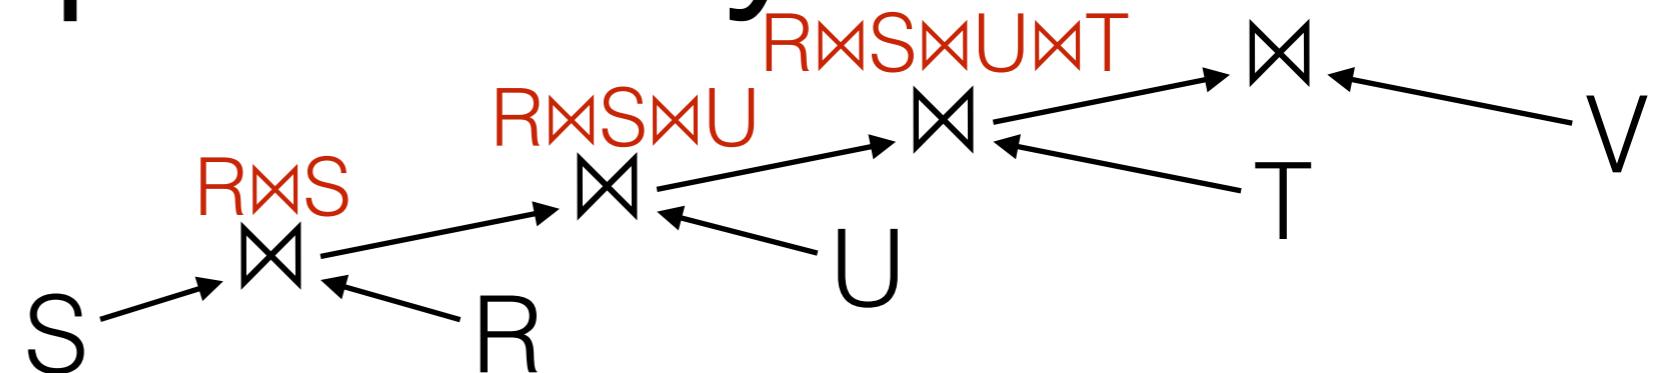
R Before S
T Before U U Before T

S Before R
T Before U U Before T

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Left-Deep Query Plans

Left-Deep Plan ~
Table Order



Partitioning via Binary Constraints

Parallelism $*2^{nrPartitionLevels}$

Time $*(3/4)^{nrPartitionLevels}$

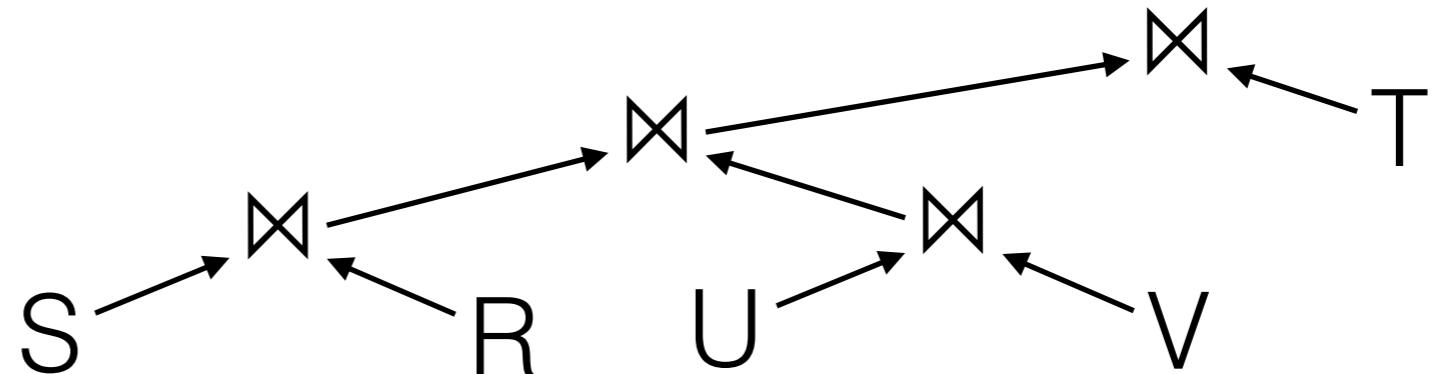
Sub-Query

| | | | | | | | | |
|-----|-----|-----|-------|-------|-------|-------|---------|-----------|
| TUV | S,V | R,V | S,U | R,S,V | R,S,U | R,S,T | R,S,T,U | R,S,T,U,V |
| TUV | T,U | R,U | S,T,U | | | | | |
| TUV | | R,V | S,V | | | | | |
| TUV | | R,U | S,U | | | | | |
| TUV | | R,T | S,T | | | | | |
| TUV | | | | R,S | | | | |

Bushy Query Plans

Plan Space

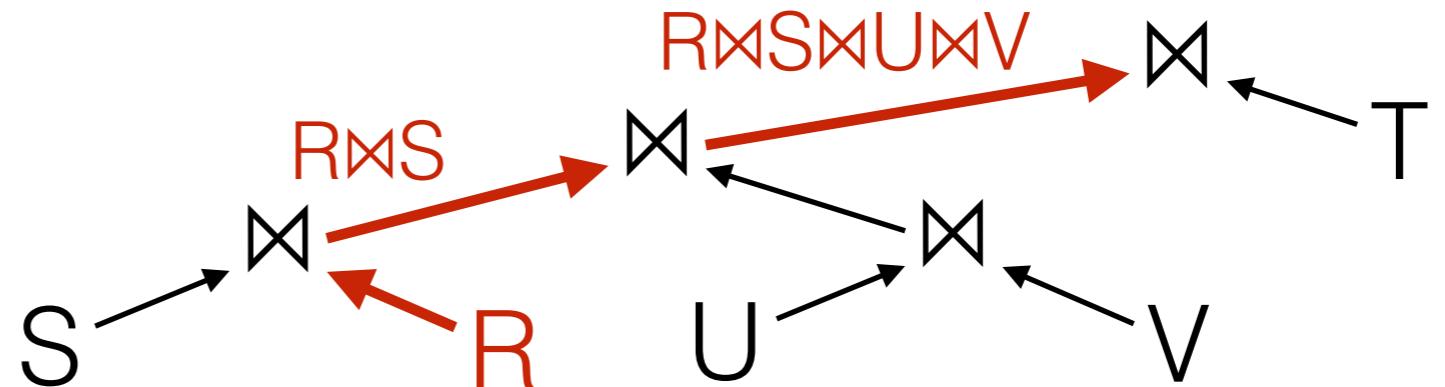
*Bushy Plan ~
Binary Tree*



Bushy Query Plans

Plan Space

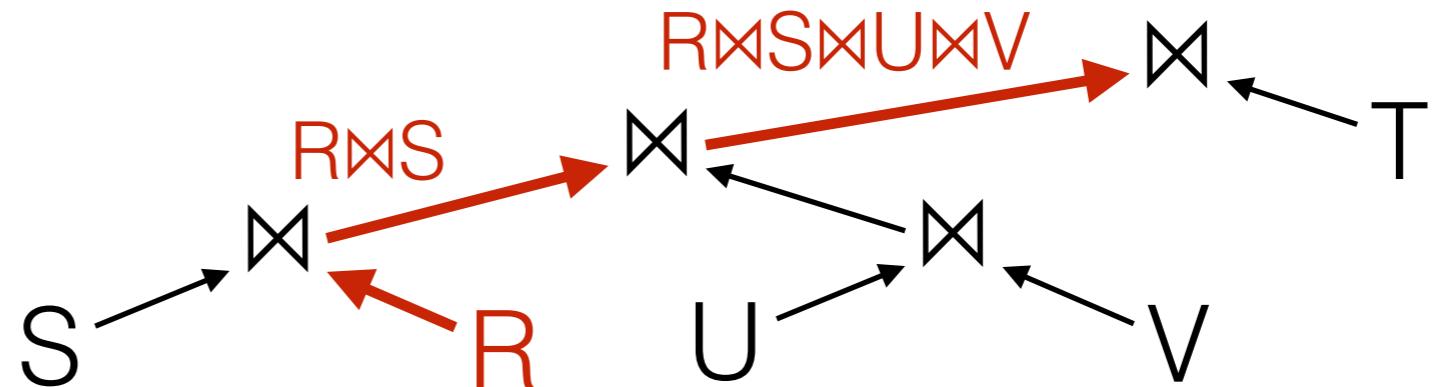
*Bushy Plan ~
Binary Tree*



Bushy Query Plans

Plan Space

*Bushy Plan ~
Binary Tree*



Partitioning via Ternary Constraints

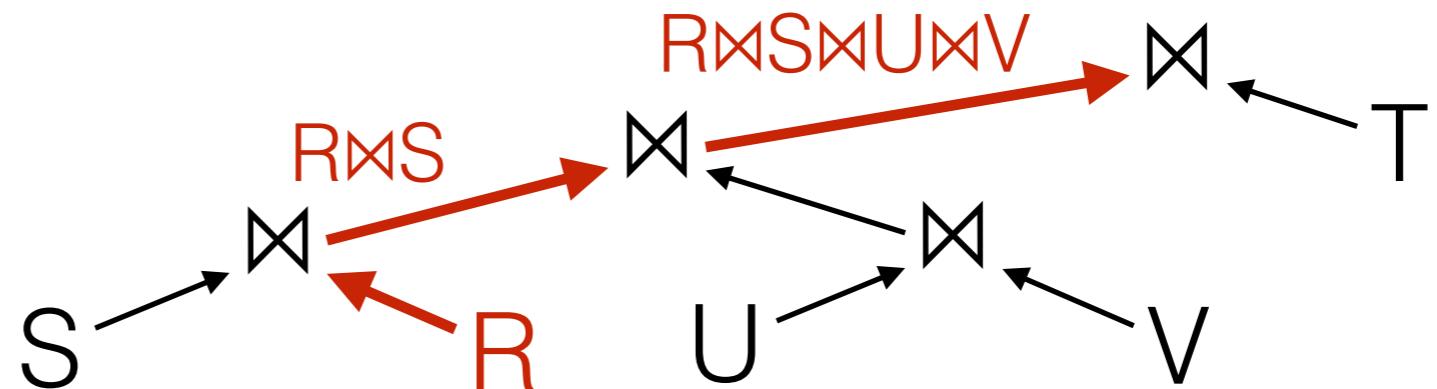
Path **R**→**Root**: **T** after **U**

Path **R**→**Root**: **T** not after **U**

Partitioning

Bushy Query Plans

Bushy Plan ~ Binary Tree



Partitioning via Ternary Constraints

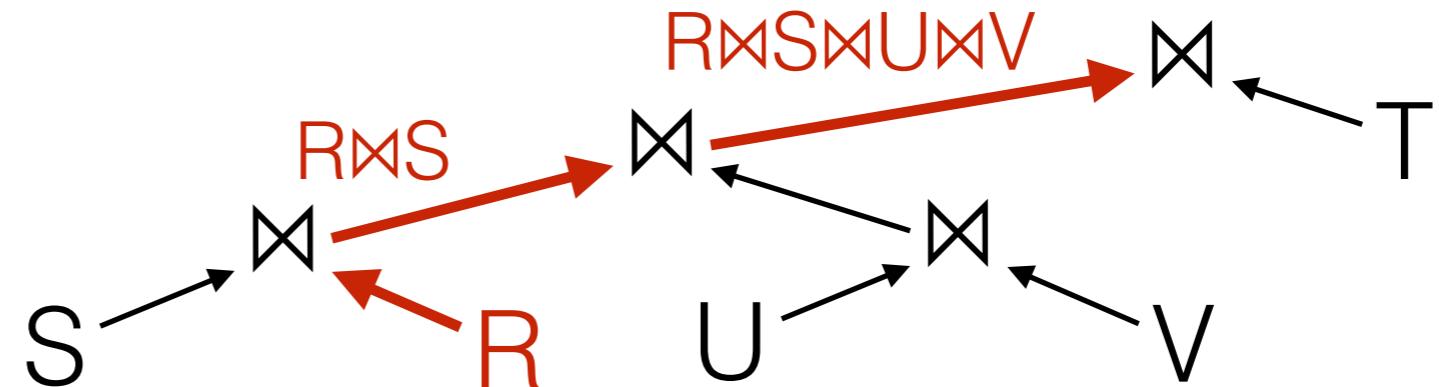
Path **R**→**Root**: **T** after **U**

Path **R**→**Root**: **T** not after **U**

Bushy Query Plans

Plan Space

*Bushy Plan ~
Binary Tree*



Partitioning

Partitioning via Ternary Constraints

Path **R**→**Root**: **T** after **U**

Path **R**→**Root**: **T** not after **U**

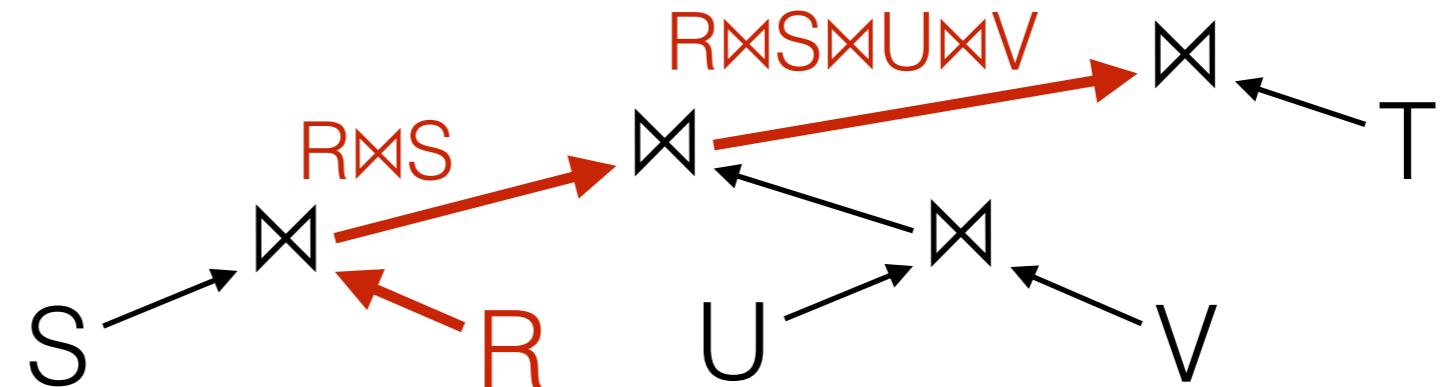
Sub-Queries

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Bushy Query Plans

Plan Space

*Bushy Plan ~
Binary Tree*



Partitioning

Partitioning via Ternary Constraints

Path $\mathbf{R} \rightarrow \mathbf{Root: T}$ after \mathbf{U}

Path $\mathbf{R} \rightarrow \mathbf{Root: T}$ not after \mathbf{U}

Sub-Queries

TUV

T,U,V

TUV

T,V

TUV

U,V

TUV

T,U

TUV

TUV

TUV

TUV

RS

R,T,U,V

R,T,V

R,U,V

R,T,U

R,V

R,U

R,T

RS

S,T,U,V

S,T,V

S,U,V

S,T,U

S,V

S,U

S,T

RS

R,S,T,U,V

R,S,T,V

R,S,U,V

R,S,T,U

R,S,V

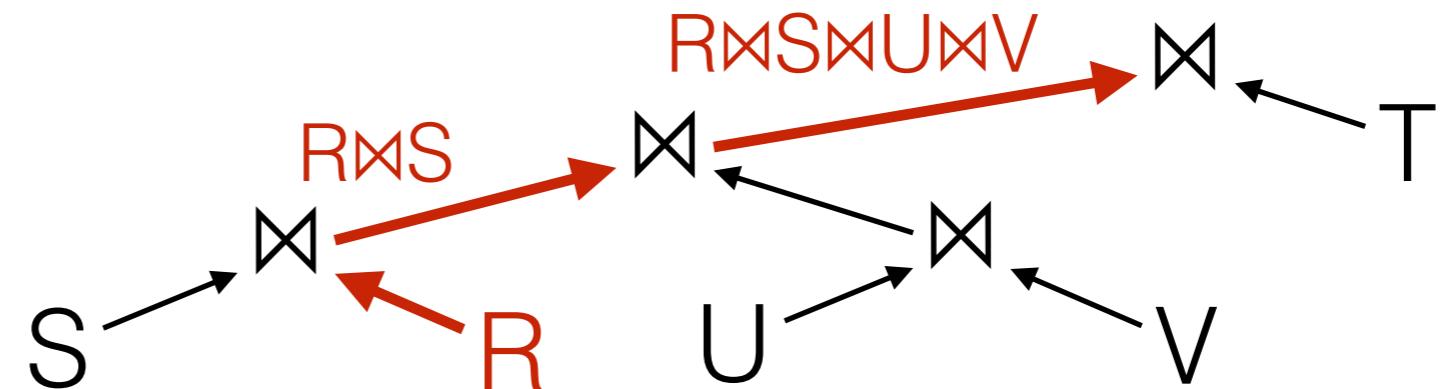
R,S,U

R,S,T

R,S

Bushy Query Plans

Bushy Plan ~ Binary Tree



Partitioning via Ternary Constraints

Path **R**→**Root**: **T** after **U**

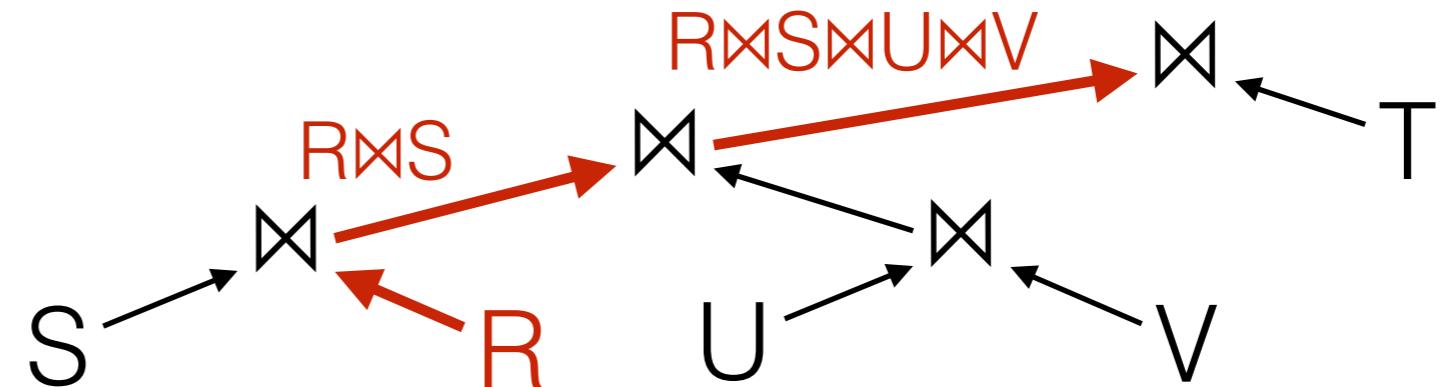
Path **R**→**Root**: **T** not after **U**

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | R,T,V | S,T,V | R,S,T,V |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | R,T | S,T | R,S,T |
| TUV | | | | R,S |

Bushy Query Plans

Plan Space

*Bushy Plan ~
Binary Tree*



Partitioning

Partitioning via Ternary Constraints

Path $R \rightarrow \text{Root: } T$ after U

Path $R \rightarrow \text{Root: } T$ not after U

Sub-Queries

| | RS | RS | RS | RS |
|-----|-------|---------|---------|-----------|
| TUV | T,U,V | R,T,U,V | S,T,U,V | R,S,T,U,V |
| TUV | T,V | ..., | S,T,V | ..., |
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | ..., | S,T | ..., |
| TUV | | | | R,S |

Bushy Query Plans

Space

Bushy Plan ~



Parallelism

nrPartitionLevels

Sub-Queries

$$*(7/8)^{nrPartitionLevels}$$

Time

*** $(21/27)^{nrPartitionLevels}$**

Sub-Querie

| | | | | |
|-----|-----|-------|-------|---------|
| TUV | U,V | R,U,V | S,U,V | R,S,U,V |
| TUV | T,U | R,T,U | S,T,U | R,S,T,U |
| TUV | | R,V | S,V | R,S,V |
| TUV | | R,U | S,U | R,S,U |
| TUV | | S,T | R,S,T | R,S,T |
| TUV | | | | R,S |

Bushy Query Plans

Space

Bushy Plan ~



Parallelism

nrPartitionLevels

Sub-Queries

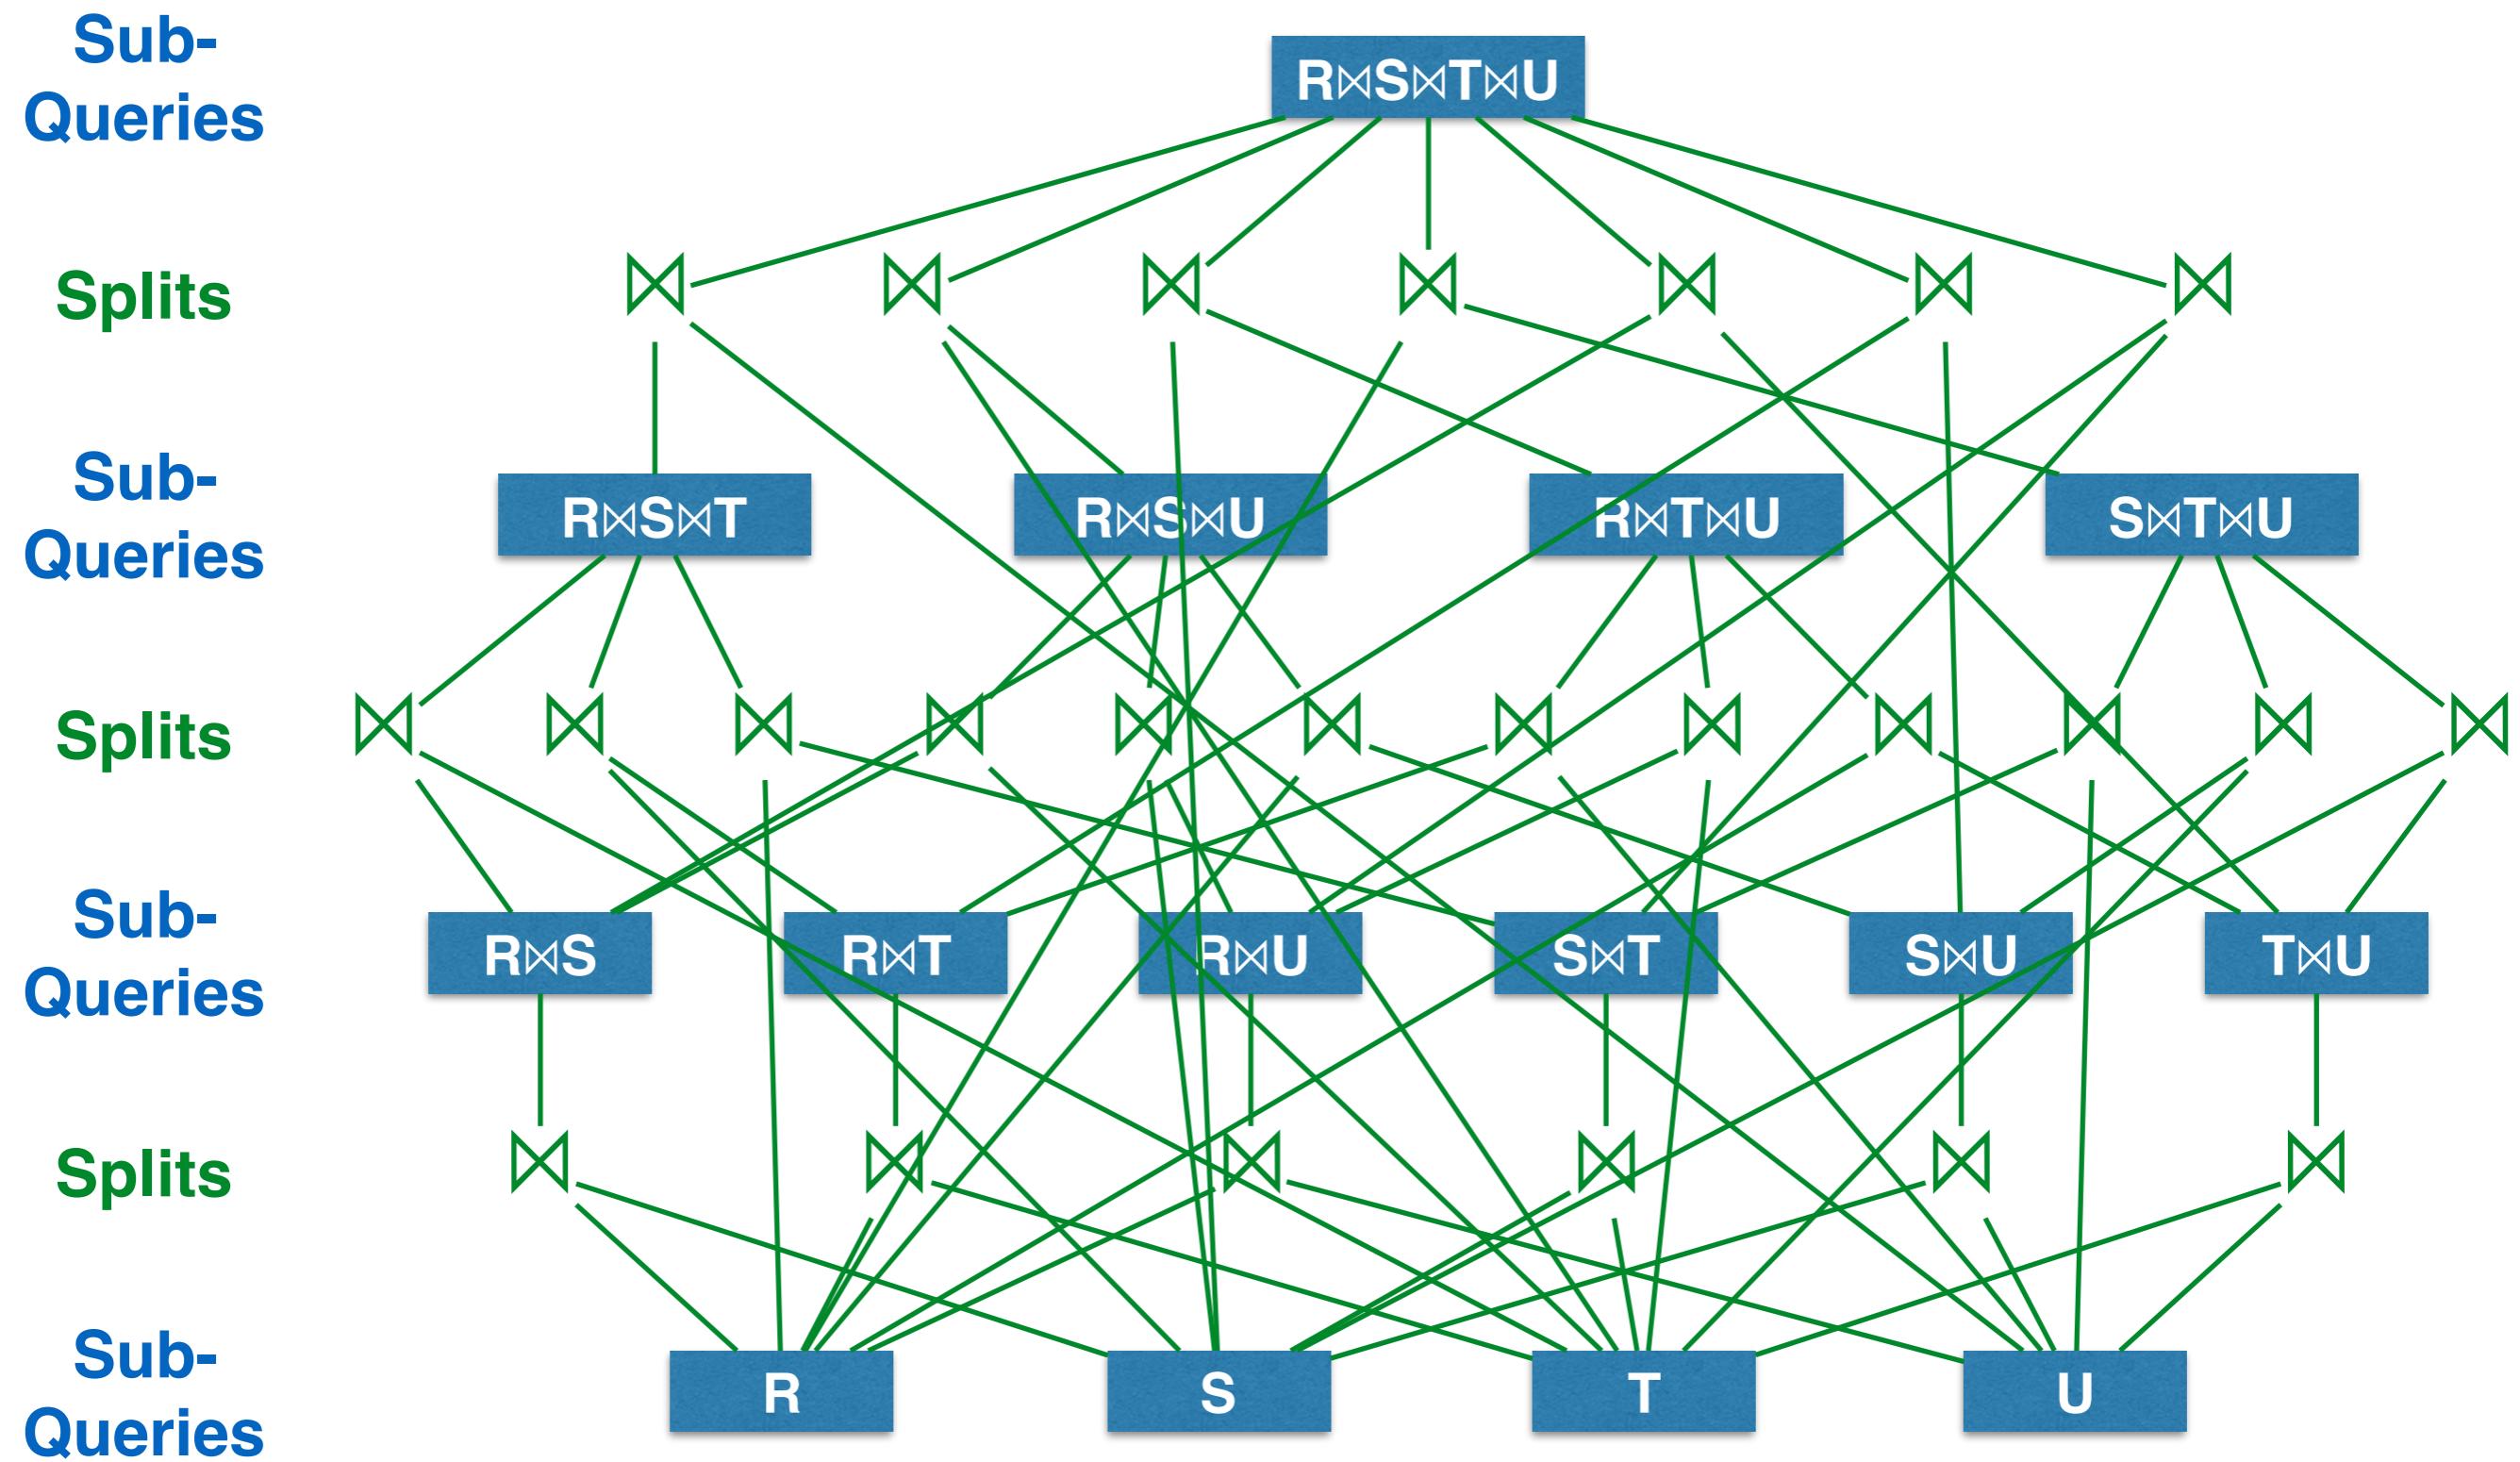
$*(7/8)^{nrPartitionLevels}$

Time

*(21/27)

A large, bold red question mark and exclamation mark graphic is centered on a white rectangular background. The graphic is set against a yellow background with black and red diagonal stripes.

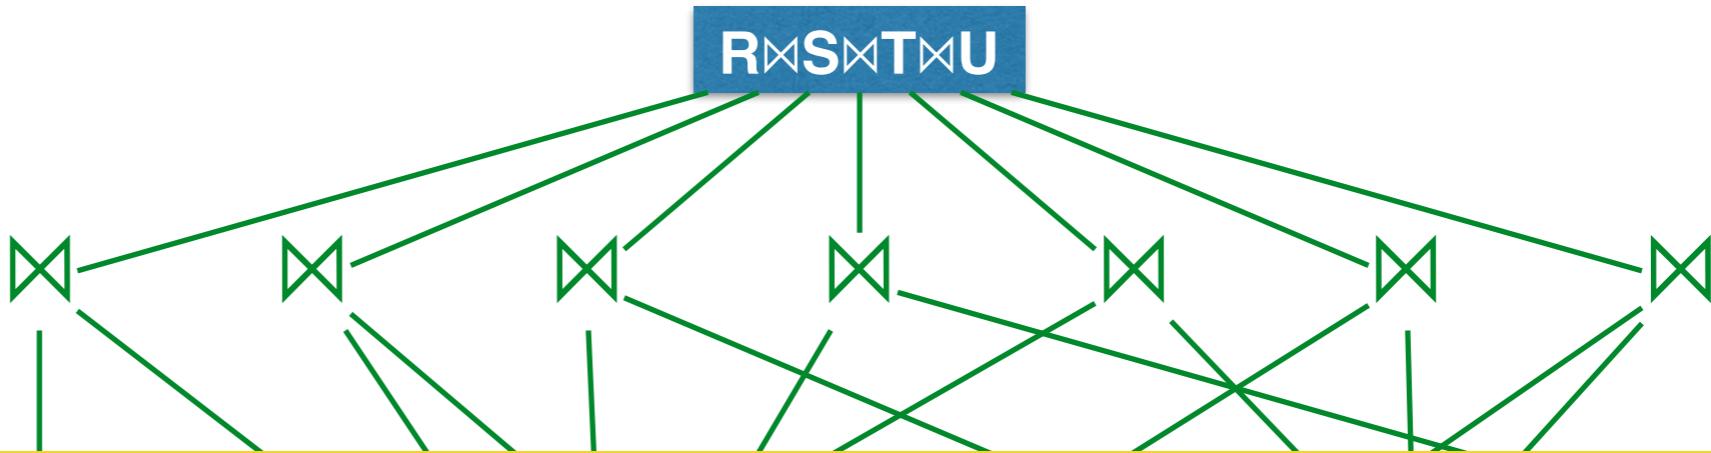
Complexity for Bushy Plans



Complexity for Bushy Plans

Sub-
Queries

Splits



Nr. Sub-Queries $\sim 2^{nrTables}$

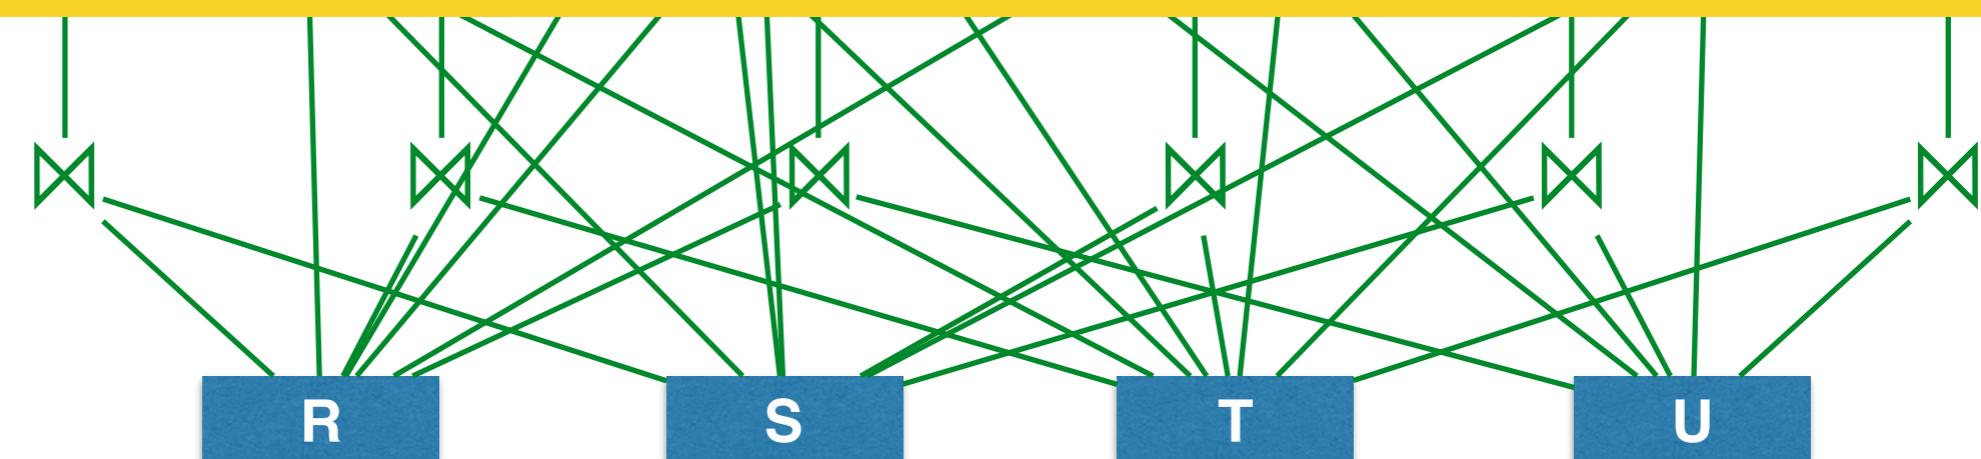
Nr. Splits

$\sim 3^{nrTables}$

Sub-
Queries

Splits

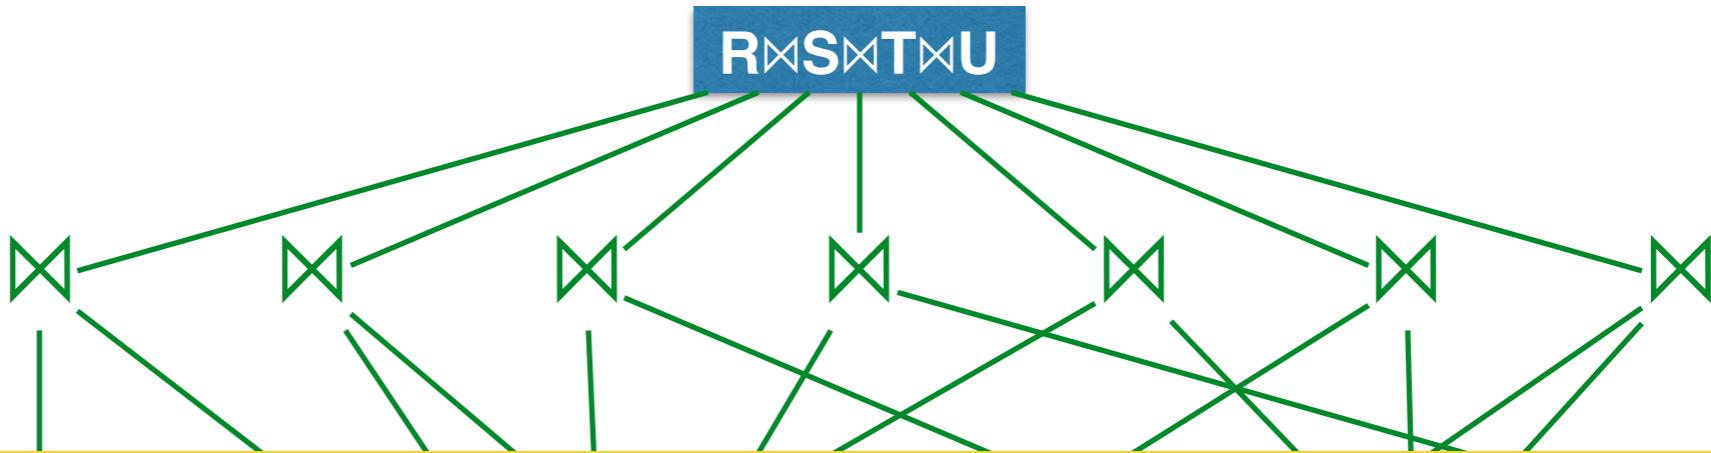
Sub-
Queries



Complexity for Bushy Plans

Sub-
Queries

Splits



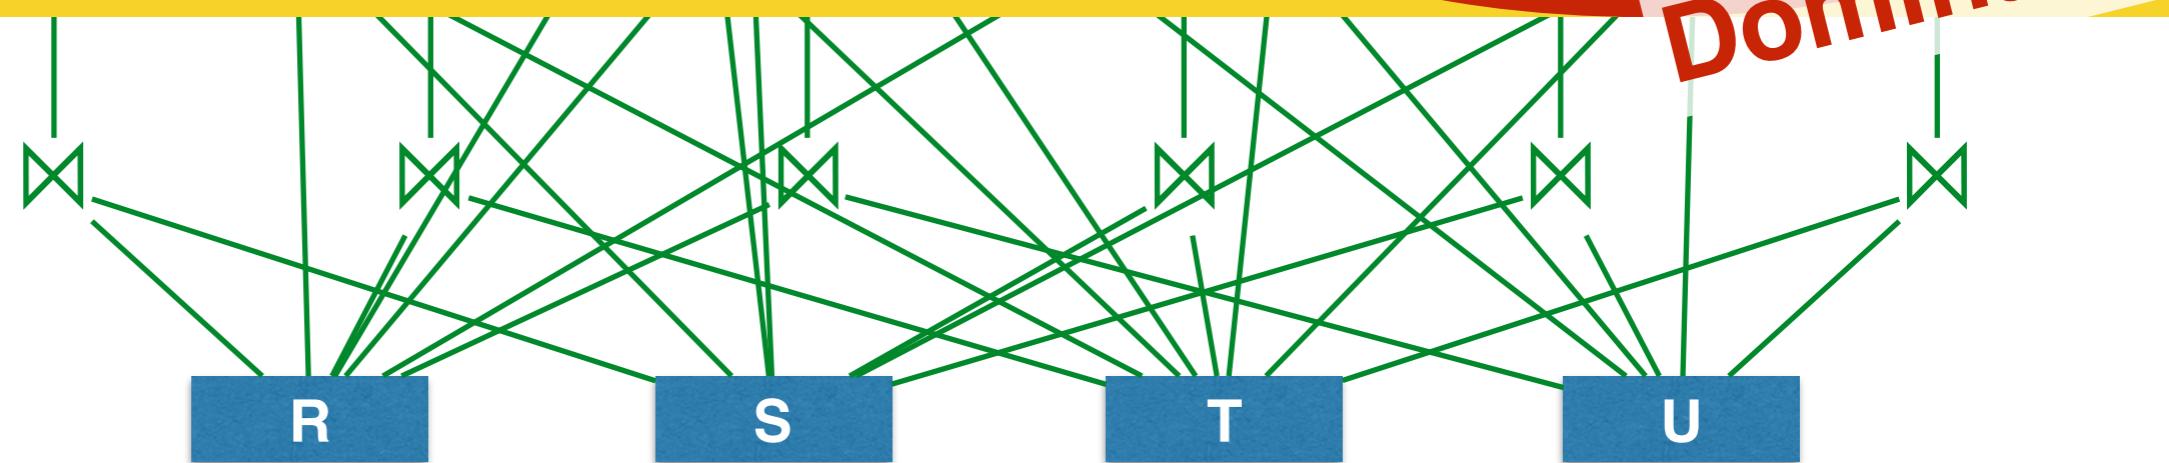
Nr. Sub-Queries $\sim 2^{nrTables}$

Nr. Splits

Sub-
Queries

Splits

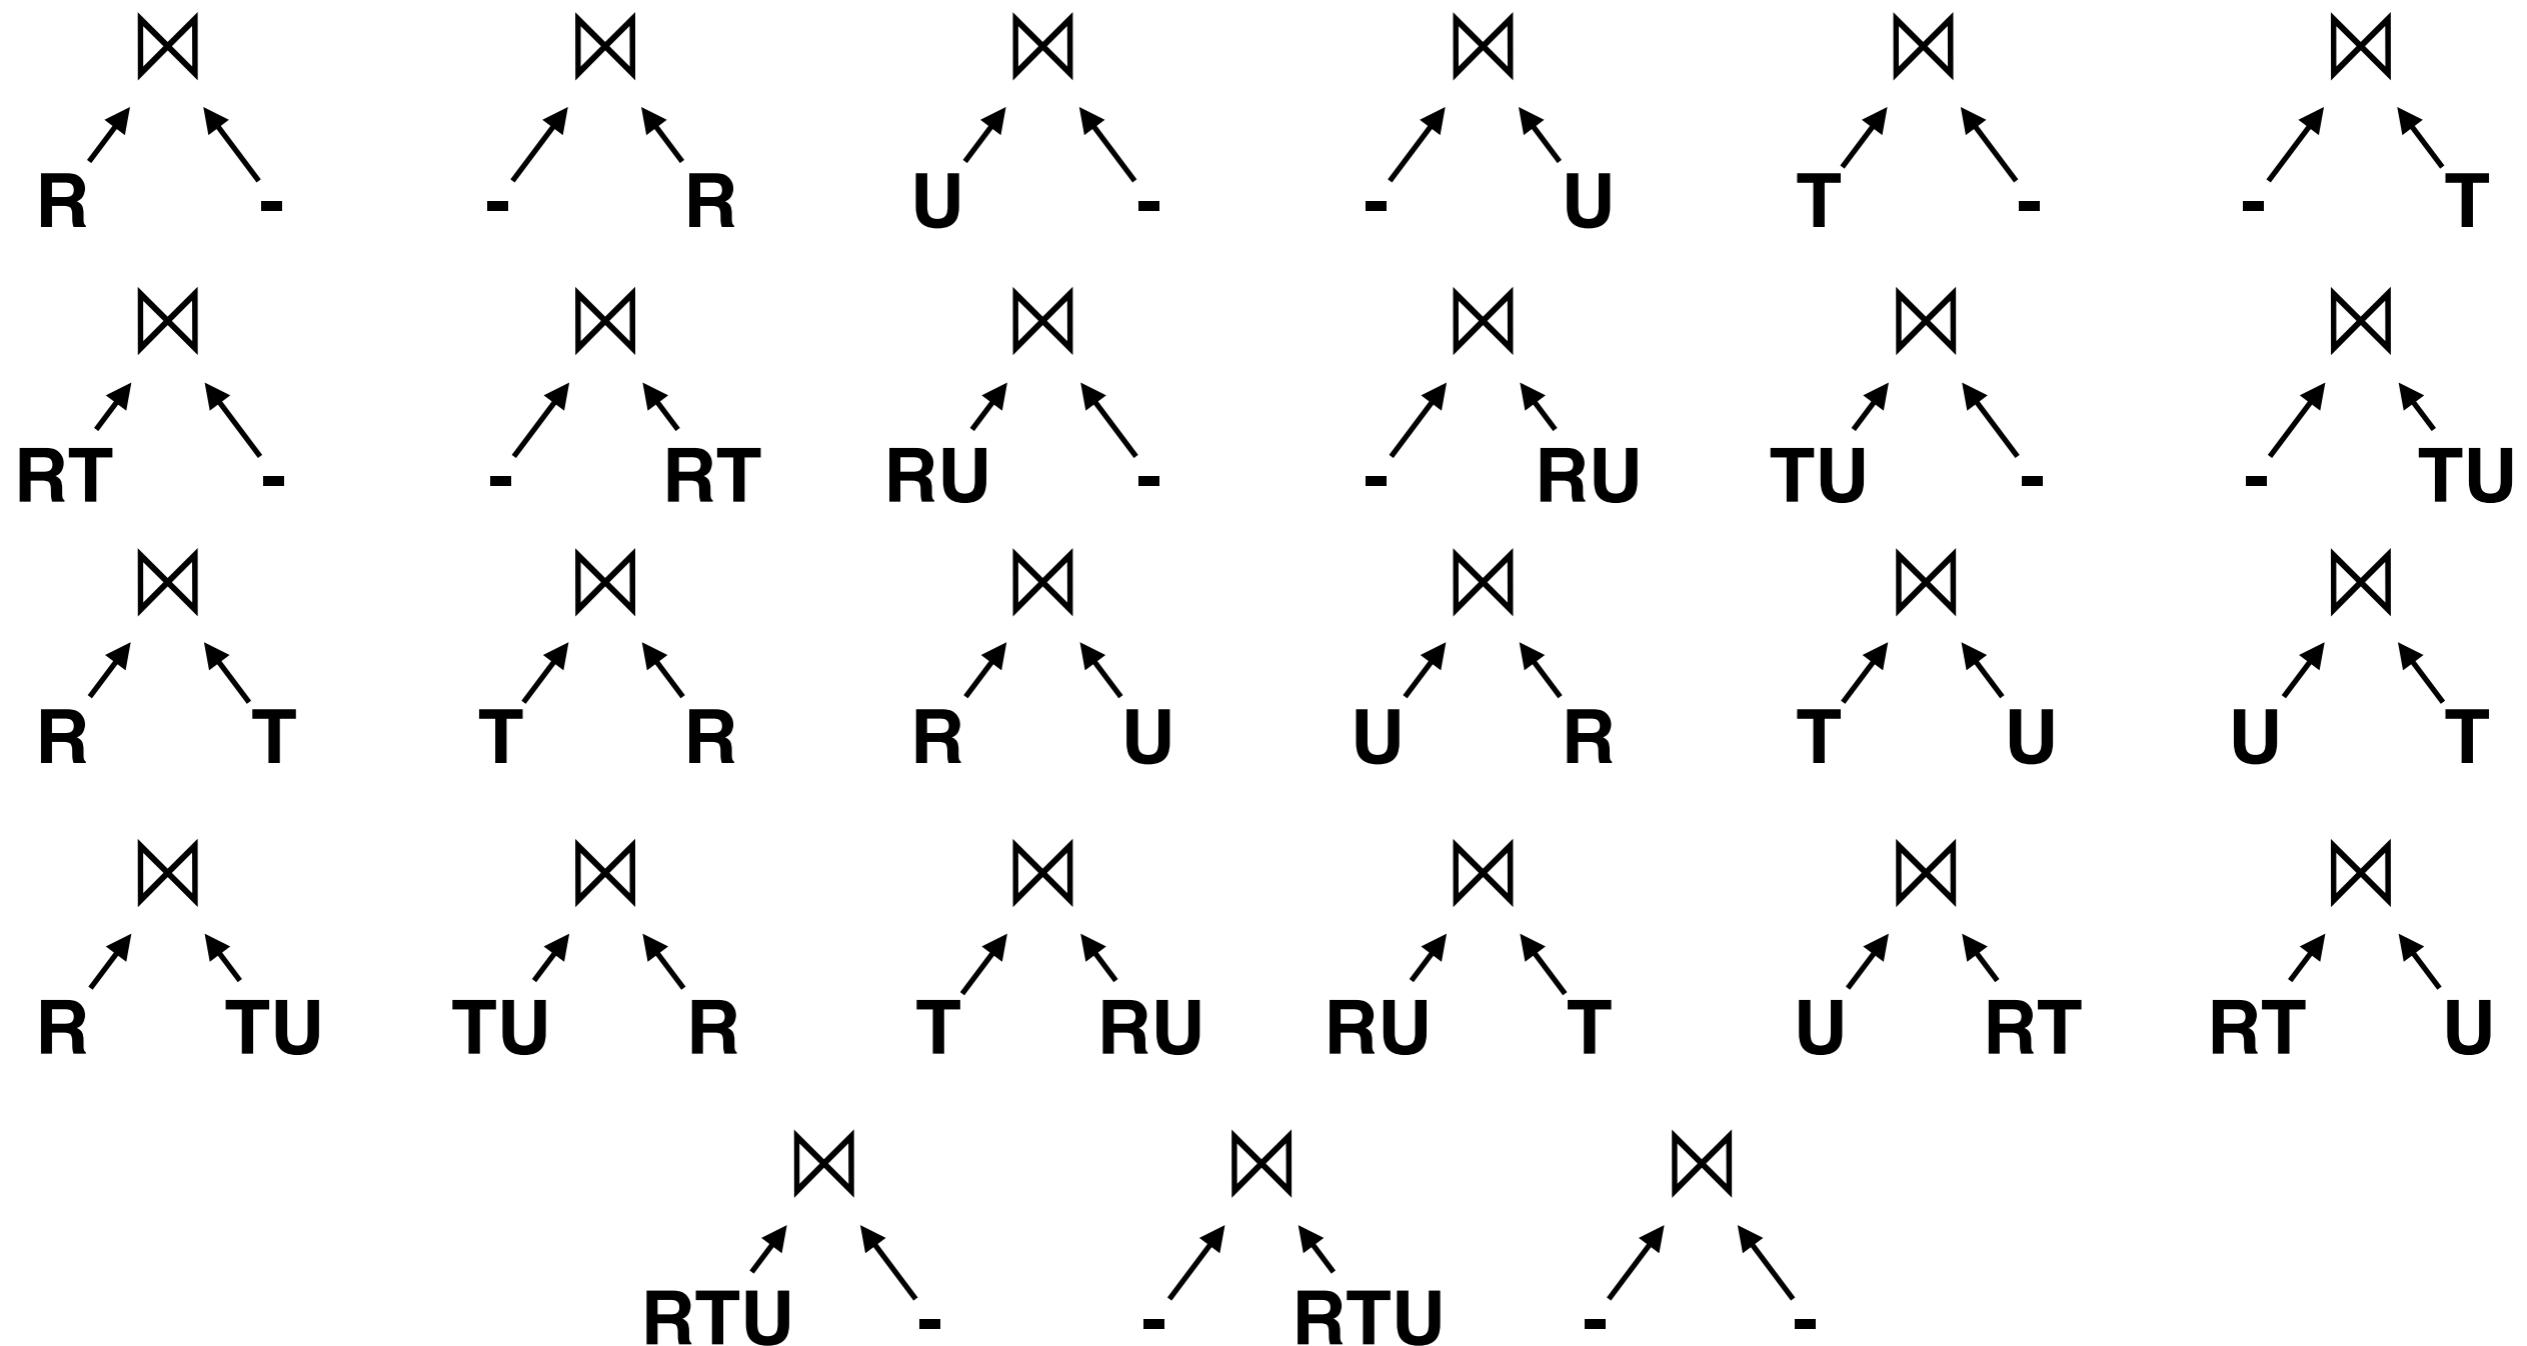
Sub-
Queries



$\sim 3^{nrTables}$

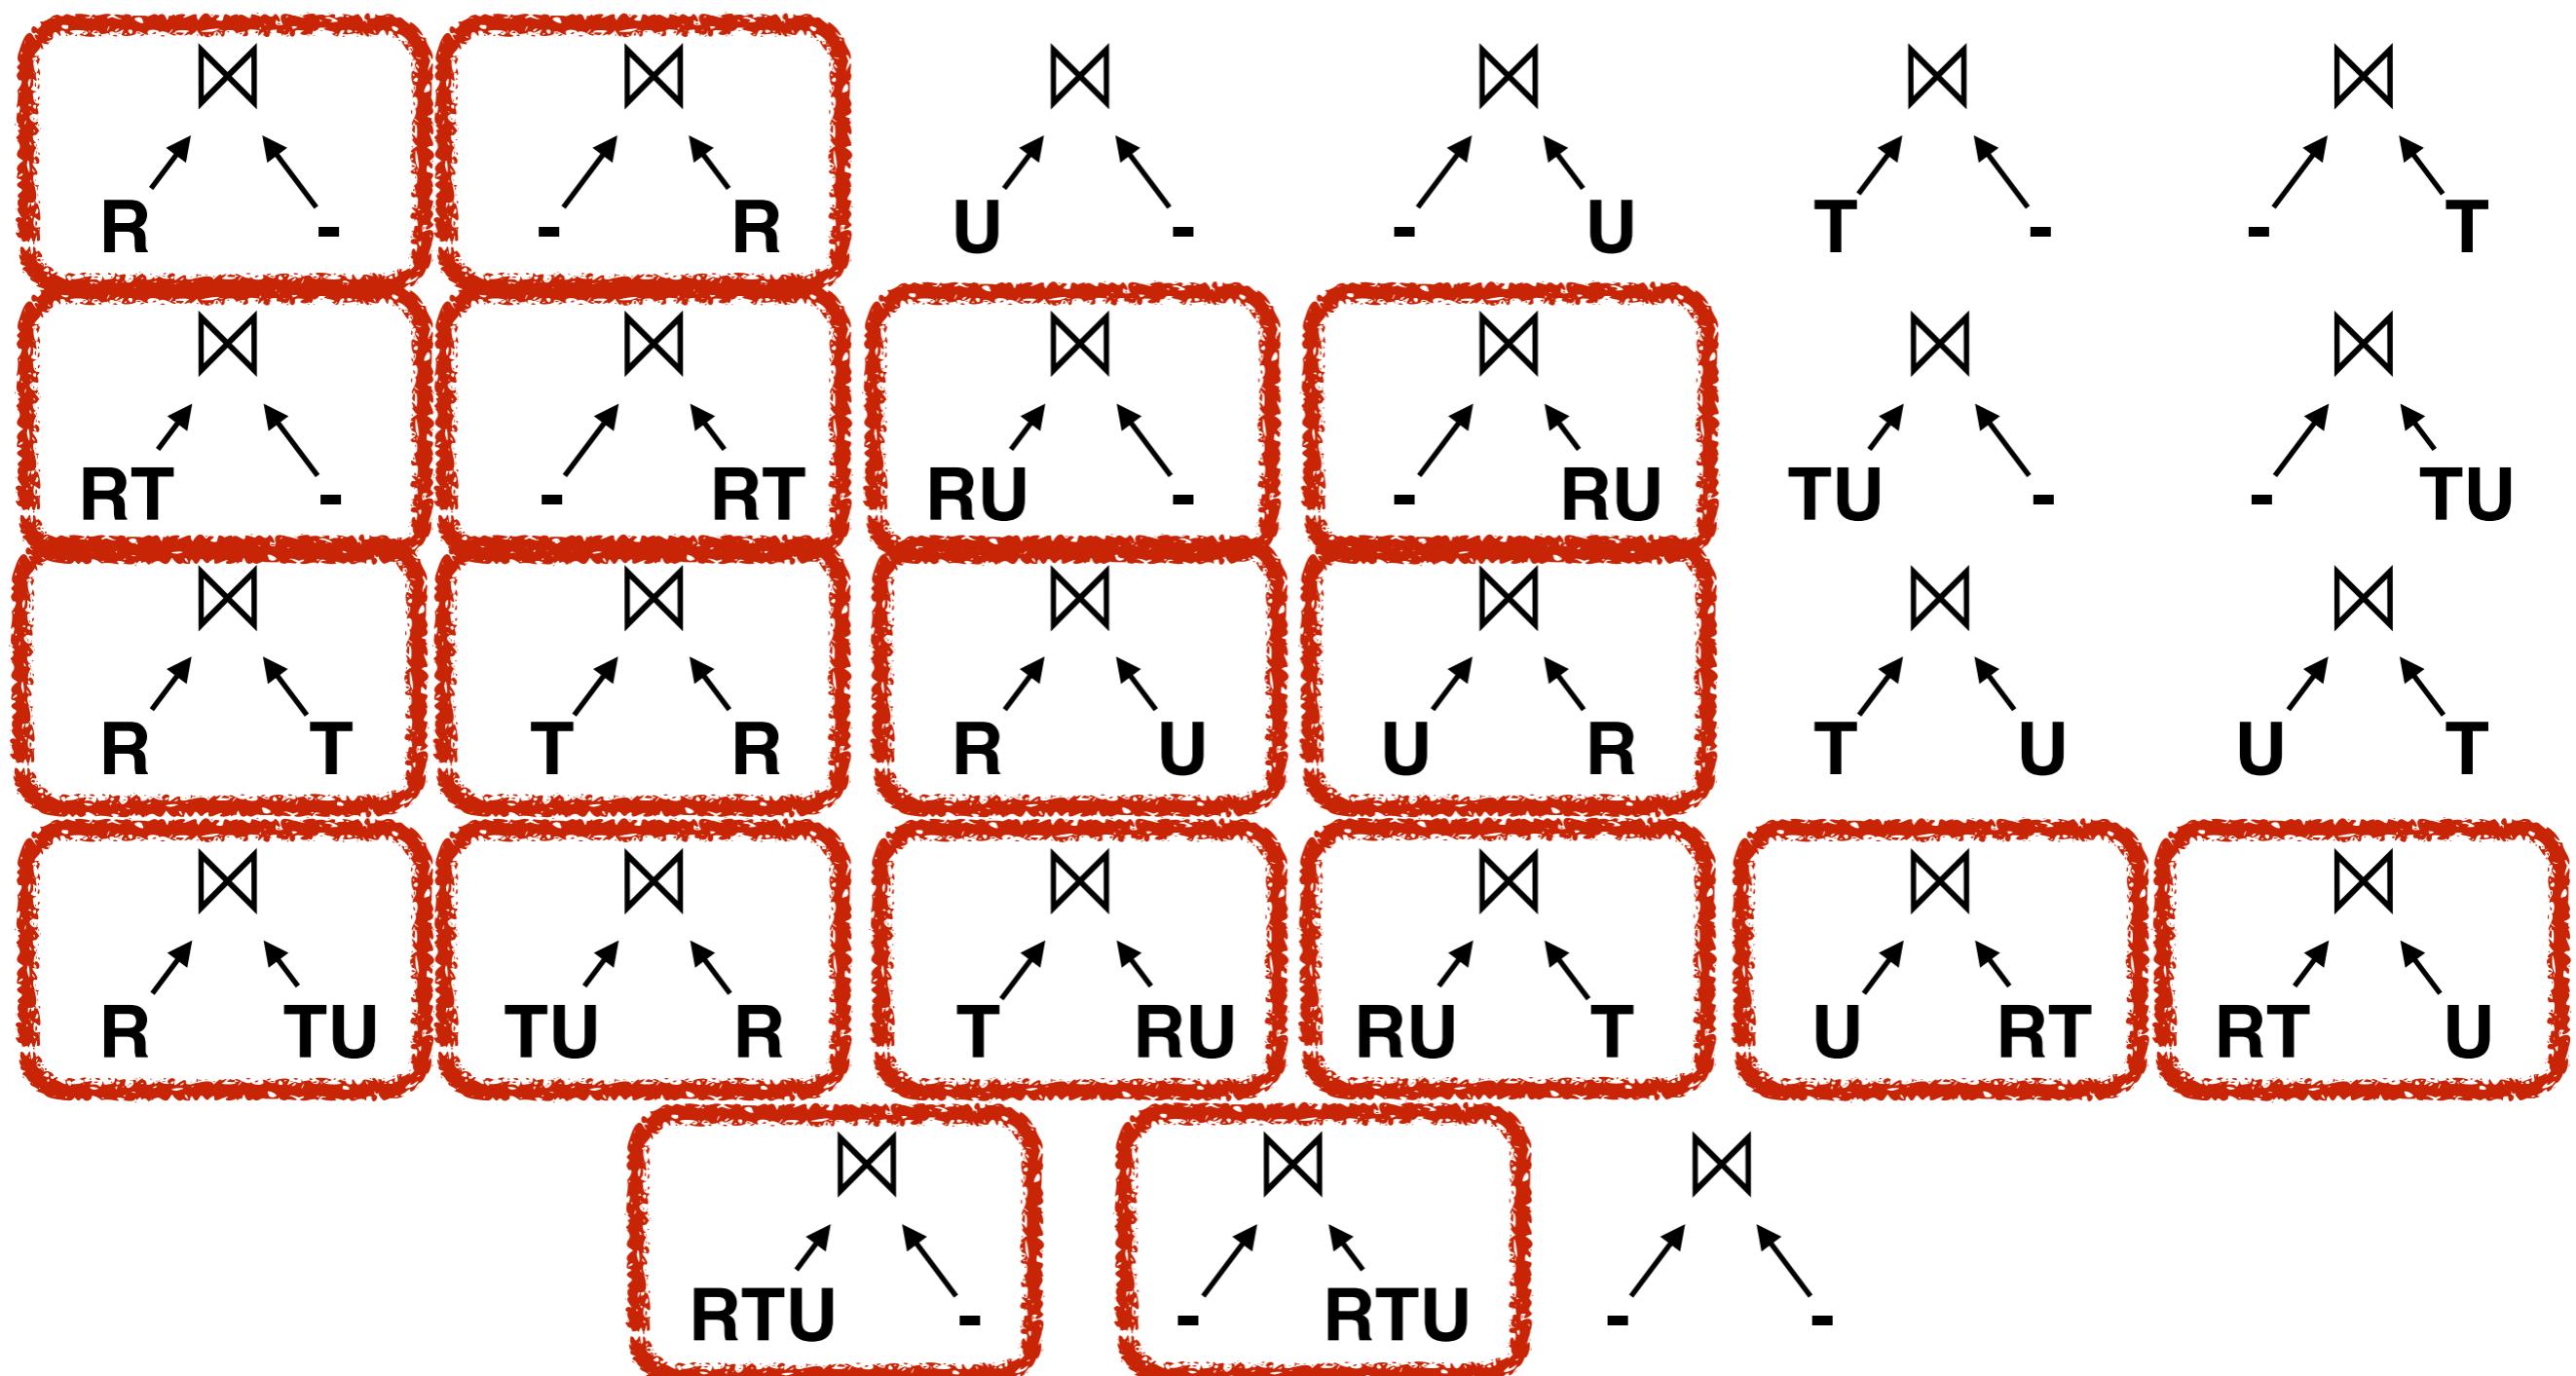
Dominant

Calculating Bushy Speedup



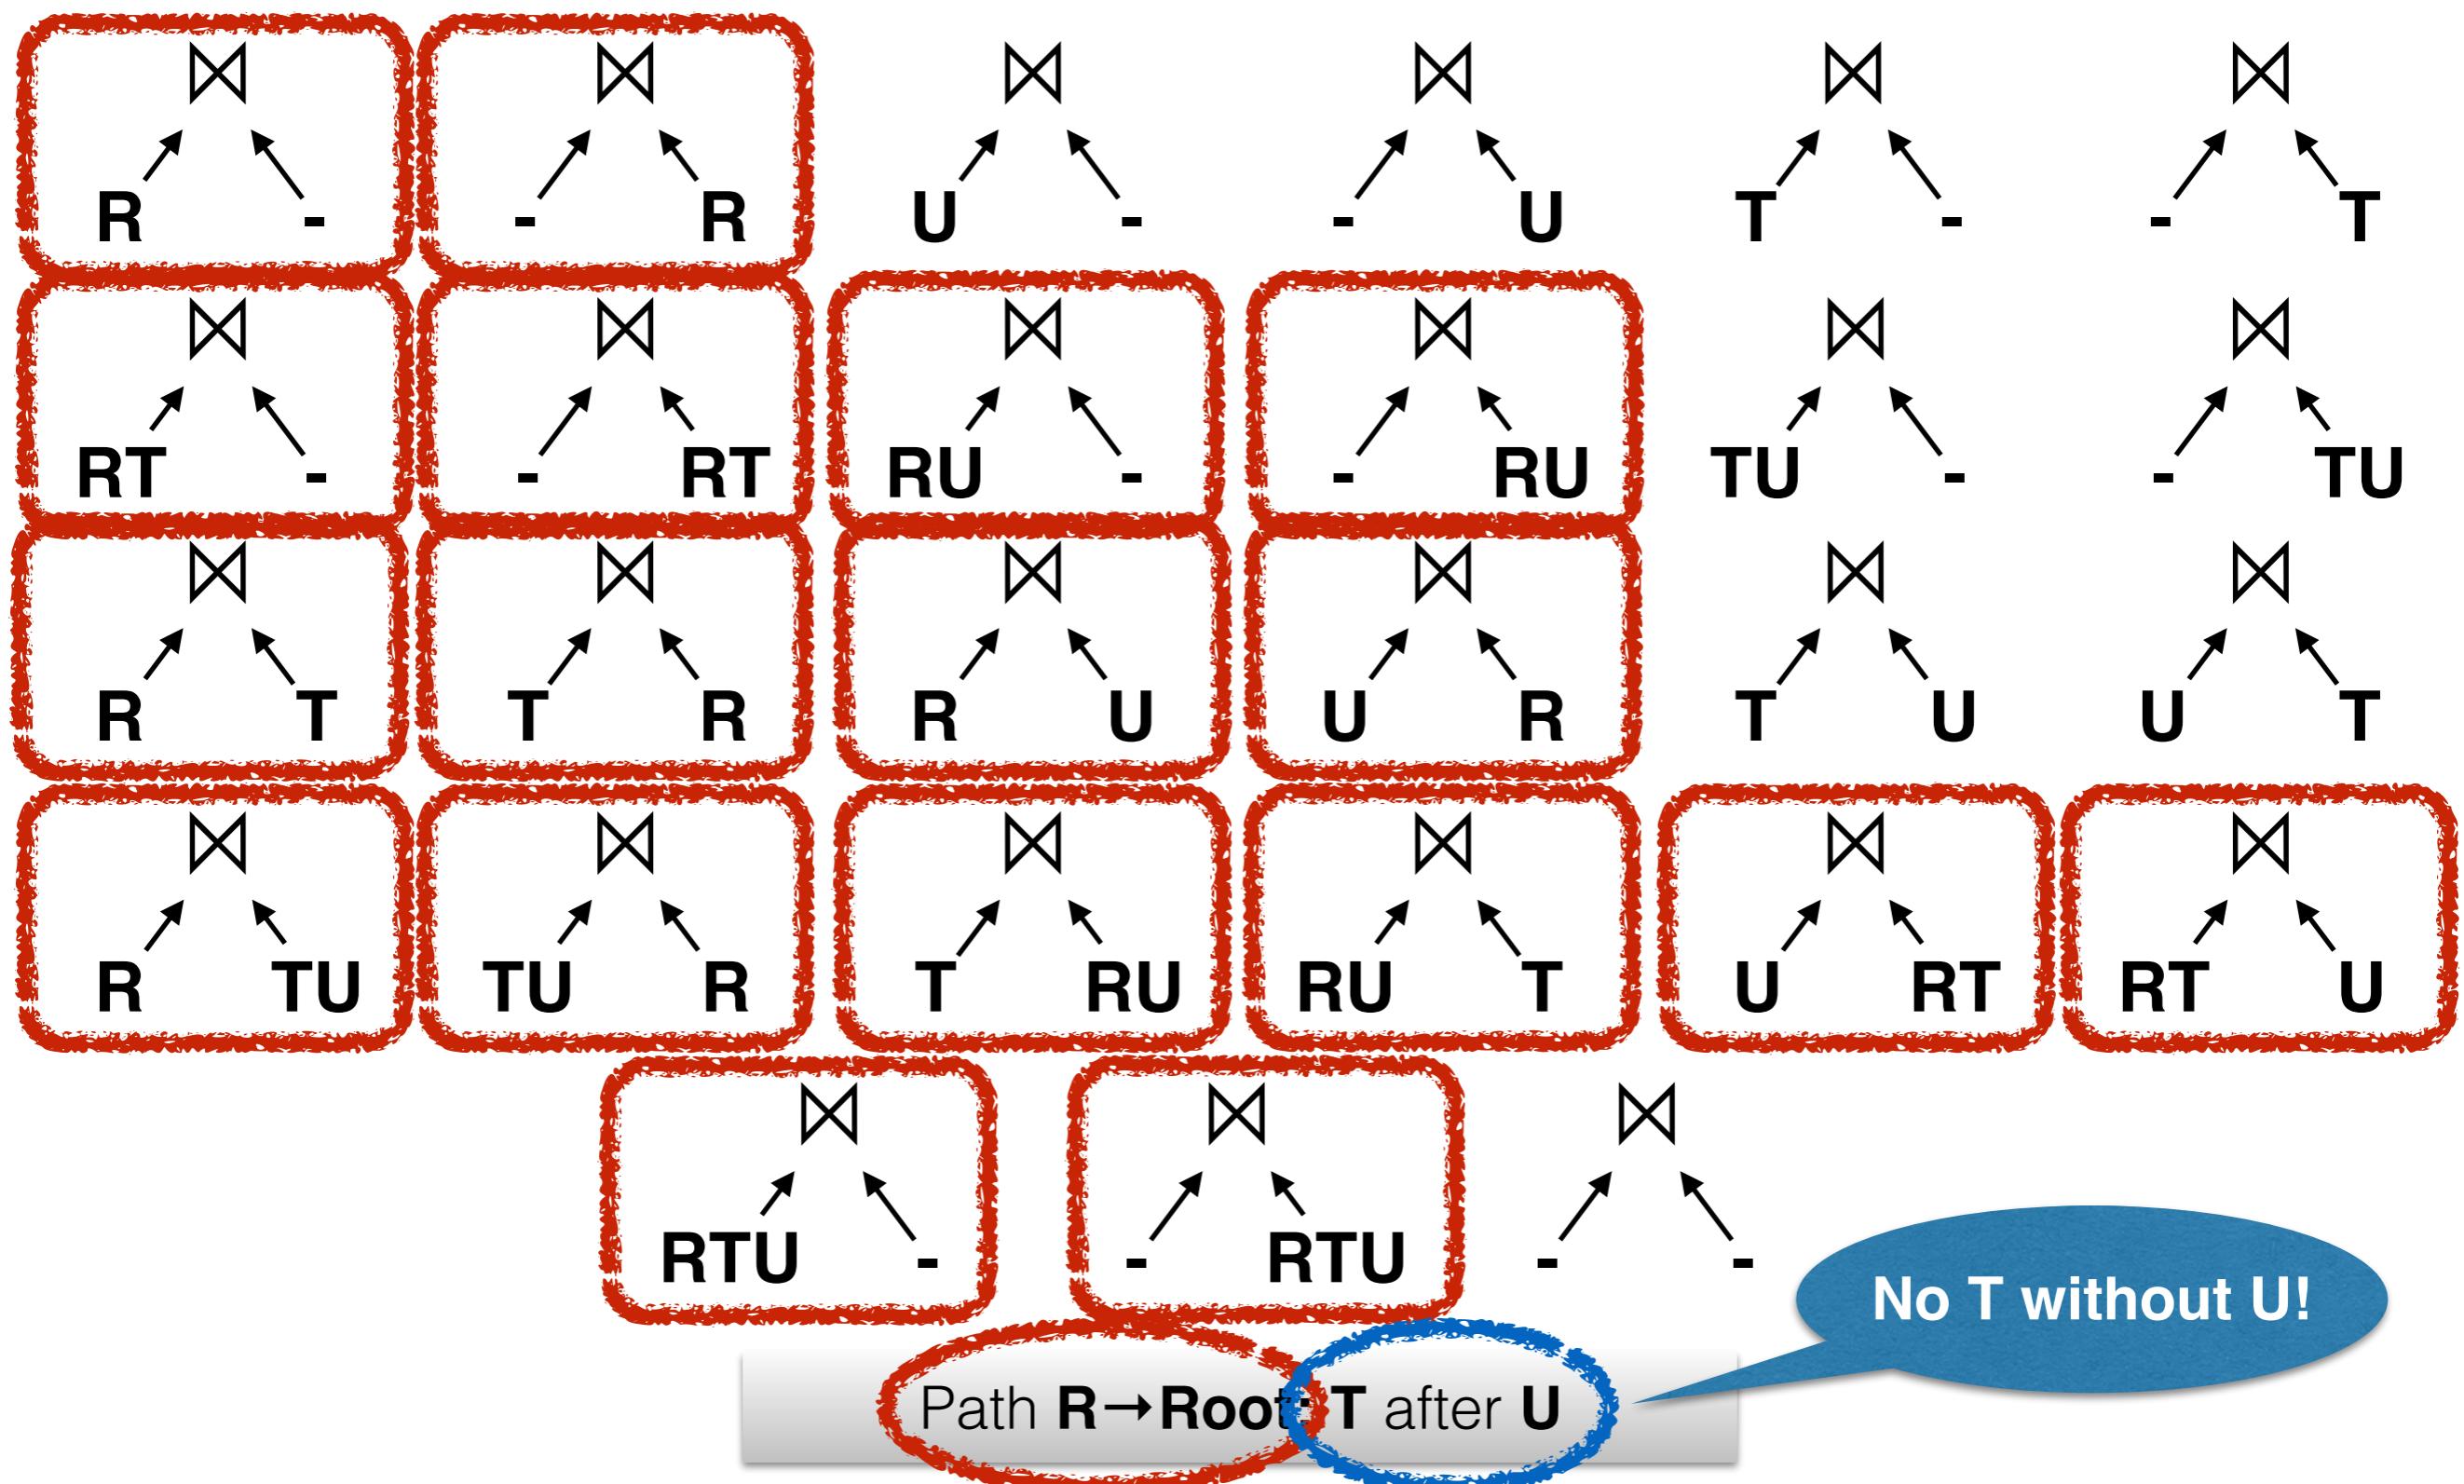
Path **R→Root: T after U**

Calculating Bushy Speedup

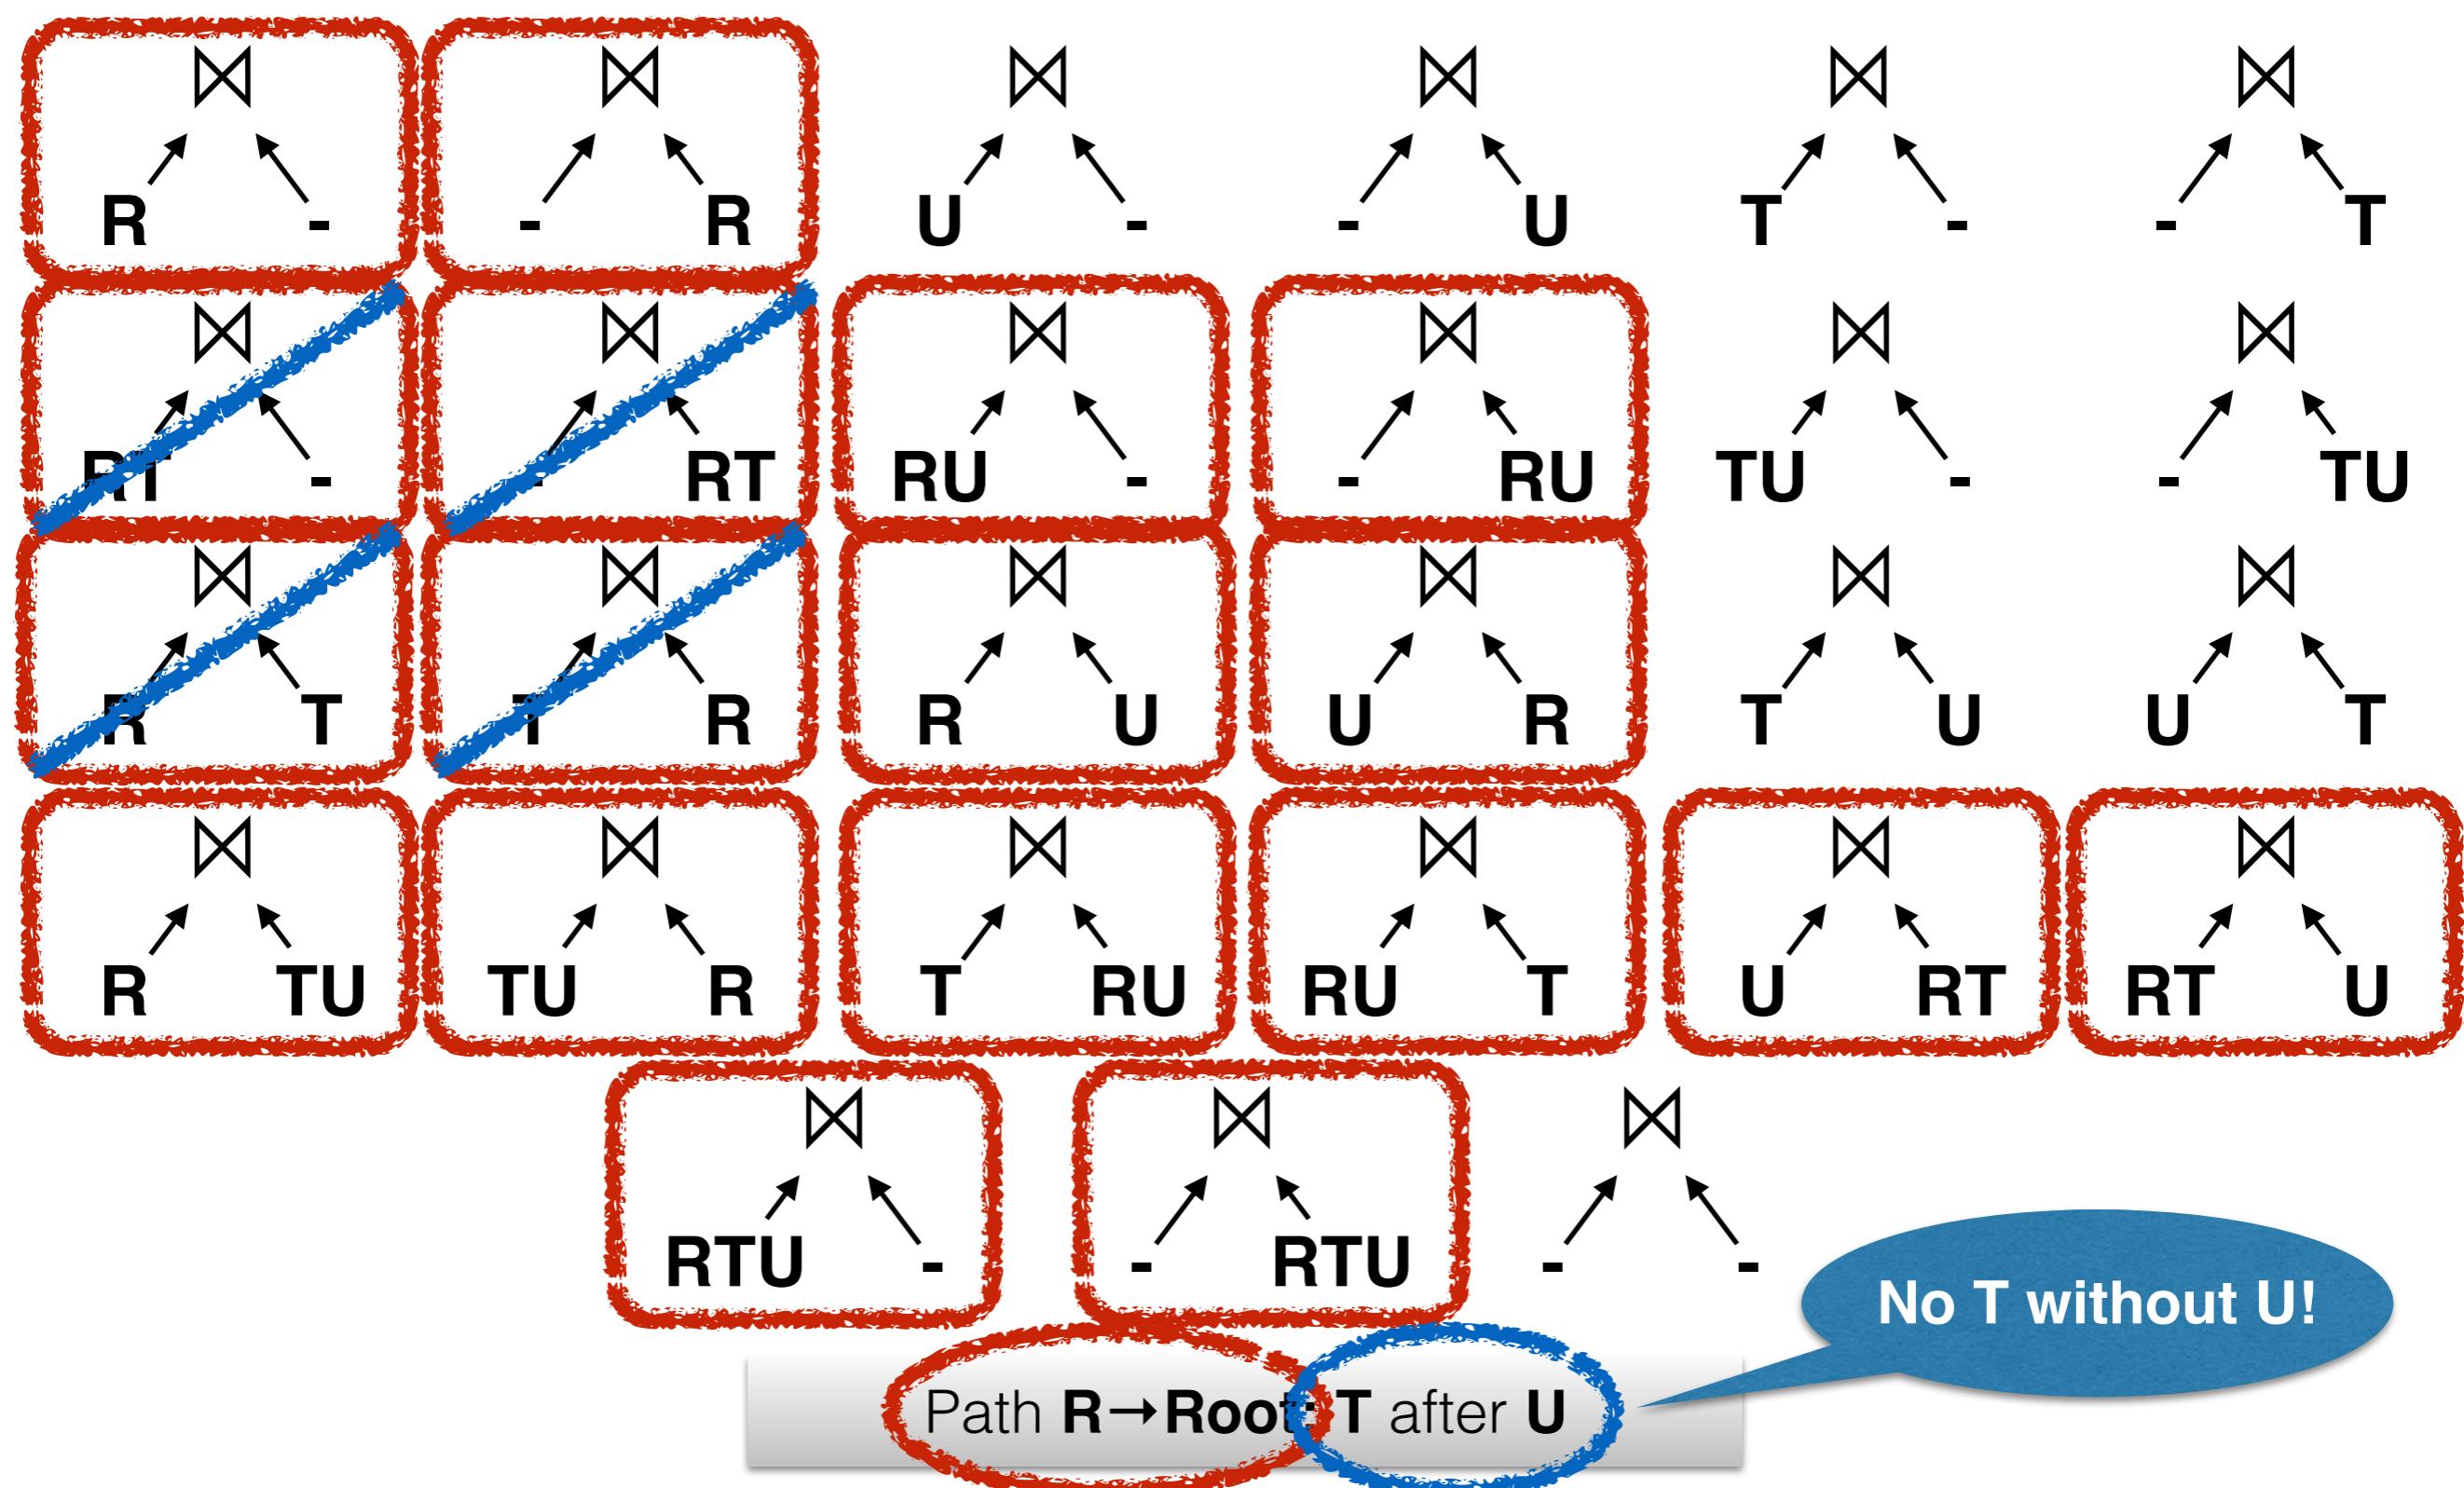


Path **R→Root: T after U**

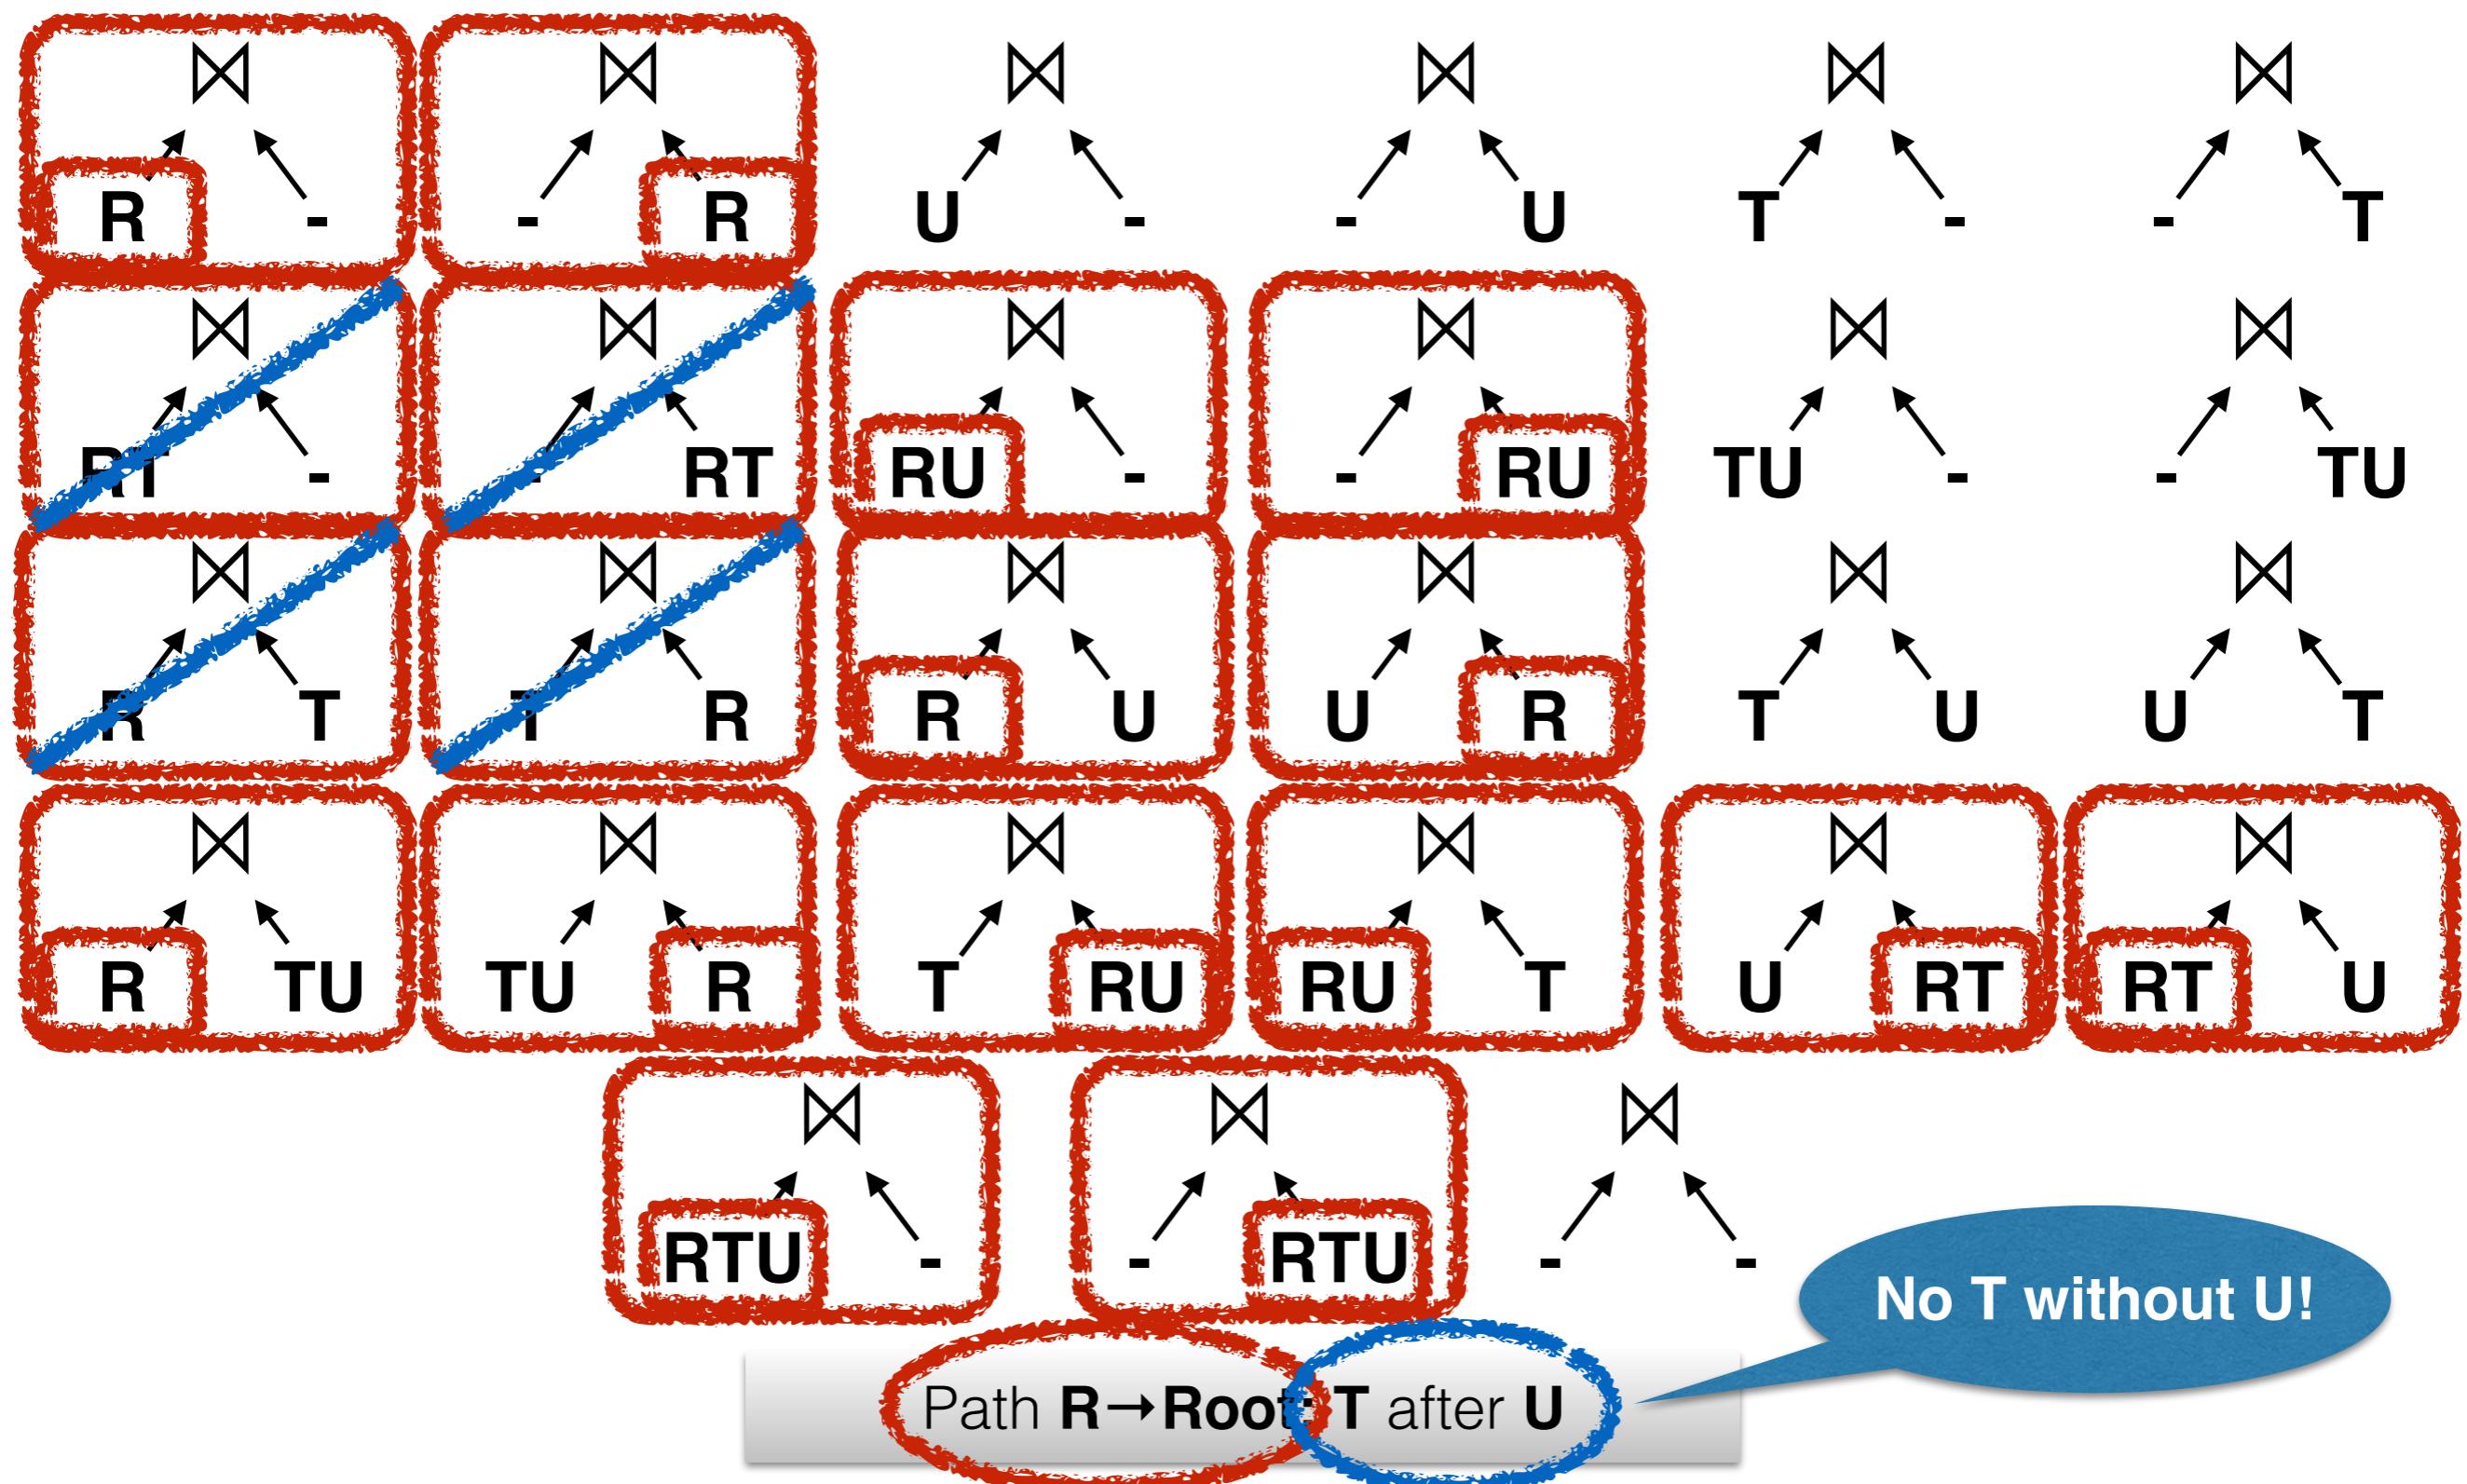
Calculating Bushy Speedup



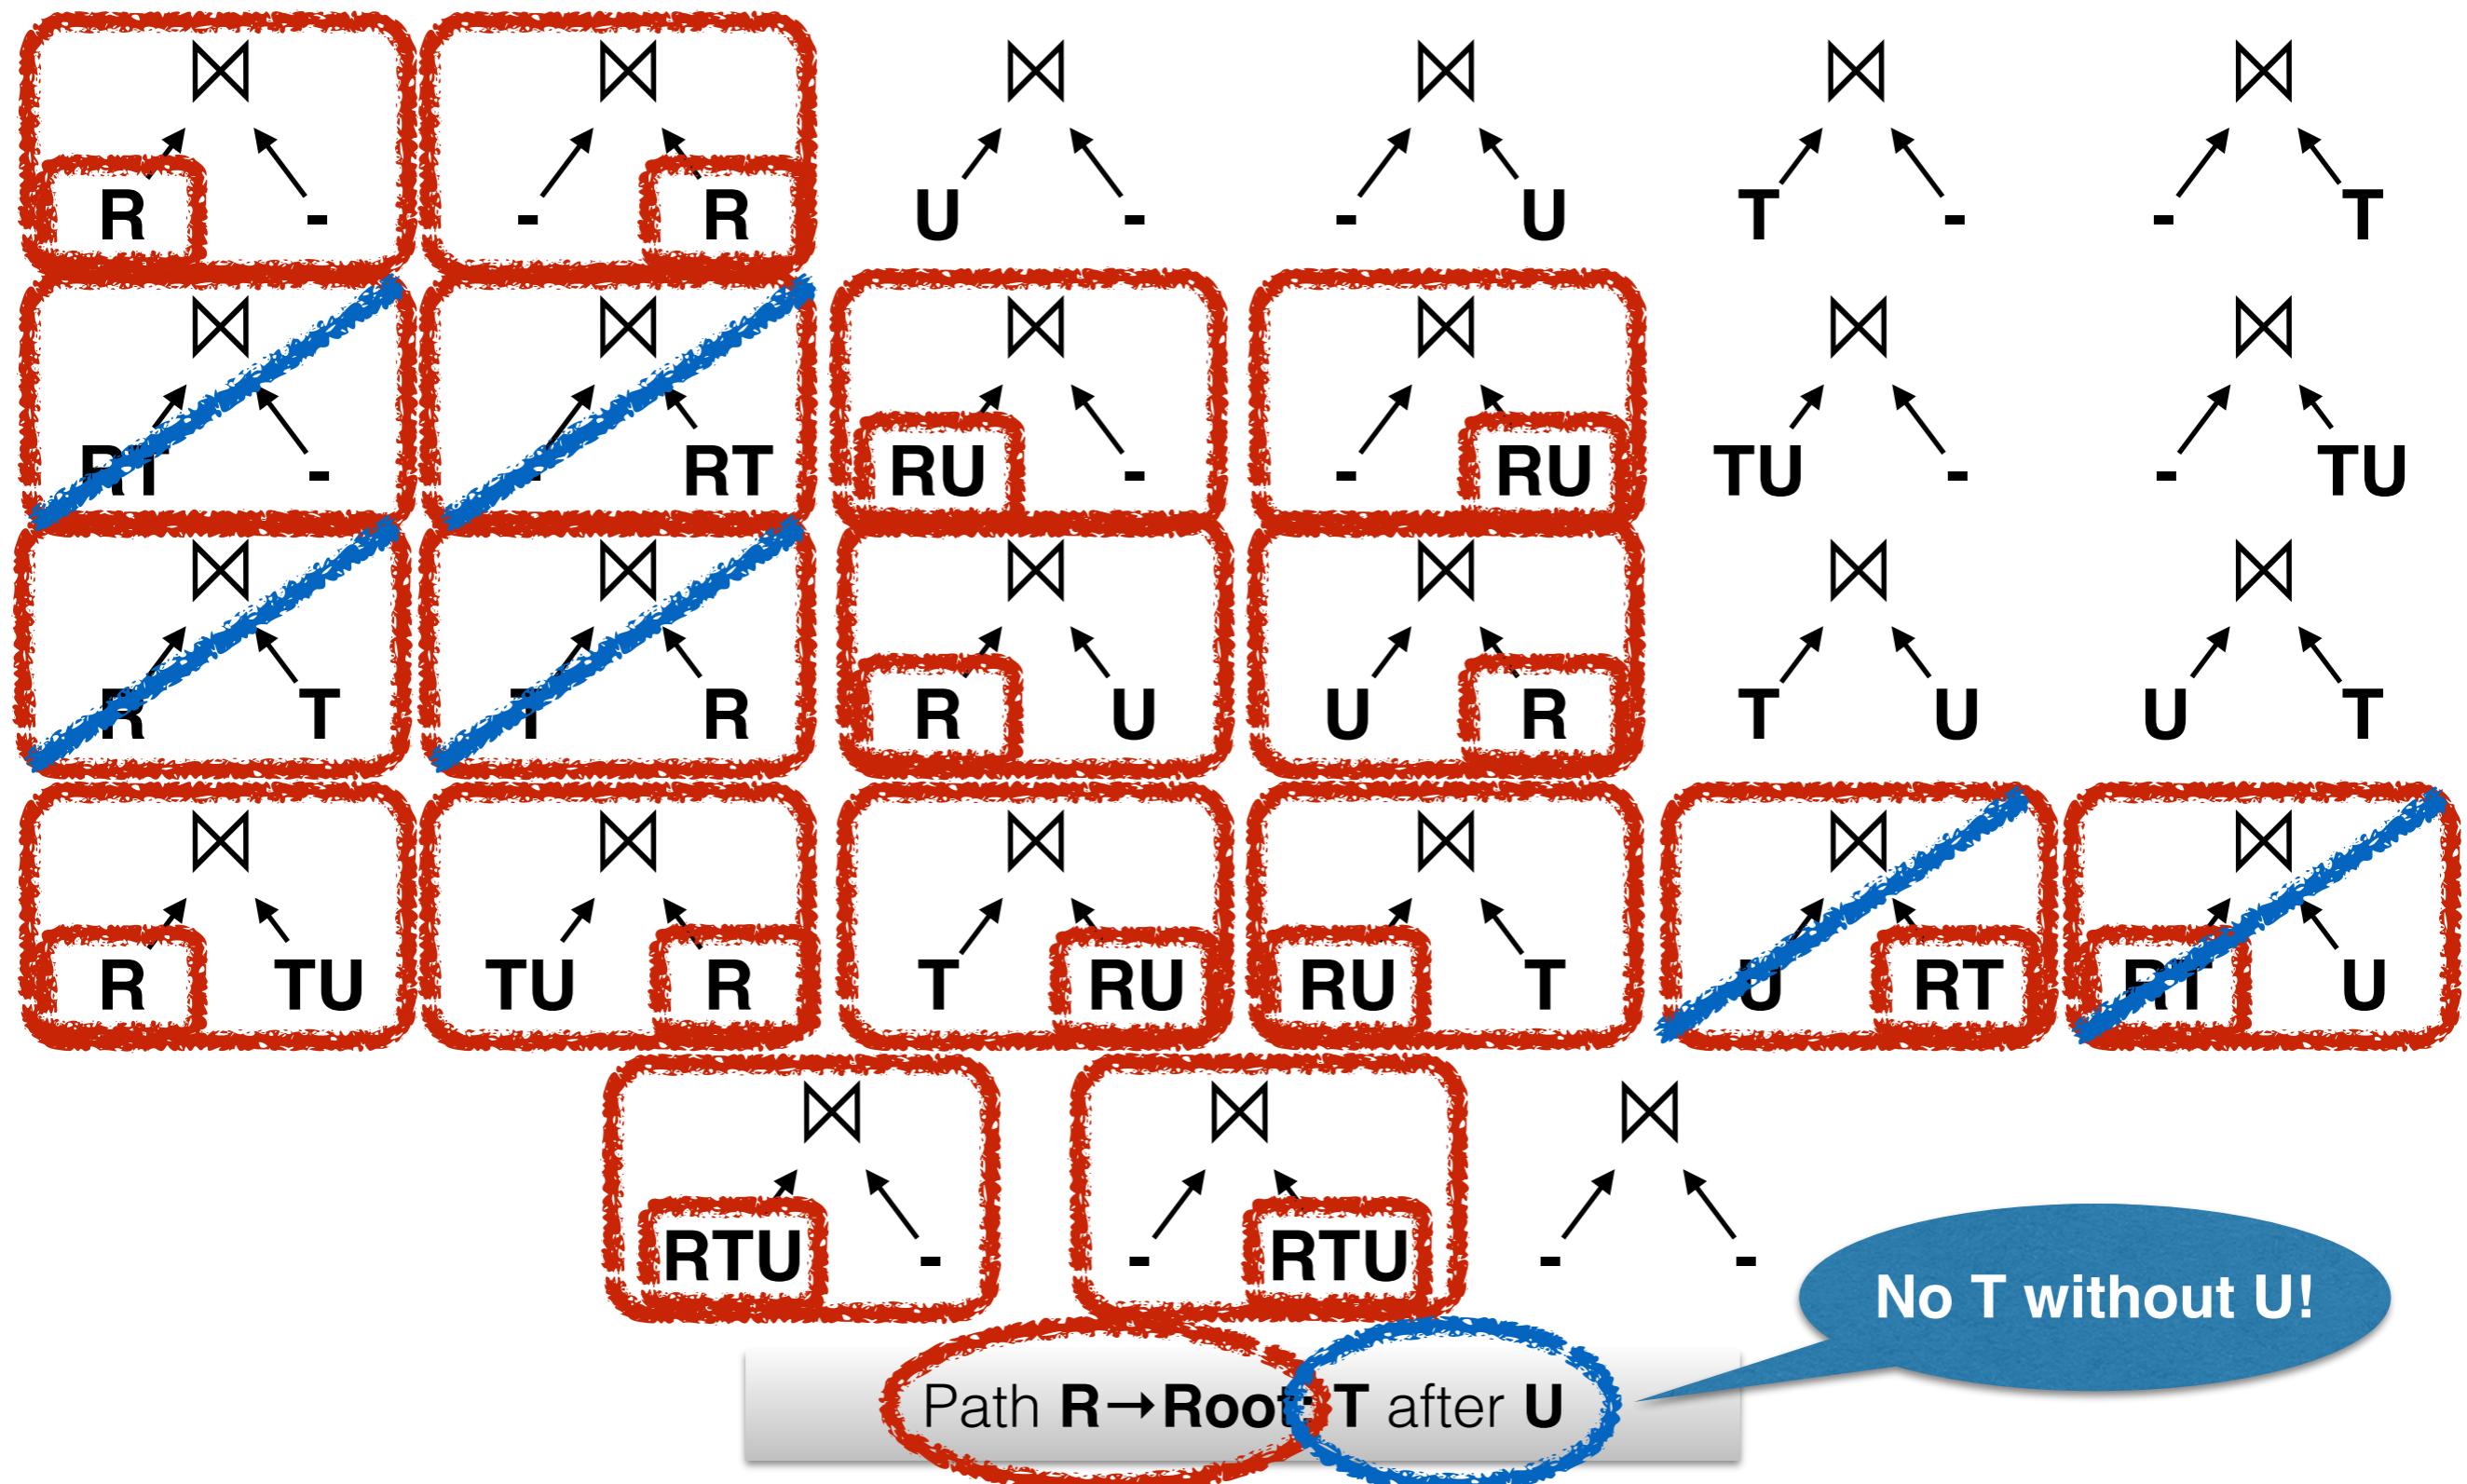
Calculating Bushy Speedup



Calculating Bushy Speedup



Calculating Bushy Speedup



Calculating Bushy Speedup

27 Cases
6 Excluded
Splits * 21/27

Path **R**→**Root** **T** after **U**

No T without U!

Complexity Results

Results for one cost metric:

| Metric \ Plan Space | Left-Deep | Bushy |
|---------------------|---|---------------------------------------|
| Bytes Sent | $nrWorkers * (querySize + planSize)$ | |
| Time - Master | | $nrWorkers^2$ |
| Memory - Worker | $2^{querySize} * (3/4)^{nrWorkers}$ | $2^{querySize} * (7/8)^{nrWorkers}$ |
| Time - Worker | $2^{querySize} * (3/4)^{nrWorkers} * querySize$ | $3^{querySize} * (21/27)^{nrWorkers}$ |

(Considering multiple metrics multiplies by polynomial factor)

Complexity Results

Results for one cost metric:

| Metric \ Plan Space | Left-Deep | Bushy |
|---------------------|--|----------------------------------|
| Bytes Sent | $nrWorkers * (querySize + planSize)$ | |
| Time - Master | | $nrWorkers^2$ |
| Memory - Worker | $2querySize * (3/4) nrWorkers$ | $2querySize * (7/8) nrWorkers$ |
| Time - Worker | $2querySize * (3/4) nrWorkers * querySize$ | $3querySize * (21/27) nrWorkers$ |

(Considering multiple metrics multiplies by polynomial factor)

Complexity Results

Results for one cost metric:

| Metric \ Plan Space | Left-Deep | Bushy |
|---------------------|---|---------------------------------------|
| Bytes Sent | $nrWorkers * (querySize + planSize)$ | |
| Time - Master | | $nrWorkers^2$ |
| Memory - Worker | $2^{querySize} * (3/4)^{nrWorkers}$ | $2^{querySize} * (7/8)^{nrWorkers}$ |
| Time - Worker | $2^{querySize} * (3/4)^{nrWorkers} * querySize$ | $3^{querySize} * (21/27)^{nrWorkers}$ |

(Considering multiple metrics multiplies by polynomial factor)

Complexity Results

Results for one cost metric:

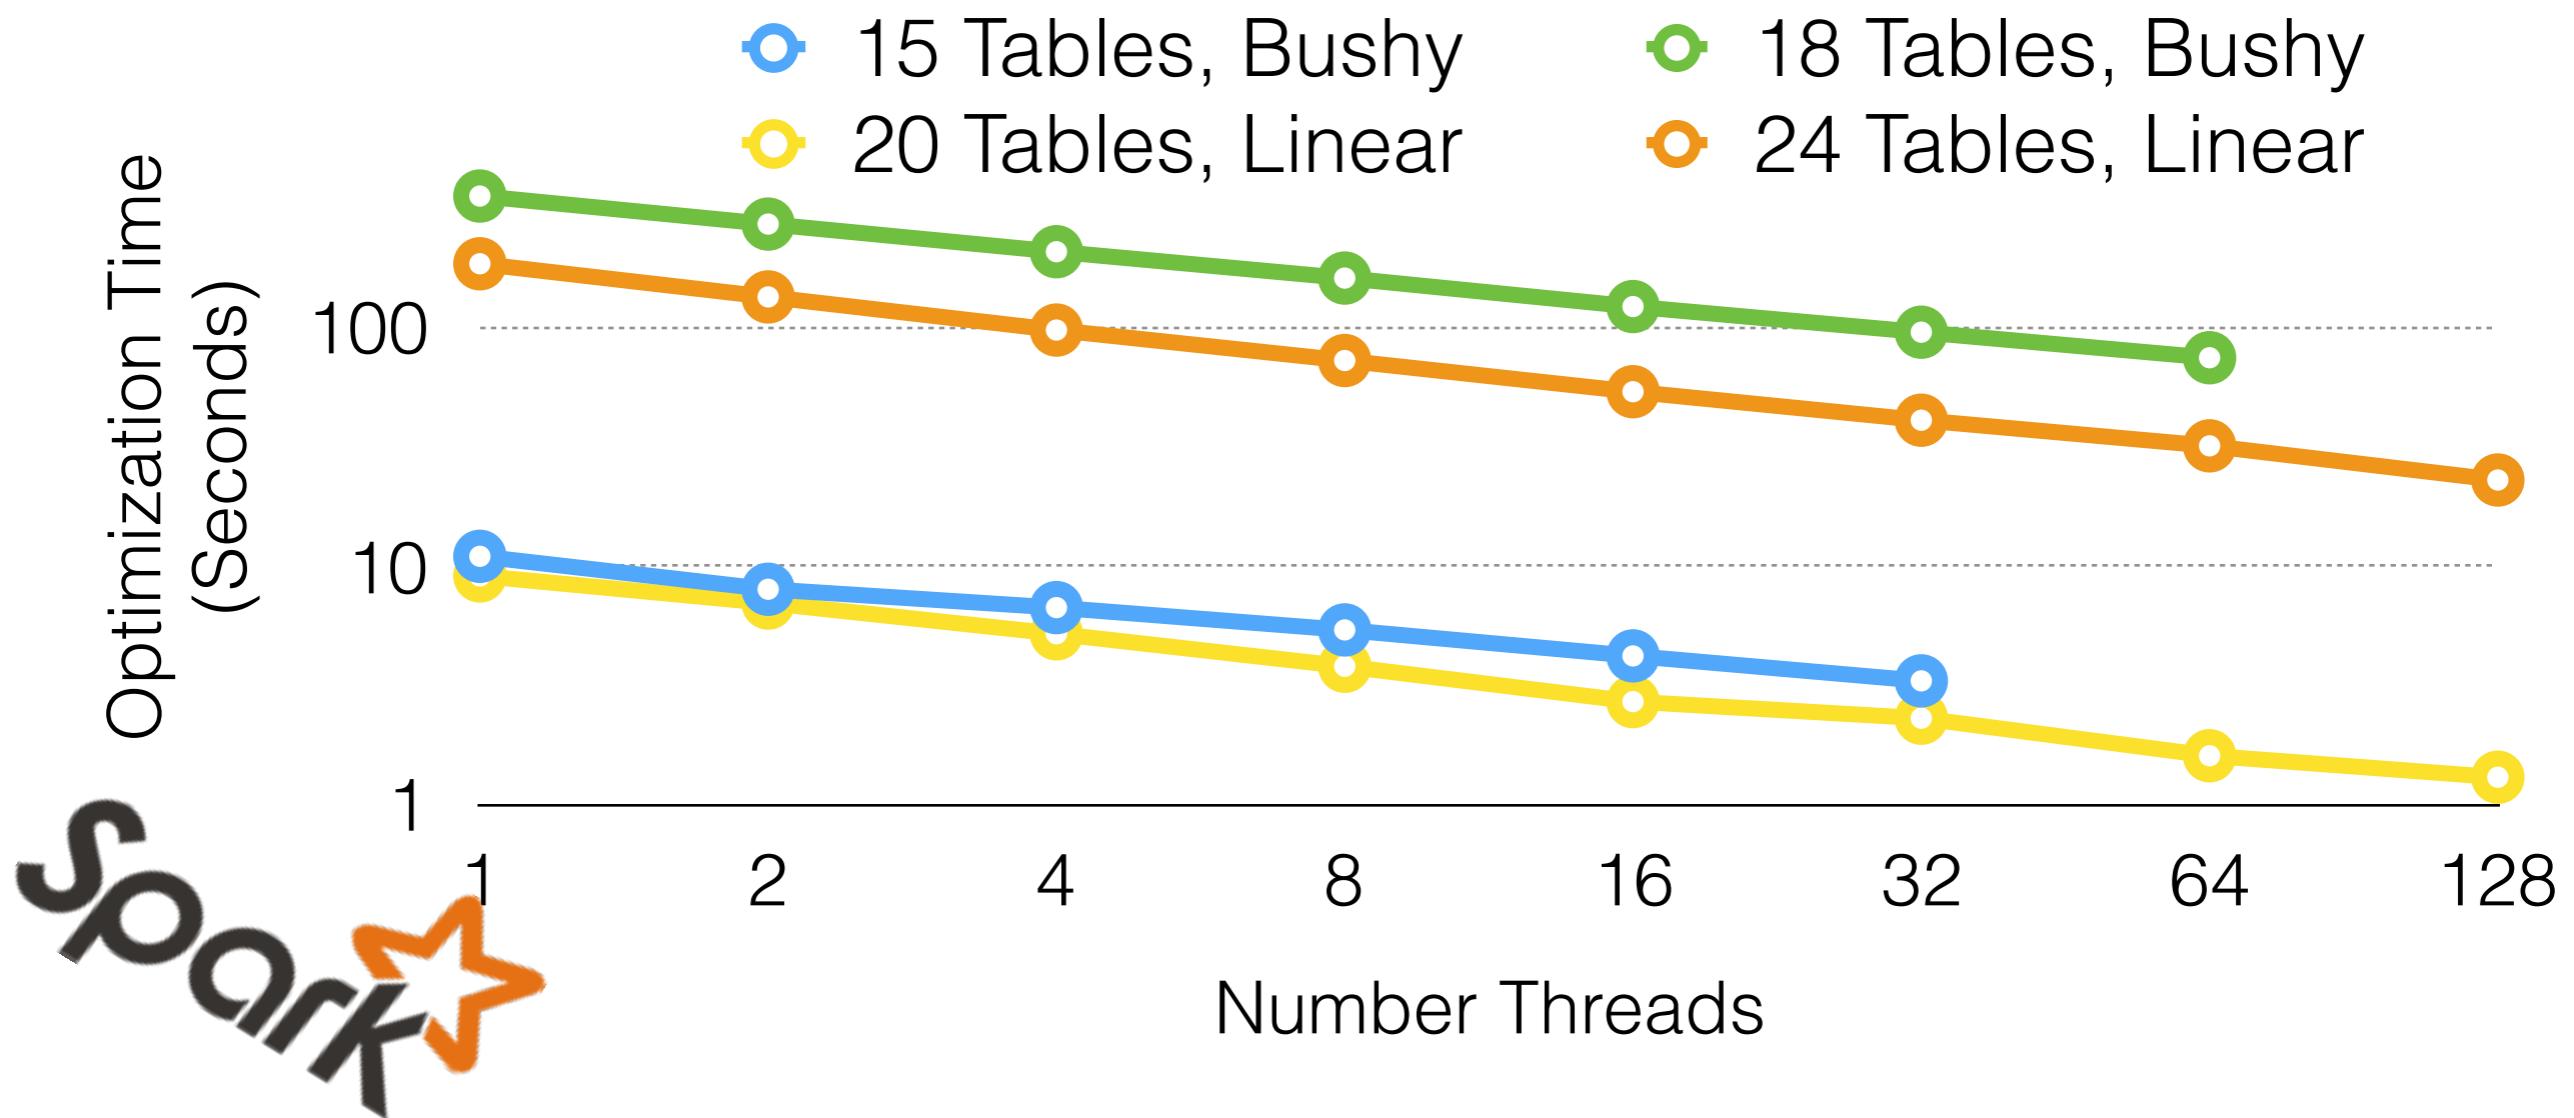
| Metric \ Plan Space | Left-Deep | Bushy |
|---------------------|---|---|
| Bytes Sent | $nrWorkers * (querySize + planSize)$ | |
| Time - Master | | $nrWorkers^2$ |
| Memory - Worker | $Optimal!$ $2^{querySize} * (3/4)^{nrWorkers}$ | $Optimal!$ $2^{querySize} * (7/8)^{nrWorkers}$ |
| Time - Worker | $2^{querySize} * (3/4)^{nrWorkers} * querySize$ | $3^{querySize} * (21/27)^{nrWorkers}$ |

(Considering multiple metrics multiplies by polynomial factor)

Experimental Results

Experimental Setup

Randomly generated queries; one execution cost metric; approximated optimization
Spark 1.5 on Yarn 2.7; 100 nodes with 2x Intel Xeon 2.6 GHz, 128 GB memory

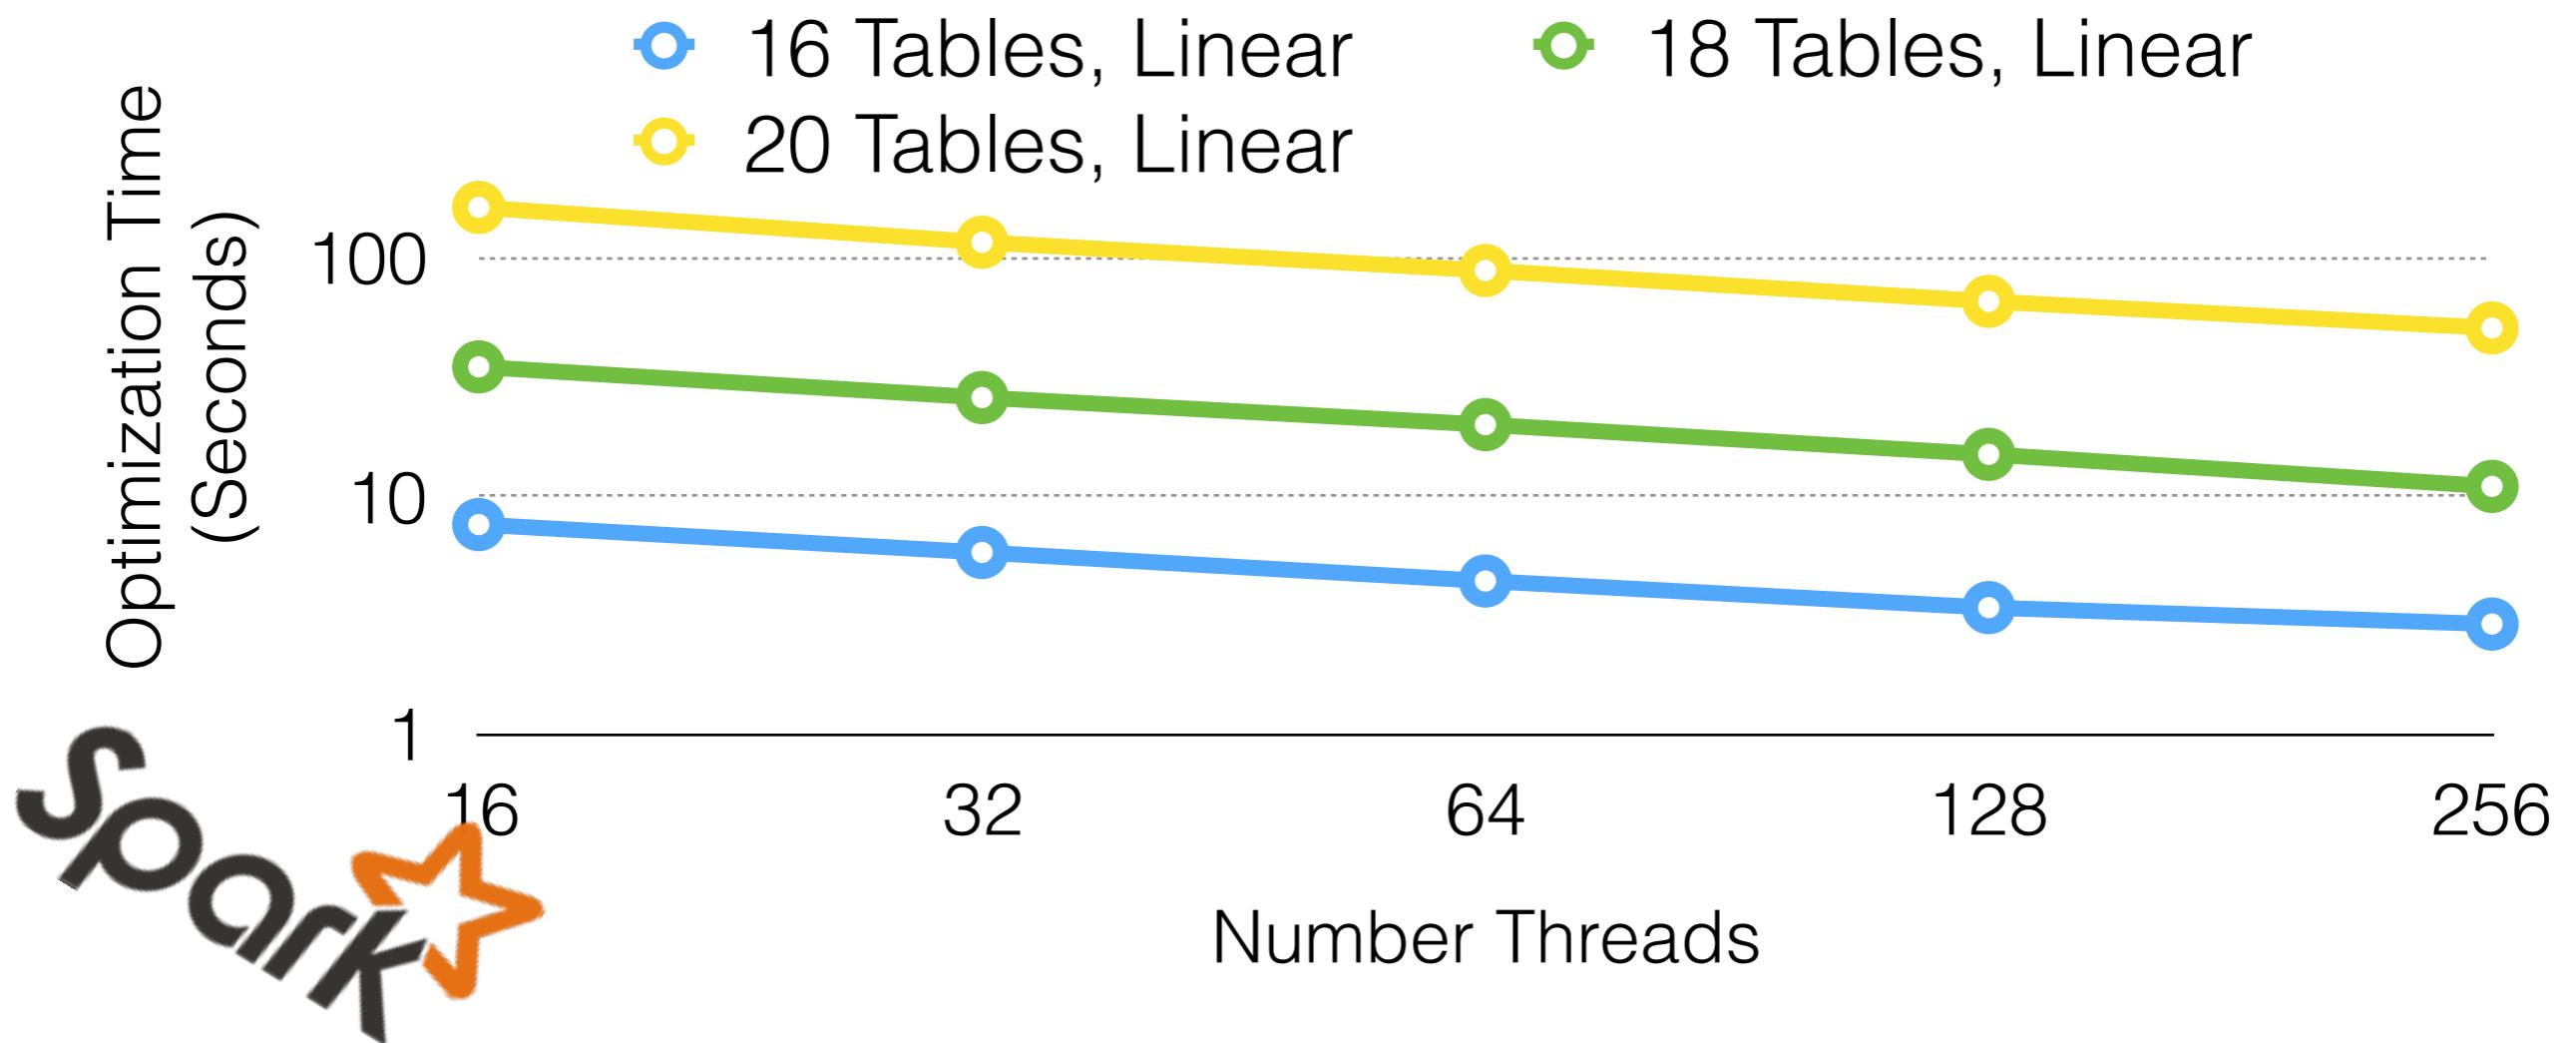


Scalability of new parallel query optimization algorithm

Experimental Results

Experimental Setup

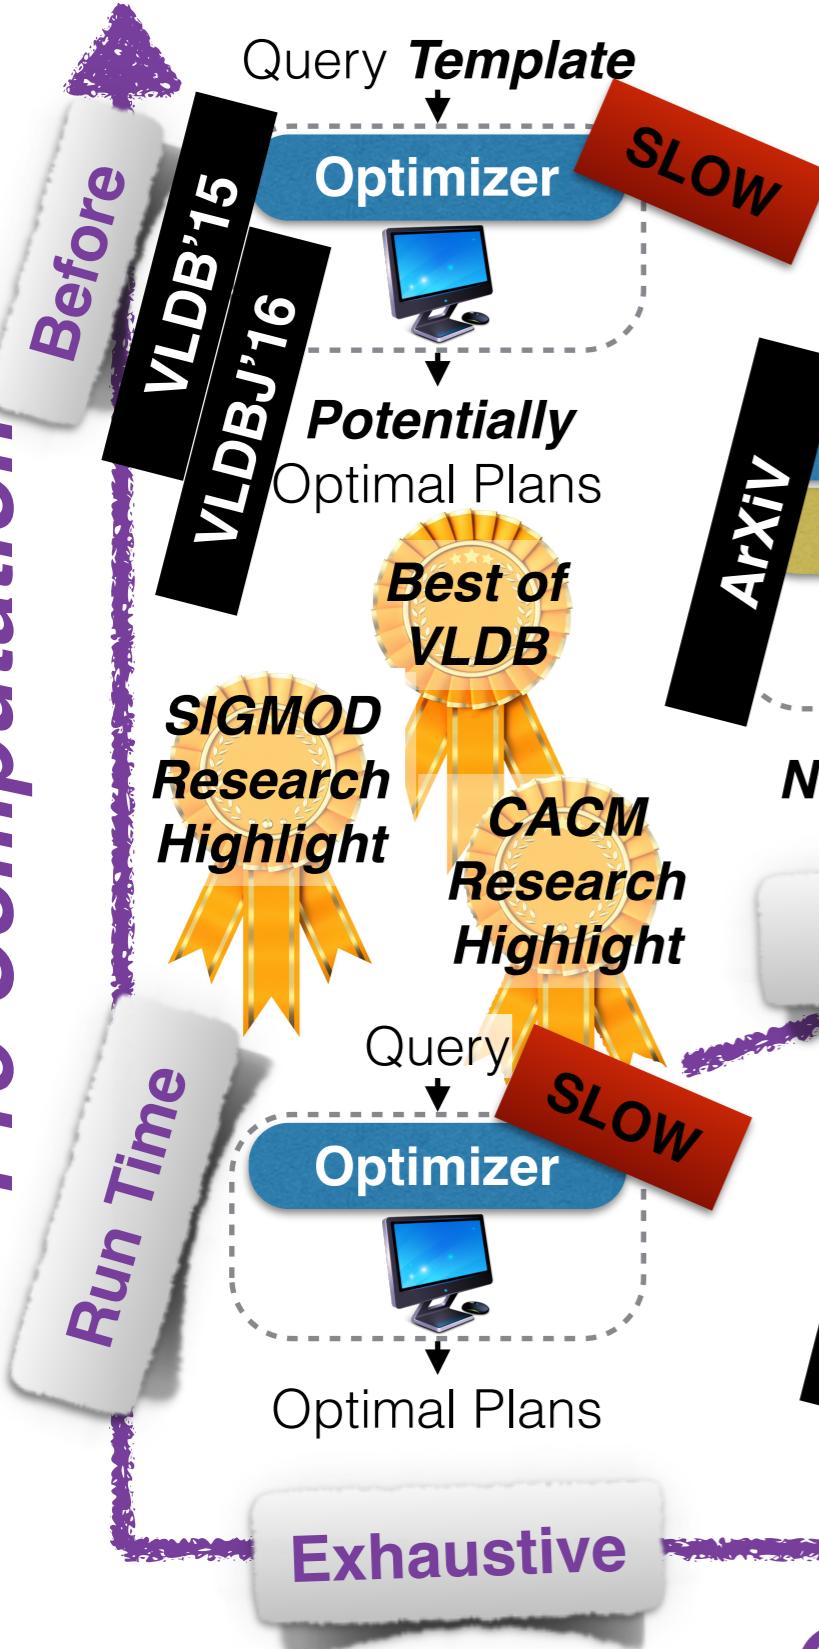
Randomly generated queries; two execution cost metrics; approximated optimization
Spark 1.5 on Yarn 2.7; 100 nodes with 2x Intel Xeon 2.6 GHz, 128 GB memory



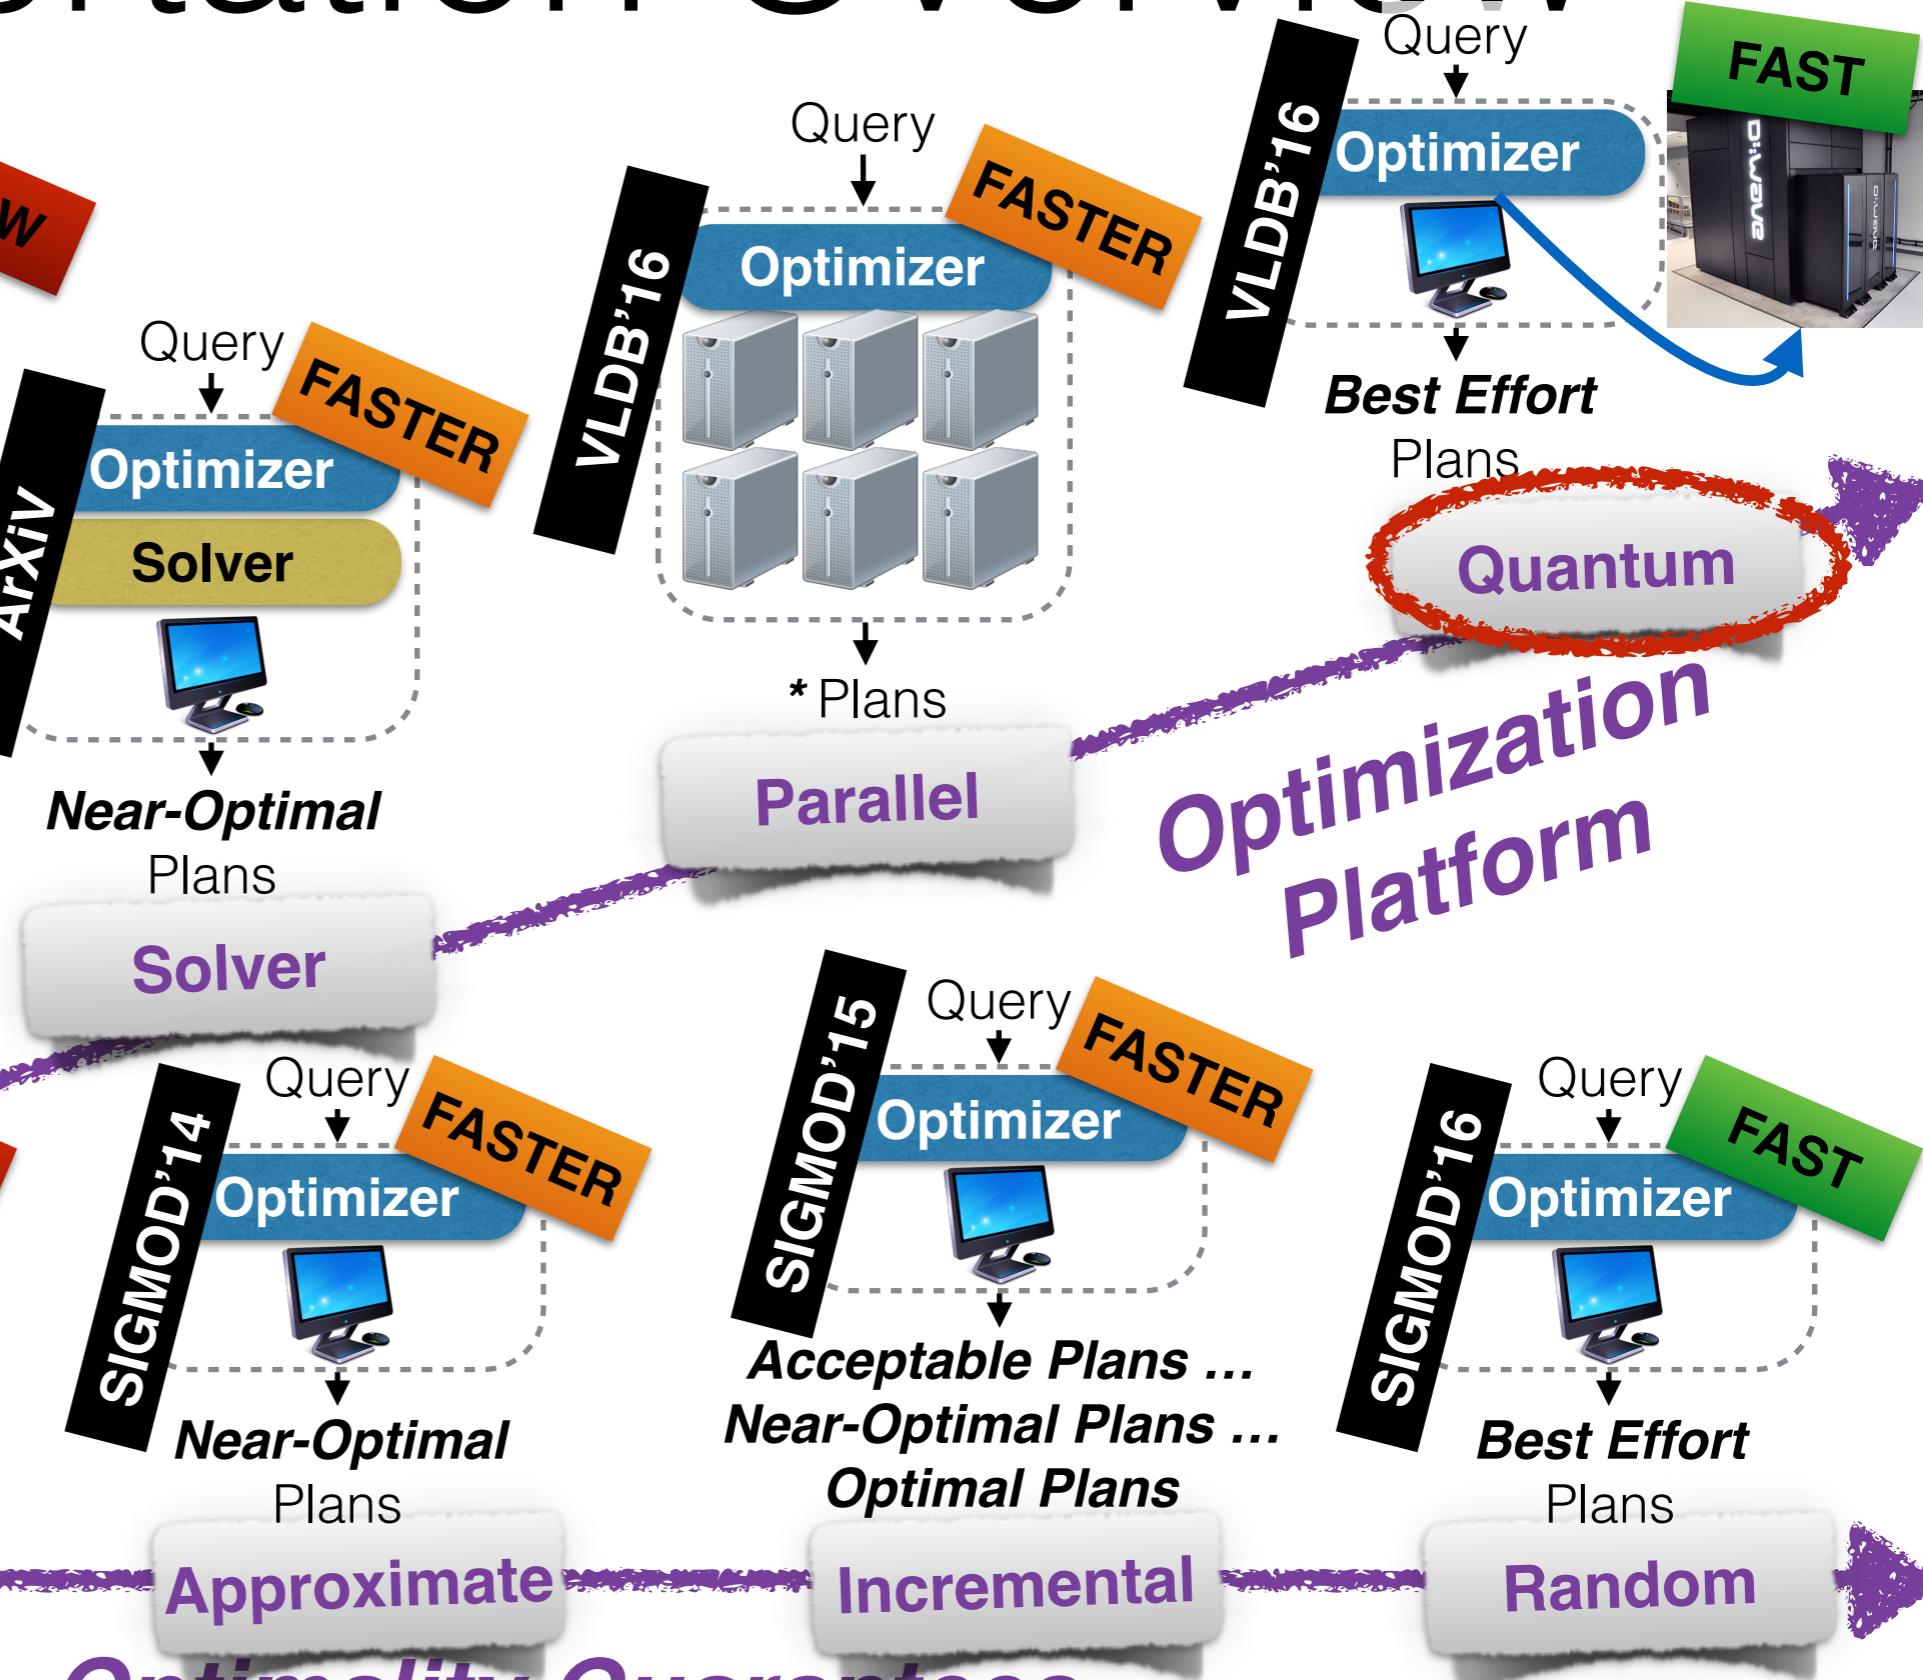
Scalability of new parallel query optimization algorithm

Dissertation Overview

Pre-Computation



Optimality Guarantees



Quantum Computing

Classical Computer

Quantum Computer

Bits

Qubits

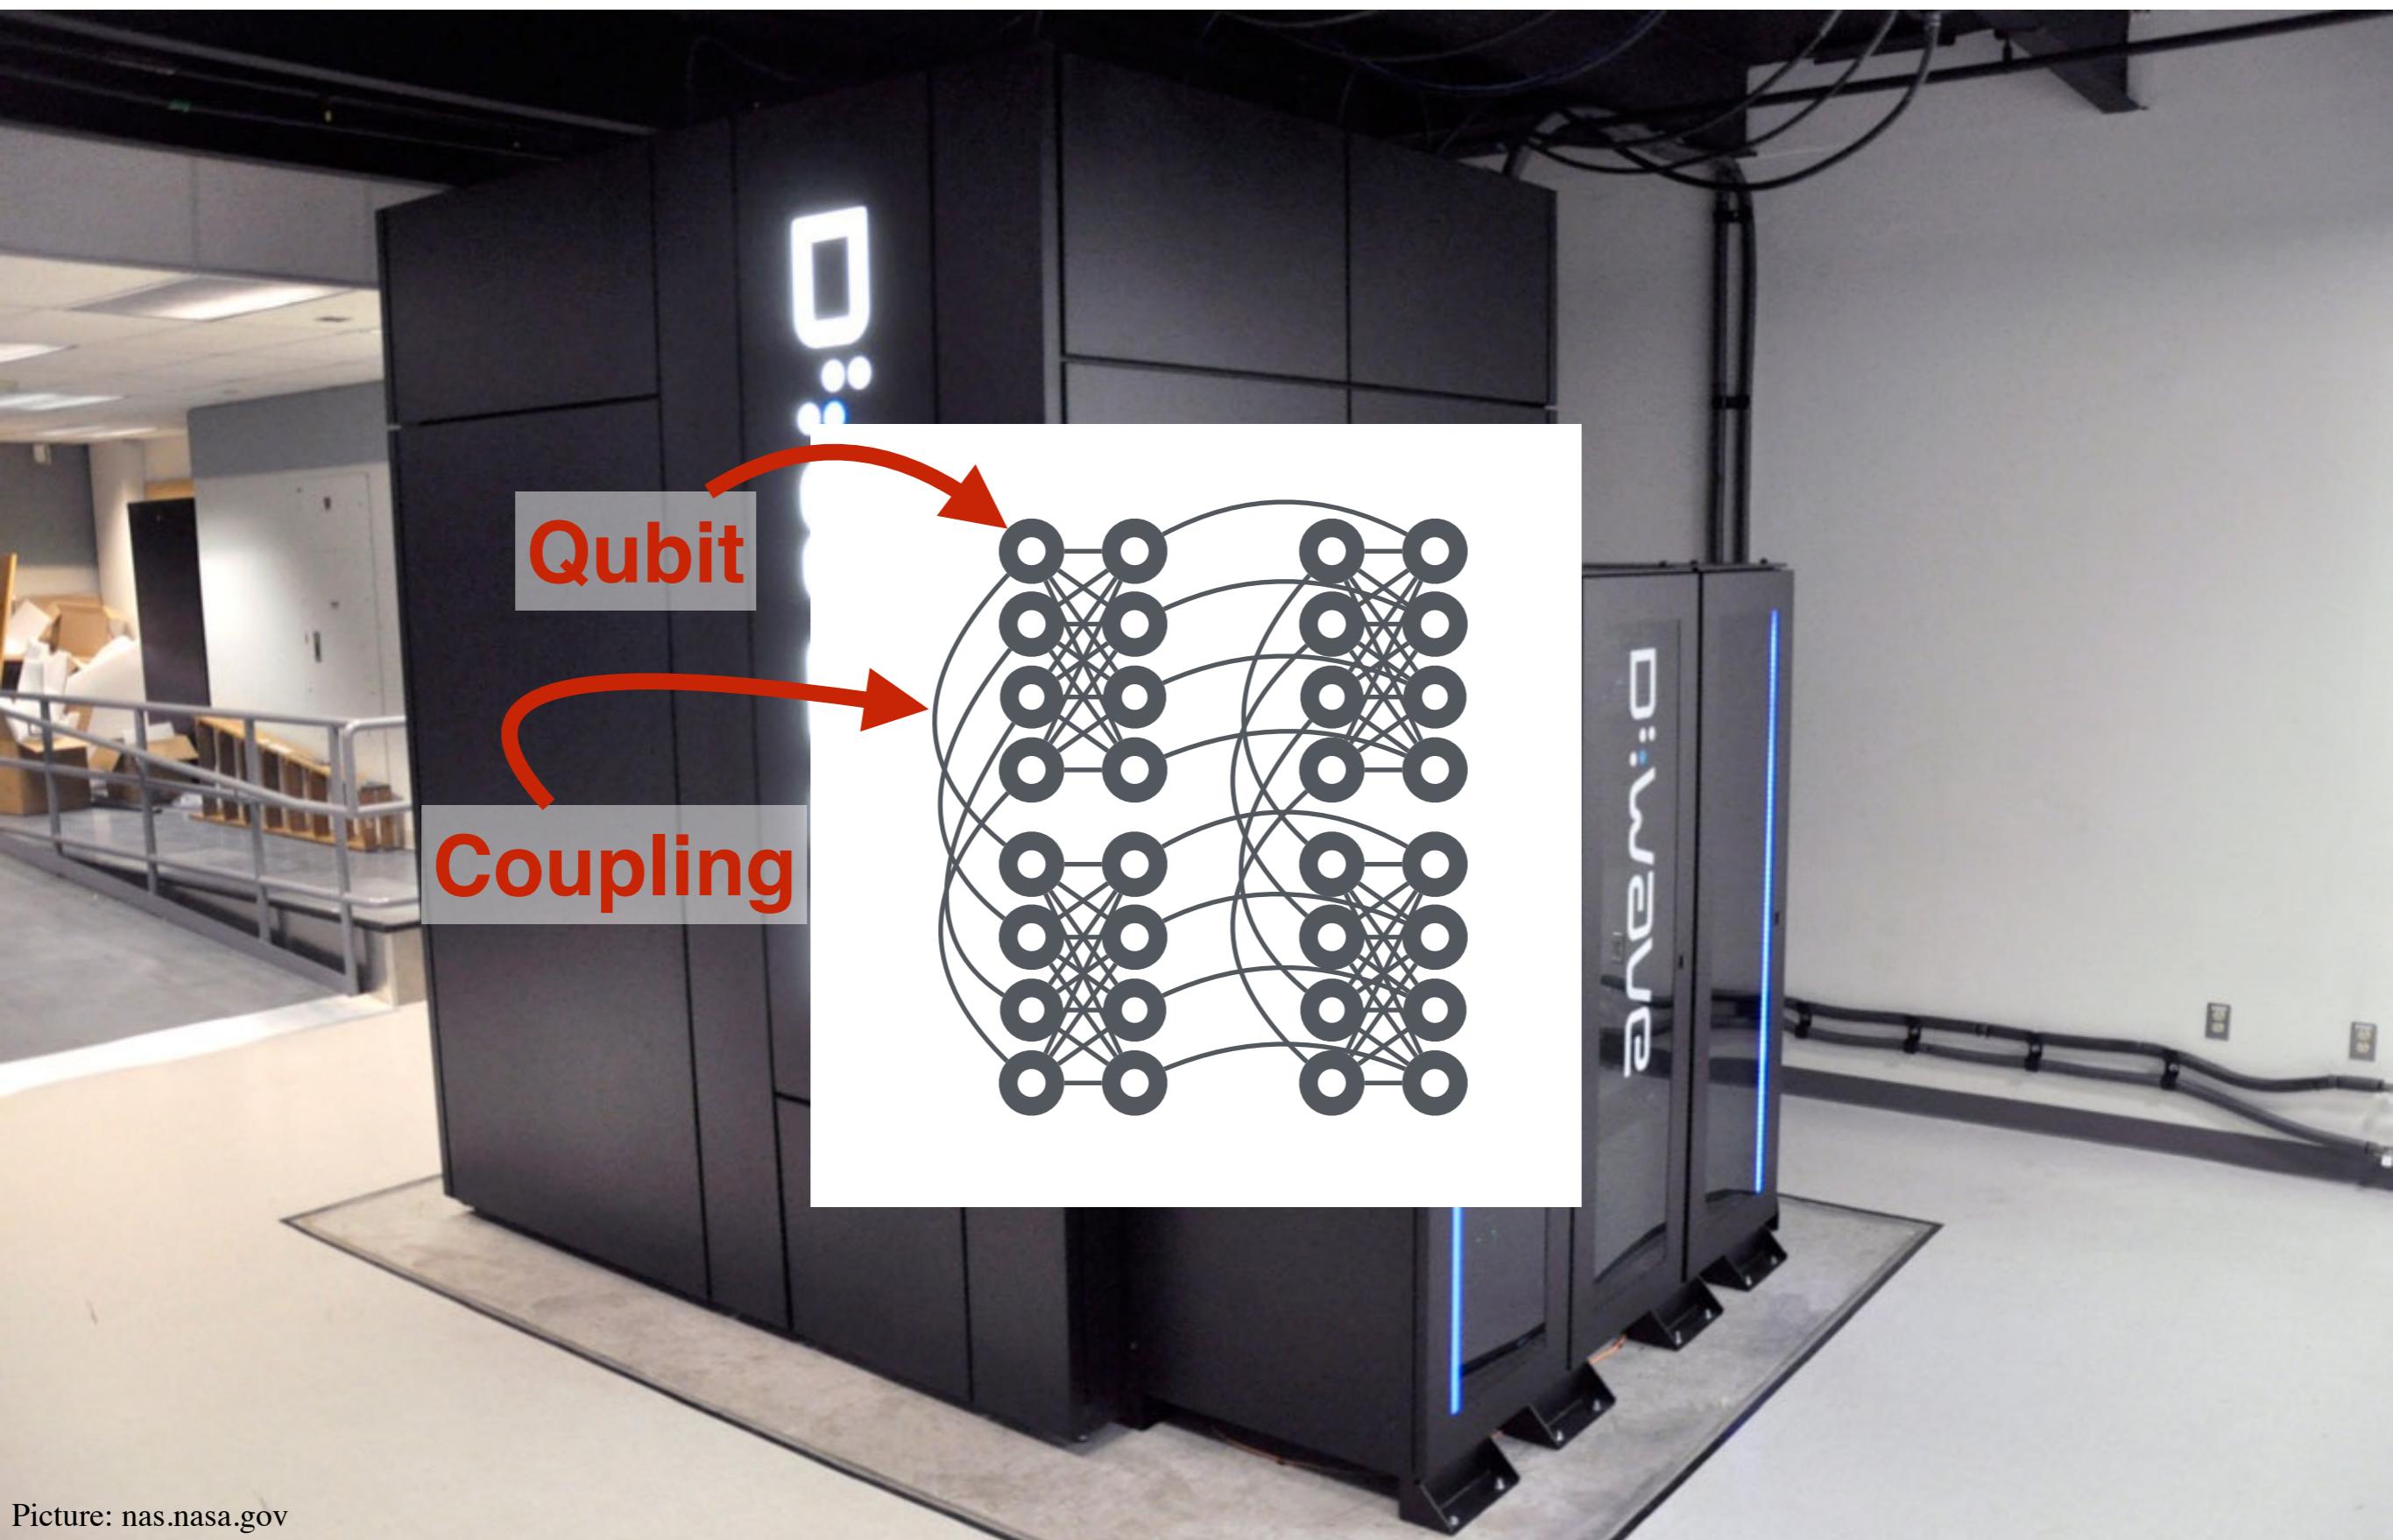
1 or 0

1 and 0

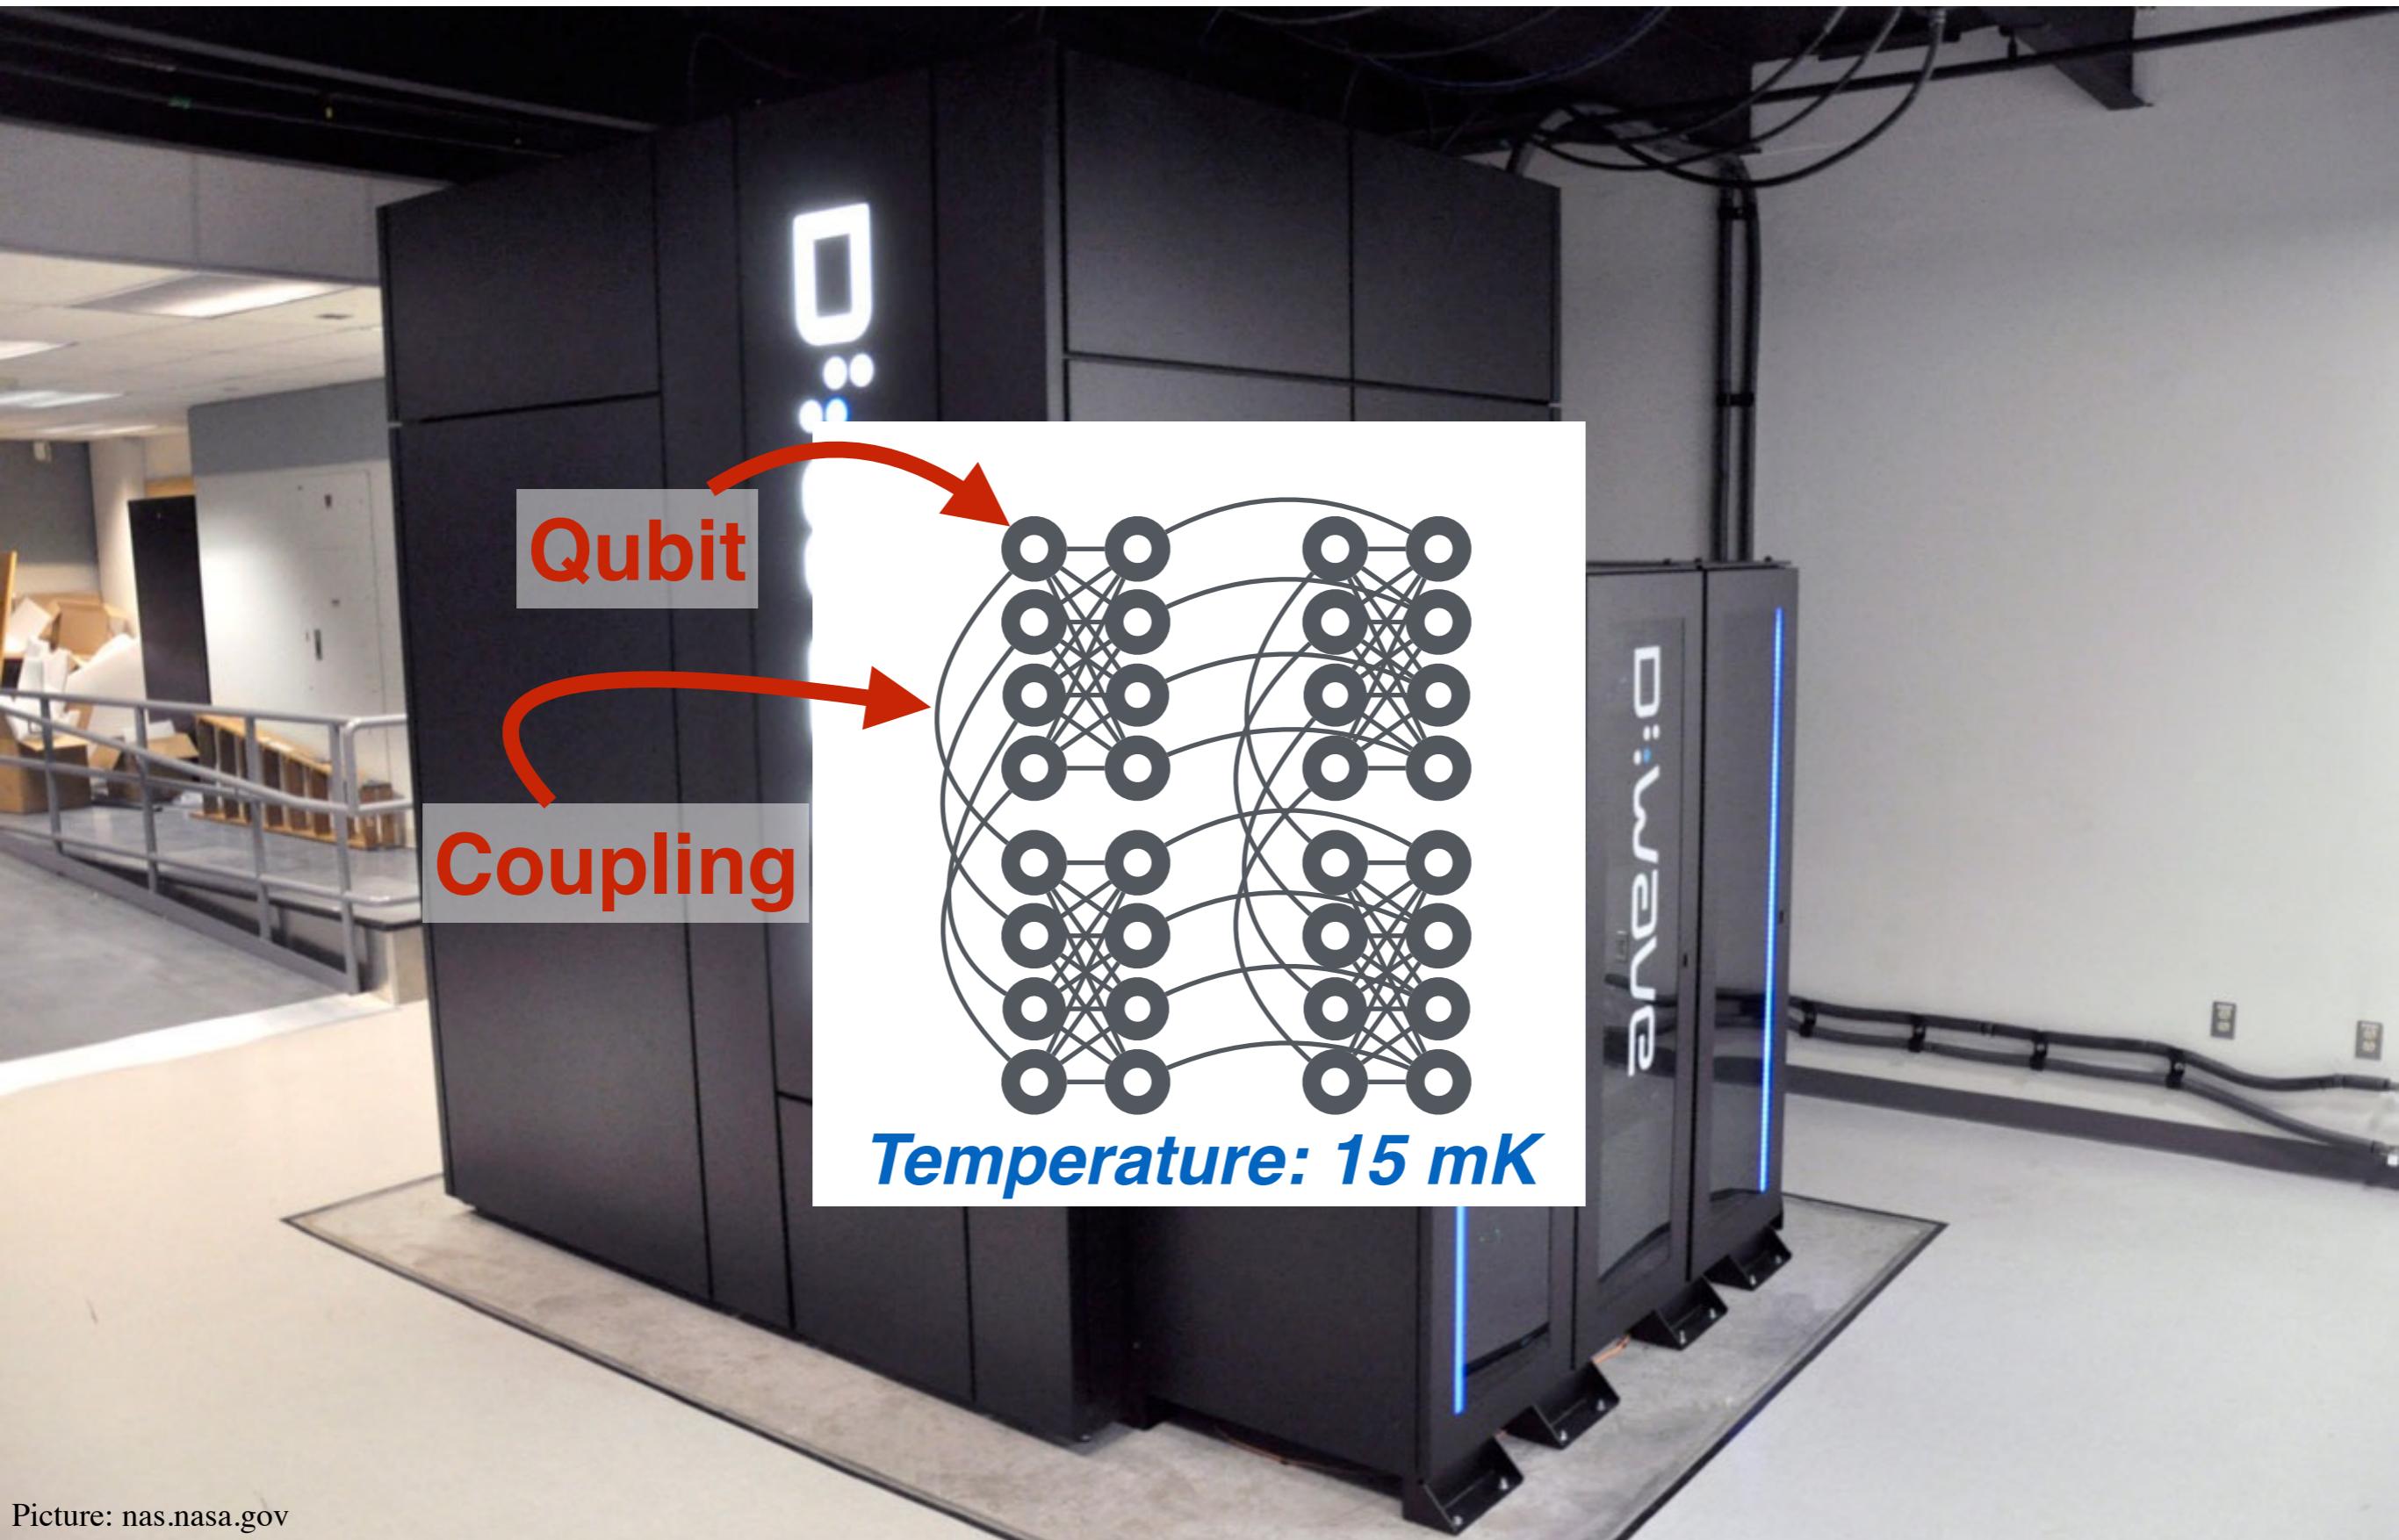
D-Wave Quantum Annealer



D-Wave Quantum Annealer



D-Wave Quantum Annealer



Comments on D-Wave

The general consensus now is that the D-Wave computers do exhibit some quantum behavior.
5.2014

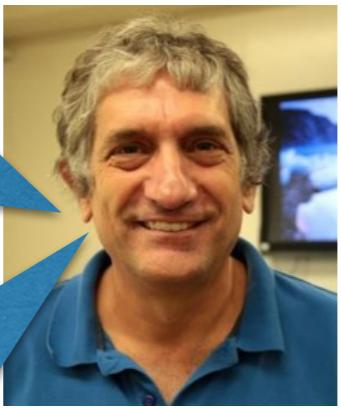
10,000,000 times faster than simulated annealing running on a single core. 12.2015

it remains uncertain whether D-Wave's approach will lead to speedups over the best known classical algorithms. 12.2015

Another challenge is the embedding of problems into the native hardware architecture 12.2015

more efficient classical optimization algorithms exist for these problems, which outperform the current D-Wave 2X device for most problem instances. 12.2015

As far as I'm concerned, this completely nails down the case for computationally-relevant collective quantum tunneling in the D-Wave machine, at least within the 8-qubit clusters 12.2015



John Martinis
Google/UCSB



Scott Aaronson
MIT



Matthias Troyer
ETH

Experiments on the D-Wave devices over the past years have indicated that the devices use quantum effects. 12.2015

The sweet spot for quantum annealers is thus quickly finding approximate solutions. 12.2015

D-Wave has performed a careful study showing that their latest device is 15x faster than the best optimized classical codes on a single core of an Intel CPU (and thus still comparable to a high-end multi-core CPU) 12.2015

if you wanted to solve a practical optimization problem, you'd first need to encode it into the Chimera graph, and that reduction entails a loss 12.2015

Comments on D-Wave

The general consensus now is that the D-Wave computers do exhibit some quantum behavior.
5.2014

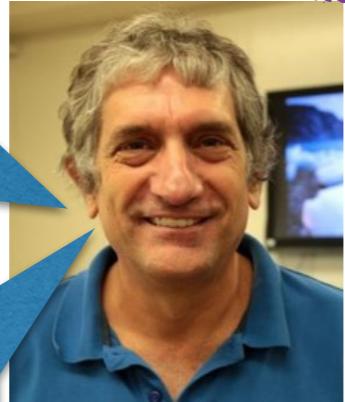
10,000,000 times faster than simulated annealing running on a single core. 12.2015

it remains uncertain whether D-Wave's approach will lead to speedups over the best known classical algorithms. 12.2015

Another challenge is the embedding of problems into the native hardware architecture 12.2015

more efficient classical optimization algorithms exist for these problems, which outperform the current D-Wave 2X device for most problem instances. 12.2015

As far as I'm concerned, this completely nails down the case for computationally-relevant collective quantum tunneling in the D-Wave machine, at least within the 8-qubit clusters 12.2015



John Martinis
Google/UCSB



Scott Aaronson
MIT



Matthias Troyer
ETH

Experiments on the D-Wave devices over the past years have indicated that the devices use quantum effects. 12.2015

The sweet spot for quantum annealers is thus quickly finding approximate solutions. 12.2015

D-Wave has performed a careful study showing that their latest device is 15x faster than the best optimized classical codes on a single core of an Intel CPU (and thus still comparable to a high-end multi-core CPU) 12.2015

if you wanted to solve a practical optimization problem, you'd first need to encode it into the Chimera graph, and that reduction entails a loss 12.2015

Comments on D-Wave

The general consensus now is that the D-Wave computers do exhibit some quantum behavior.
5.2014

As far as I'm concerned, this completely nails down the case for computationally-relevant collective quantum tunneling in the D-Wave machine, at least within the 8 qubit cluster. 12.2015



Experiments on the D-Wave devices over the past years have indicated that the devices use quantum effects. 12.2015

10,000,000 times faster than simulated annealing running on a single core. 12.2015

John Martinis
Google/UCSB

Scott Aaronson
MIT

Matthias Troyer
ETH

The sweet spot for quantum annealers is thus quickly finding approximate solutions. 12.2015

it remains uncertain whether D-Wave's approach will lead to speedups over the best known classical algorithms. 12.2015

Another challenge is the embedding of problems into the native hardware architecture 12.2015

more efficient classical optimization algorithms exist for these problems, which outperform the current D-Wave 2X device for most problem instances. 12.2015

D-Wave has performed a careful study showing that their latest device is 15x faster than the best optimized classical codes on a single core of an Intel CPU (and thus still comparable to a high-end multi-core CPU) 12.2015

if you wanted to solve a practical optimization problem, you'd first need to encode it into the Chimera graph, and that reduction entails a loss 12.2015

Comments on D-Wave

The general consensus now is that the D-Wave computers do exhibit some quantum behavior.

5.2014

As far as I'm concerned, this completely nails down the case for computationally-relevant collective quantum tunneling in the D-Wave machine, at least within the 8 qubit cluster. 12.2015

Quantum Effects

Experiments on the D-Wave devices over the past years have indicated that the devices use quantum effects. 12.2015

10,000,000 times faster than simulated annealing running on a single core. 12.2015

it remains uncertain whether D-Wave's approach will lead to speedups over the best known classical algorithms. 12.2015

Another challenge is the embedding of problems into the native hardware architecture. 12.2015

more efficient classical optimization algorithms exist for these problems, which outperform the current D-Wave 2X device for most problem instances. 12.2015



John Martinis

Scott Aaronson

MIT

Matthias Troyer

ETH

D-Wave has performed a careful study showing that their latest device is 15x faster than the best optimized classical code on a single core of an Intel CPU (and thus still comparable to a high-end multi-core CPU) 12.2015

if you wanted to solve a practical optimization problem, you'd first need to encode it into the Chimera graph, and that reduction entails a loss 12.2015

Often Beaten at Finding Optimum Faster than Single-Core

Comments on D-Wave

The general consensus now is that the D-Wave computers do exhibit some quantum behavior.
5.2014

10,000,000 times faster than simulated annealing running on a single core. 12.2015

it remains uncertain whether D-Wave's approach will lead to speedups over best known classical algorithms. 12.2015

Another interesting embedding of problems into the D-Wave hardware research. 12.2015

more efficient classical optimization algorithms exist for these problems, which outperform the current D-Wave 2X device for most problem instances. 12.2015

As far as I'm concerned, this completely nails down the case for computationally-relevant collective quantum tunneling in the D-Wave machine, at least within the 8 qubit cluster. 12.2015

Experiments on the D-Wave devices over the past years have indicated that the devices use quantum effects. 12.2015

The sweet spot for quantum annealing thus quickly approaches. 12.2015

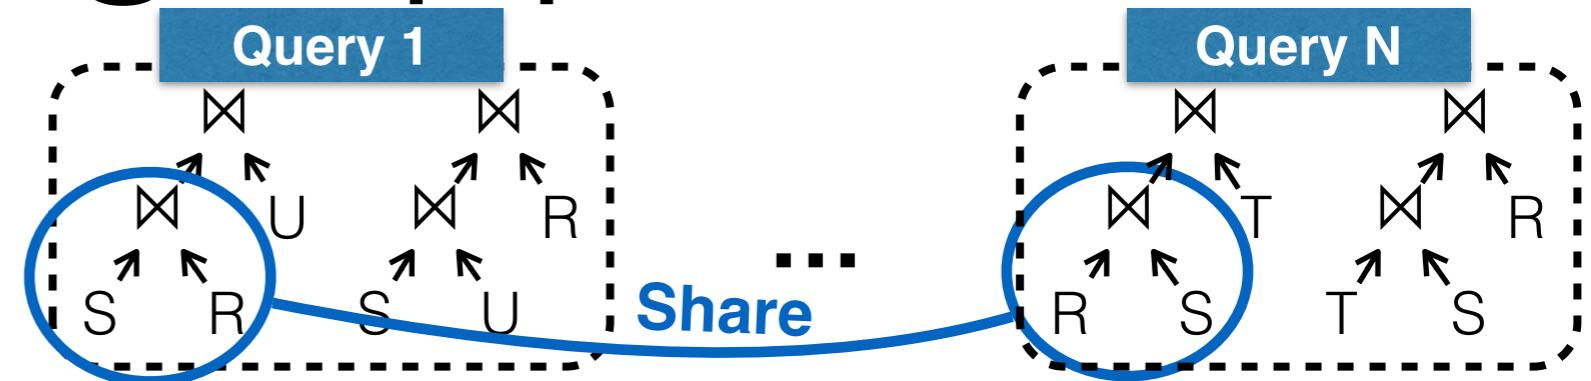
D-Wave has performed a careful study showing that their best device is 15x faster than on a single core of an Intel CPU thus still comparable to a high-end multi-core CPU) 12.2015

if you wanted to solve a practical optimization problem, you'd first need to encode it into the Chimera graph, and that reduction entails a loss 12.2015

Off Finding Beaten Quantum Effects
How About Relevant Problems??
Faster at Approximation

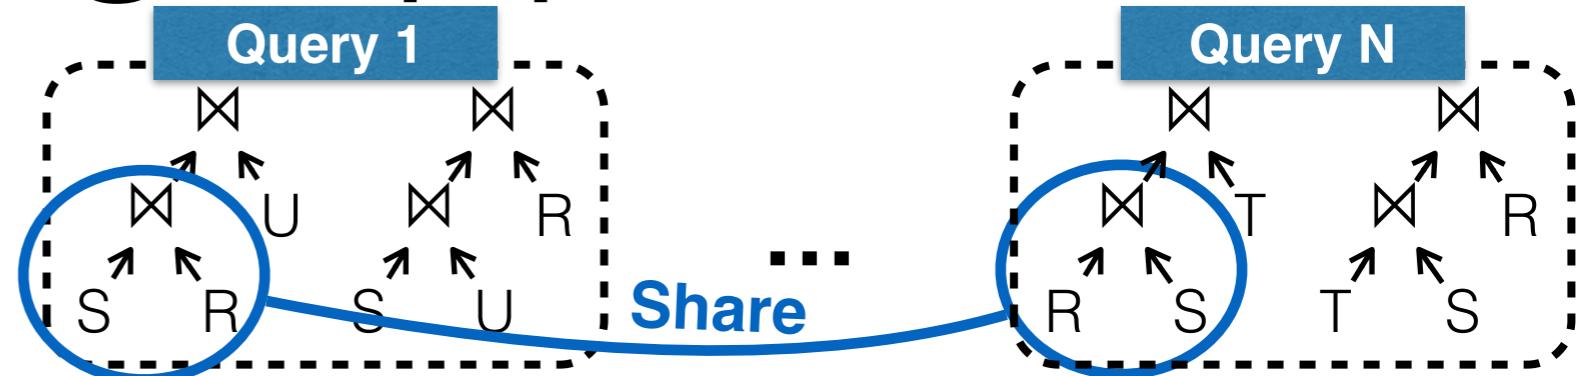
Mapping Approach

*Find Optimal
Plan Combination*



Mapping Approach

*Find Optimal
Plan Combination*



1

Transform → Quadratic Formula

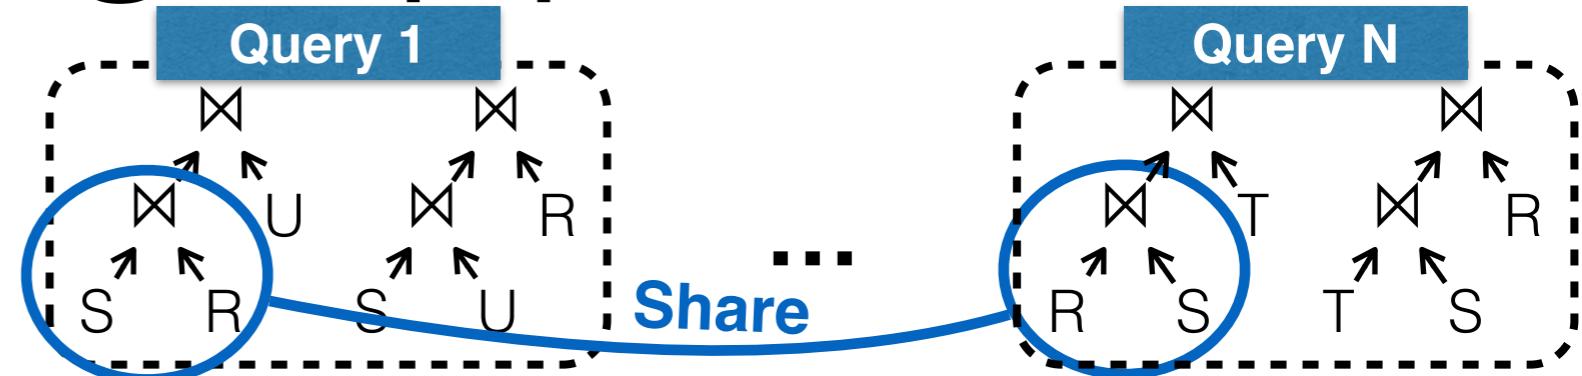
Binary Variables:

selectPlan₁ ... *selectPlan_n*

Cost Formula:
costTerms *savingsTerms*
constraintTerms

Mapping Approach

*Find Optimal
Plan Combination*



1

Transform → Quadratic Formula

Binary Variables:

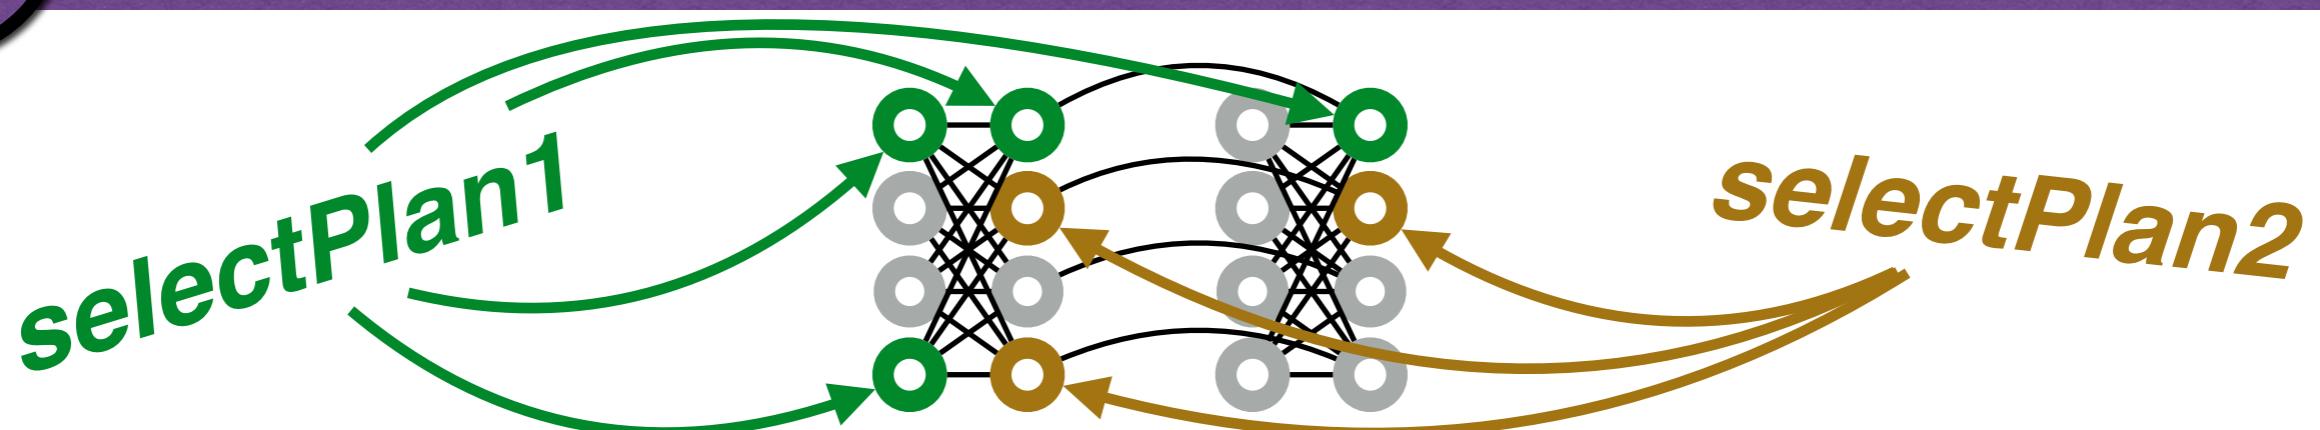
selectPlan₁ ... *selectPlan_n*

Cost Formula:

costTerms *savingsTerms*
constraintTerms

2

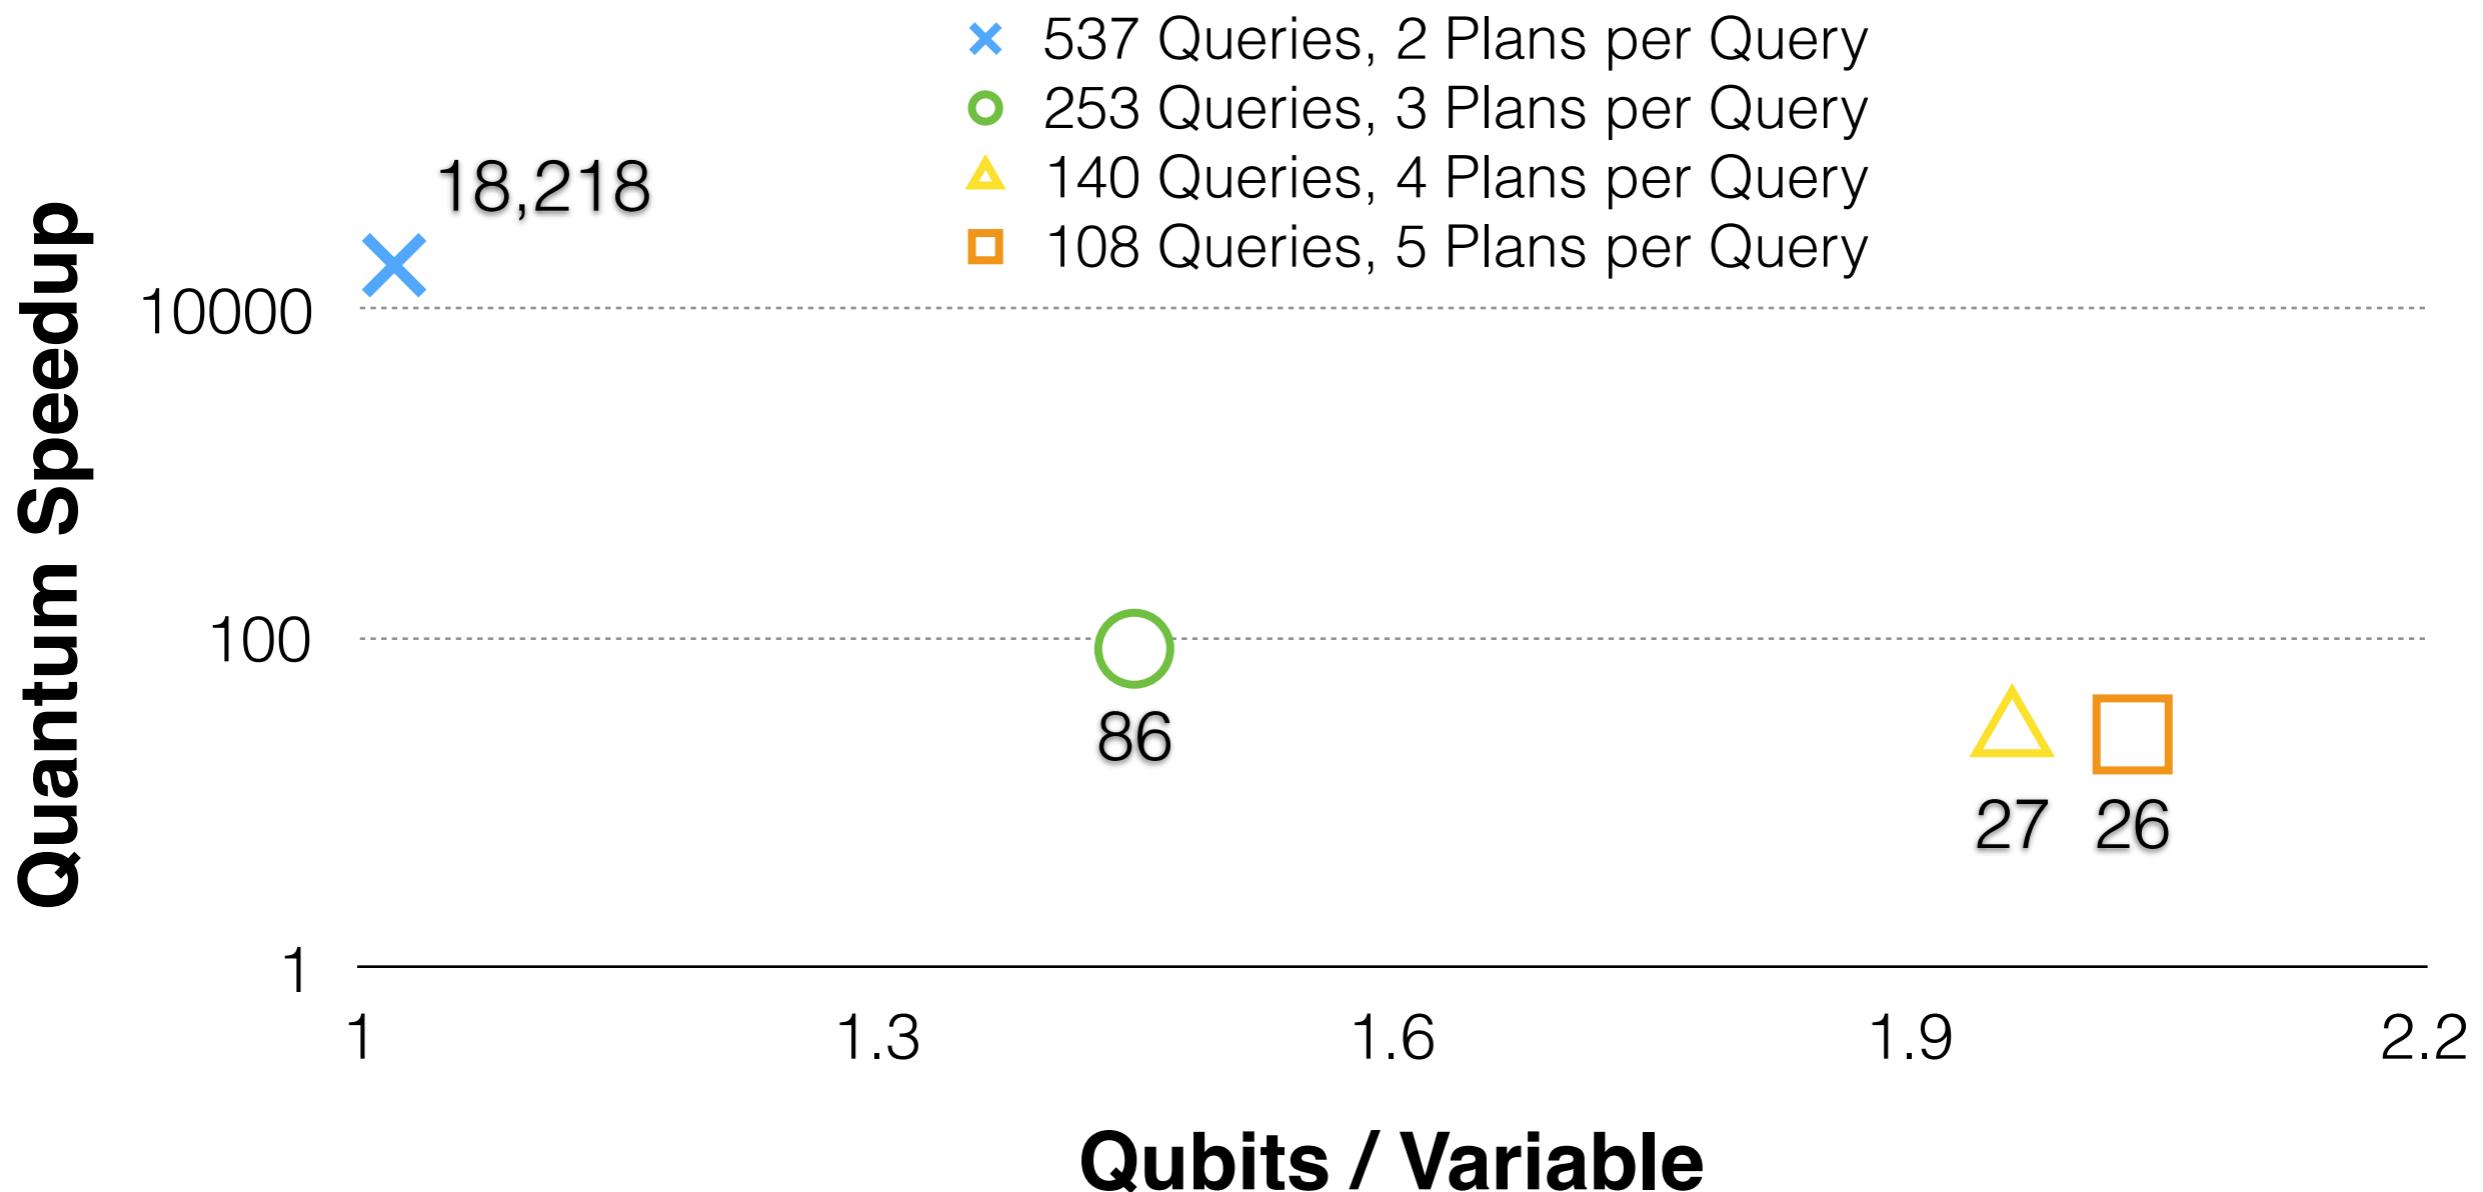
Transform → Weights on Qubits



Experimental Results

Experimental Setup

Compared D-Wave 2X vs. Integer Programming, Genetic Algorithms, Heuristics
Compared approximation time on randomly generated problem instances



More Details

Multiple Query Optimization on the D-Wave 2X Adiabatic Quantum Computer

Immanuel Trummer and Christoph Koch

{firstname}.{lastname}@epfl.ch

École Polytechnique Fédérale de Lausanne

VLDB '16

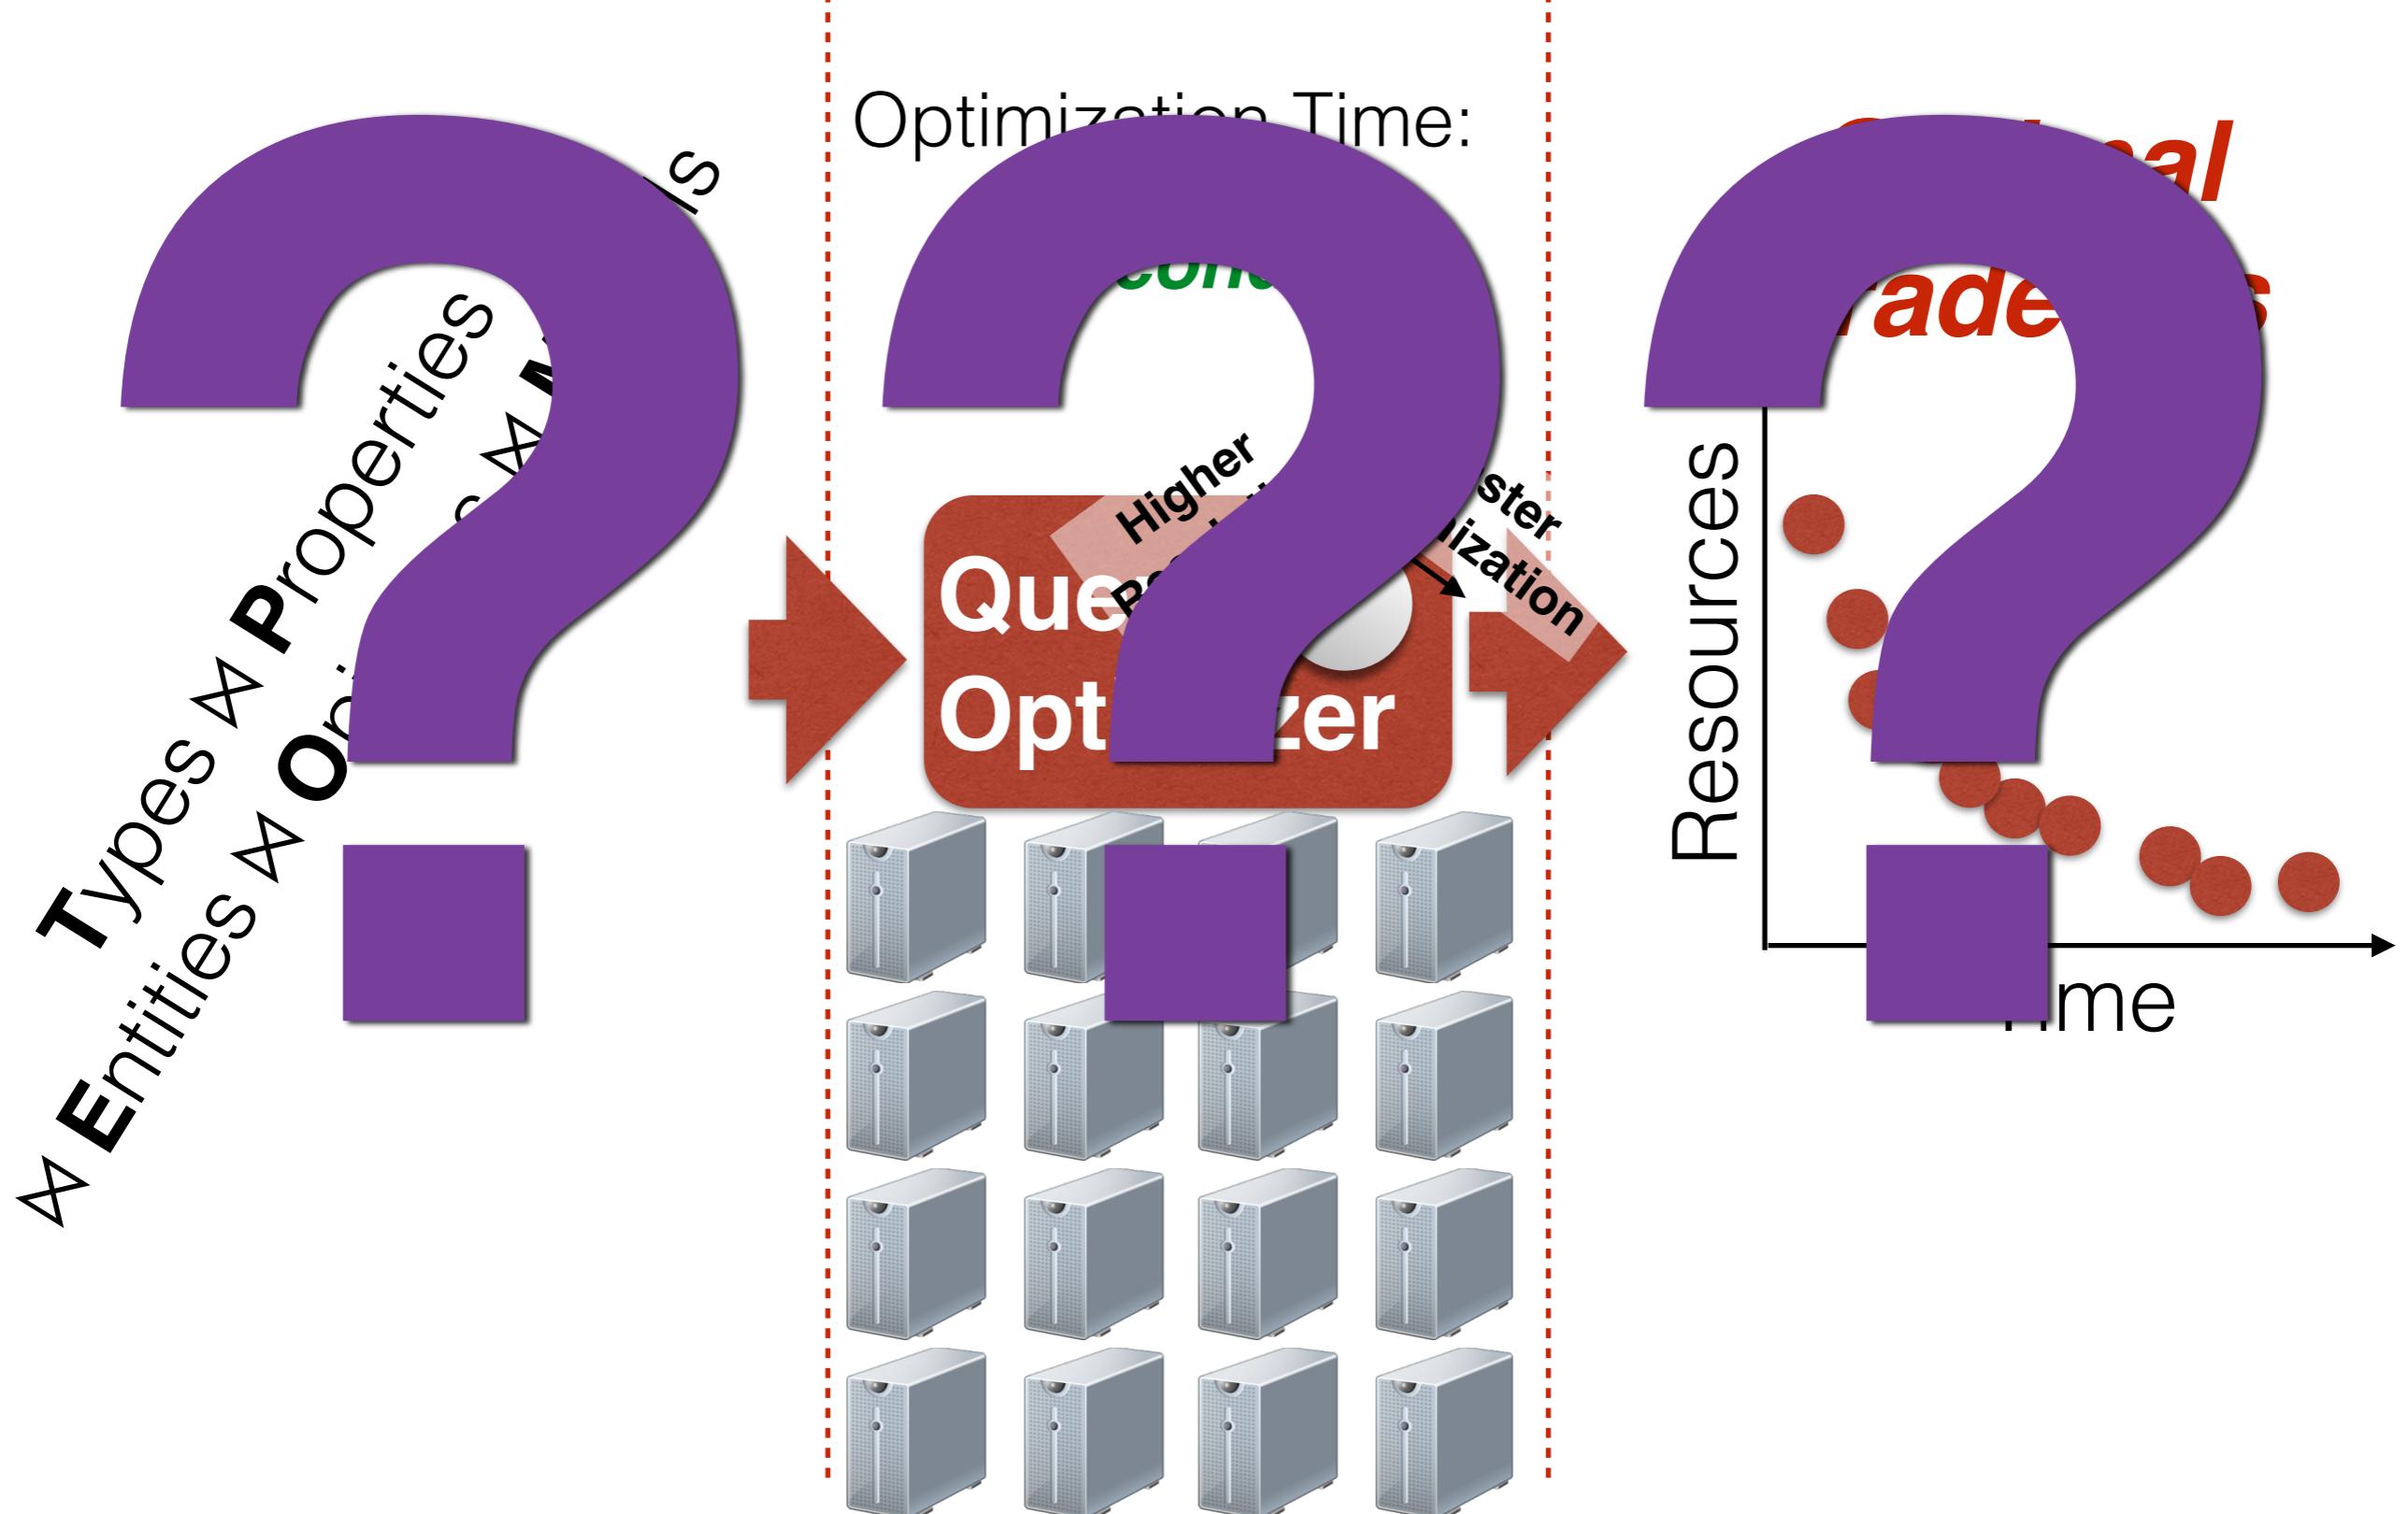
ABSTRACT

The D-Wave adiabatic quantum annealer solves hard combinatorial optimization problems leveraging quantum physics. The newest version features over 1000 qubits and was released in August 2015. We were given access to such a machine, currently hosted at NASA Ames Research Center in California, to explore the potential for hard optimization problems that arise in the context of databases.

In this paper, we tackle the problem of multiple query optimization (MQO). We show how an MQO problem instance can be transformed into a mathematical formula that

in California. This device is claimed to exploit the laws of quantum physics [6] in the hope to solve NP-hard optimization problems faster than traditional approaches. The machine supports a very restrictive class of optimization problems while it is for instance not capable of running Shor's algorithm [40] for factoring large numbers¹. We will show how instances of the multiple query optimization problem can be brought into a representation that is suitable as input to the quantum annealer. We also report results of an experimental evaluation that compares the time it takes to solve MQO problems on the quantum annealer to the time taken by algorithms that run on a traditional computer. We

Example



Mining Subjective Associations

Mining Subjective Associations

The Web

Mining Subjective Associations

I find kittens cute

Kittens-Cute

kittens
are cute

Palo Alto is not big

Palo Alto is a
big city

Palo Alto-Big

Palo Alto is a
small city

I don't find
jogging boring

Jogging-Boring

Jogging is so boring...

1

Extract Opinion Statements

The Web



Mining Subjective Associations

Cute Animals

Majority Opinion
 $\Pr(\dots)=0.9$
User Opinion
 $\Pr(\dots)=0.8$
User Post

Big Cities

Majority Opinion
 $\Pr(\dots)=0.85$
User Opinion
 $\Pr(\dots)=0.98$
User Post

Boring Sports

Majority Opinion
 $\Pr(\dots)=0.8$
User Opinion
 $\Pr(\dots)=0.95$
User Post

2

Learn User Behavior

I find kittens cute

Kittens-Cute

kittens
are cute

Palo Alto is not big

Palo Alto is a
big city

Palo Alto-Big

Palo Alto is a
small city

I don't find
jogging boring

Jogging-Boring

Jogging is so boring...

1

Extract Opinion Statements

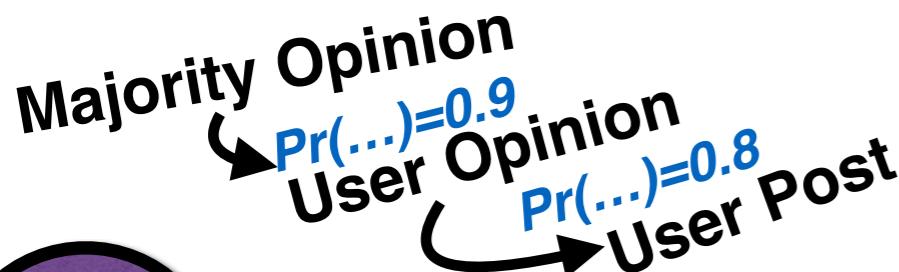
The Web

Mining Subjective Associations

3

Infer Majority Opinion

Cute Animals



Big Cities



Boring Sports



2

Learn User Behavior

I find kittens cute

Kittens-Cute

kittens are cute

Palo Alto is not big

Palo Alto is a big city

Palo Alto-Big

Palo Alto is a small city

I don't find jogging boring

Jogging-Boring

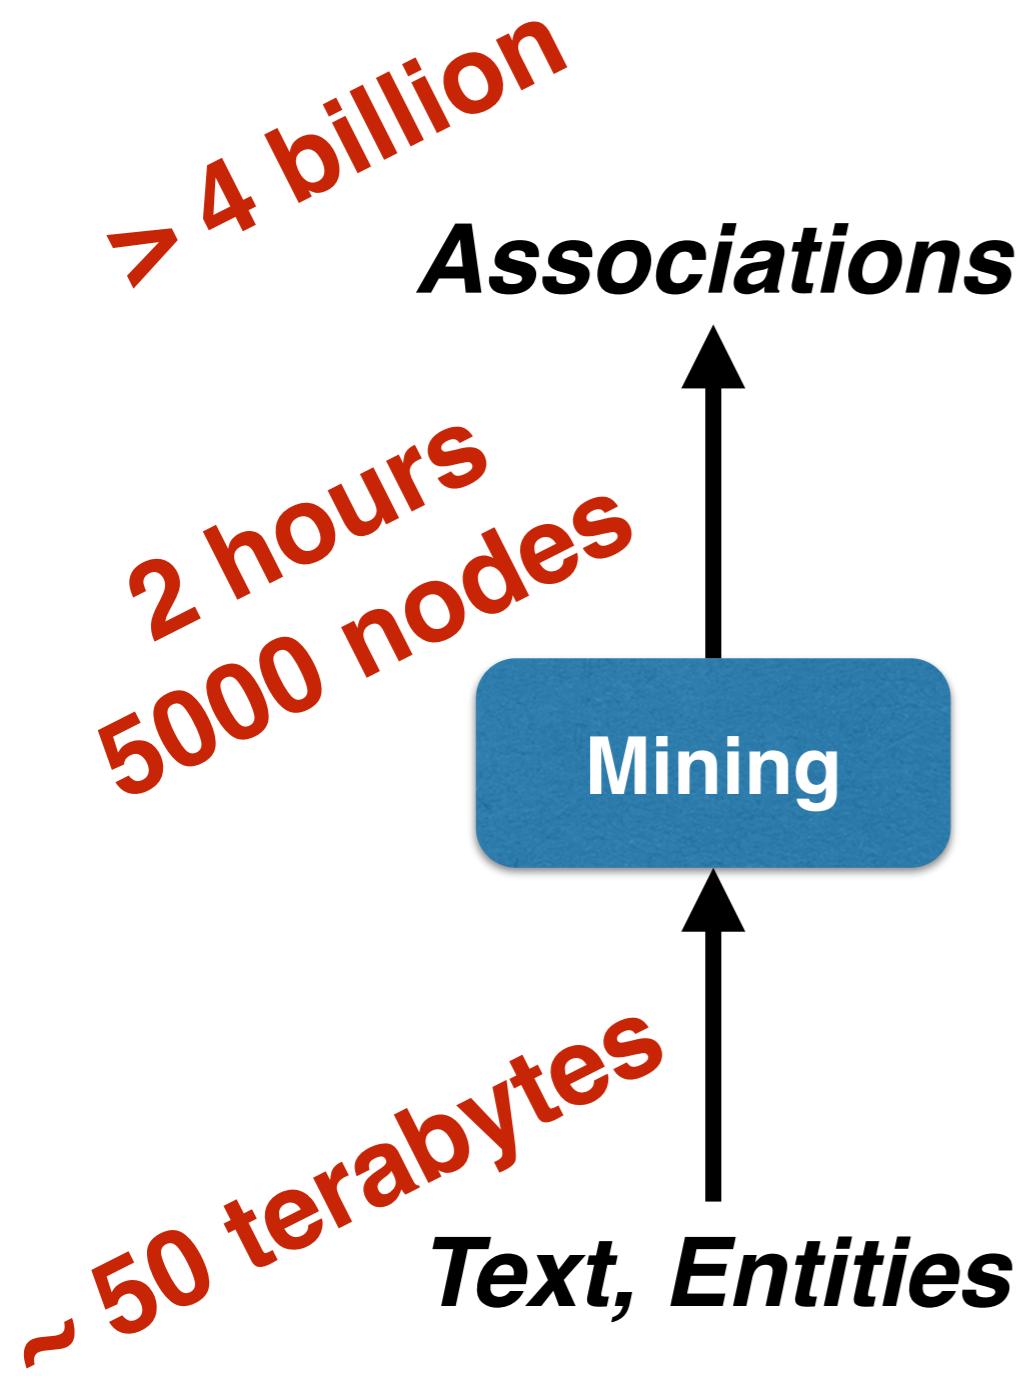
Jogging is so boring...

1

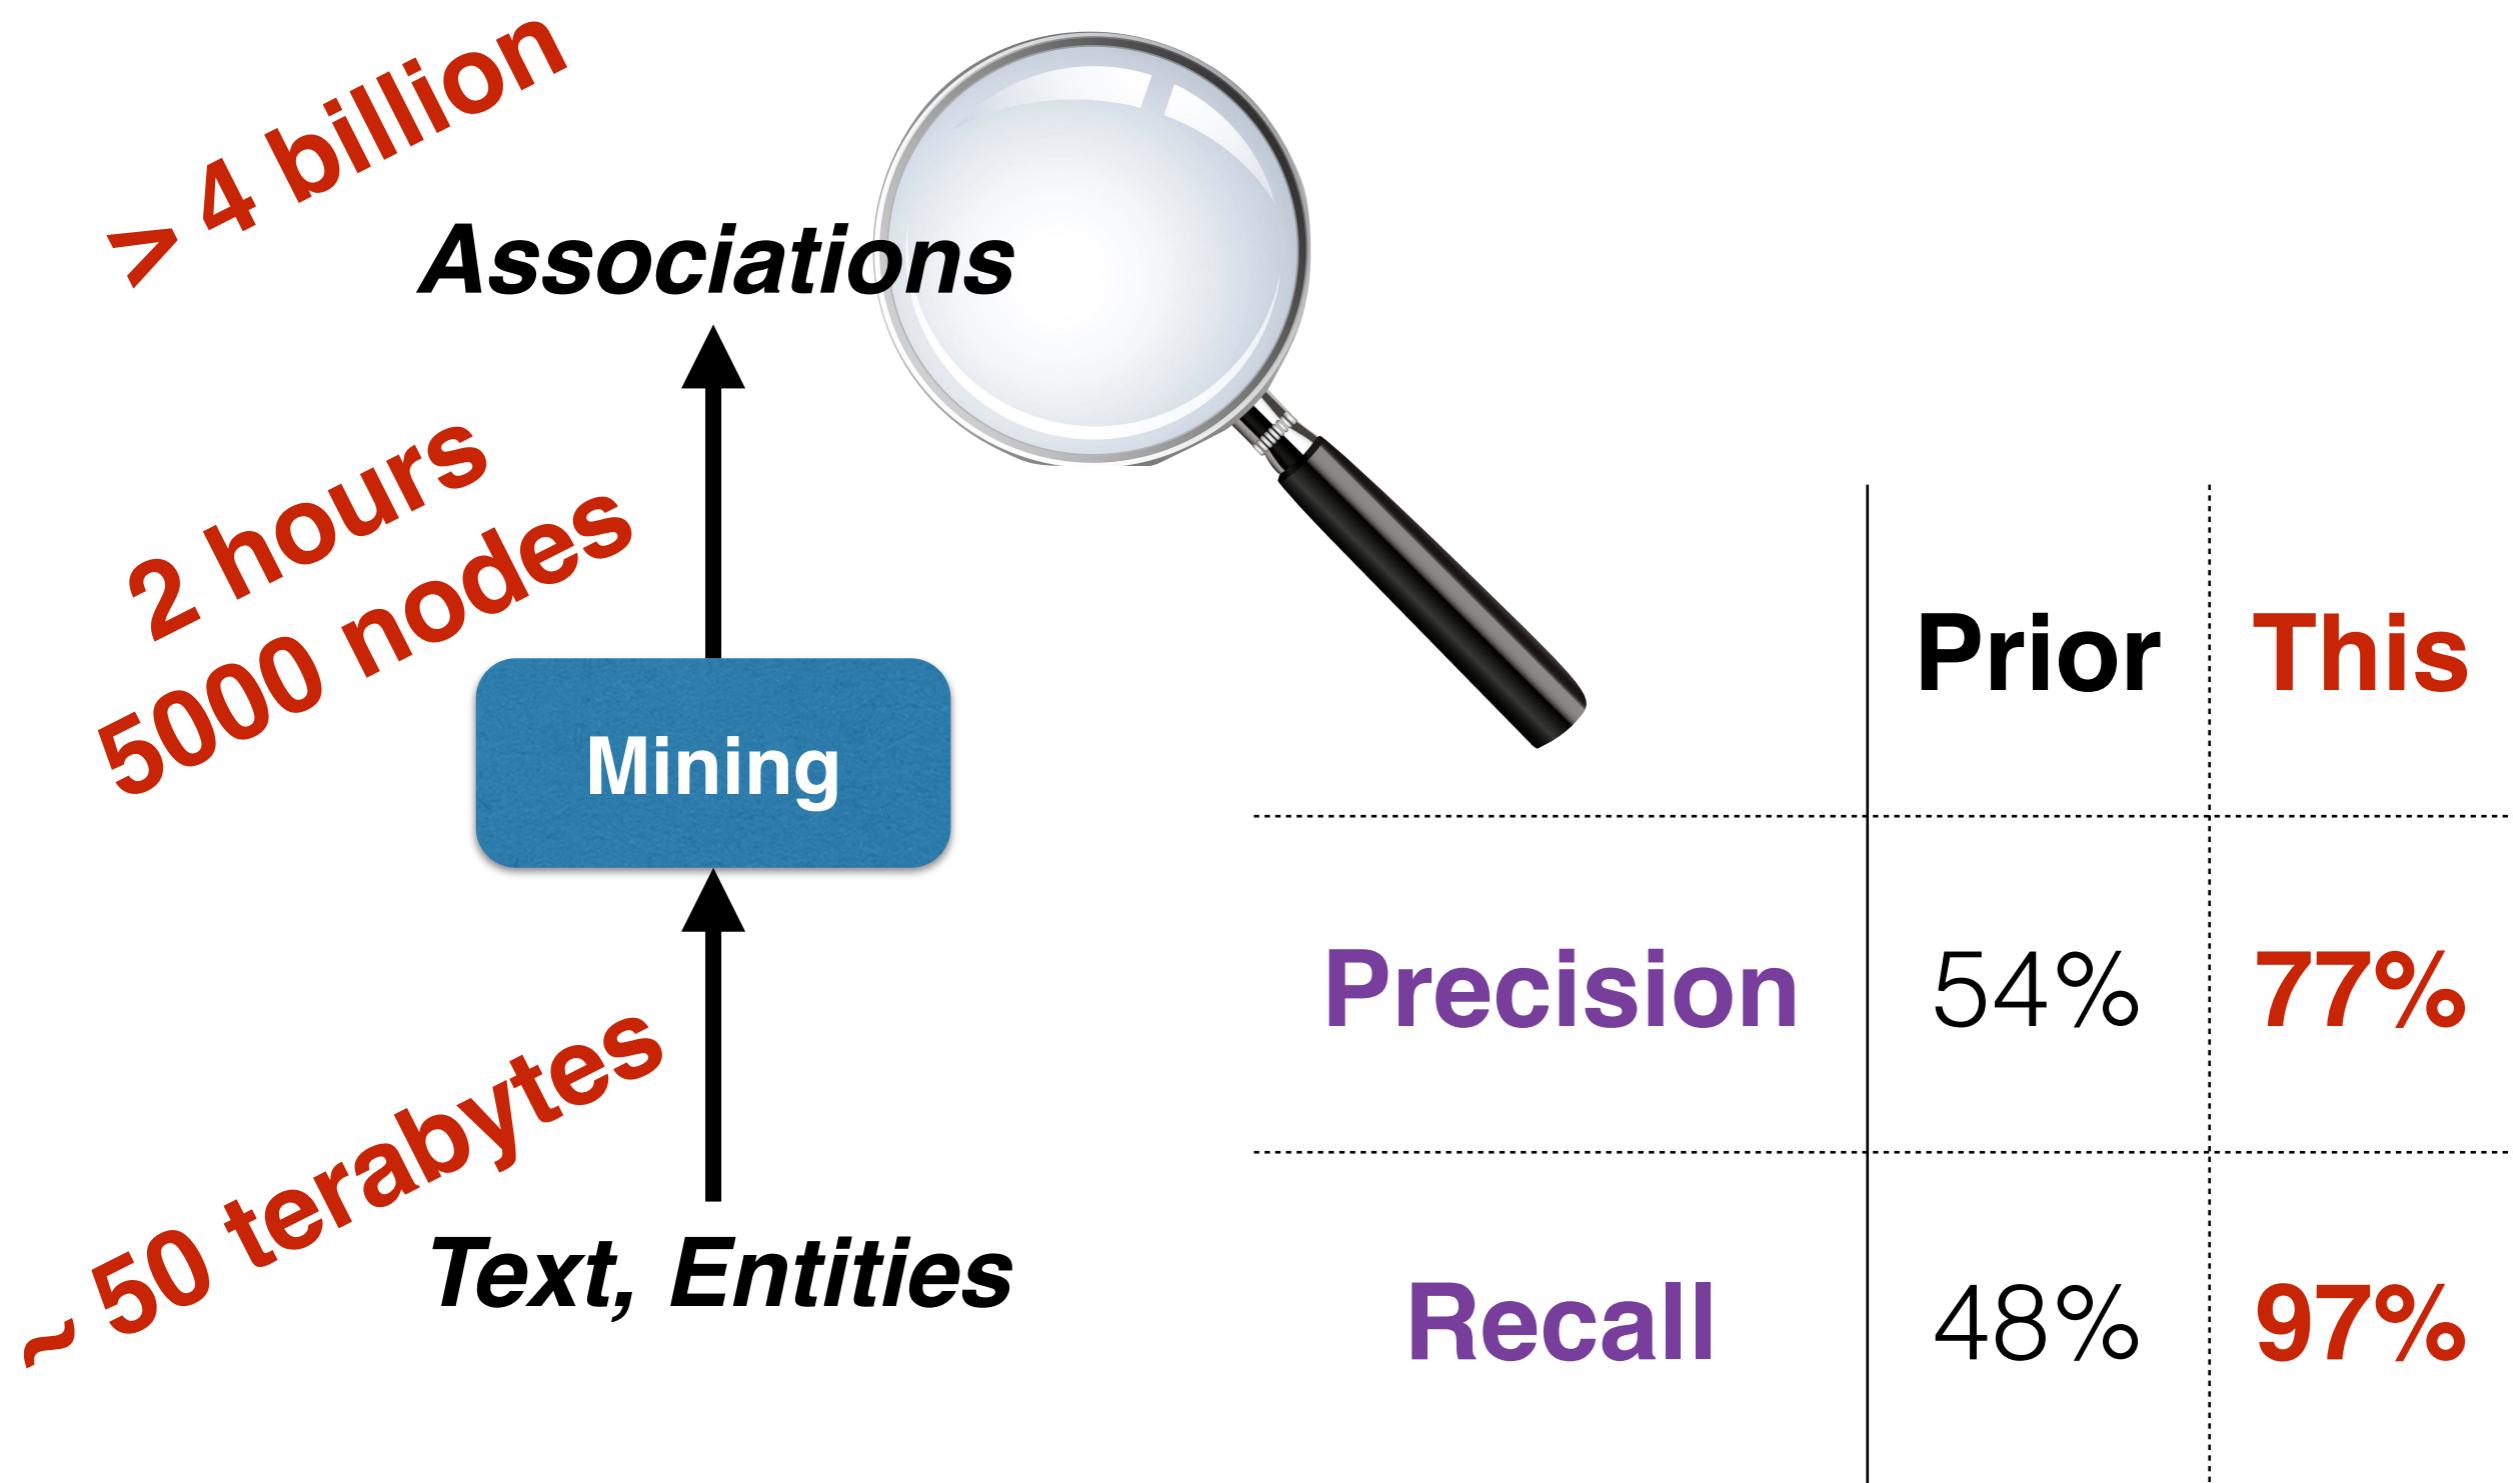
Extract Opinion Statements



Statistics and Results



Statistics and Results



More Details

SIGMOD '15

Mining Subjective Properties on the Web

| | | |
|--|--|--|
| Immanuel Trummer* EPFL Lausanne, Switzerland immanuel.trummer@epfl.ch | Alon Halevy Google, Inc. Mountain View, USA halevy@google.com | Hongrae Lee Google, Inc. Mountain View, USA hrlee@google.com |
| Sunita Sarawagi Google, Inc. and IIT Bombay Mountain View, USA/Bombay, India sarawagi@google.com | Rahul Gupta Google, Inc. Mountain View, USA grahul@google.com | |

ABSTRACT

Even with the recent developments in Web search of answering queries from structured data, search engines are still limited to queries with an objective answer, such as EUROPEAN CAPITALS or WOODY ALLEN MOVIES. However, many queries are subjective, such as SAFE CITIES, or CUTE ANIMALS. The underlying knowledge bases of search engines do not contain answers to these queries because they do not have a ground truth. We describe the SURVEYOR system that mines the dominant opinion held by authors of Web content about whether a subjective property applies to a given entity. The evidence on which SURVEYOR relies are statements extracted from Web text that either support the property or claim its

Categories and Subject Descriptors

H.4 [Information Systems Applications]: Miscellaneous

Keywords

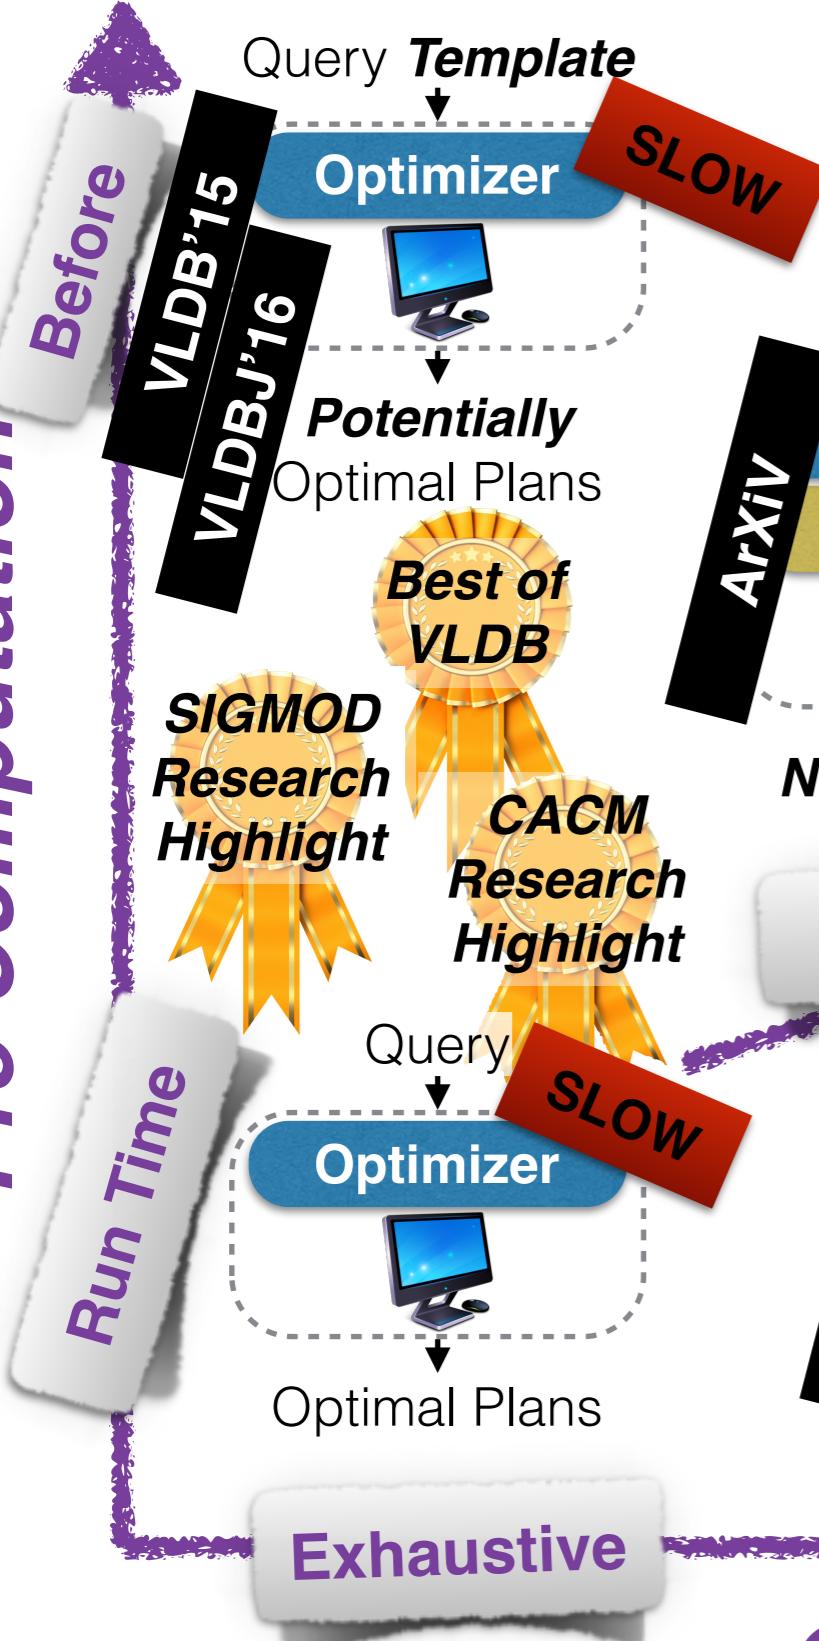
Text mining; subjective properties; user behavior model

1. INTRODUCTION

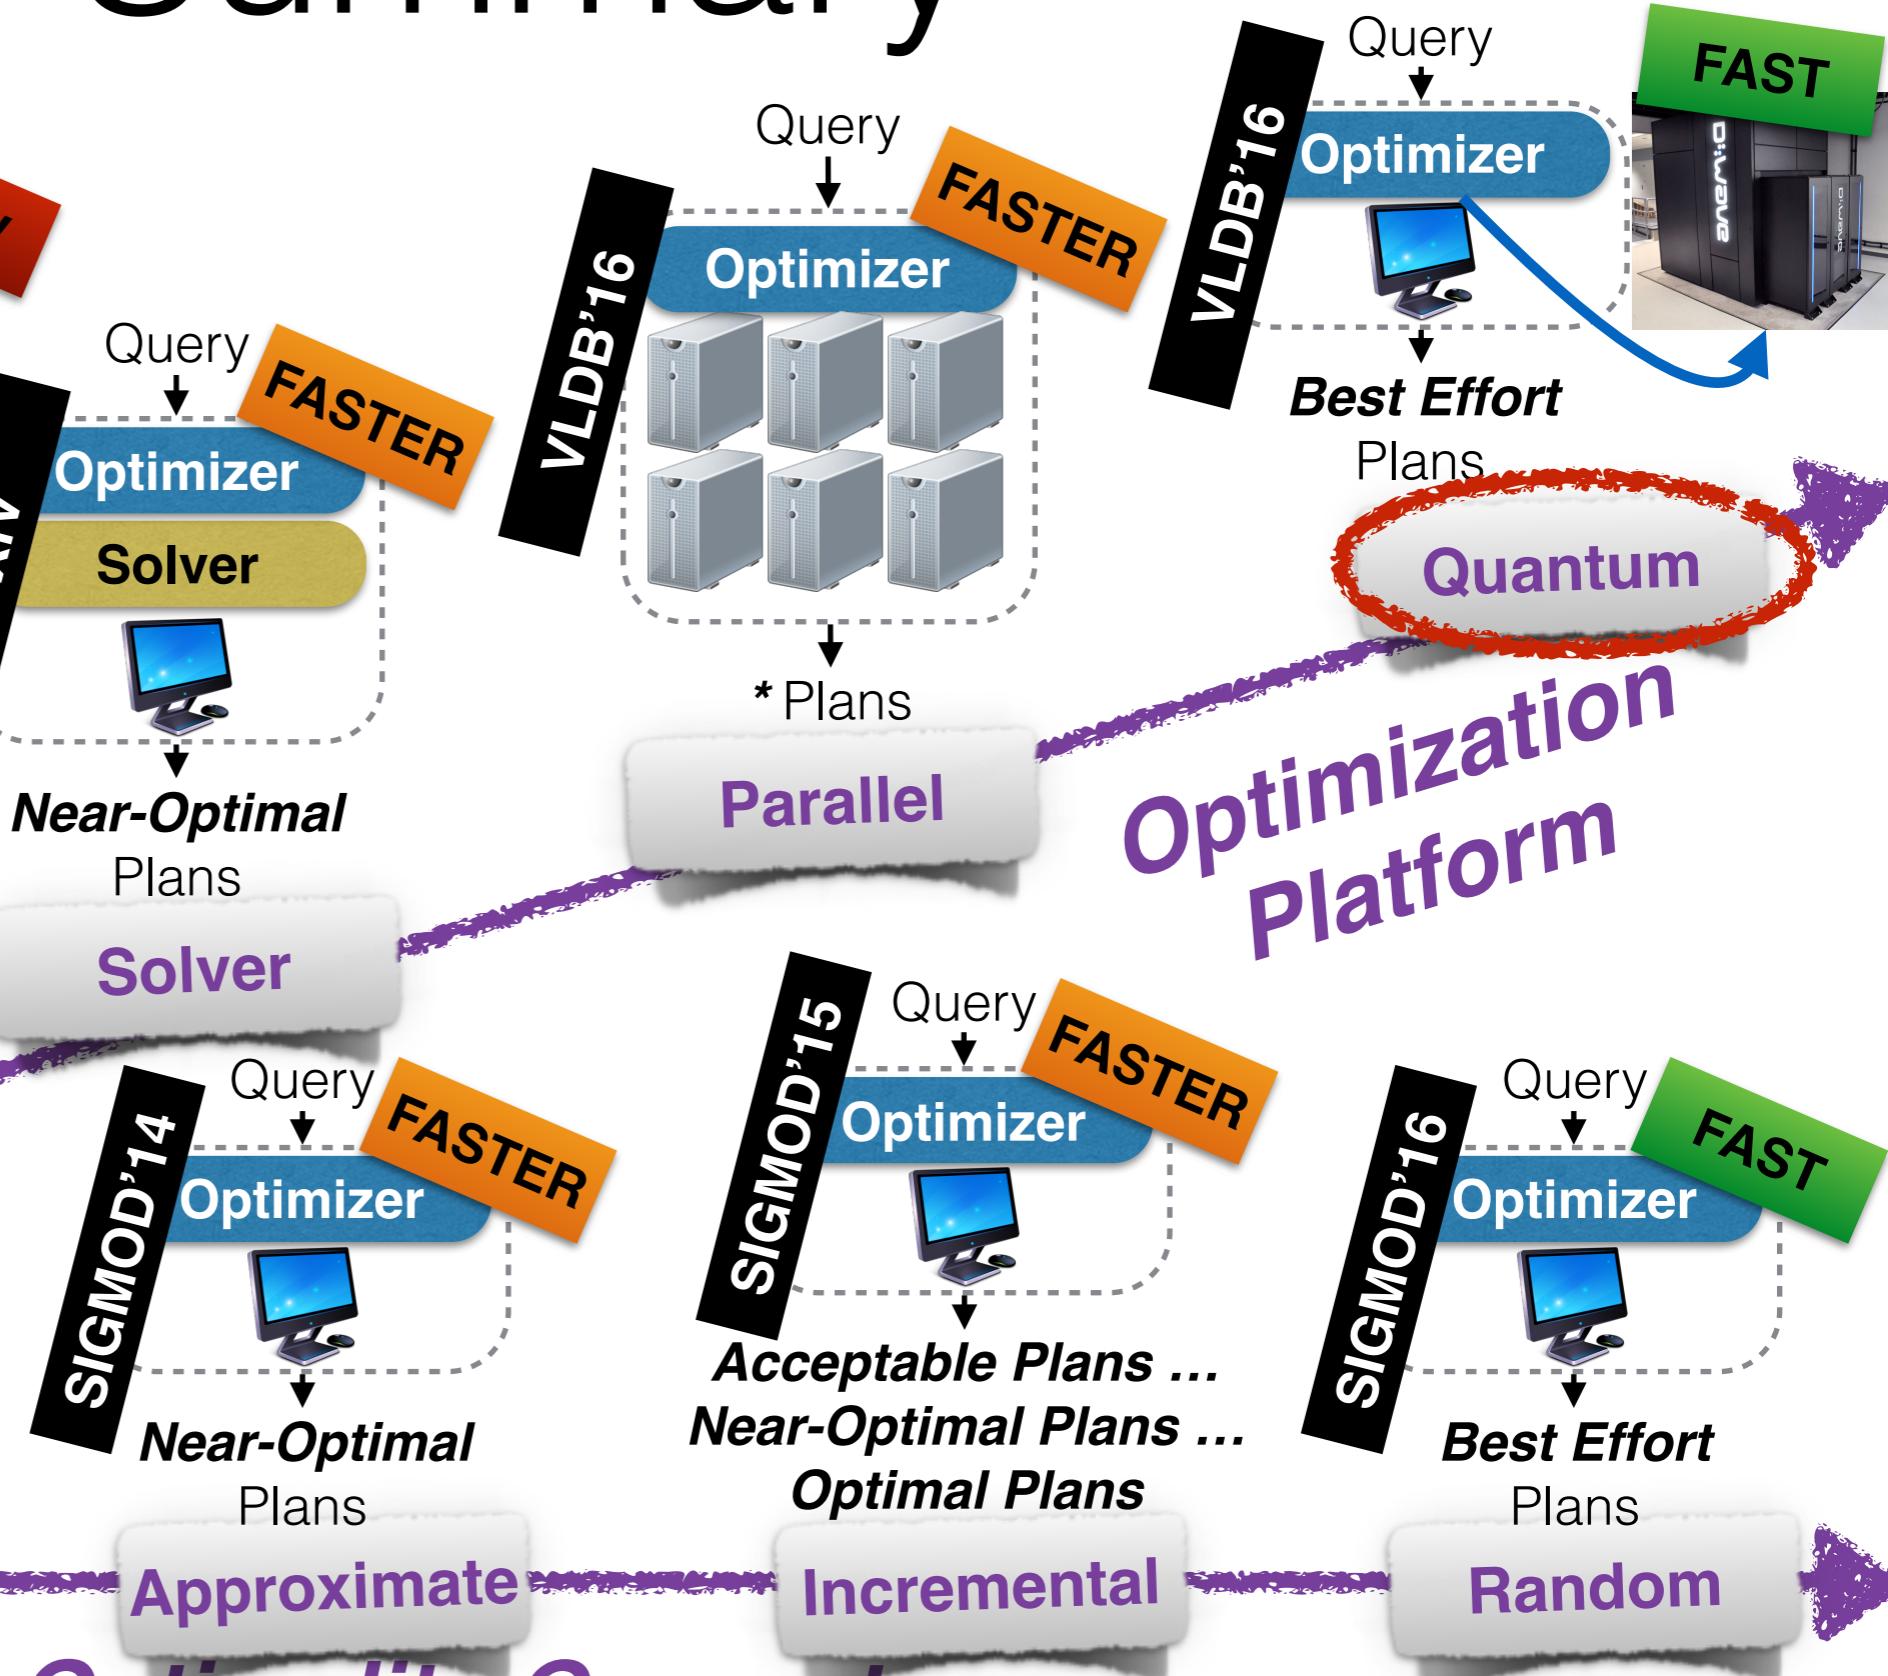
In recent years, Web search engines have invested heavily in answering queries with structured data. For example, queries such as WOODY ALLEN MOVIES or AMERICAN PRESIDENTS will yield a display of the appropriate entities. These queries are enabled by large knowledge bases about impor-

Summary

Pre-Computation



Optimality Guarantees





www.itrummer.org